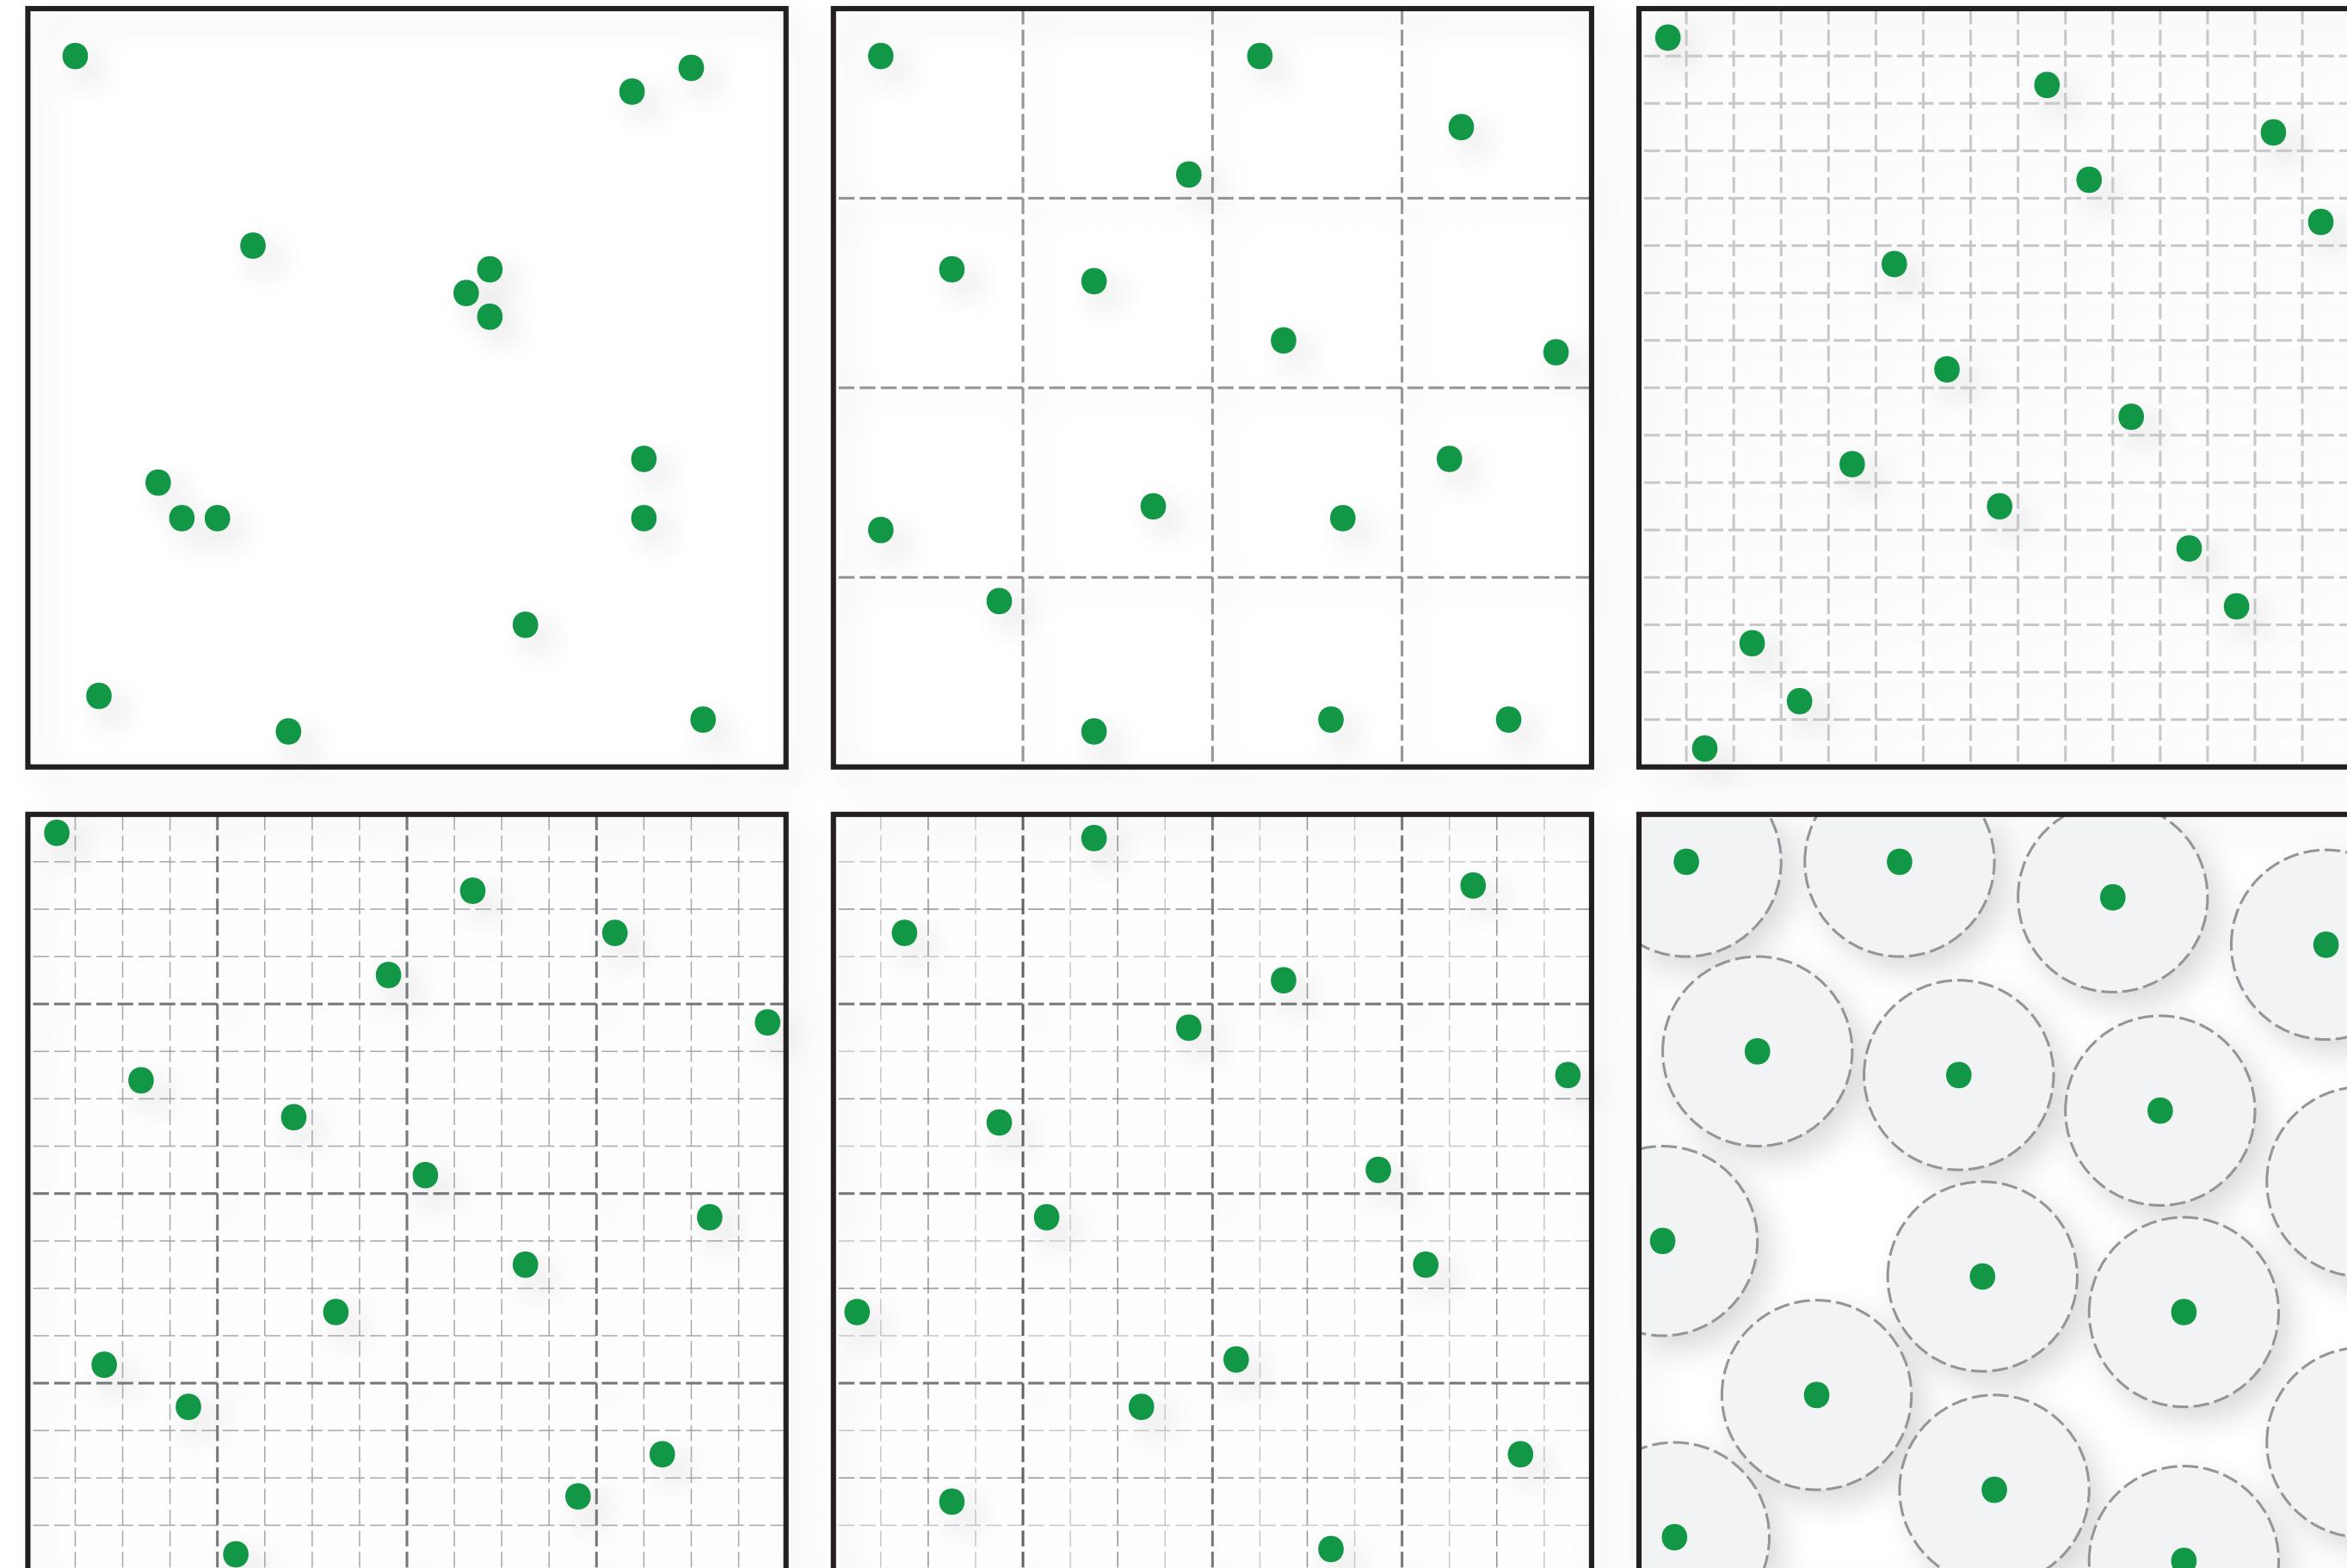


POPULAR SAMPLING PATTERNS

Fourier Analysis of Numerical Integration in Monte Carlo Rendering



Dartmouth

Wojciech Jarosz
wjarosz@dartmouth.edu

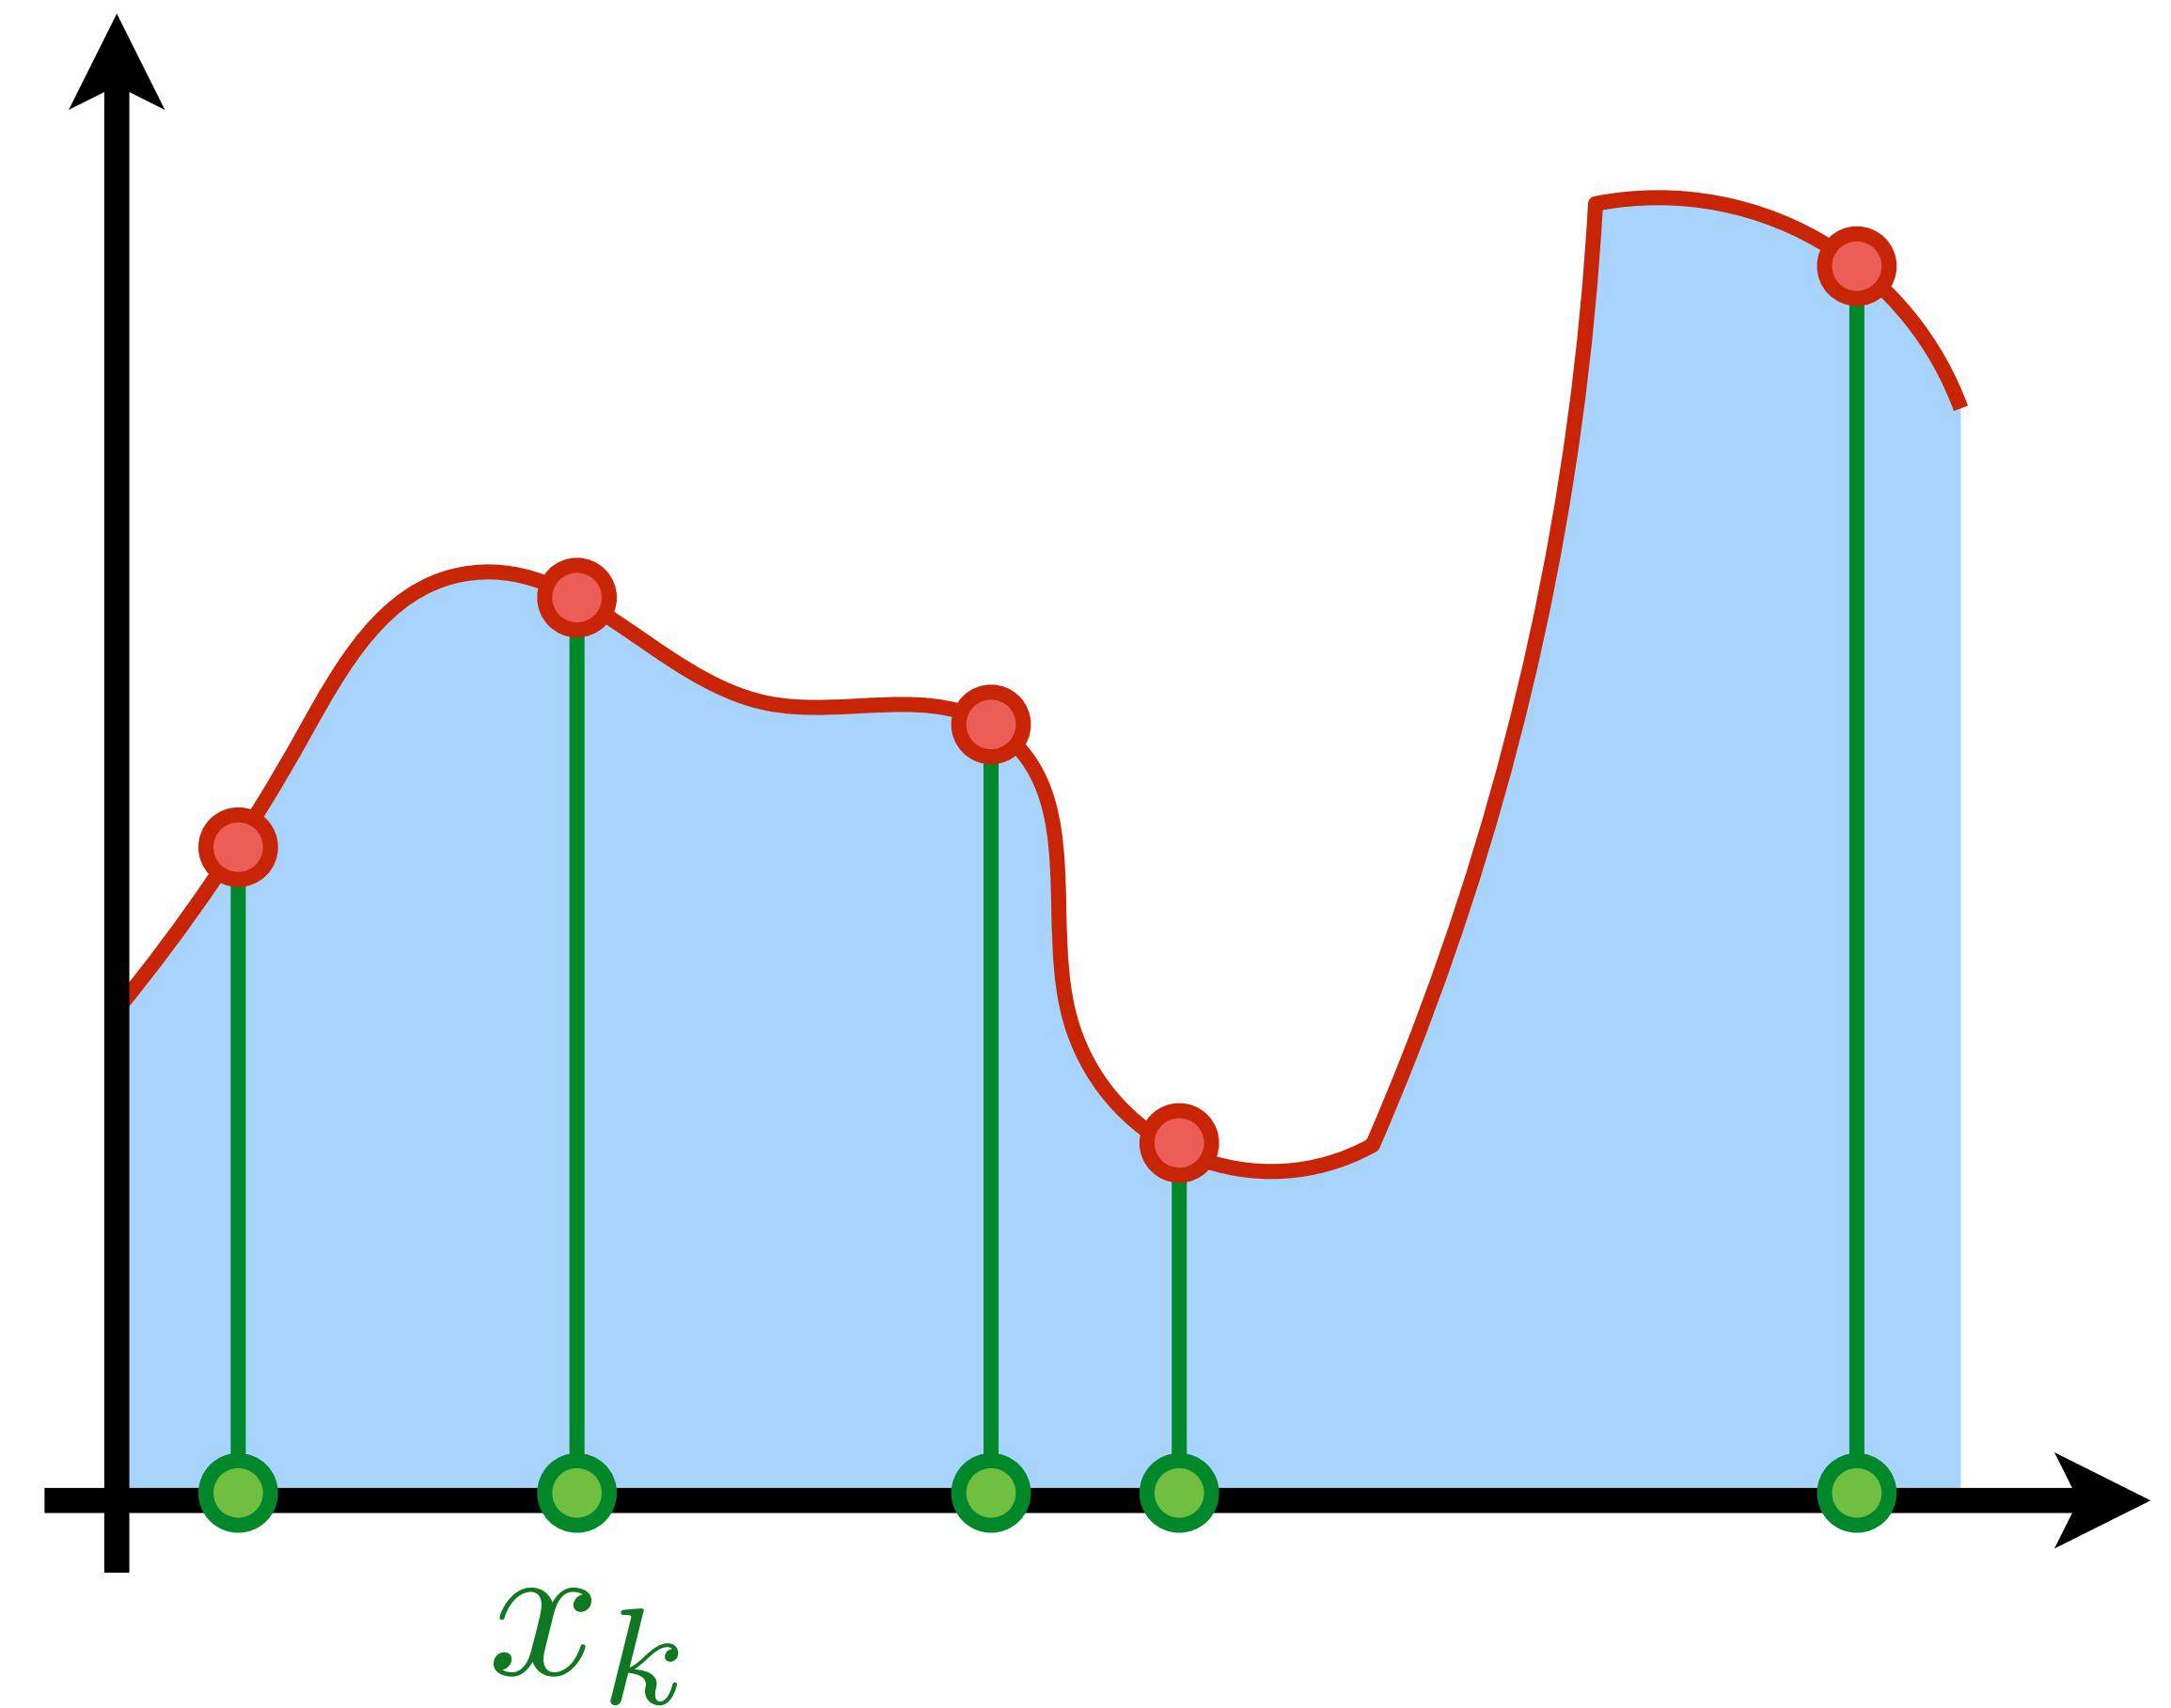
✓ CL

Recall: Monte Carlo Integration

$$I \approx \int_D f(x) S(x) dx$$

$$S(x) = \frac{1}{N} \sum_{k=1}^N \delta(x - x_k)$$

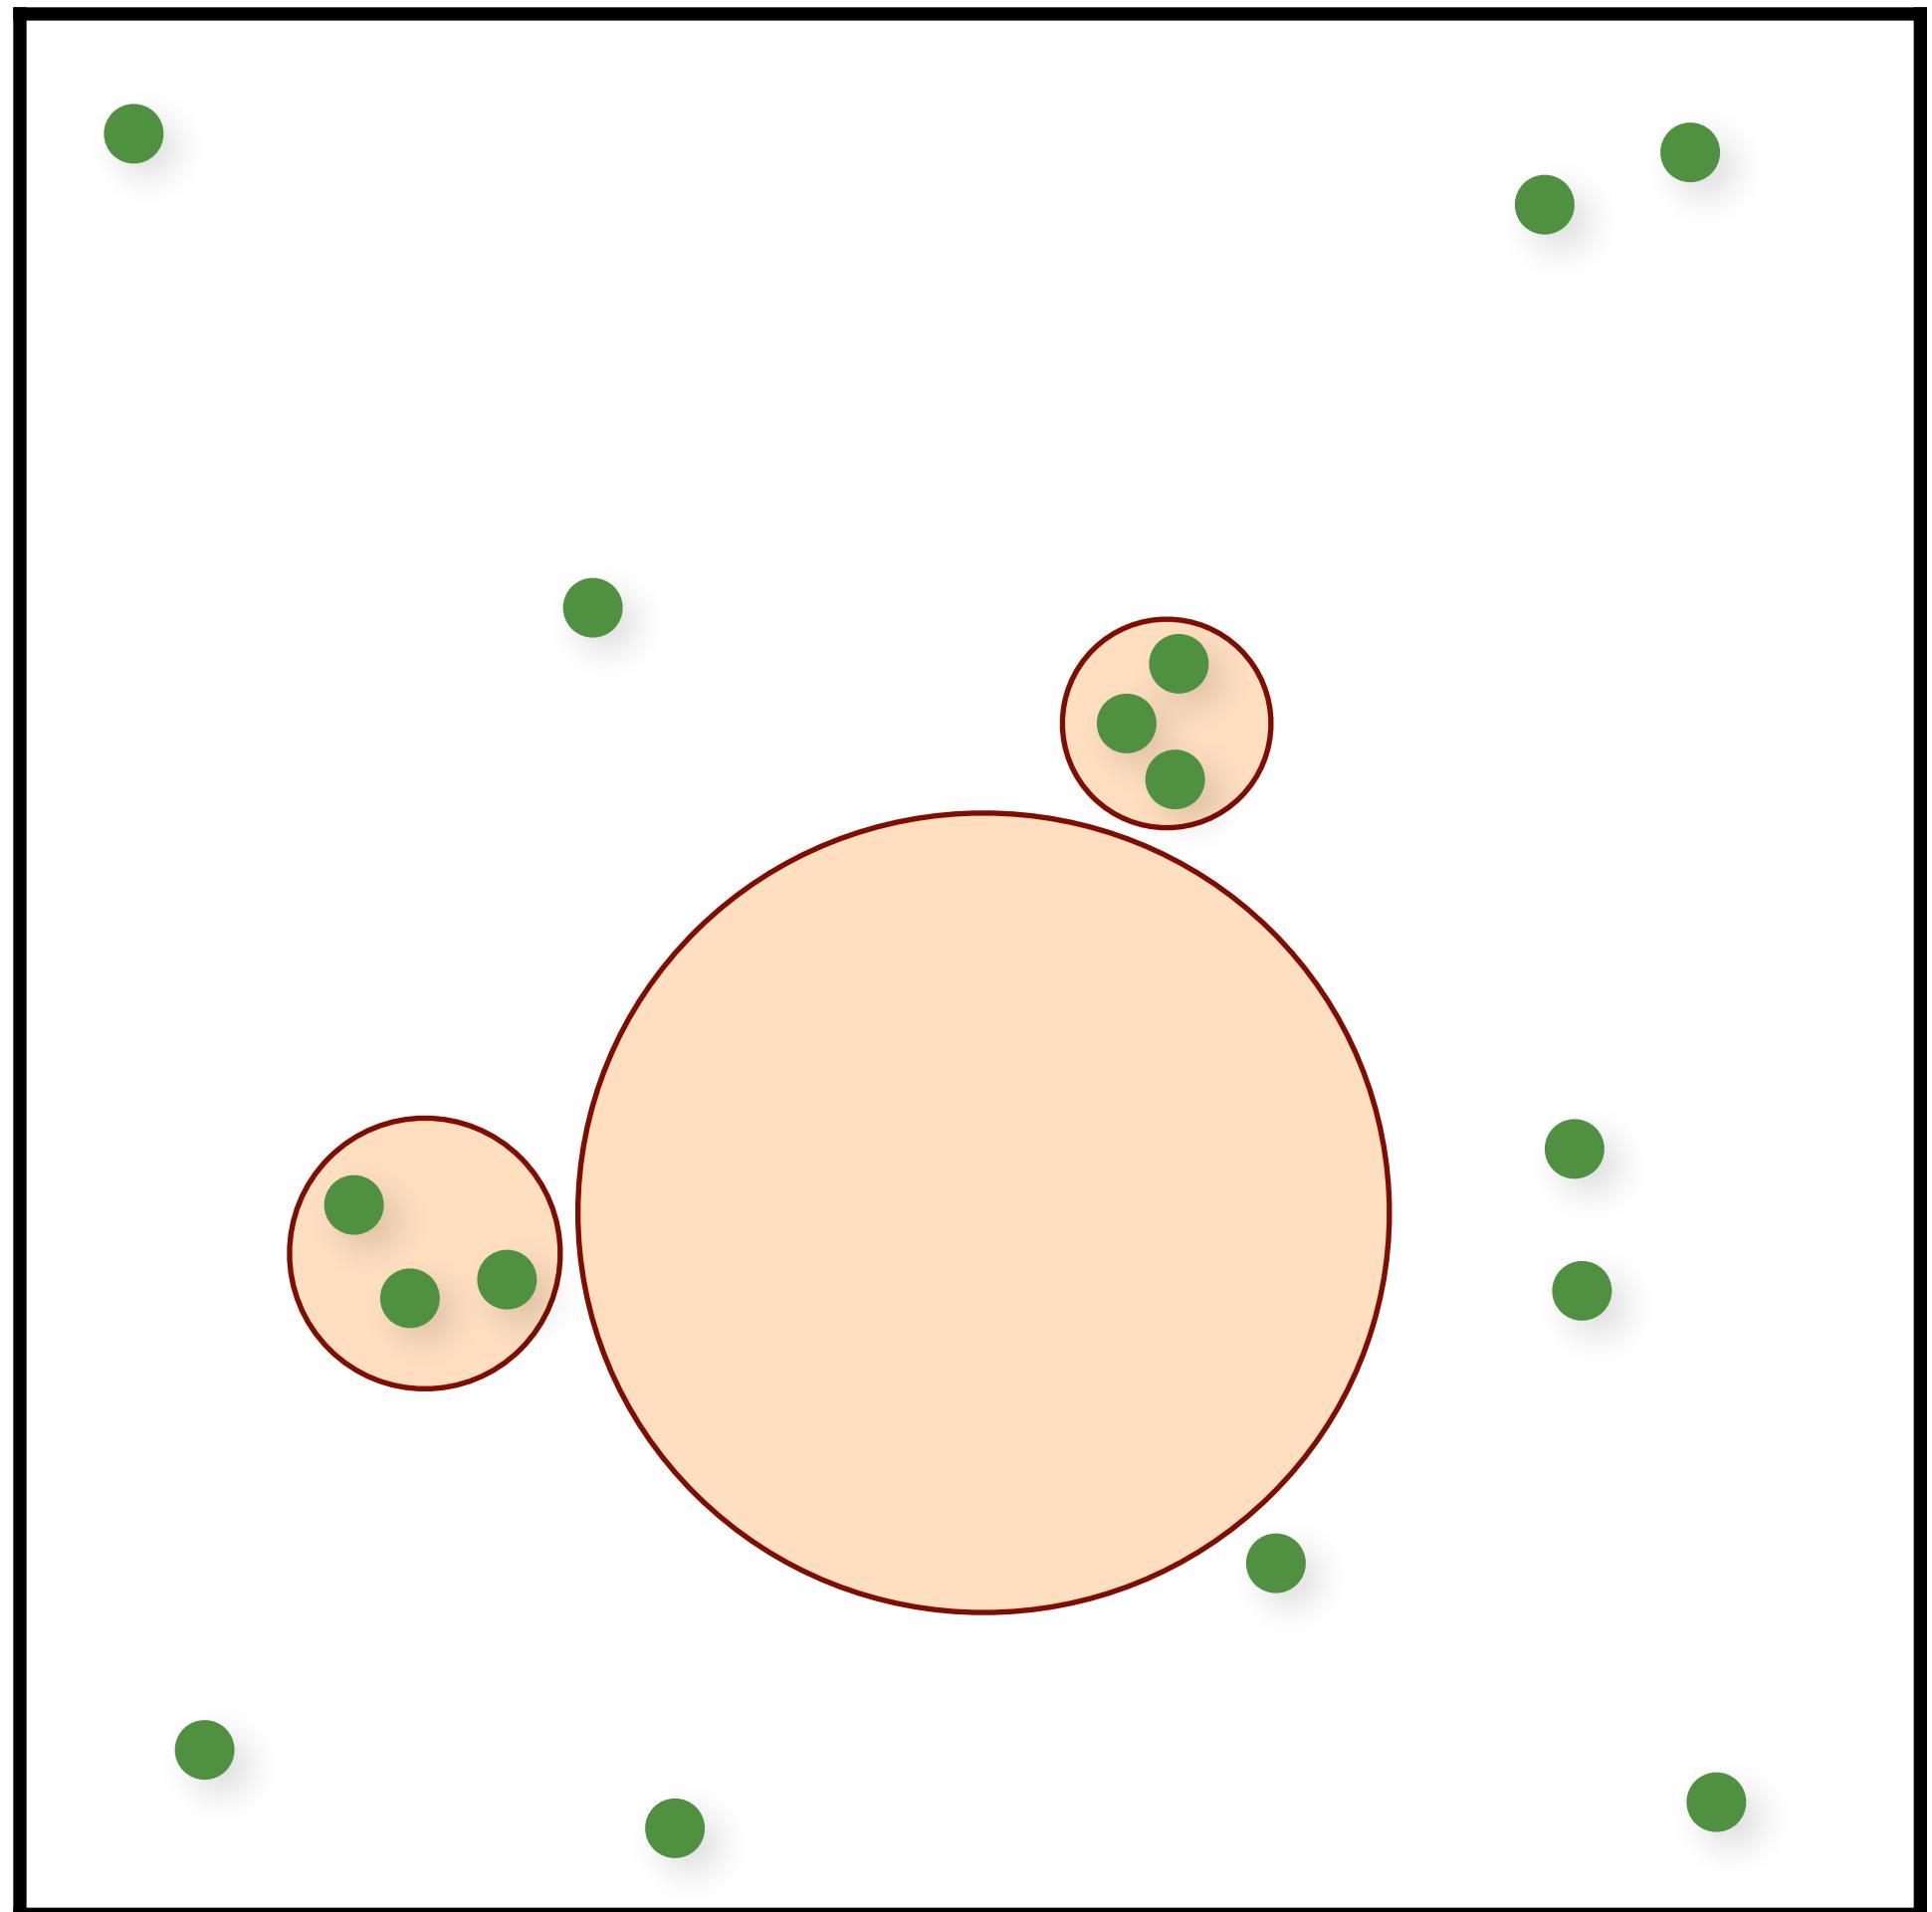
How to generate the locations x_k ?



Independent Random Sampling

```
for (uint i = 0; i < num; i++)  
{  
    samples(i).x = randf();  
    samples(i).y = randf();  
}
```

- ✓ Trivially extends to higher dimensions
- ✓ Trivially progressive and memory-less
- ✗ Big gaps
- ✗ Clumping

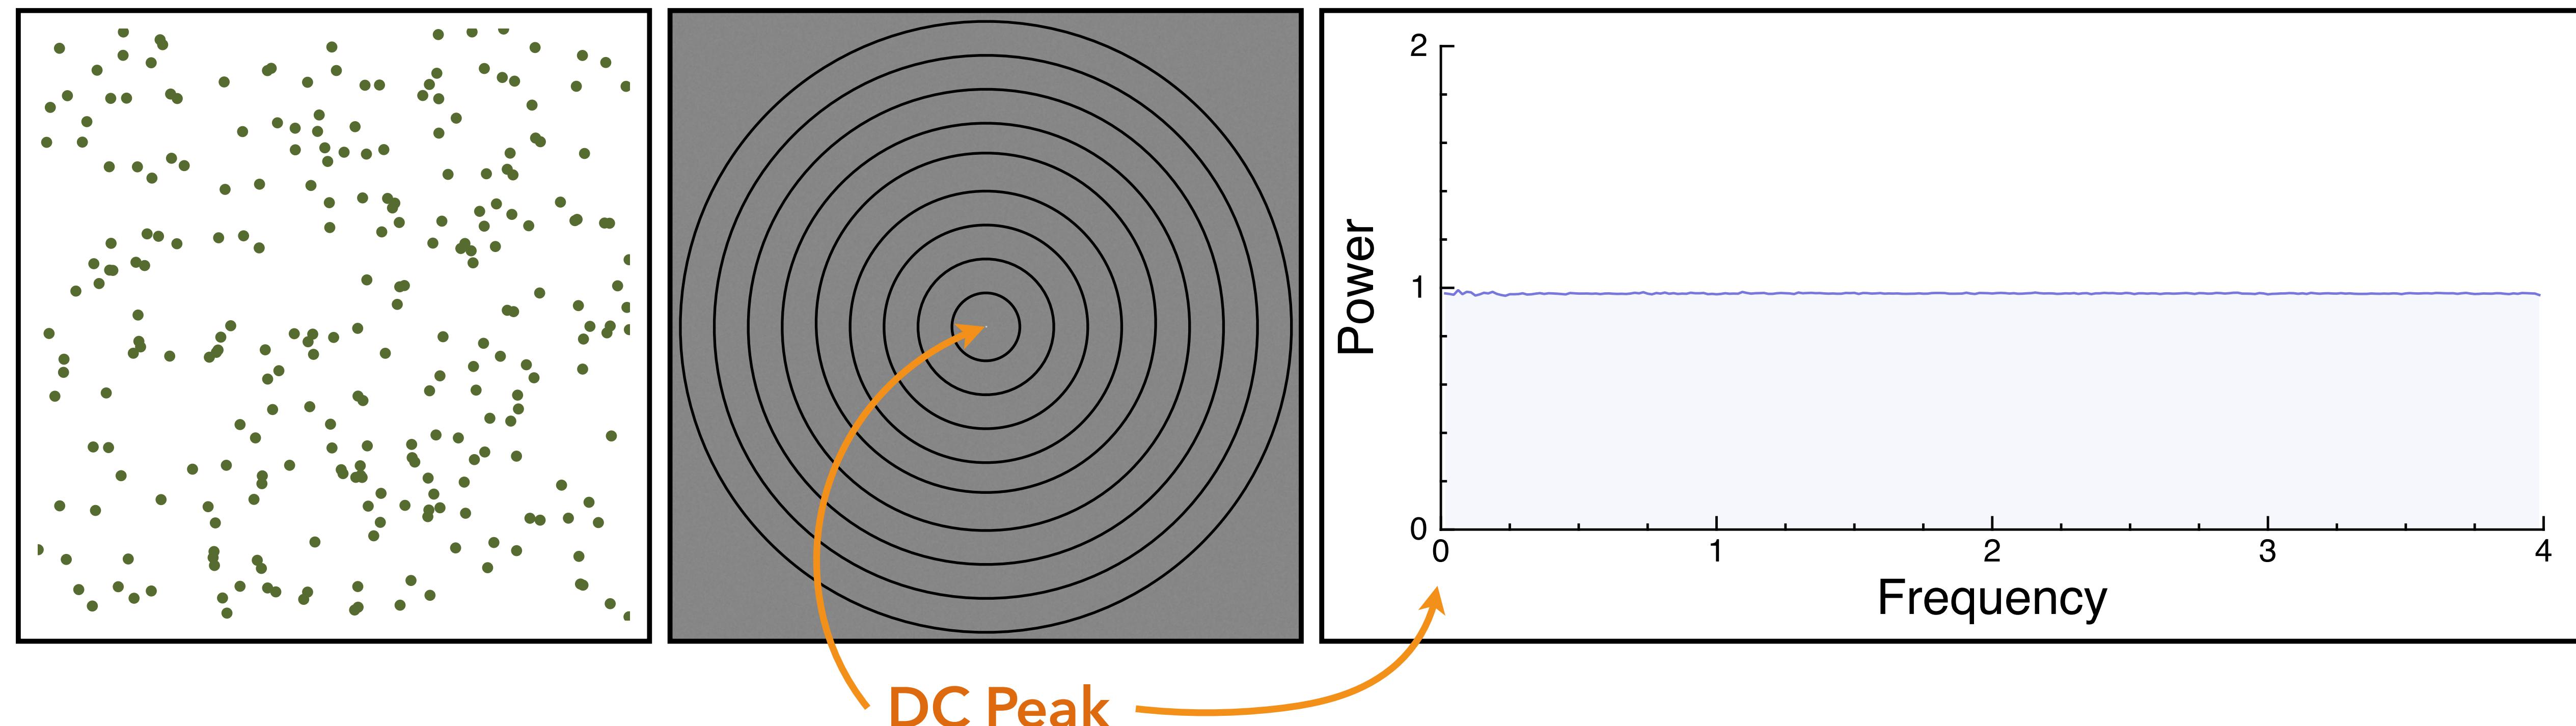


Independent Random Sampling

Samples

Power spectrum

Radial mean



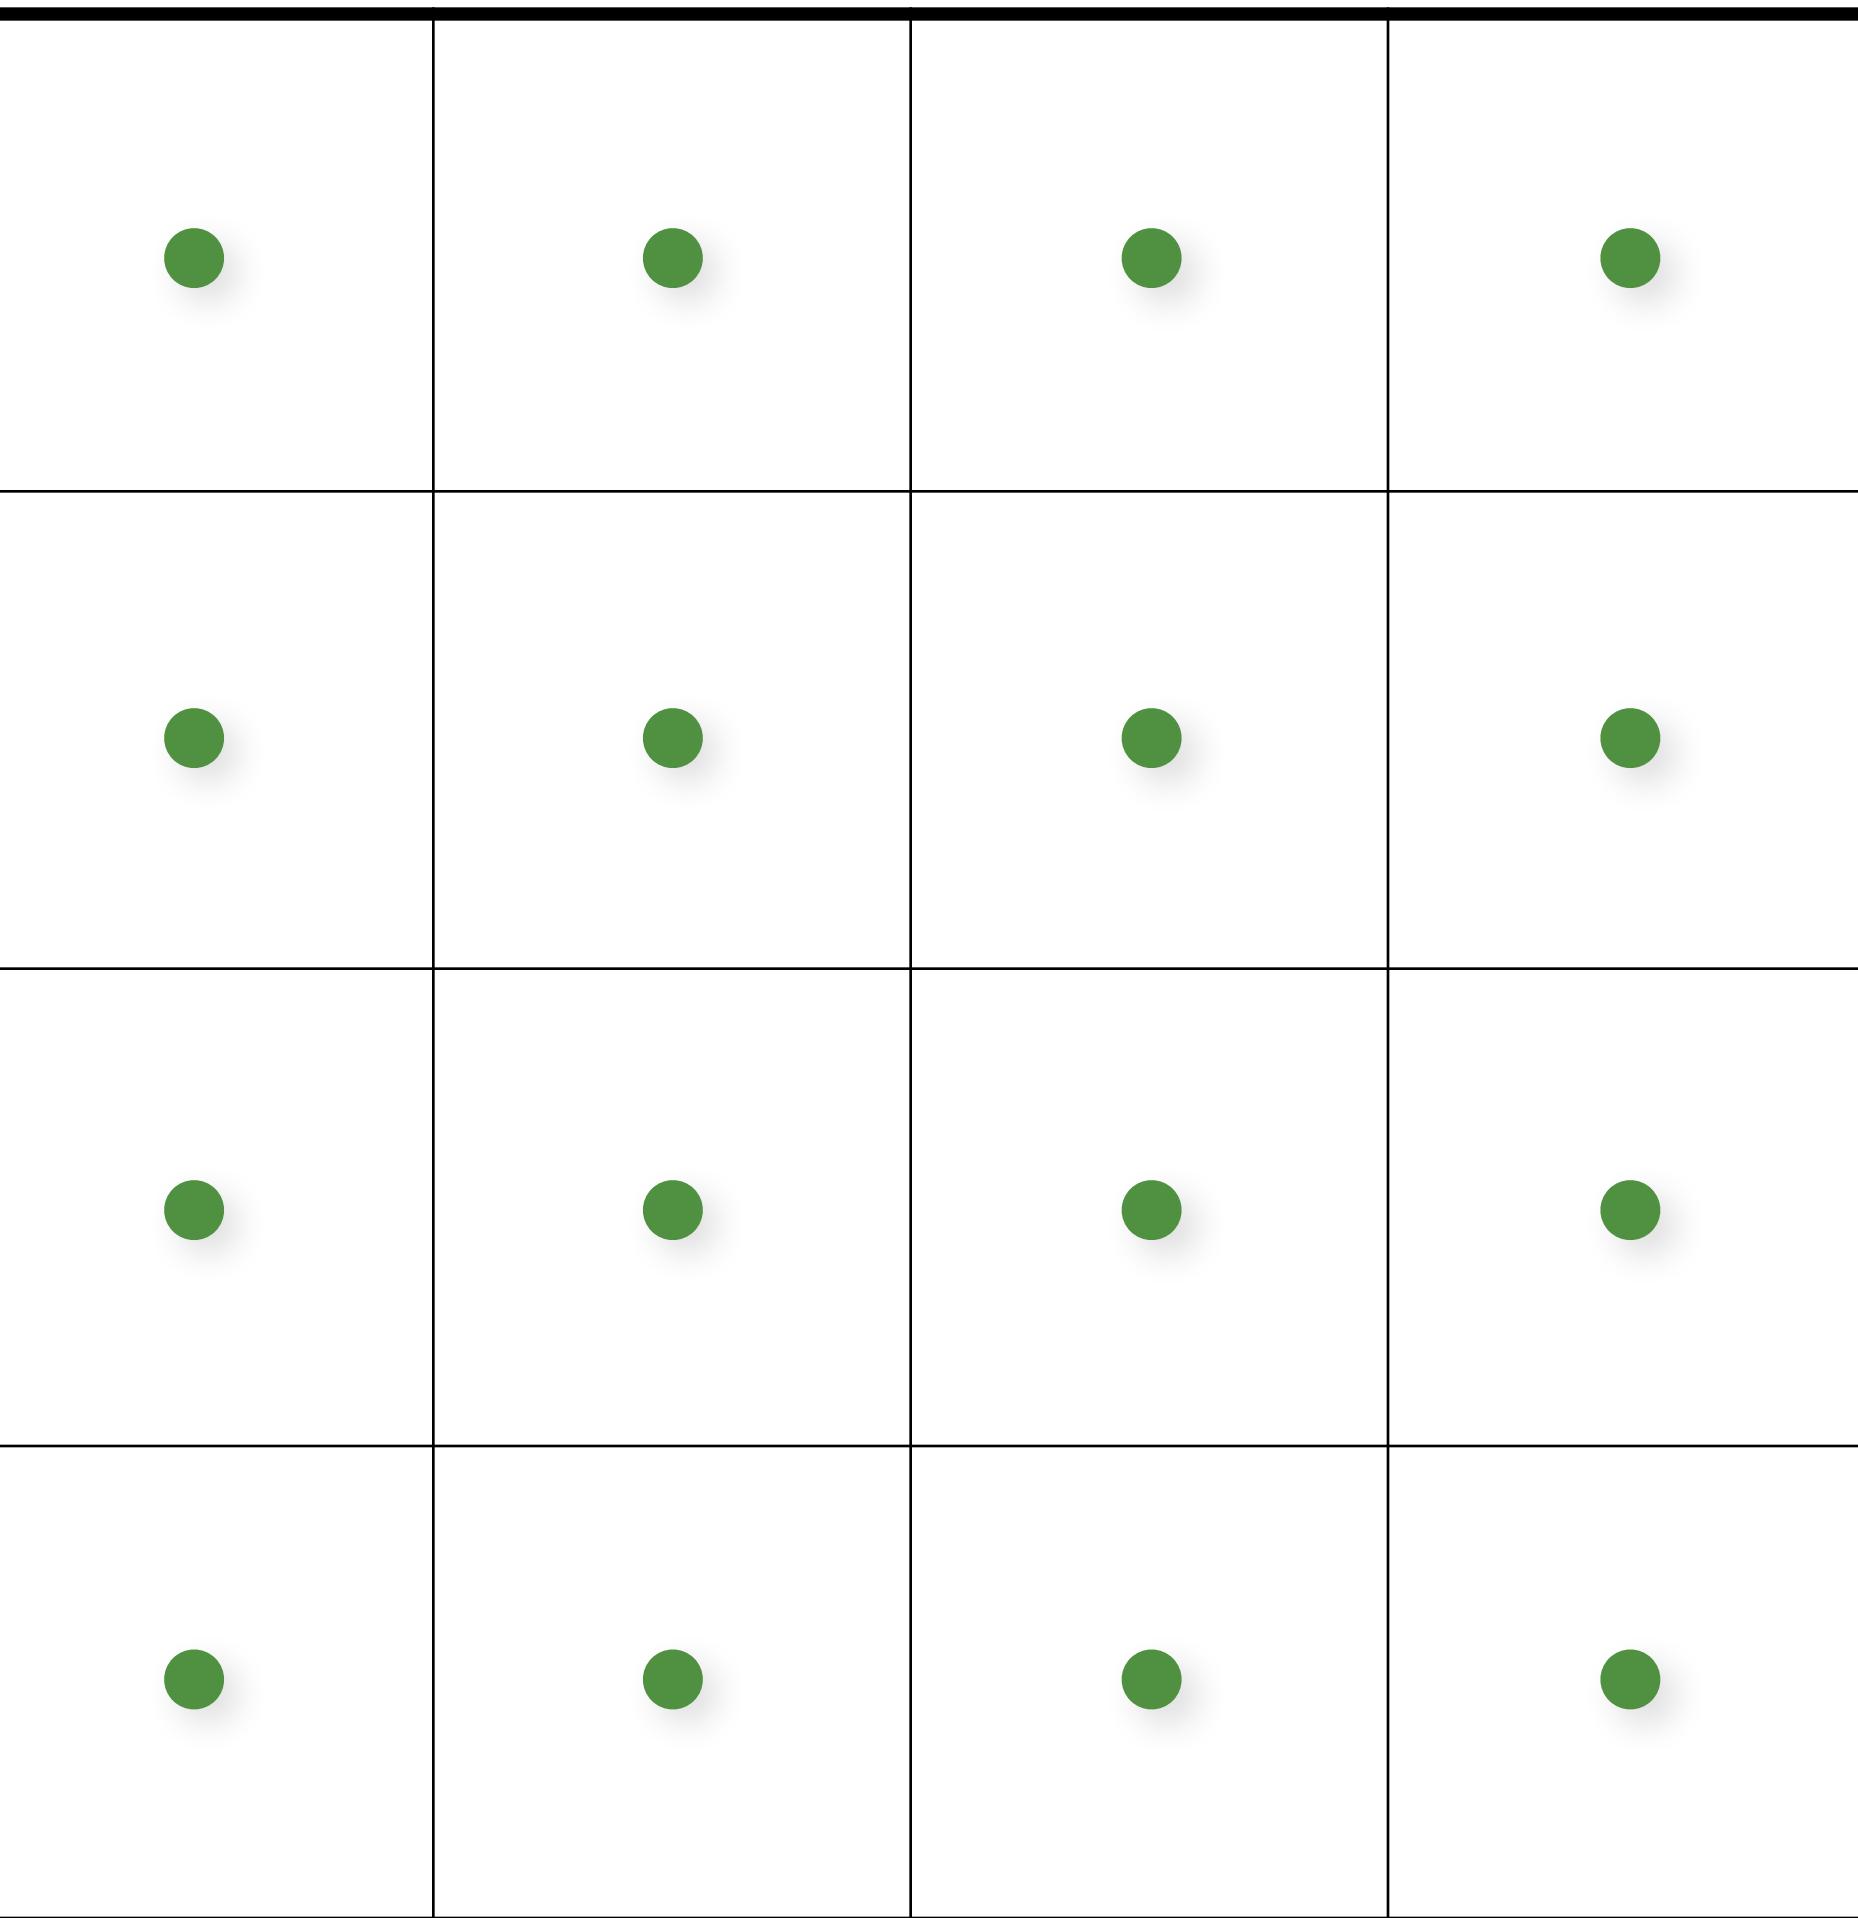
Regular Sampling

```
for (uint i = 0; i < numX; i++)  
    for (uint j = 0; j < numY; j++)  
    {  
        samples(i,j).x = (i + 0.5)/numX;  
        samples(i,j).y = (j + 0.5)/numY;  
    }
```

✓ Extends to higher dimensions, but...

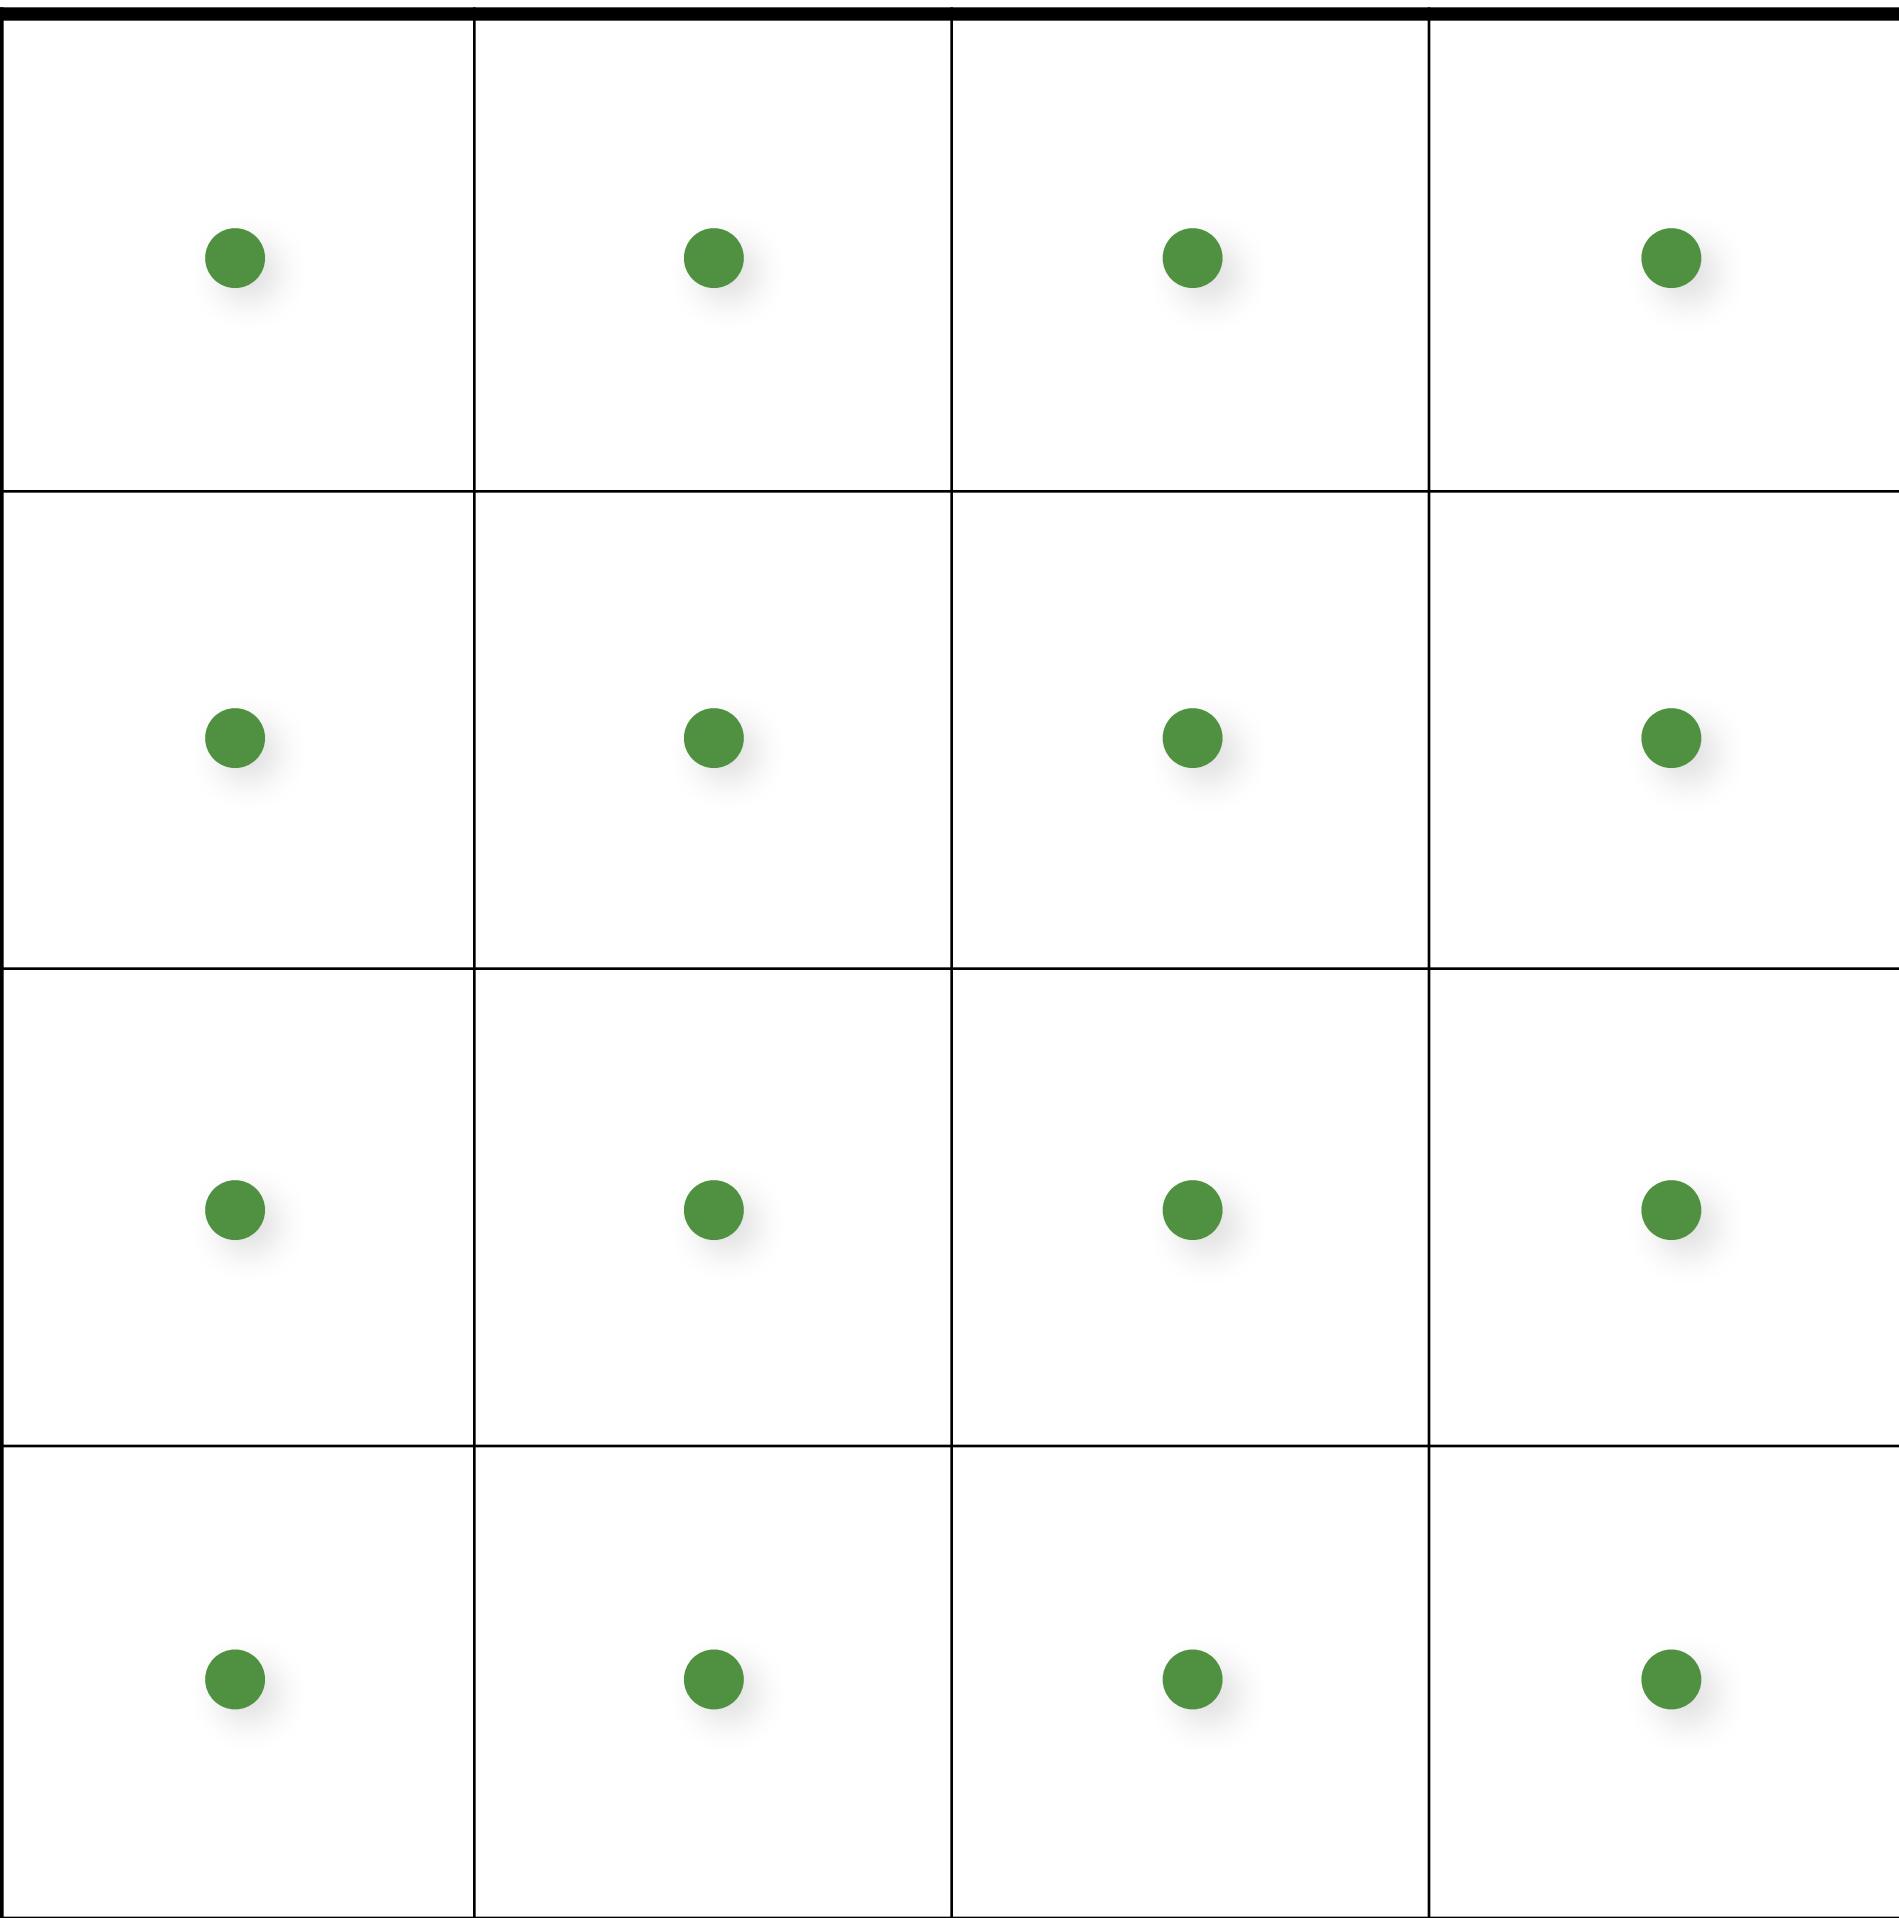
✗ Curse of dimensionality

✗ Aliasing



Regular Sampling

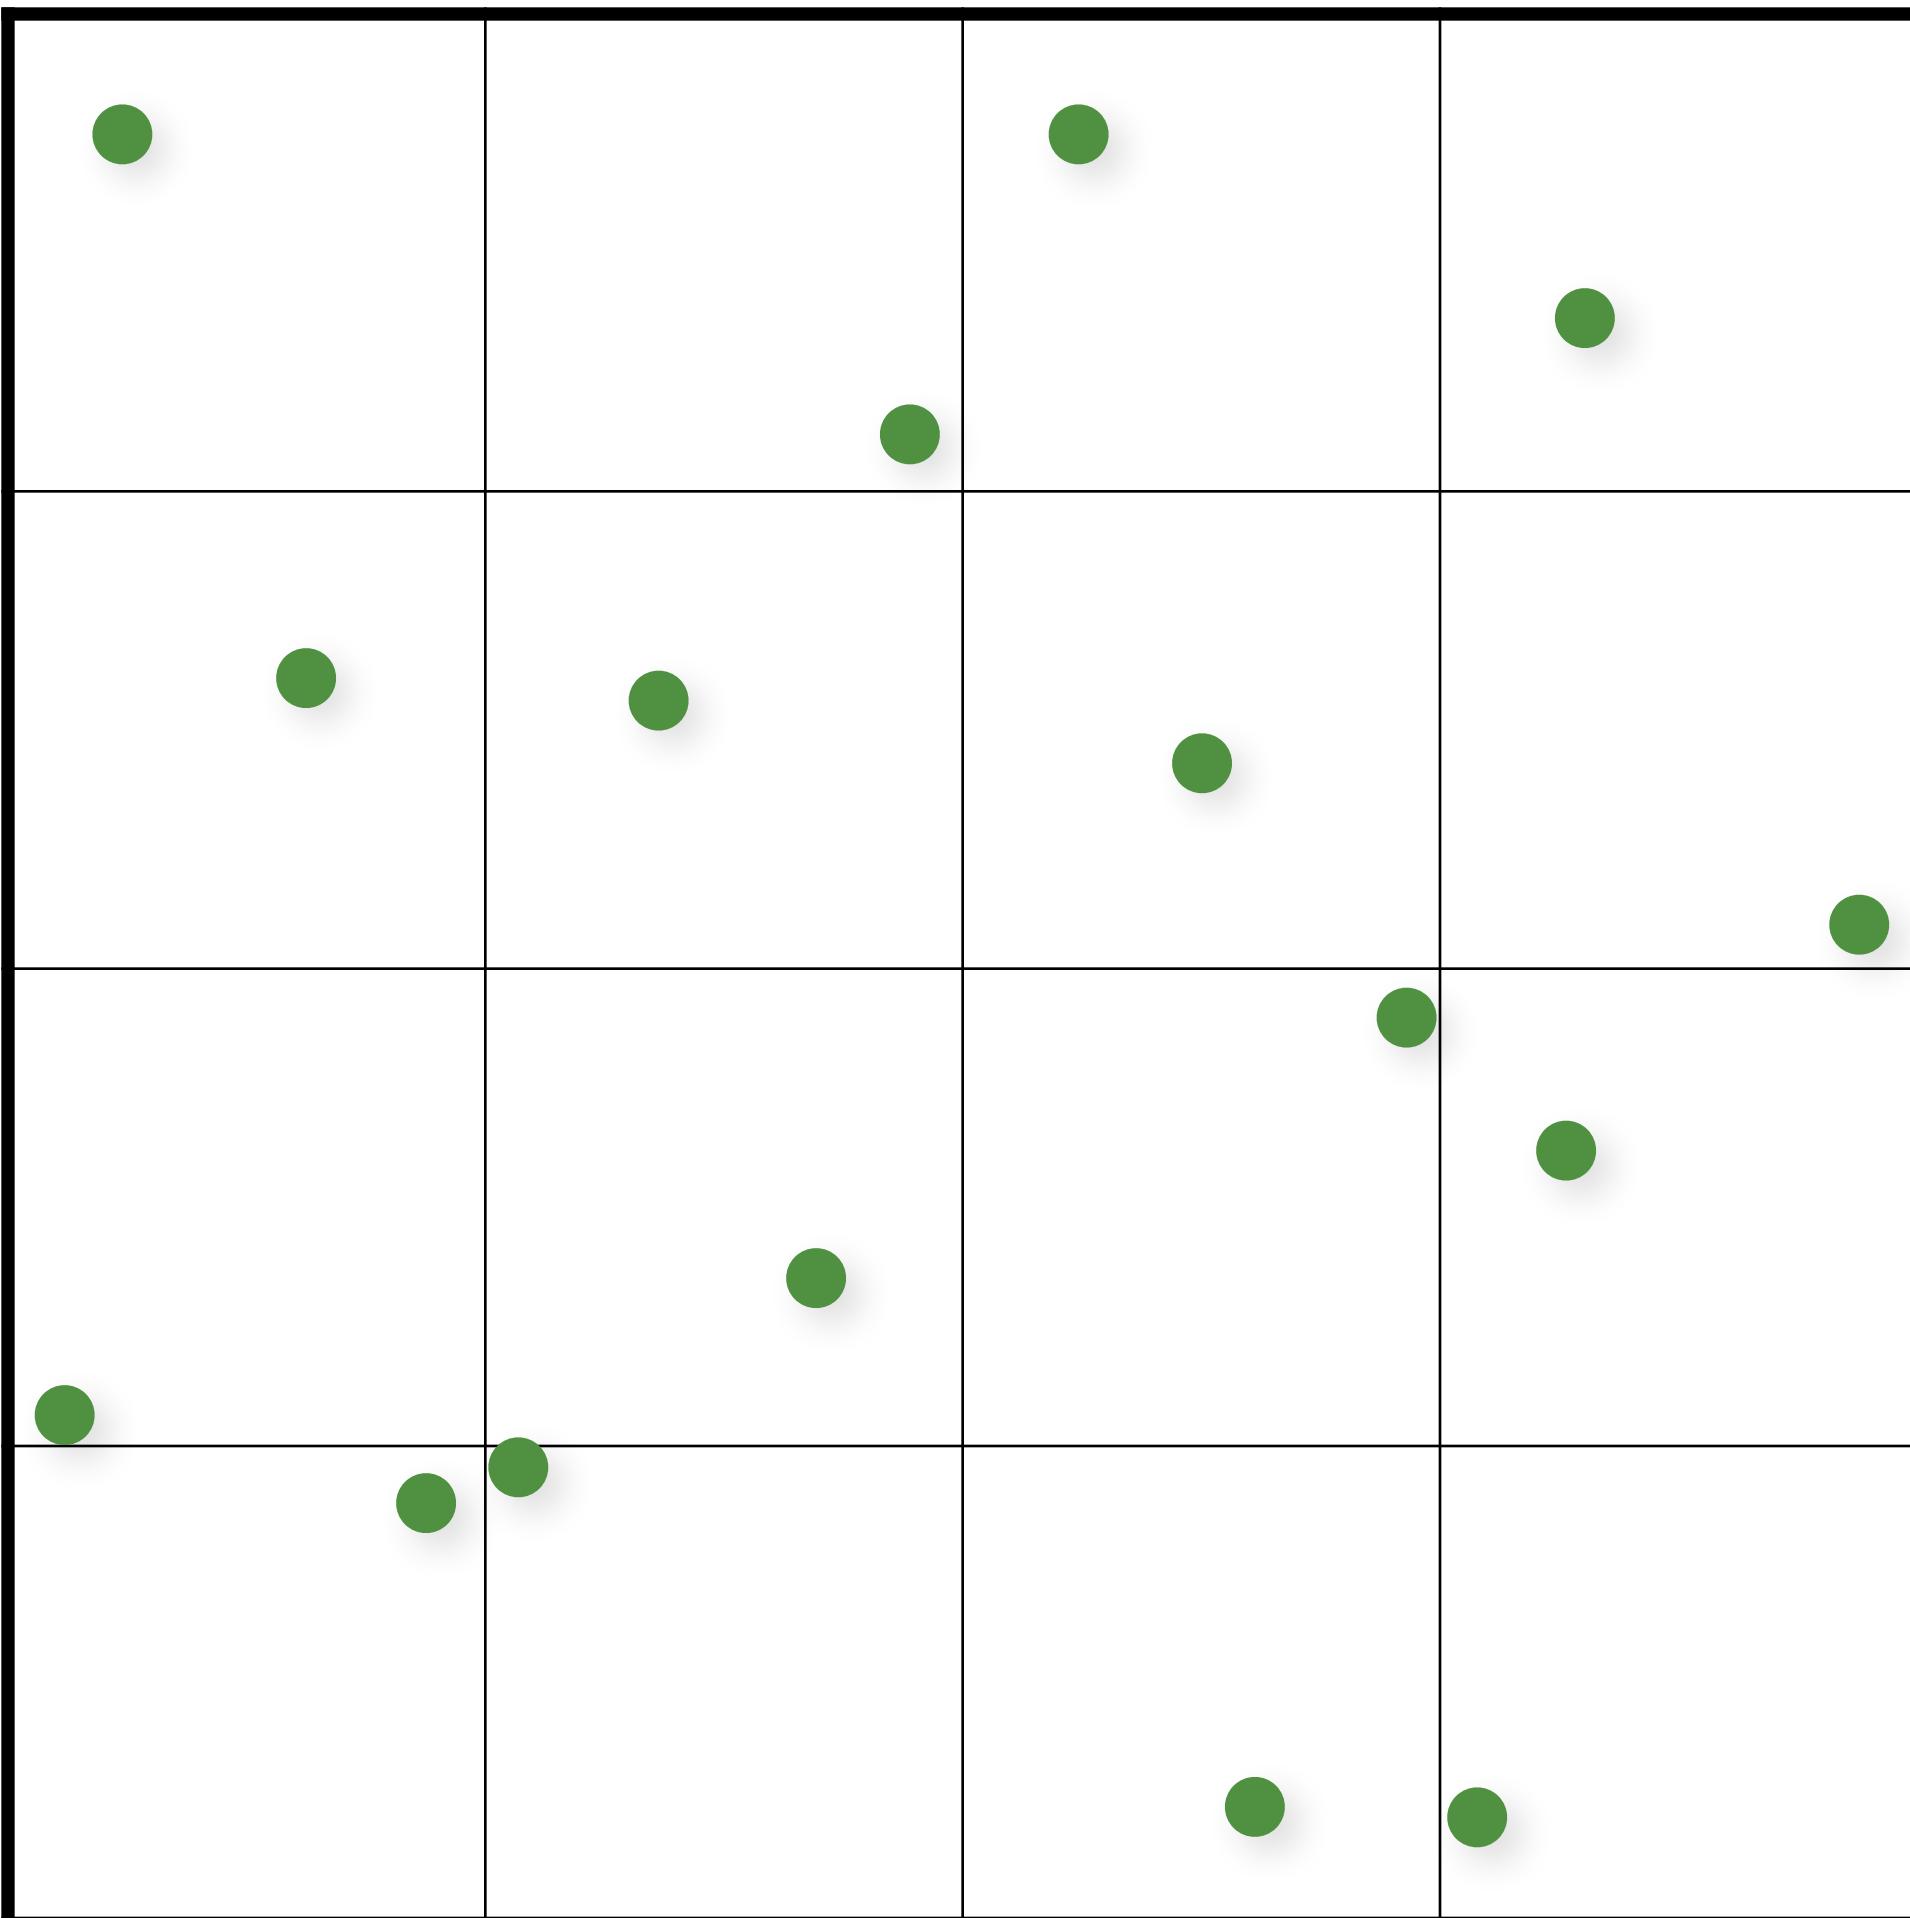
```
for (uint i = 0; i < numX; i++)  
    for (uint j = 0; j < numY; j++)  
    {  
        samples(i,j).x = (i + 0.5)/numX;  
        samples(i,j).y = (j + 0.5)/numY;  
    }
```



Jittered/Stratified Sampling

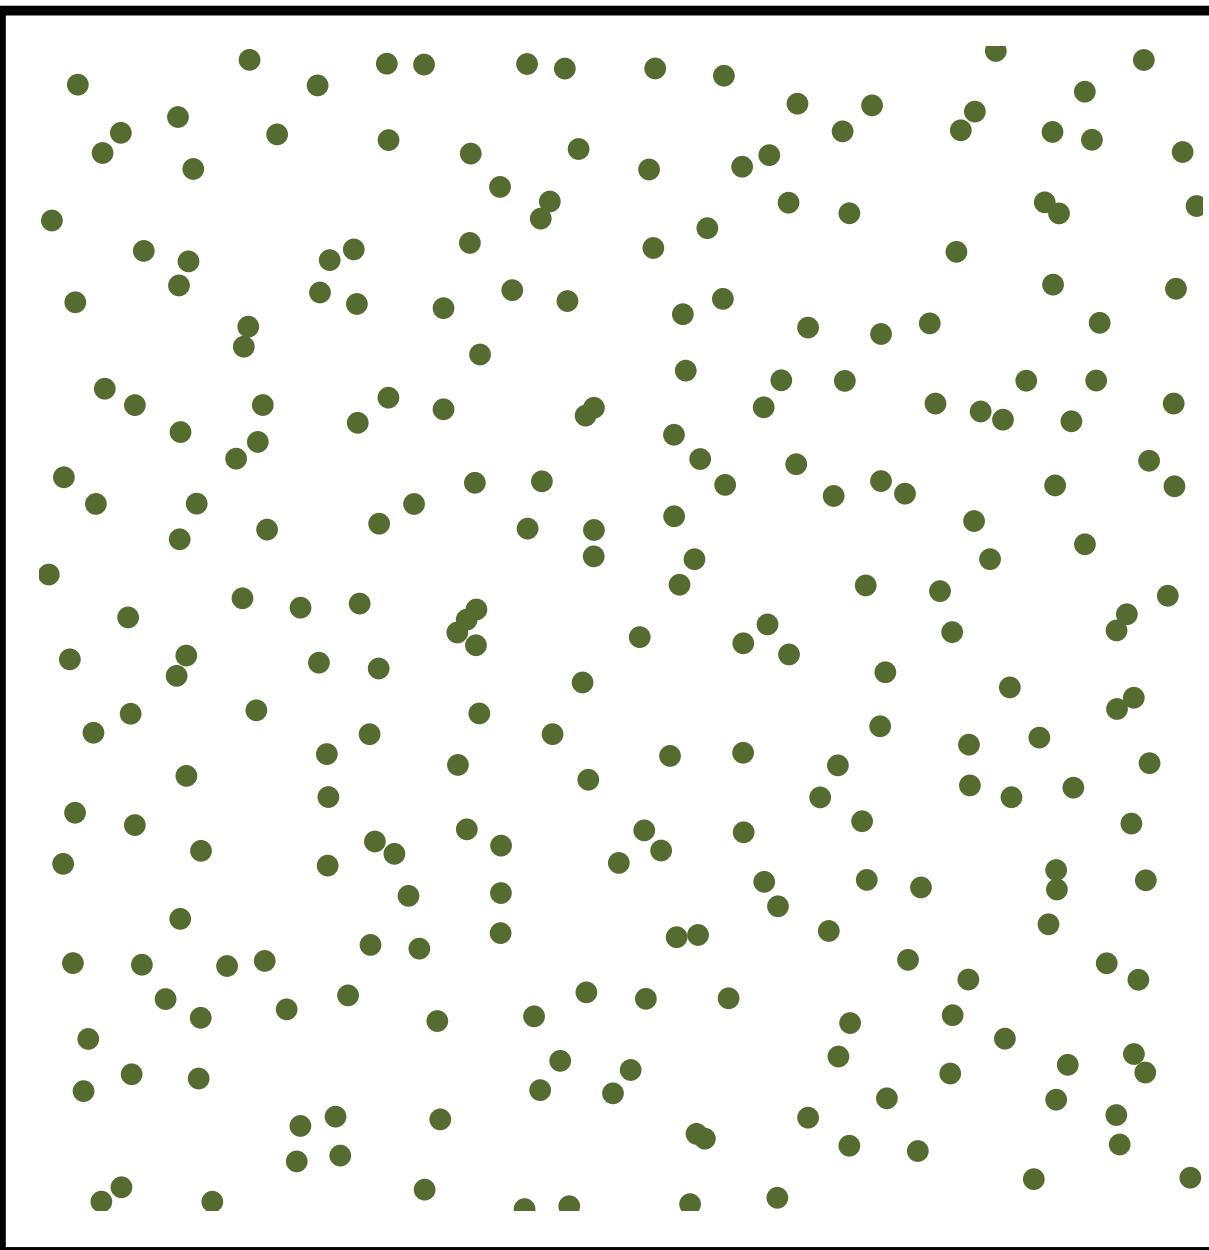
```
for (uint i = 0; i < numX; i++)  
    for (uint j = 0; j < numY; j++)  
    {  
        samples(i,j).x = (i + randf())/numX;  
        samples(i,j).y = (j + randf())/numY;  
    }
```

- ✓ Provably cannot increase variance
- ✓ Extends to higher dimensions, but...
- ✗ Curse of dimensionality
- ✗ Not progressive

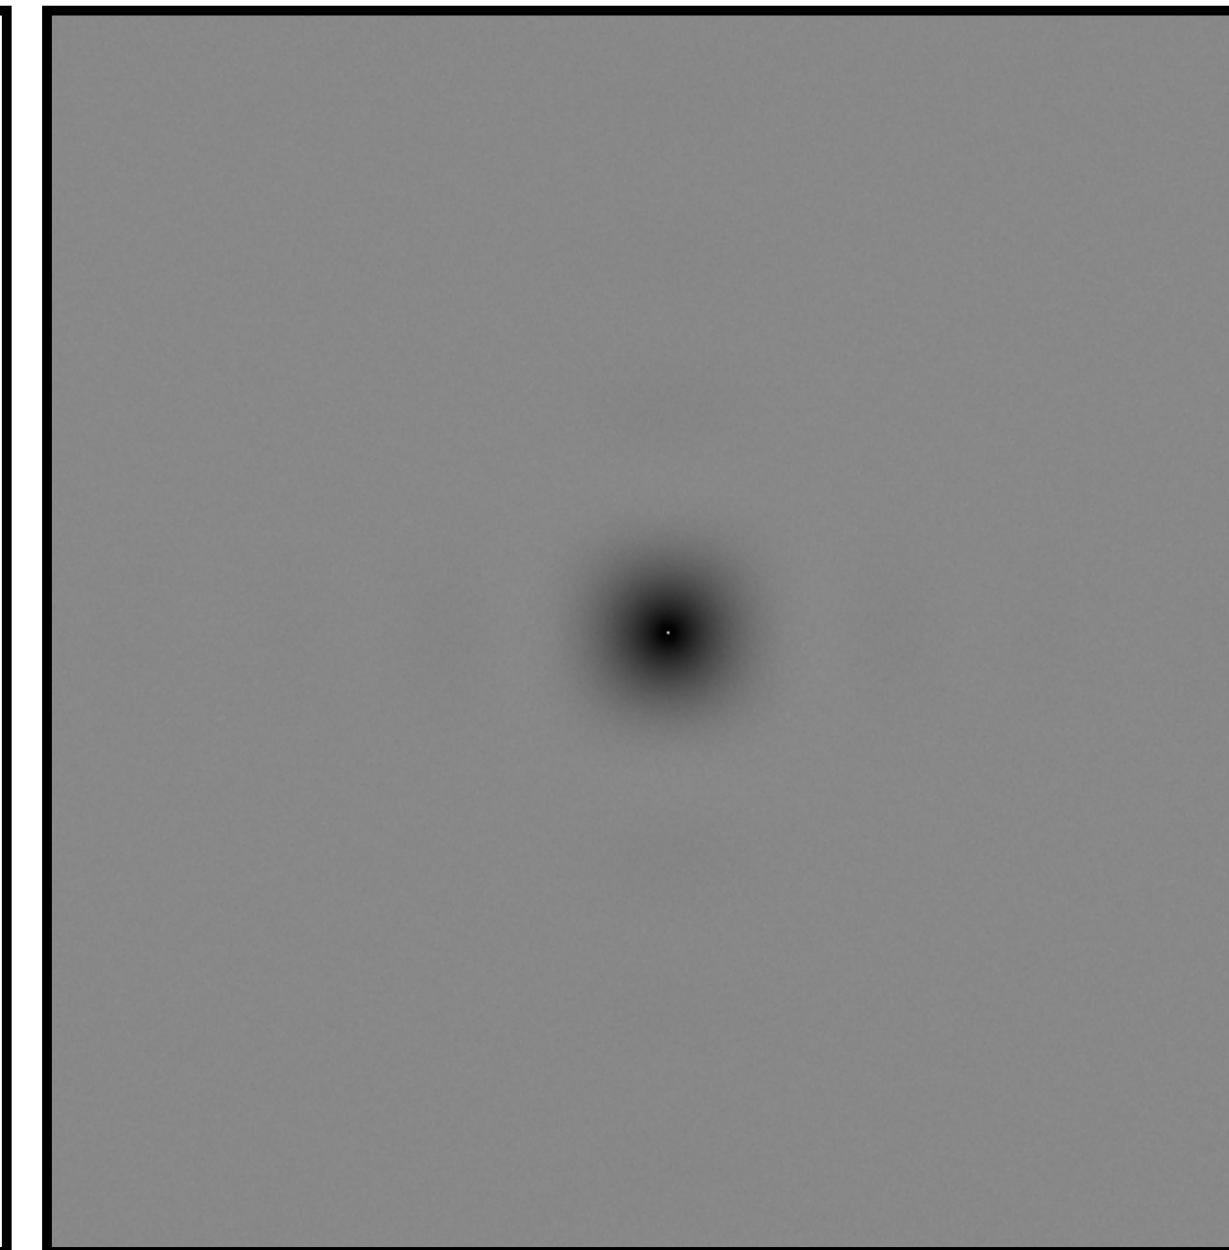


Jittered Sampling

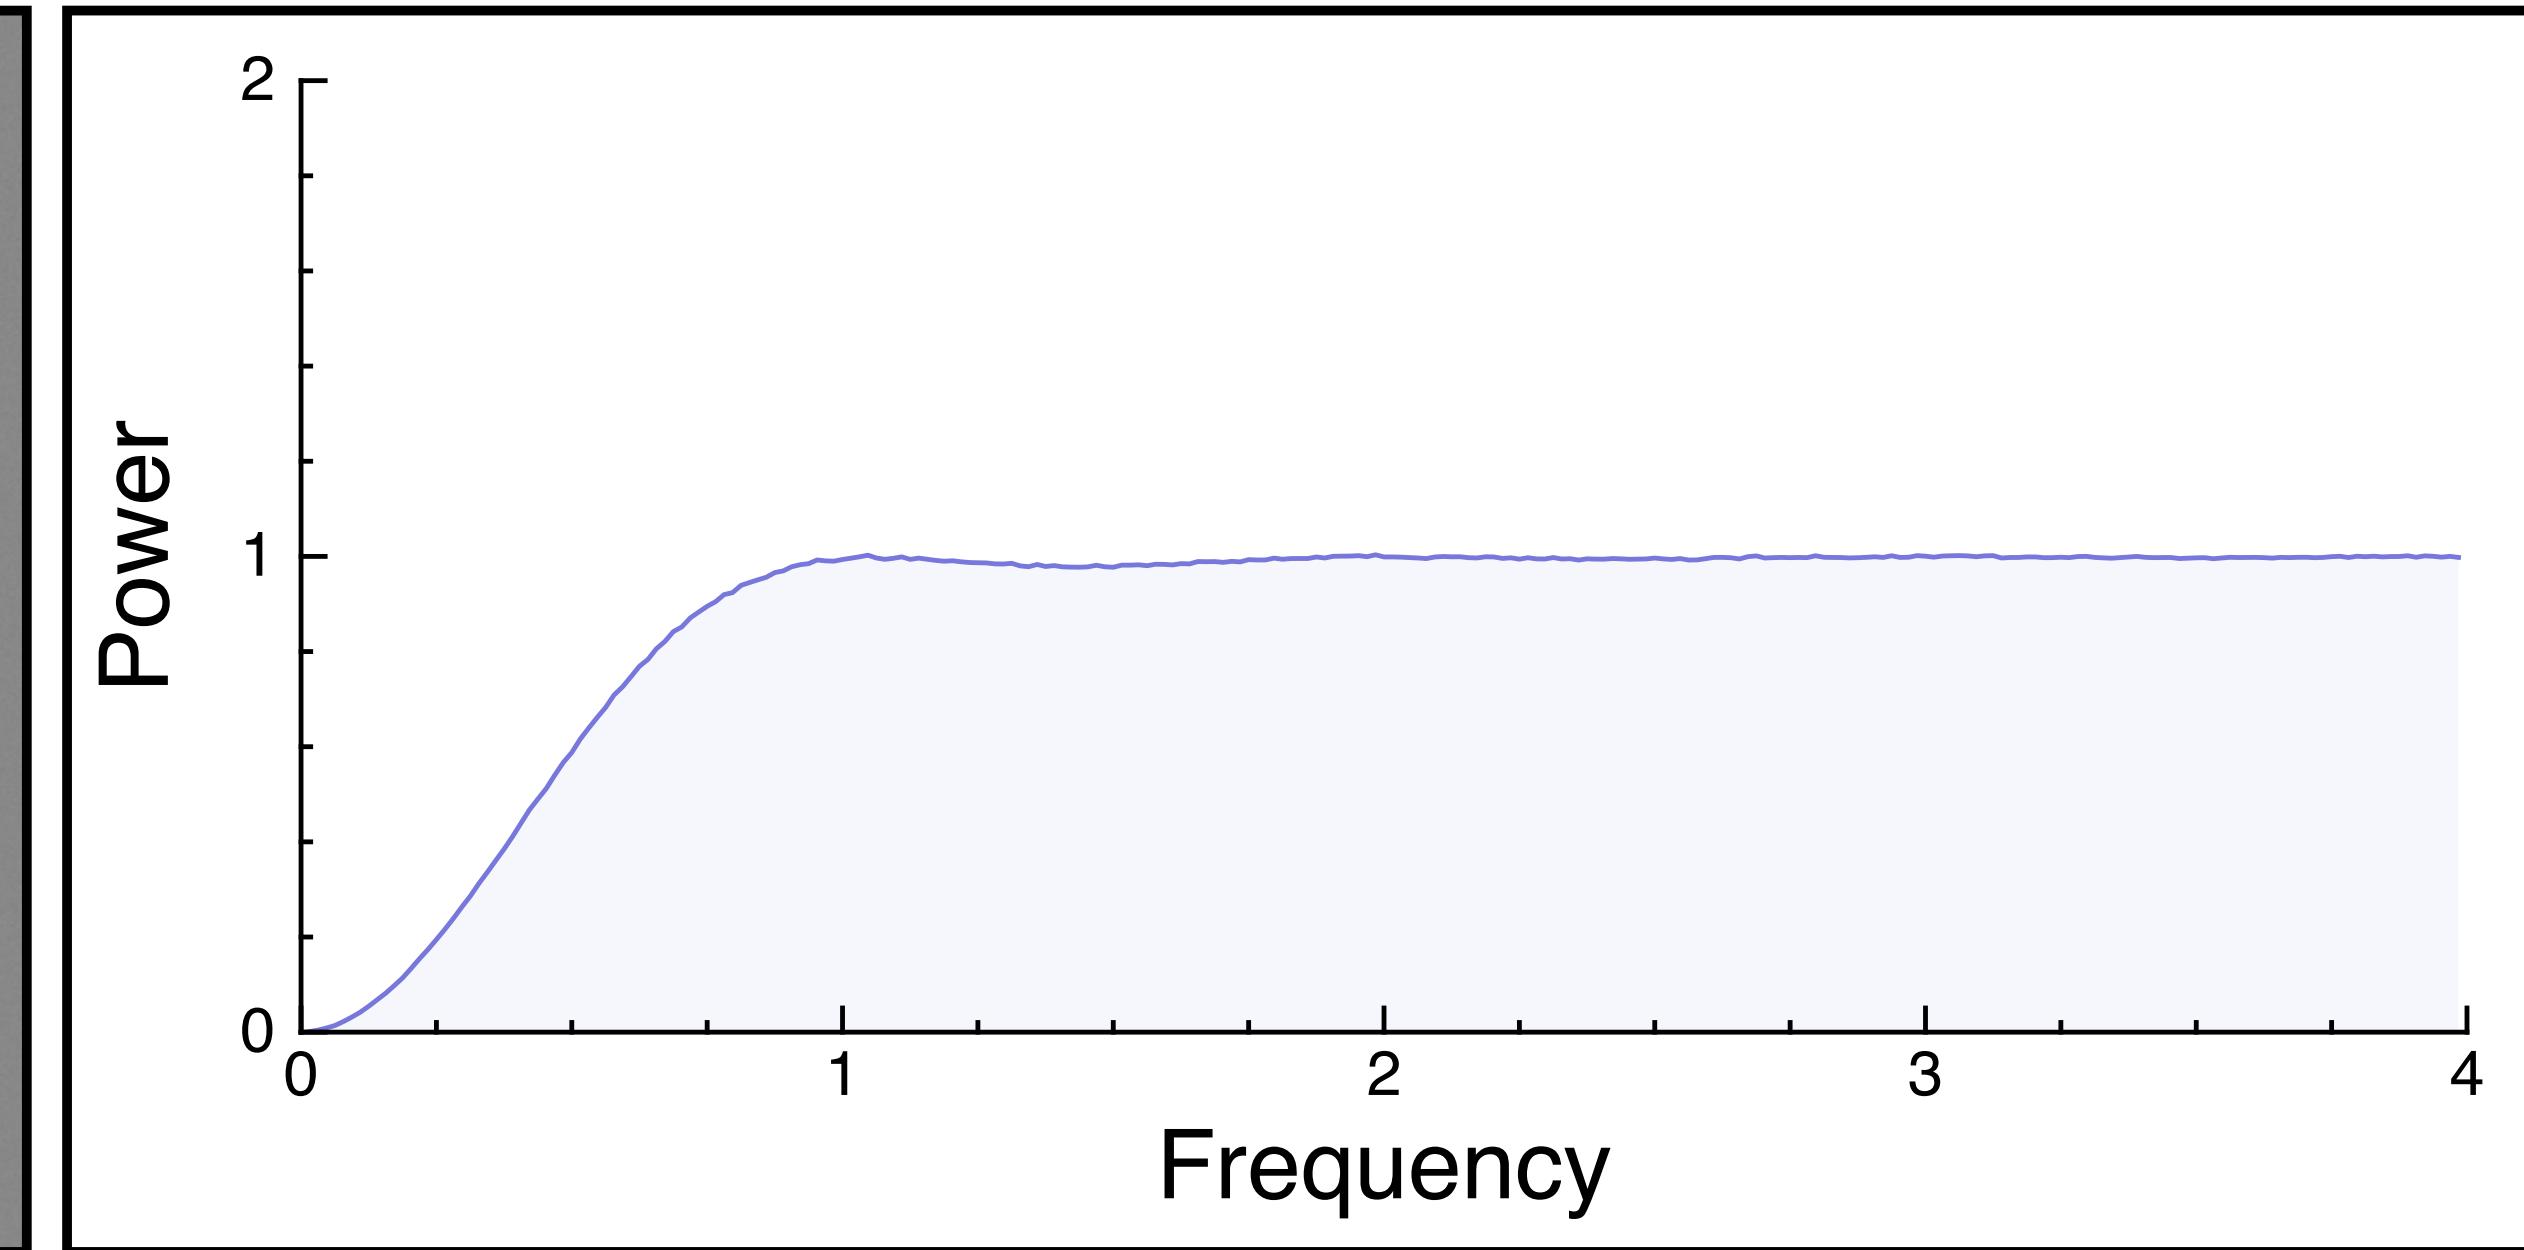
Samples



Power spectrum

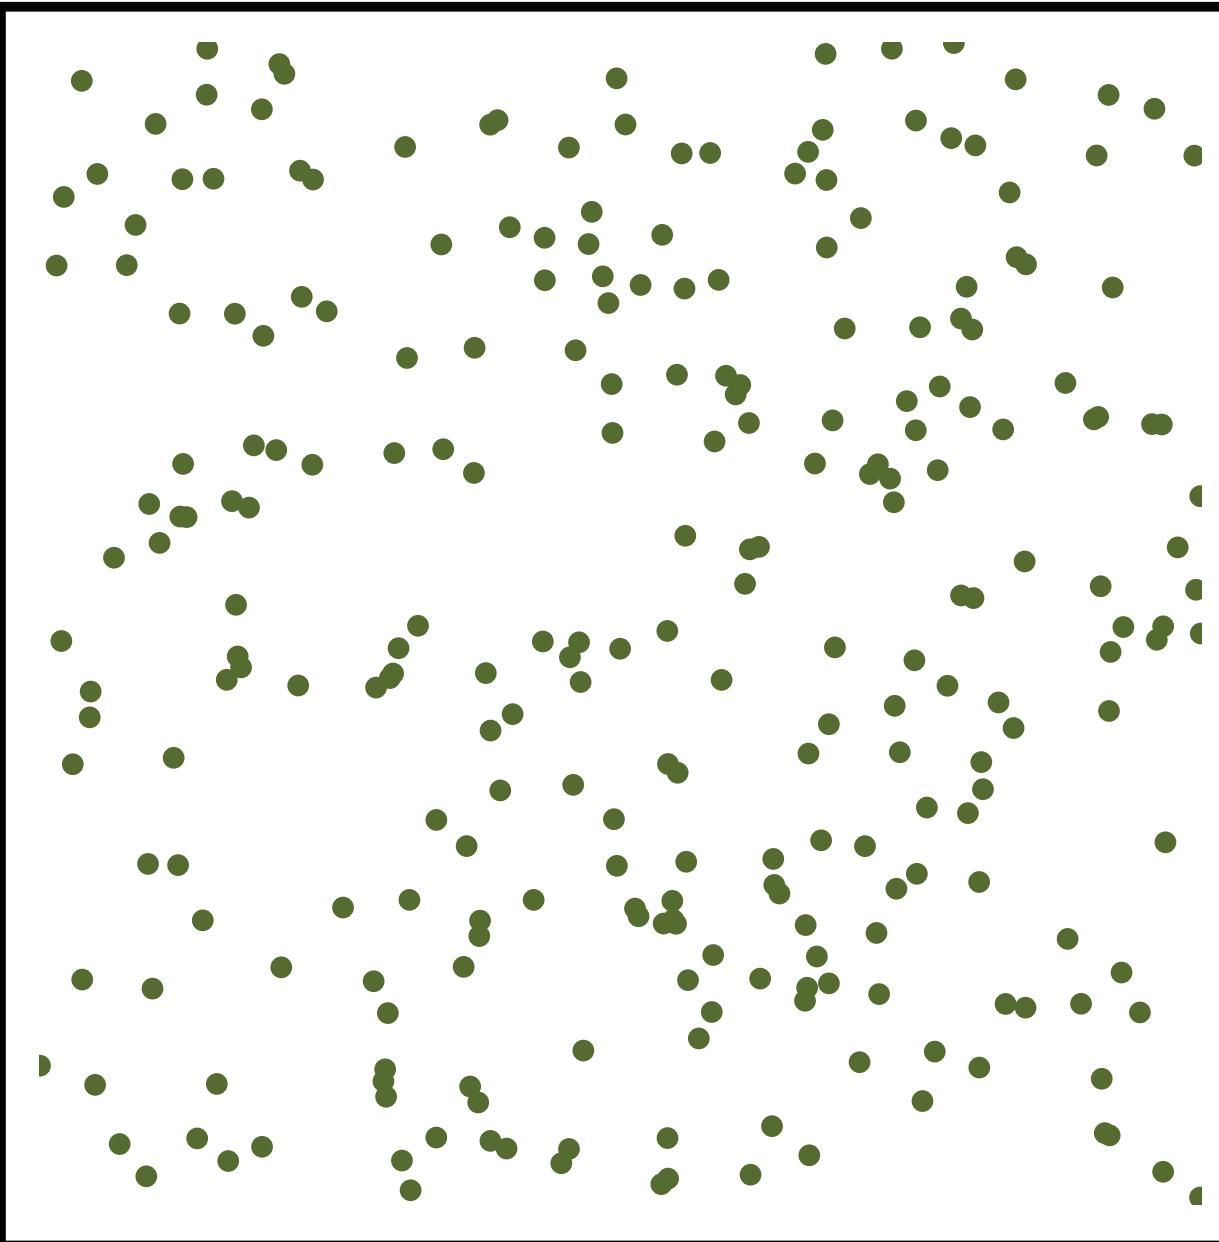


Radial mean



Independent Random Sampling

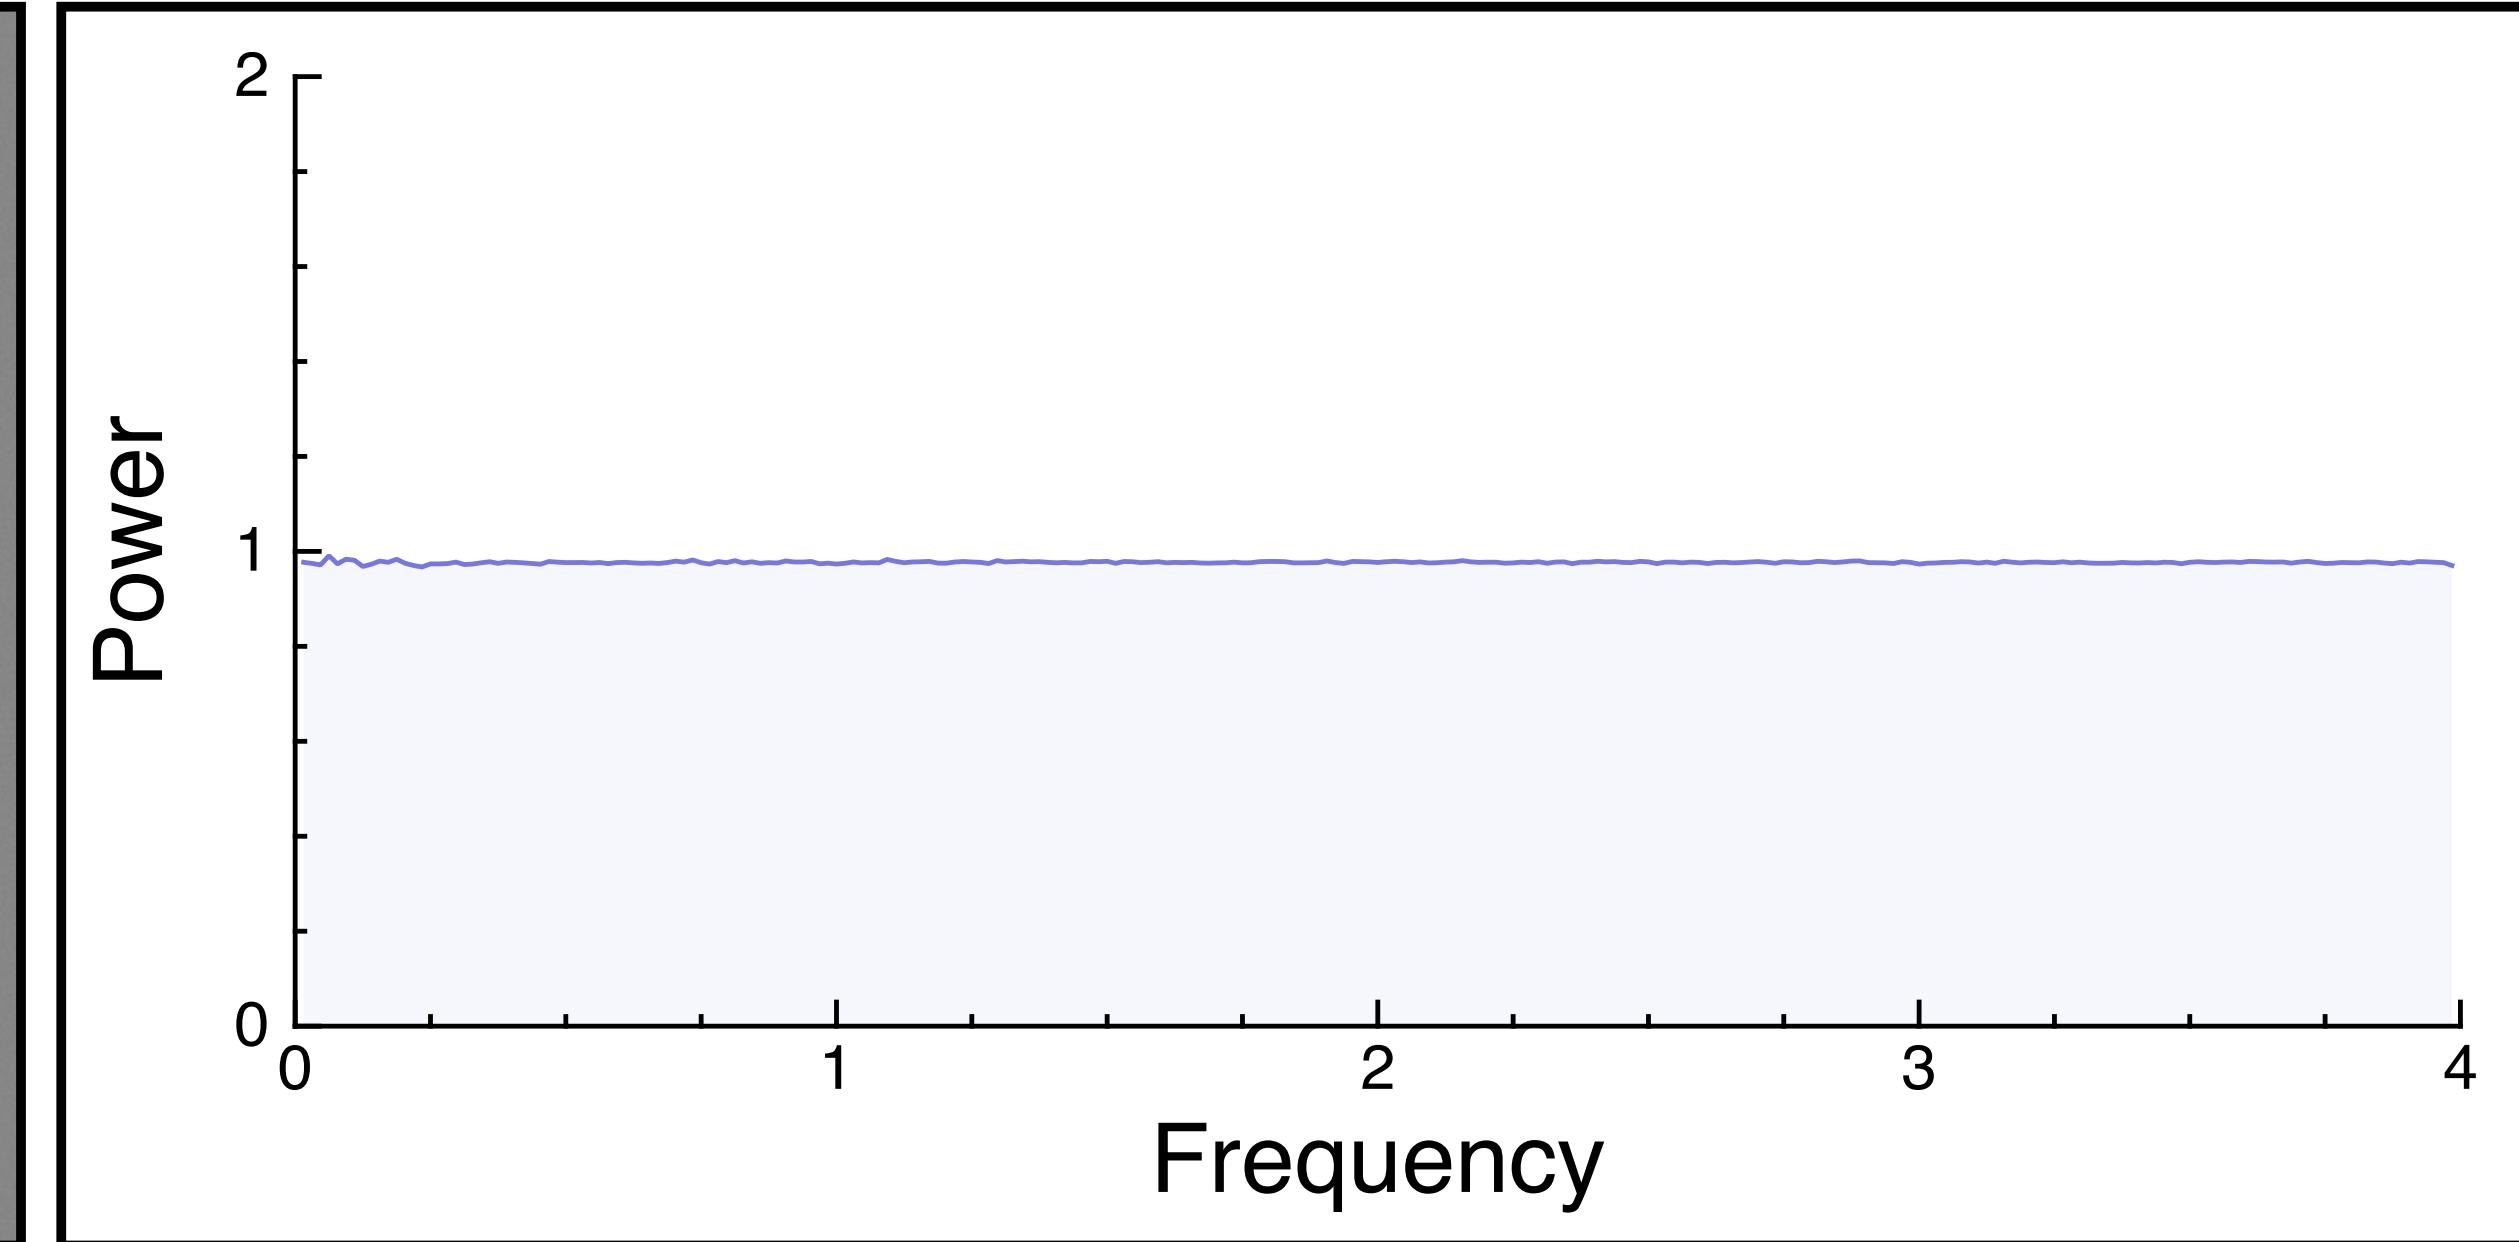
Samples



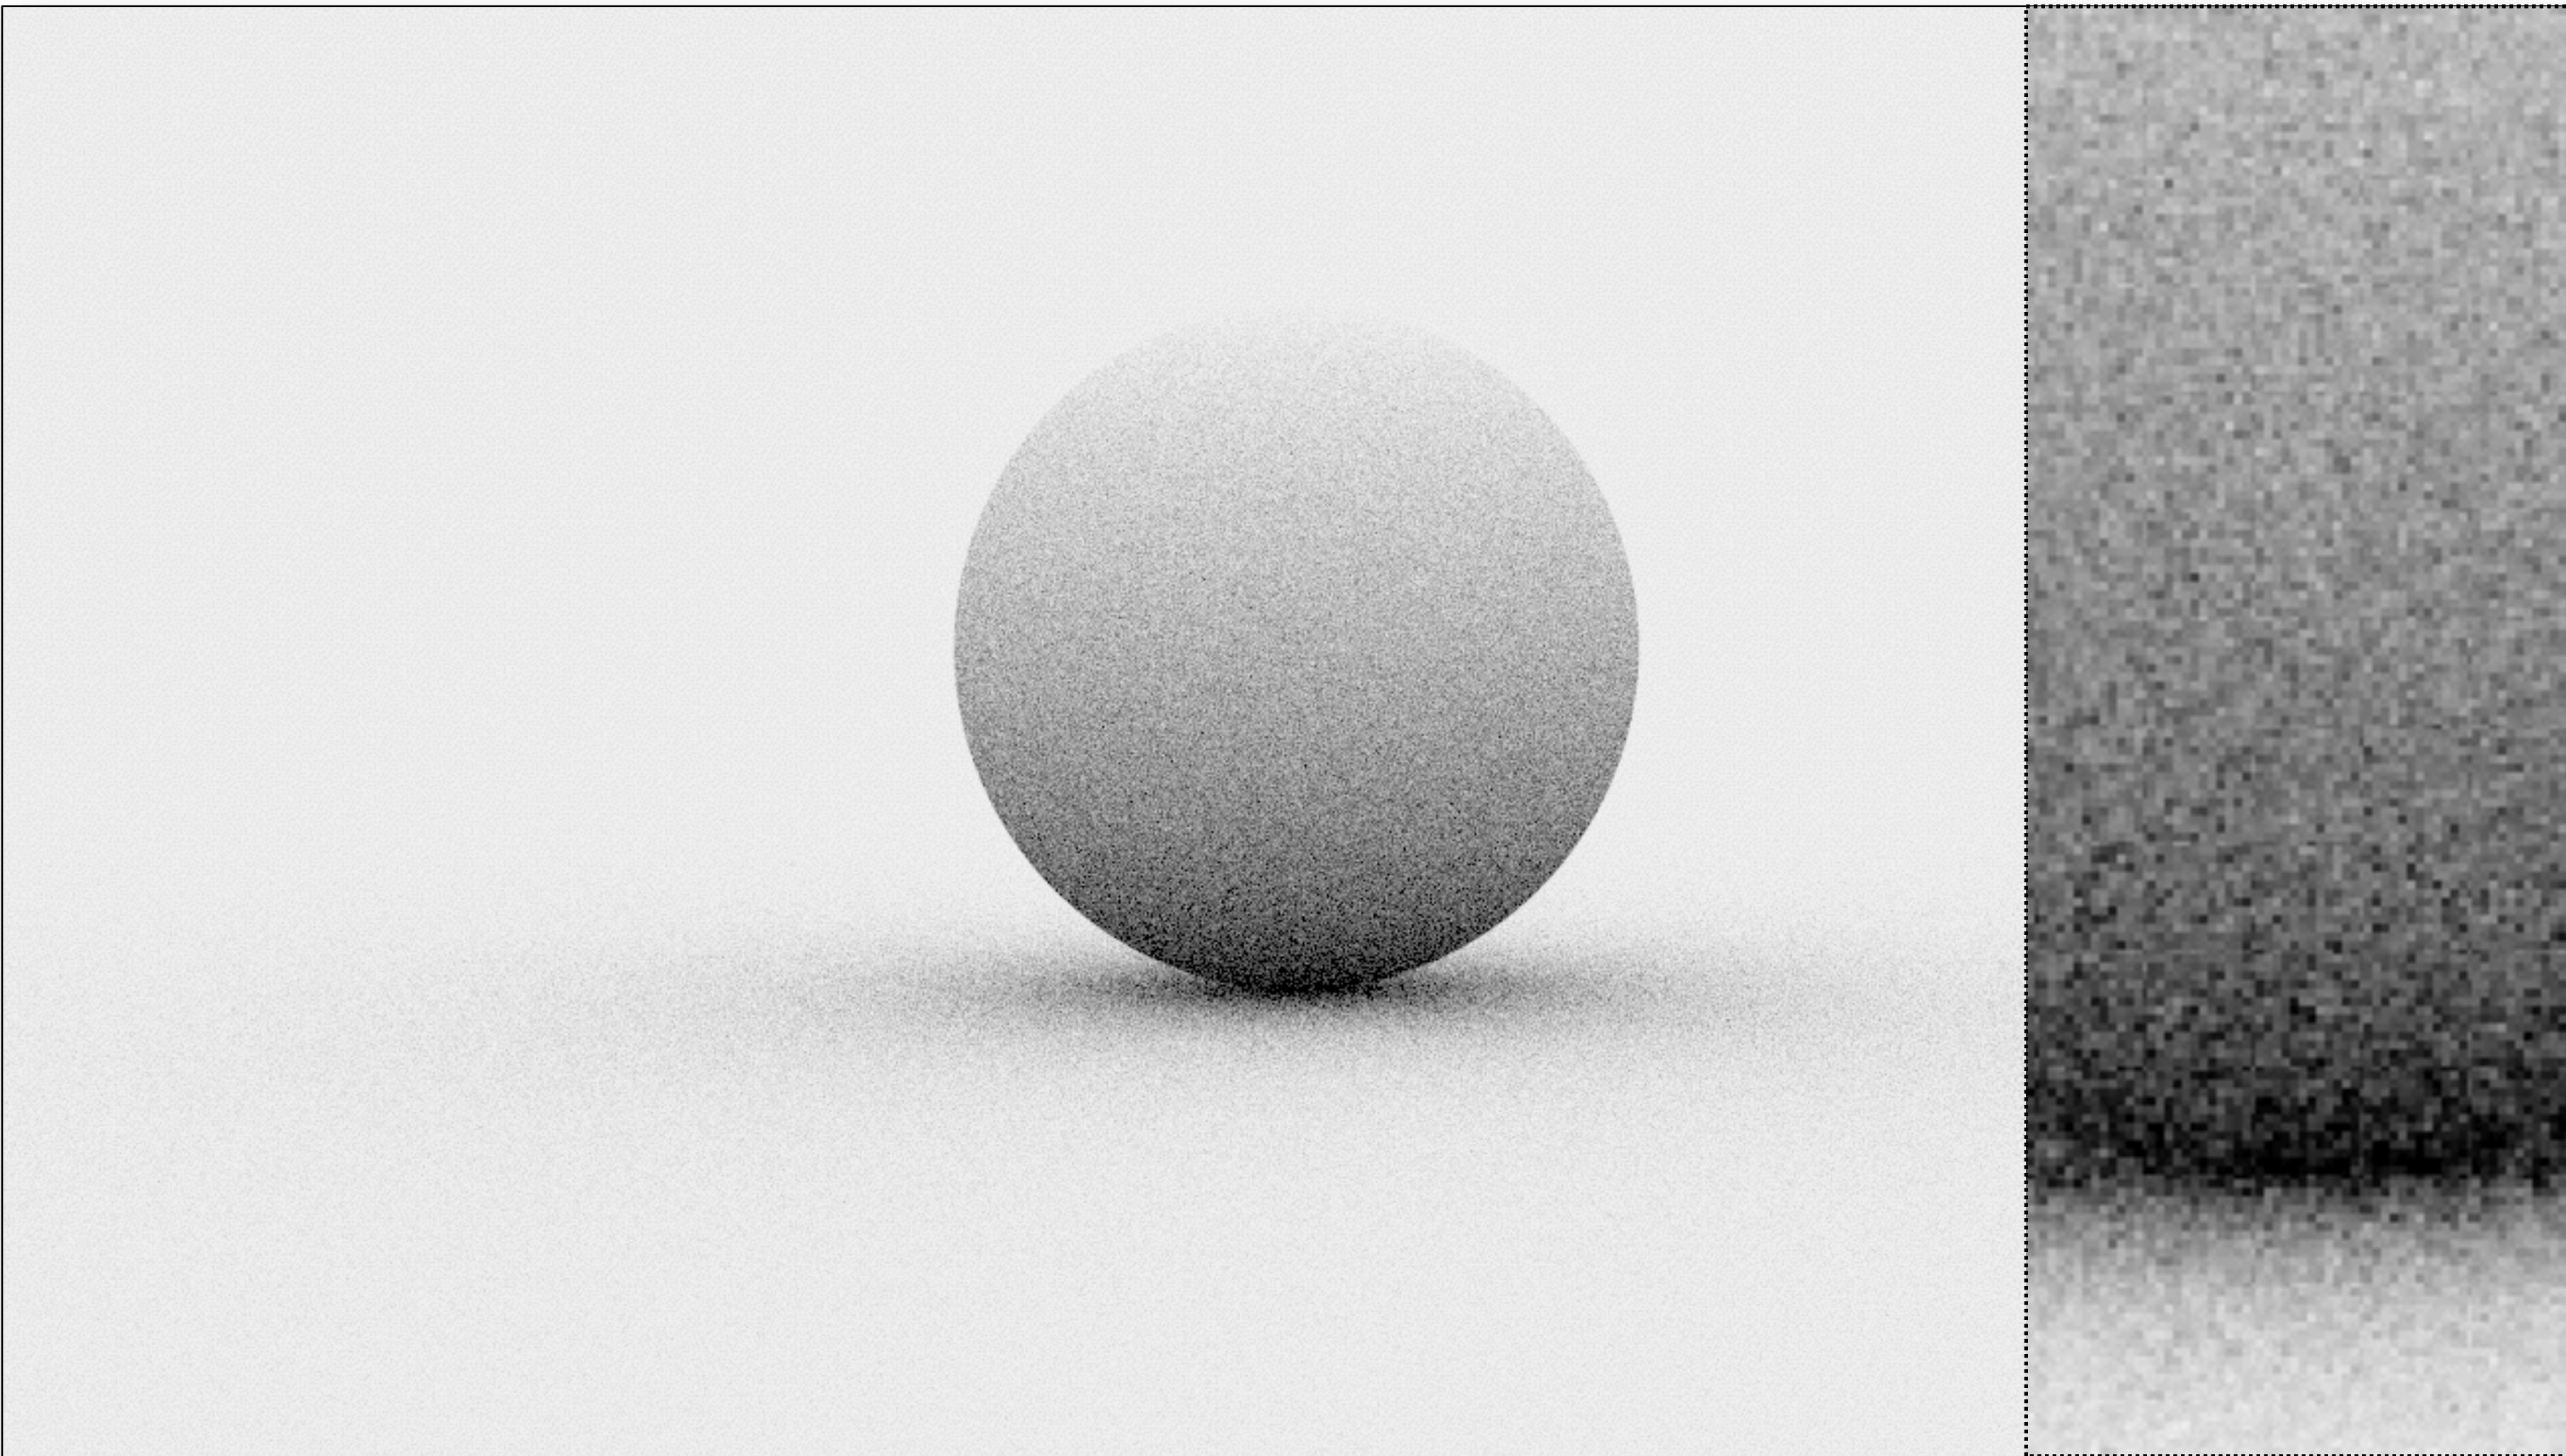
Power spectrum



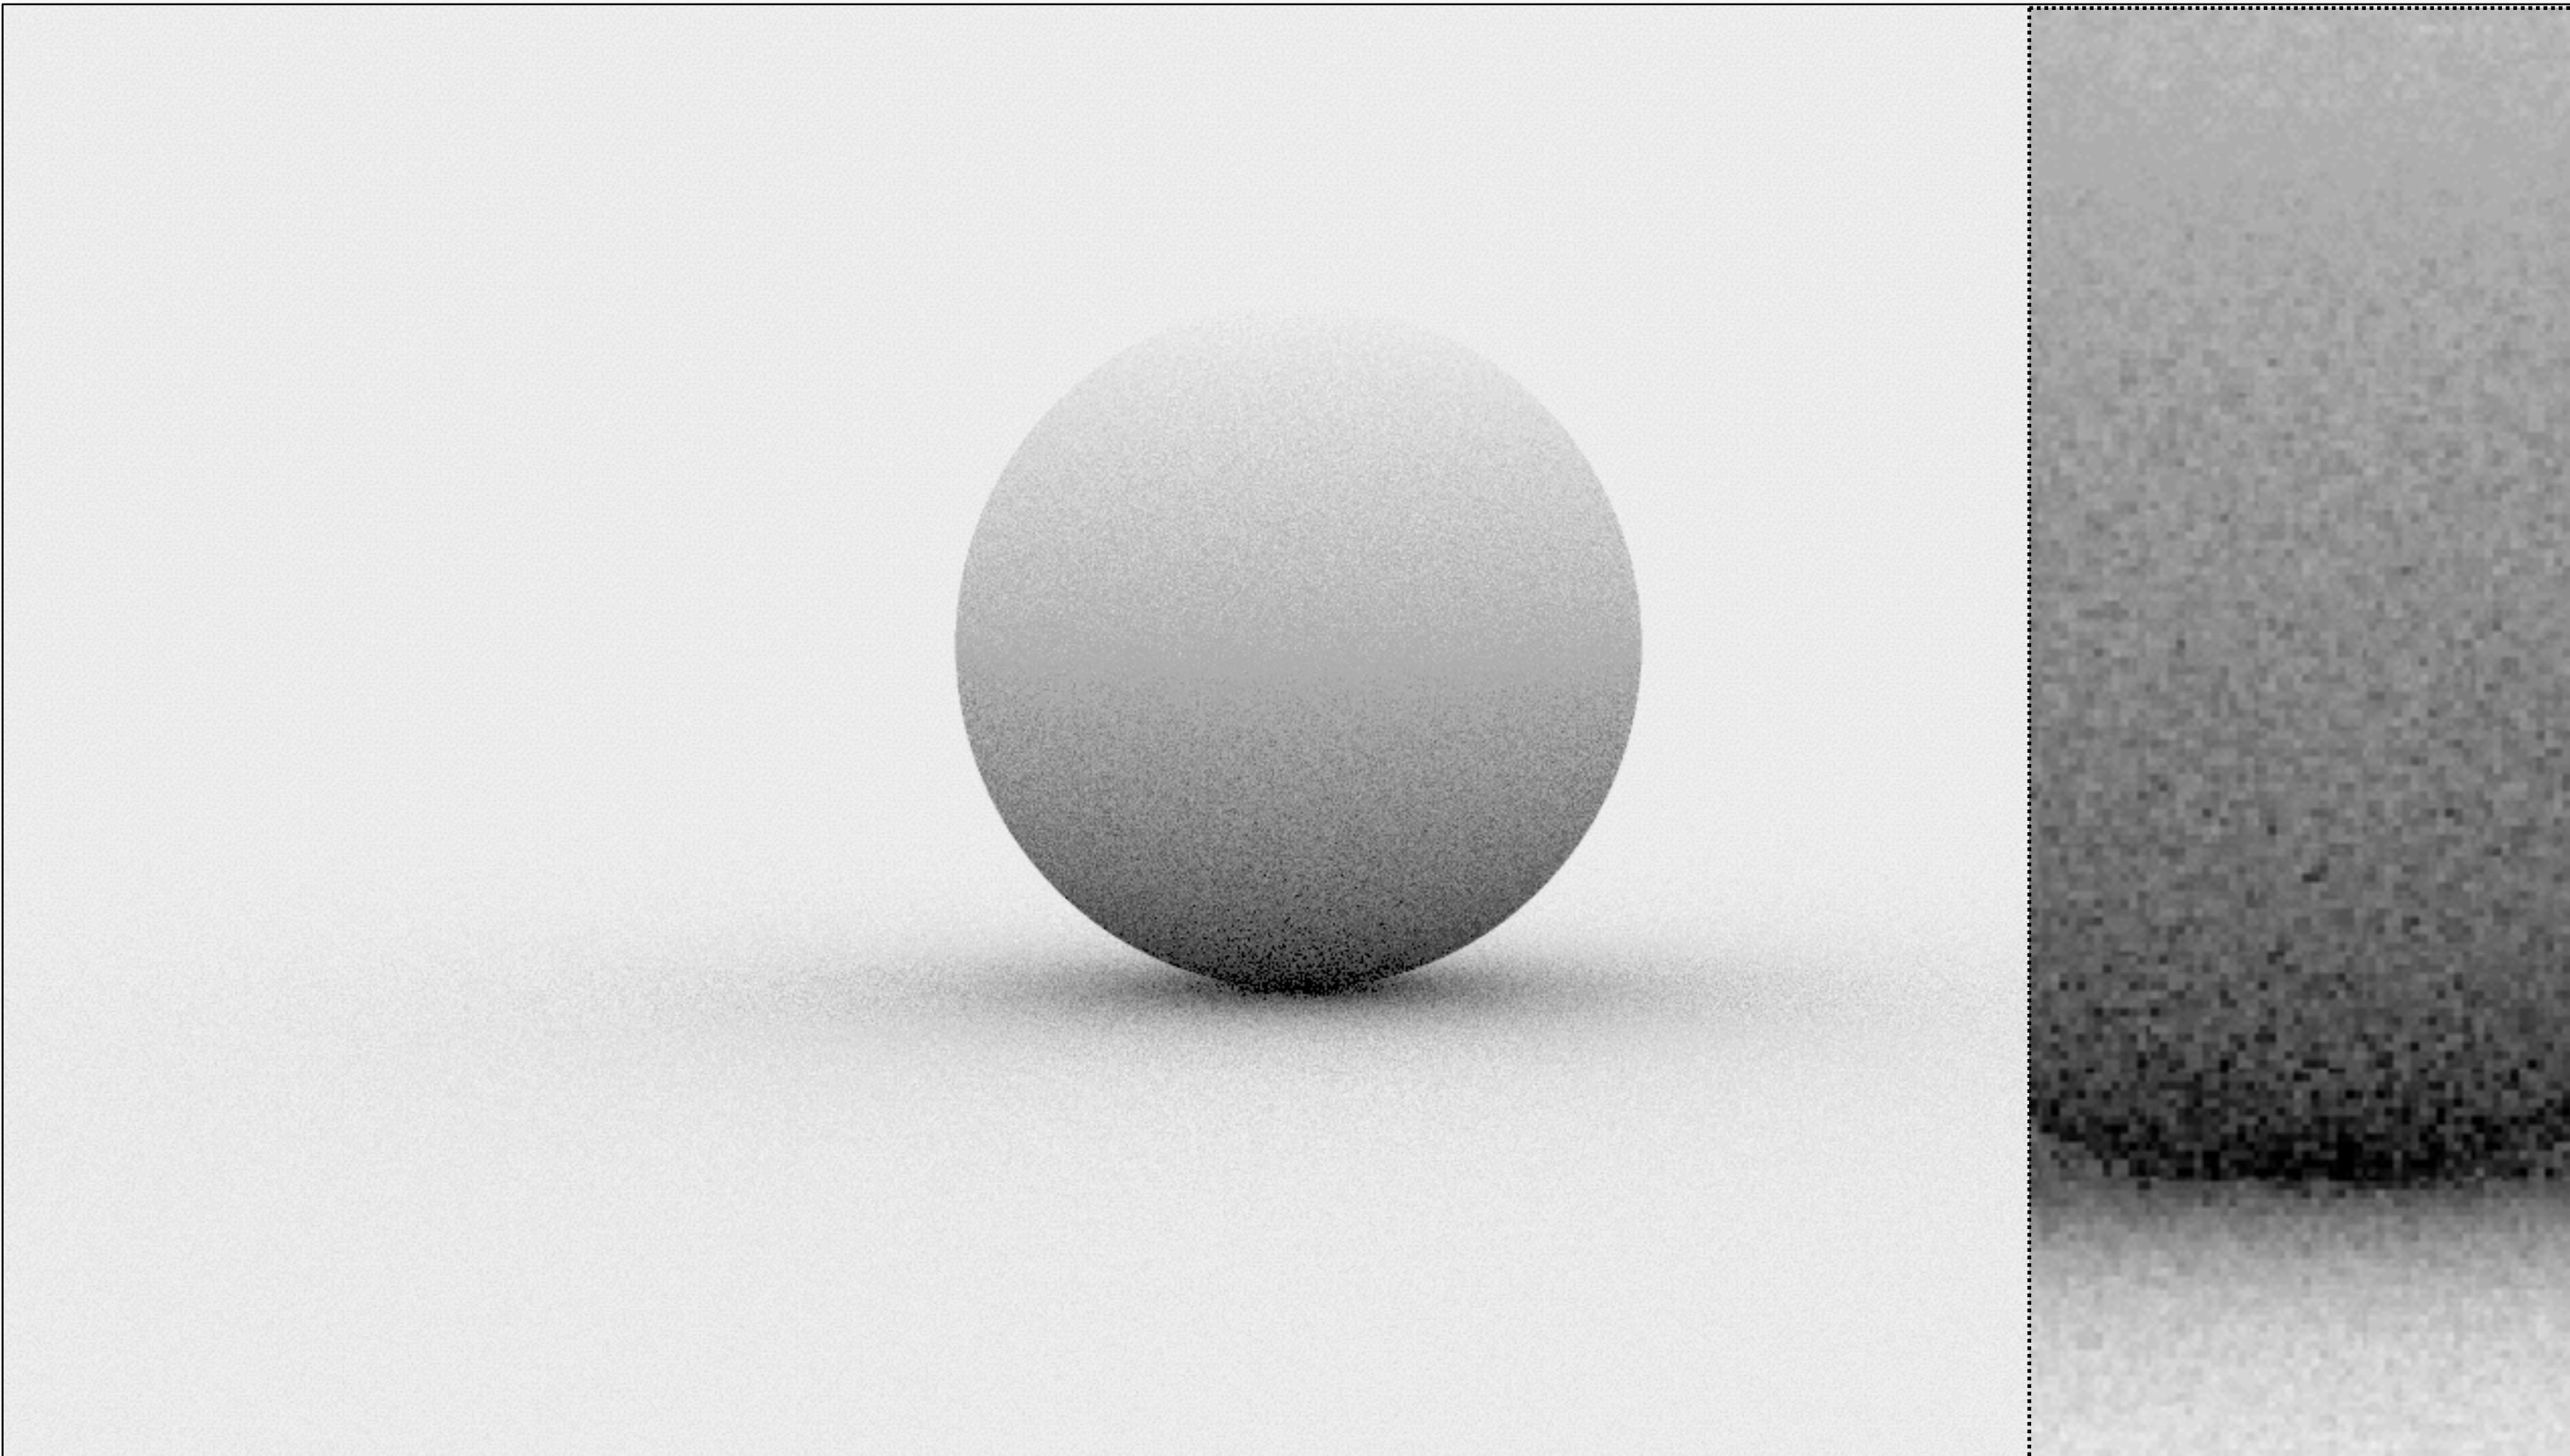
Radial mean



Monte Carlo (16 random samples)



Monte Carlo (16 jittered samples)



Stratifying in Higher Dimensions

Stratification requires $O(N^d)$ samples

- e.g. pixel (2D) + lens (2D) + time (1D) = 5D
 - splitting 2 times in 5D = $2^5 = 32$ samples
 - splitting 3 times in 5D = $3^5 = 243$ samples!

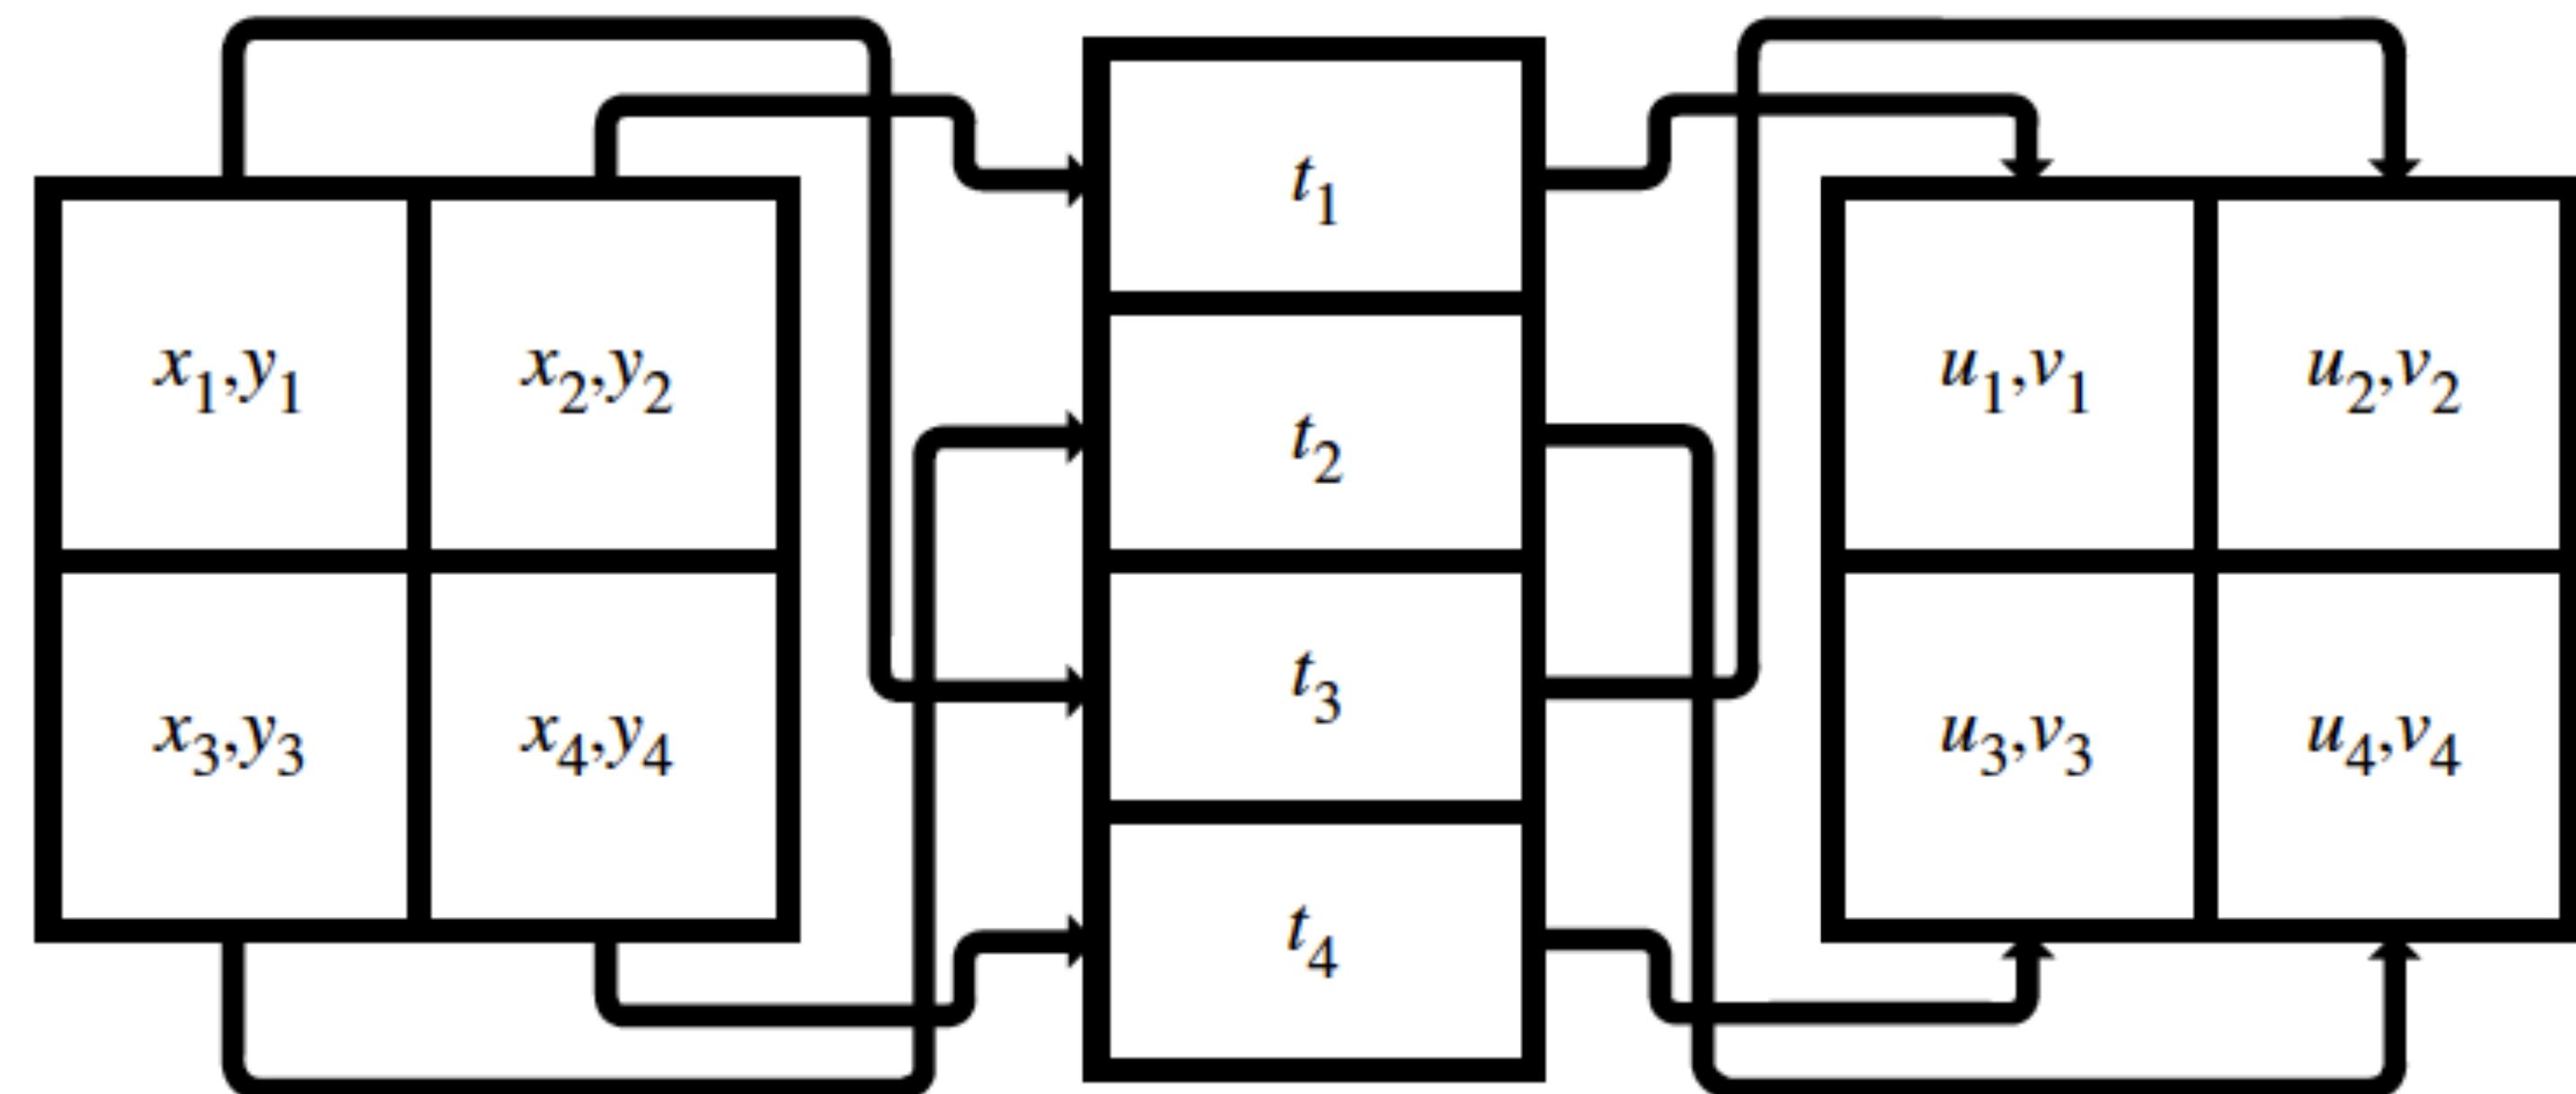
Inconvenient for large d

- cannot select sample count with fine granularity

Uncorrelated Jitter [Cook et al. 84]

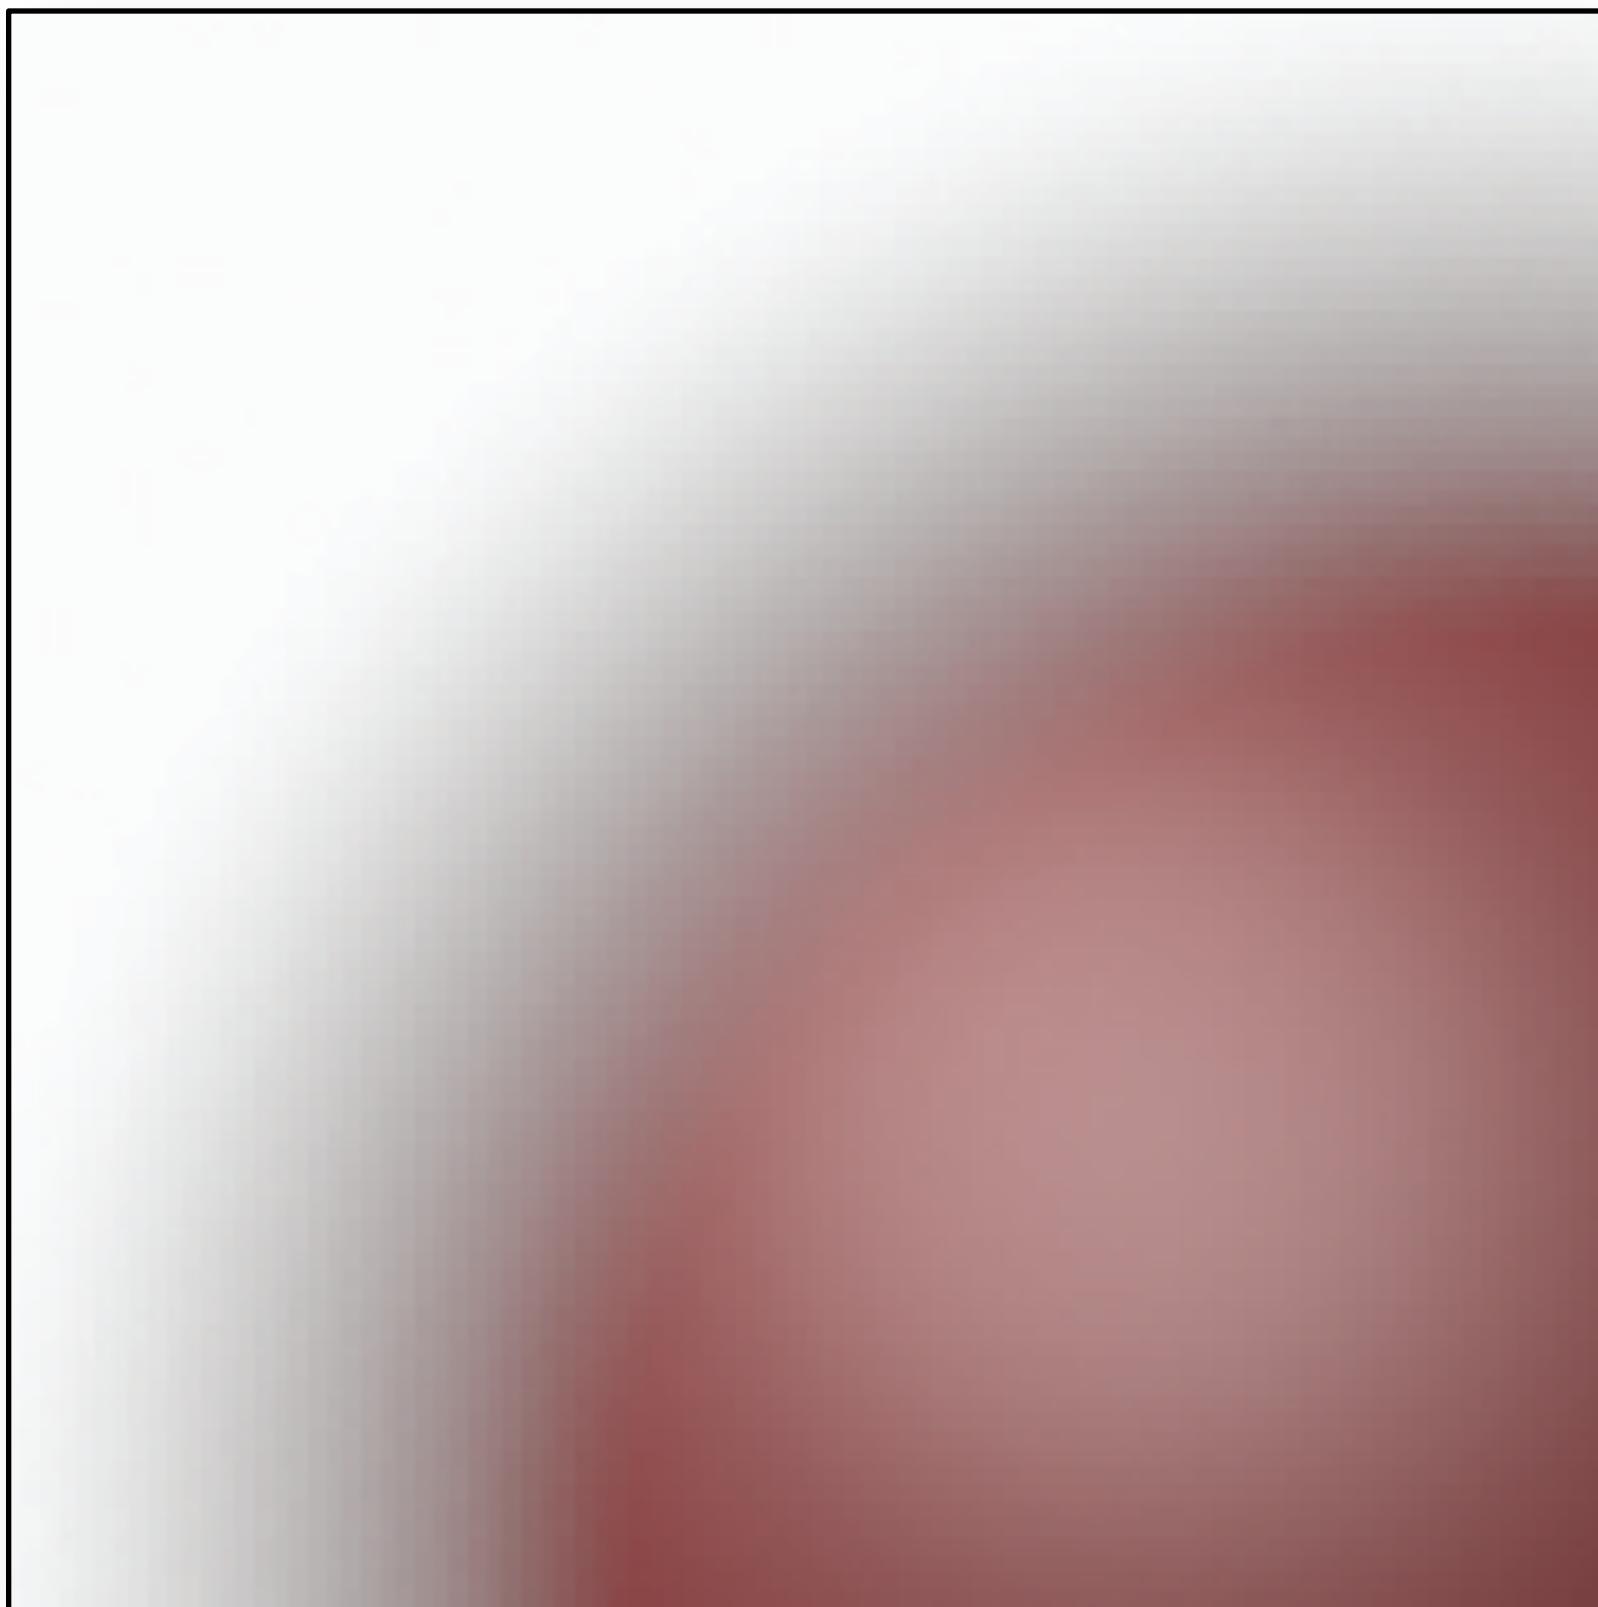
Compute stratified samples in sub-dimensions

- 2D jittered (x,y) for pixel
- 2D jittered (u,v) for lens
- 1D jittered (t) for time
- combine dimensions in random order



Depth of Field (4D)

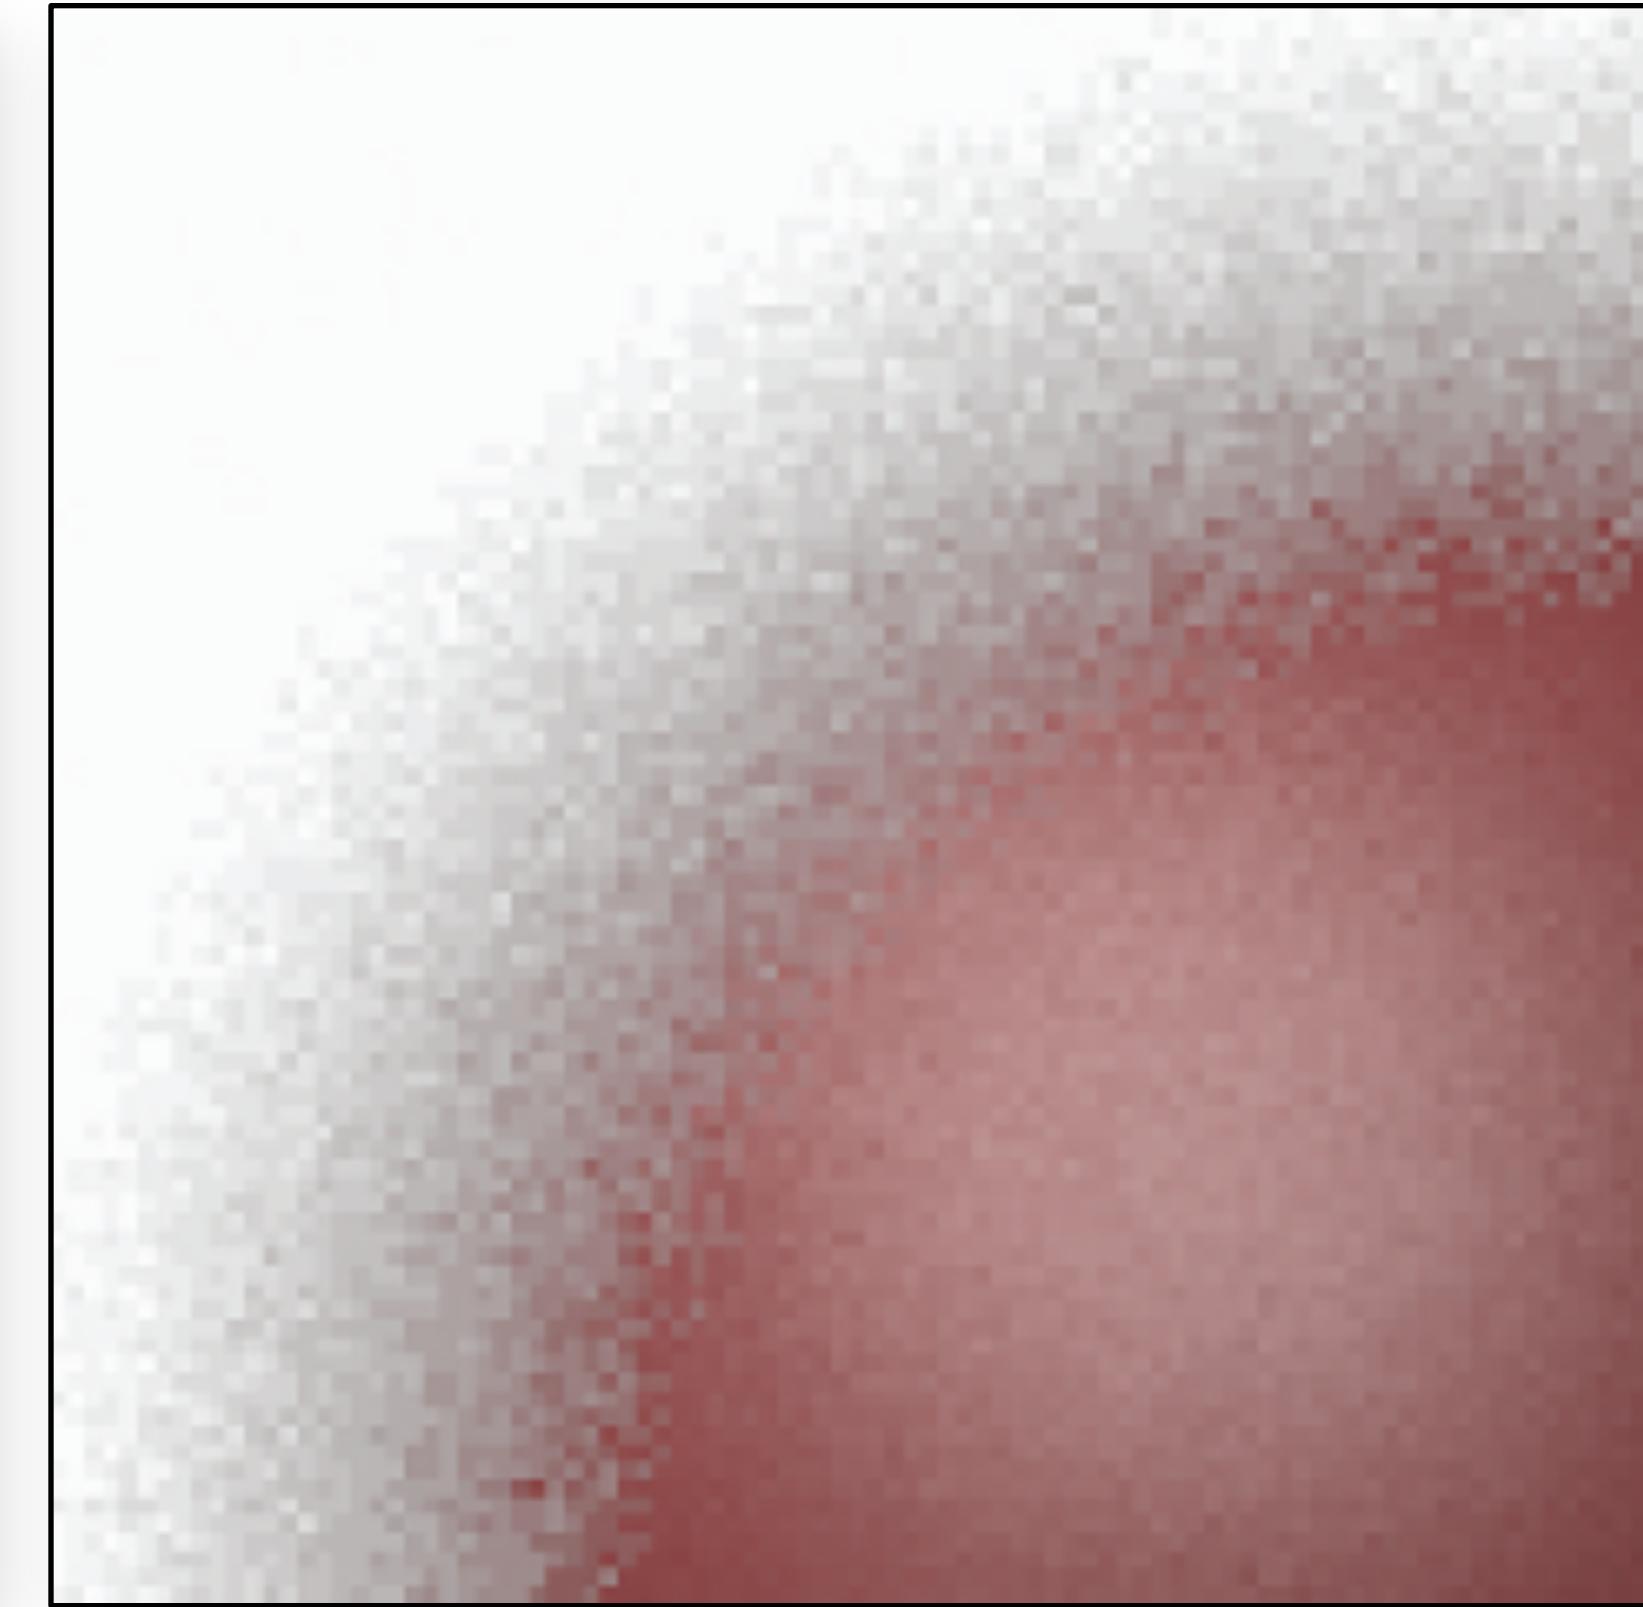
Reference



Random Sampling



Uncorrelated Jitter

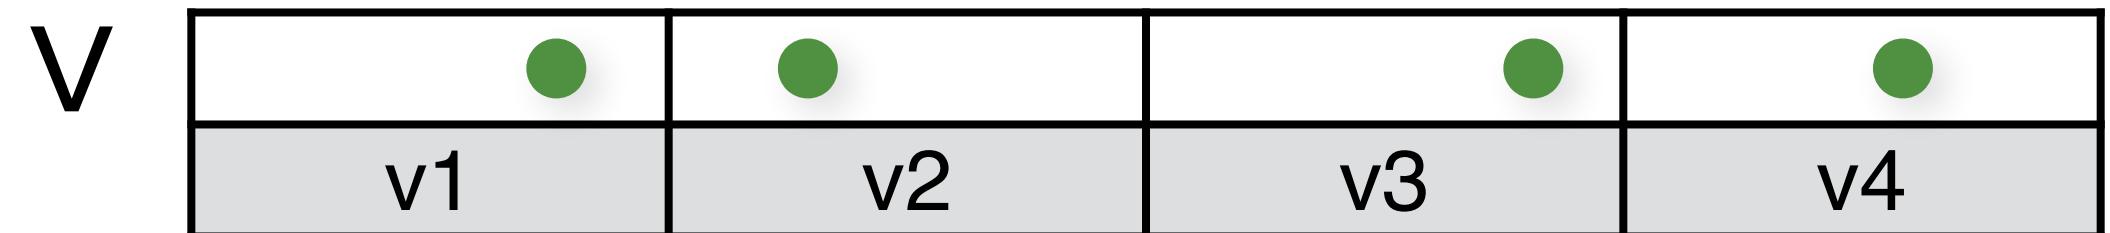
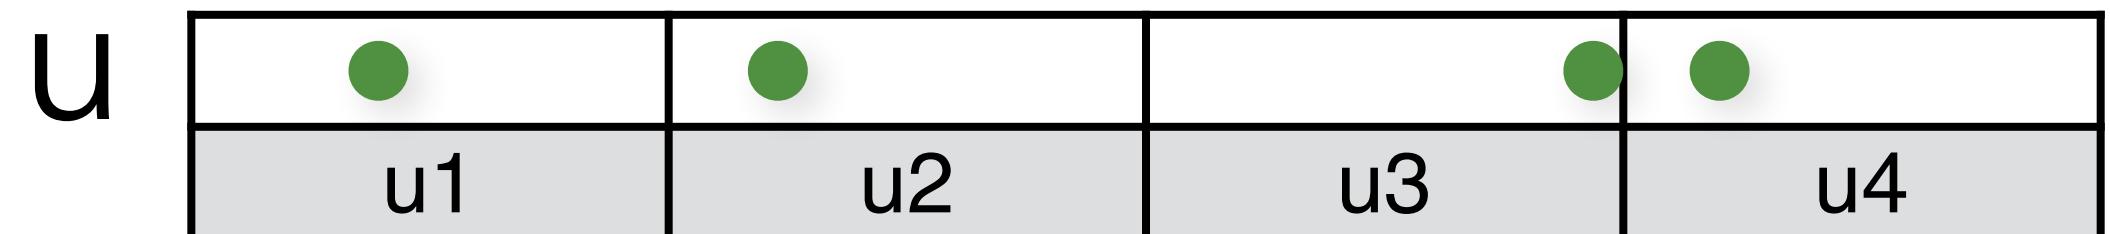
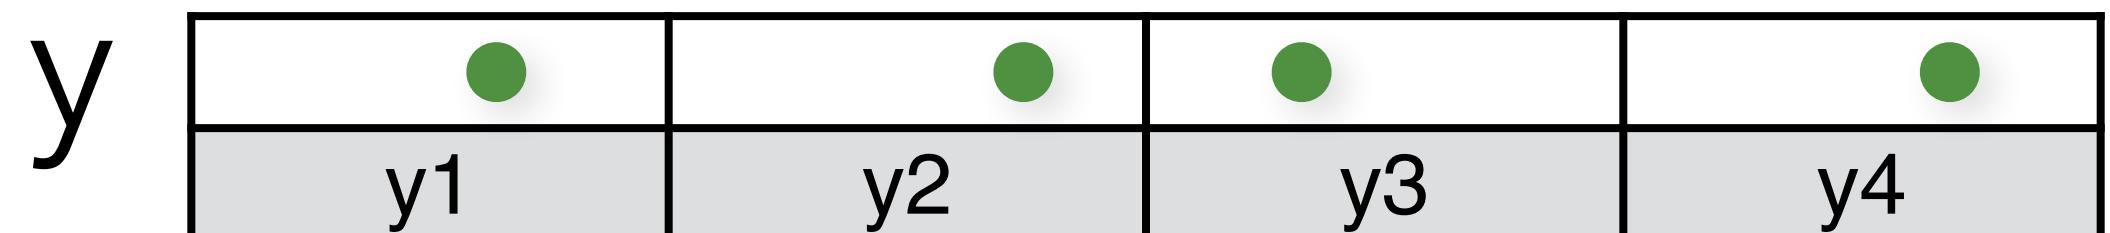
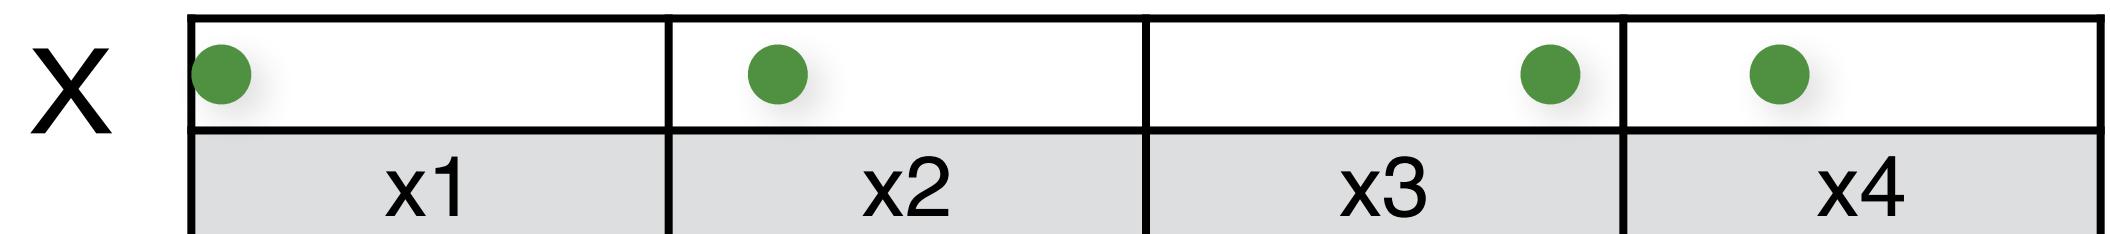


Uncorrelated Jitter → Latin Hypercube

Stratify samples in each dimension separately

- for 5D: 5 separate 1D jittered point sets

- combine dimensions
in random order

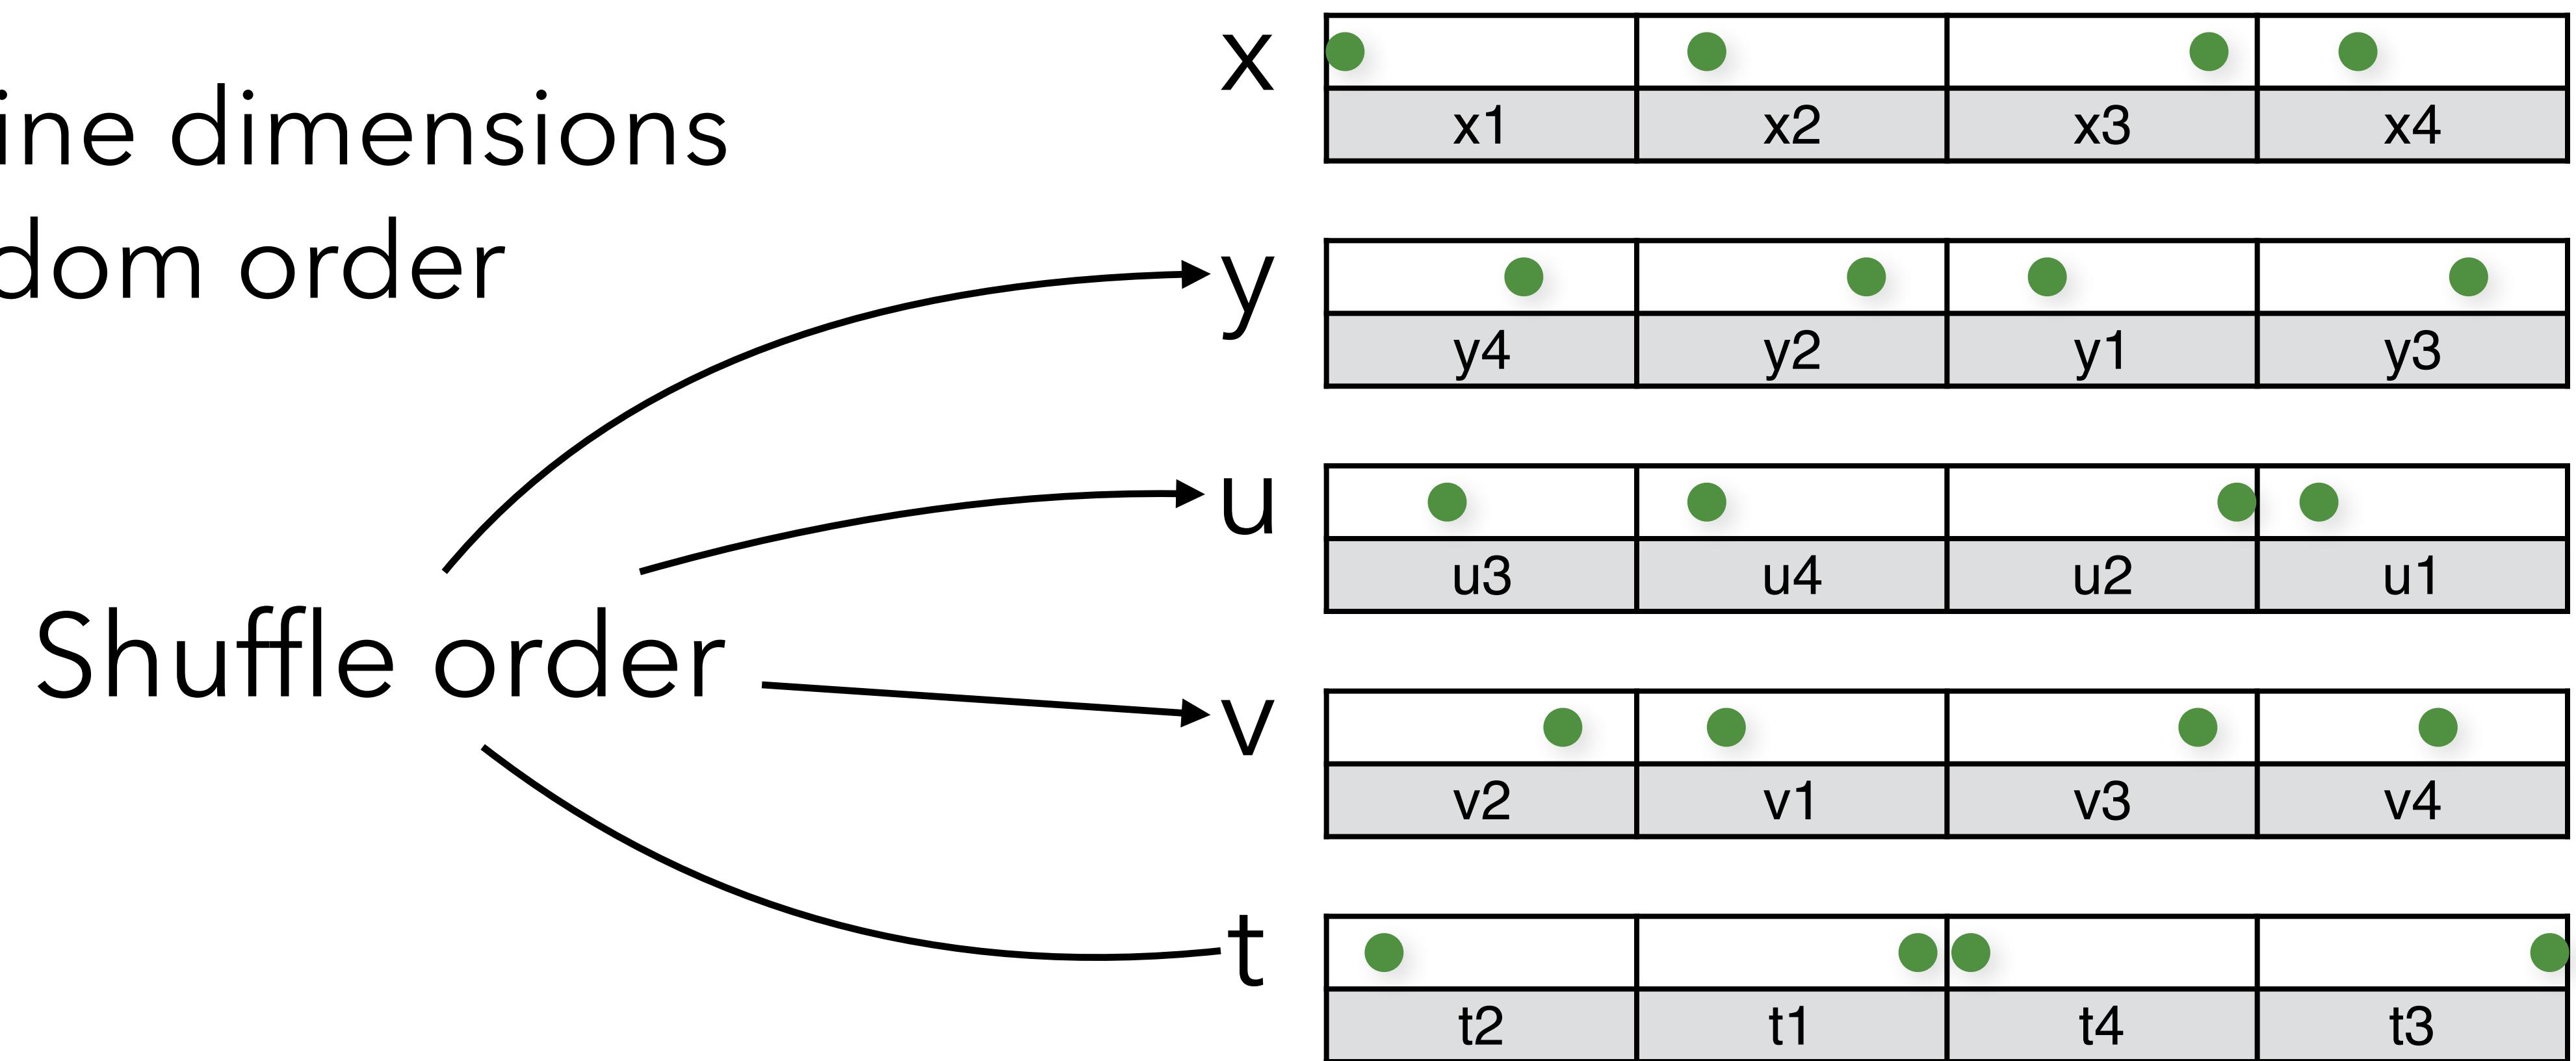


Uncorrelated Jitter → Latin Hypercube

Stratify samples in each dimension separately

- for 5D: 5 separate 1D jittered point sets

- combine dimensions
in random order

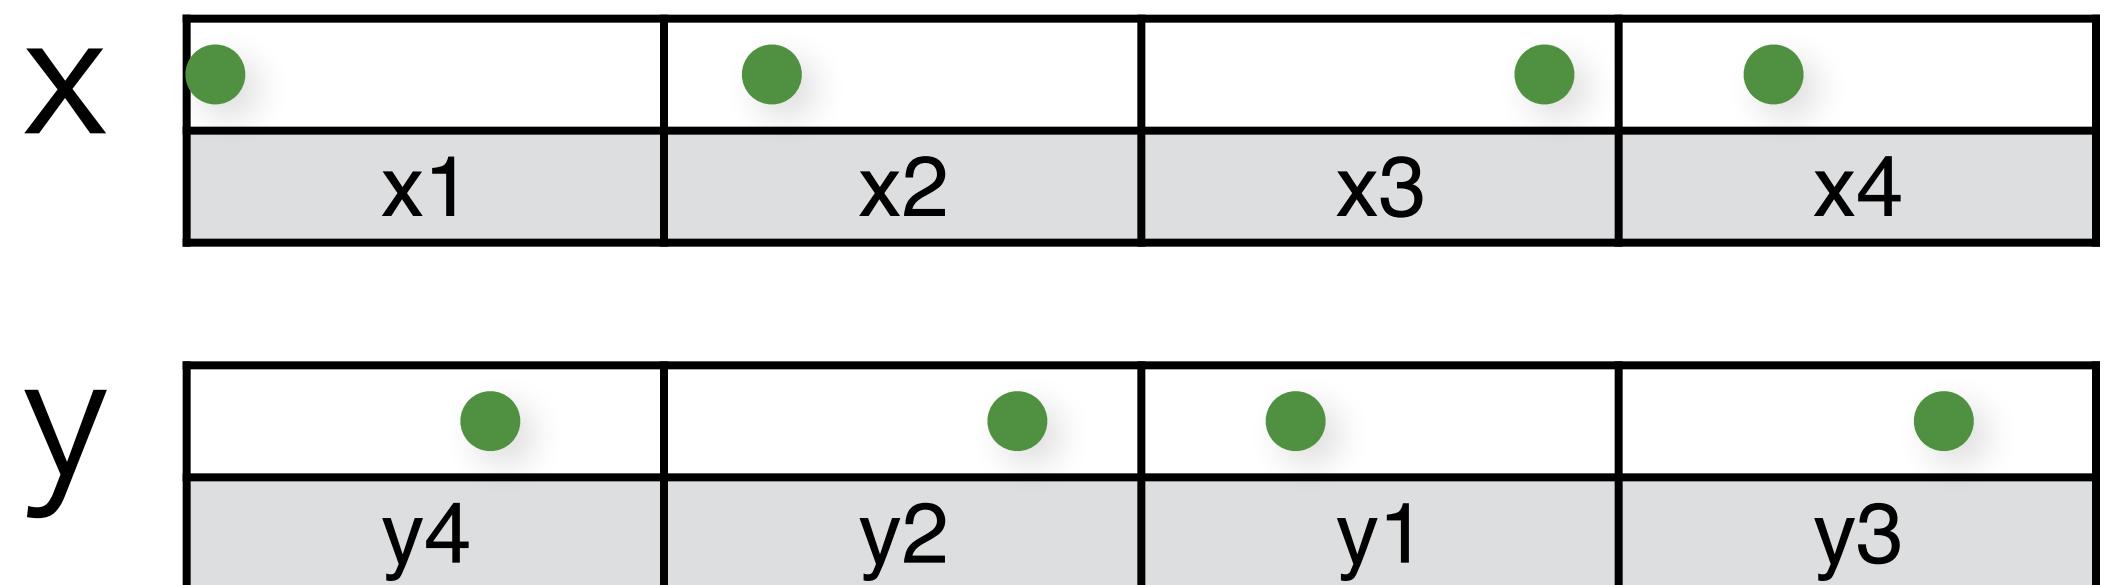


N-Rooks = 2D Latin Hypercube [Shirley 91]

Stratify samples in each dimension separately

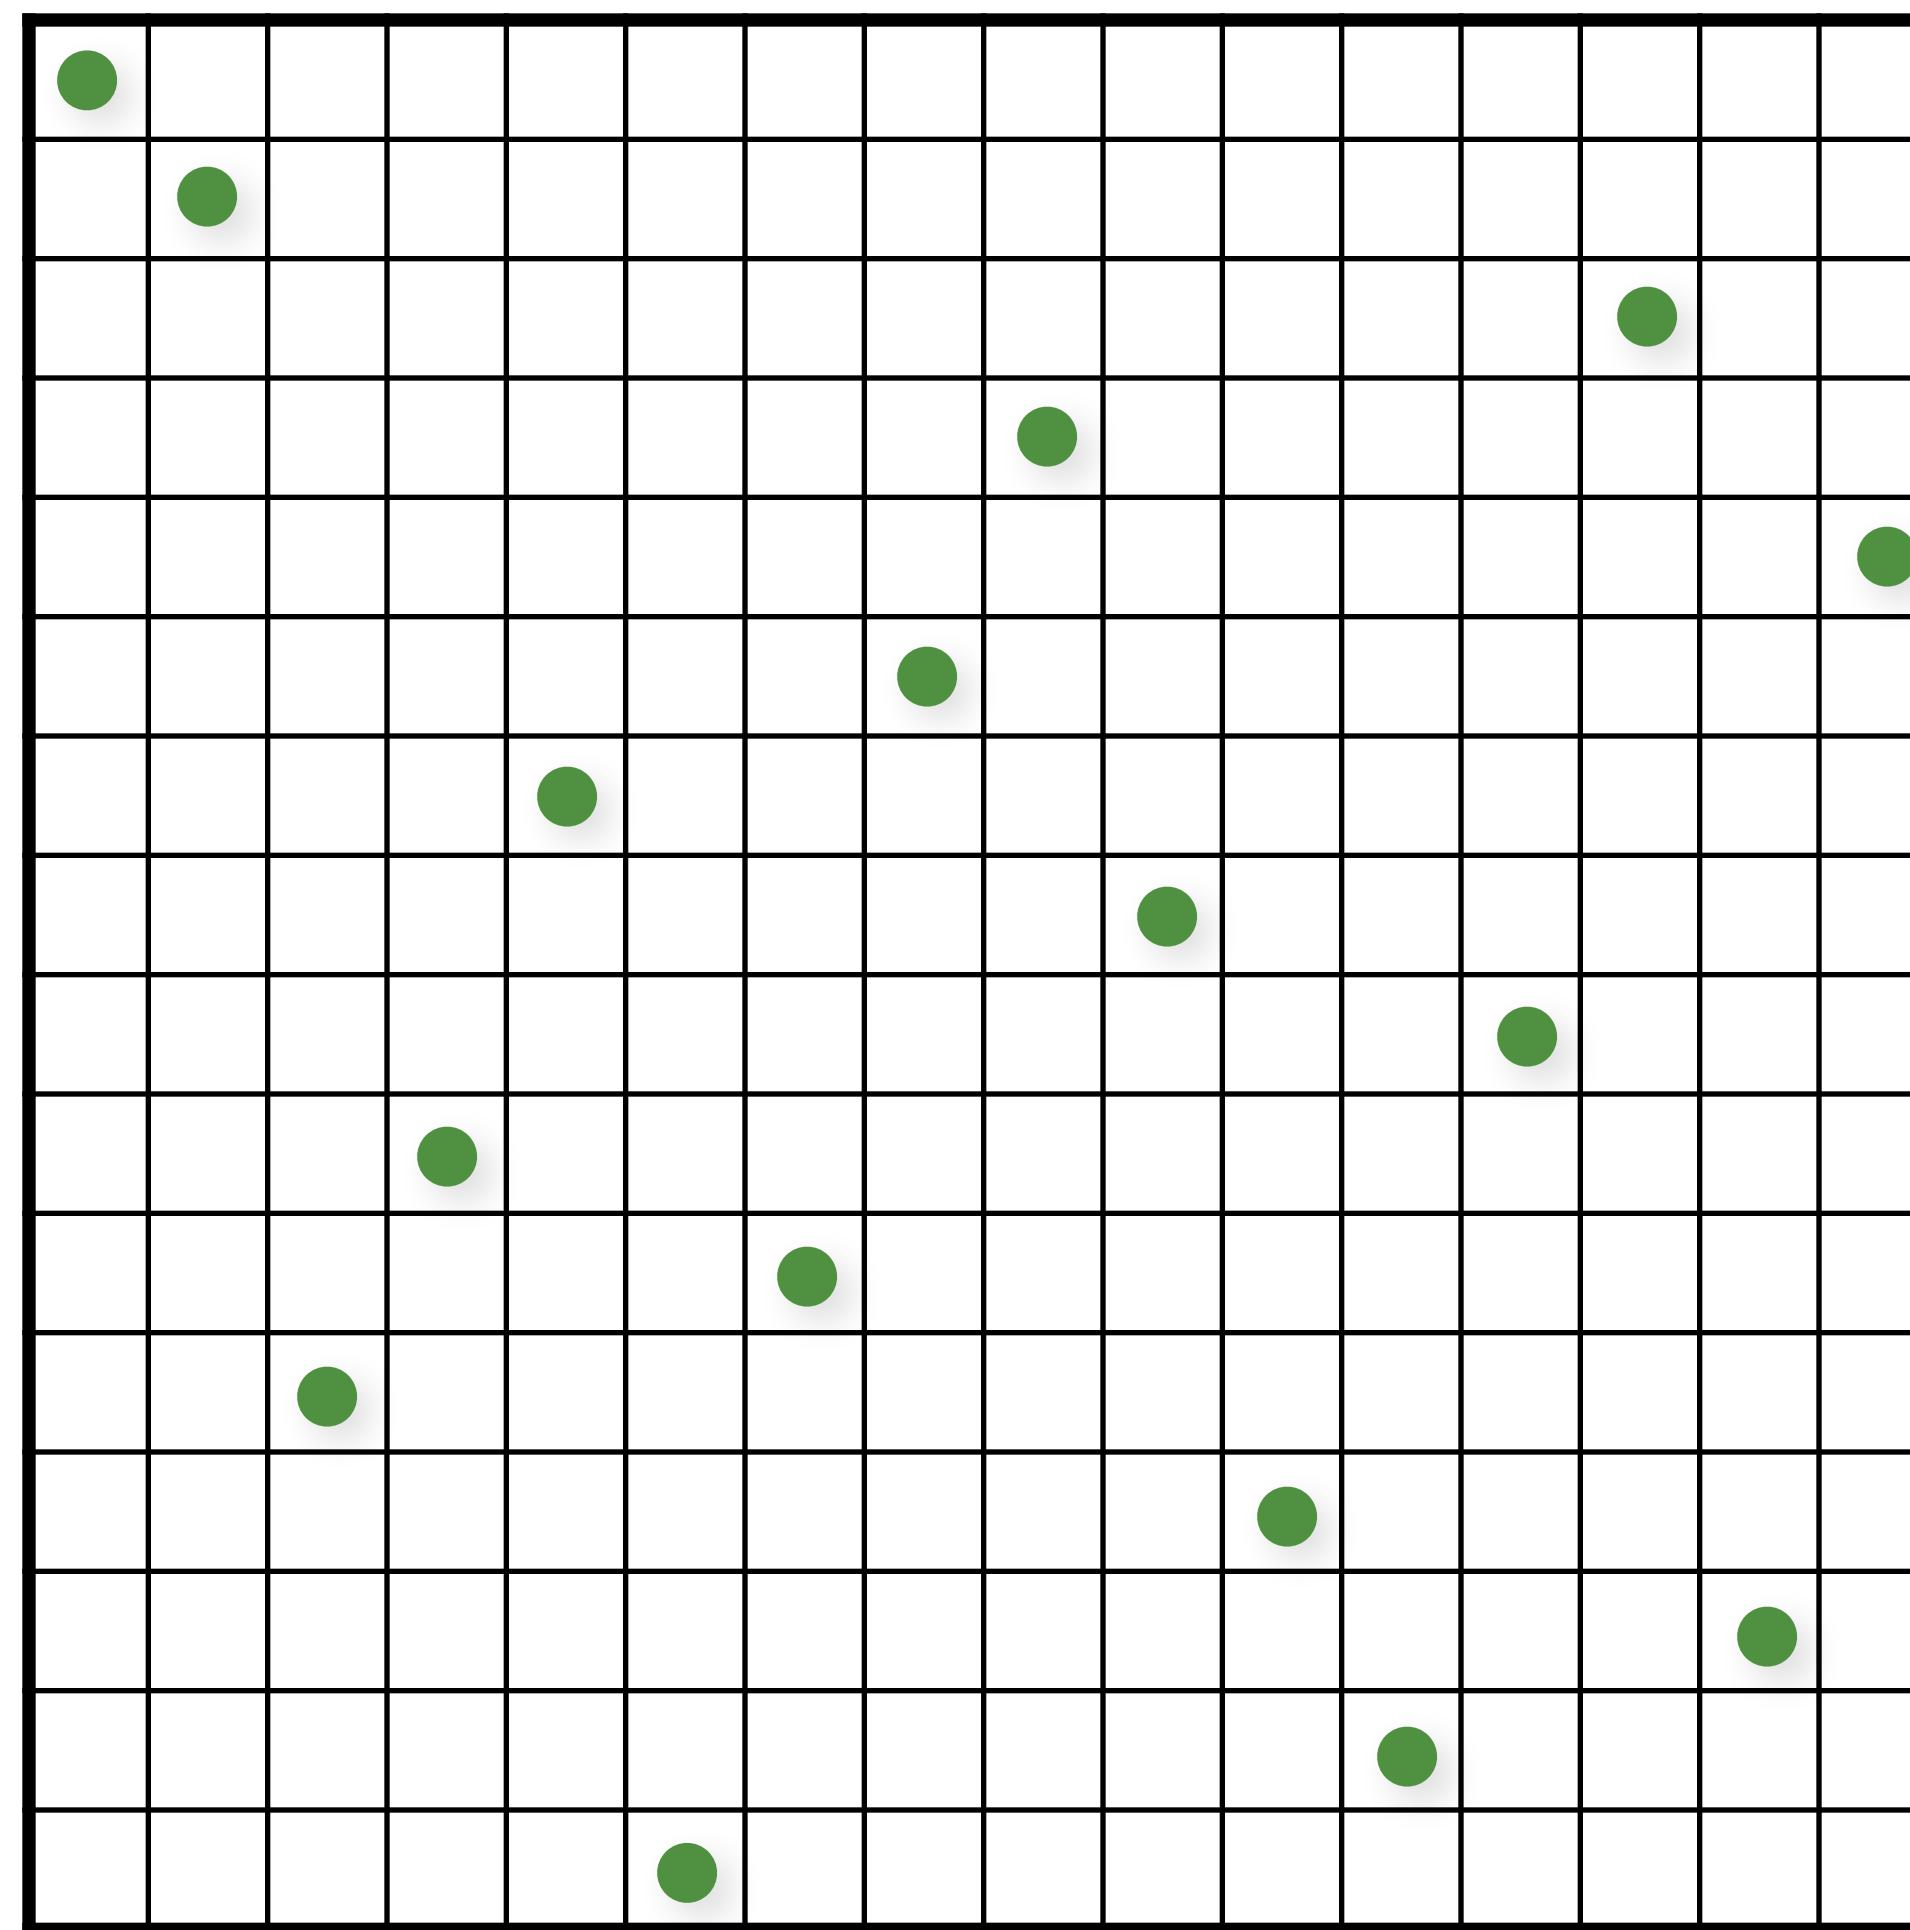
- for **2D**: 2 separate 1D jittered point sets

- combine dimensions
in random order



Latin Hypercube (N-Rooks) Sampling

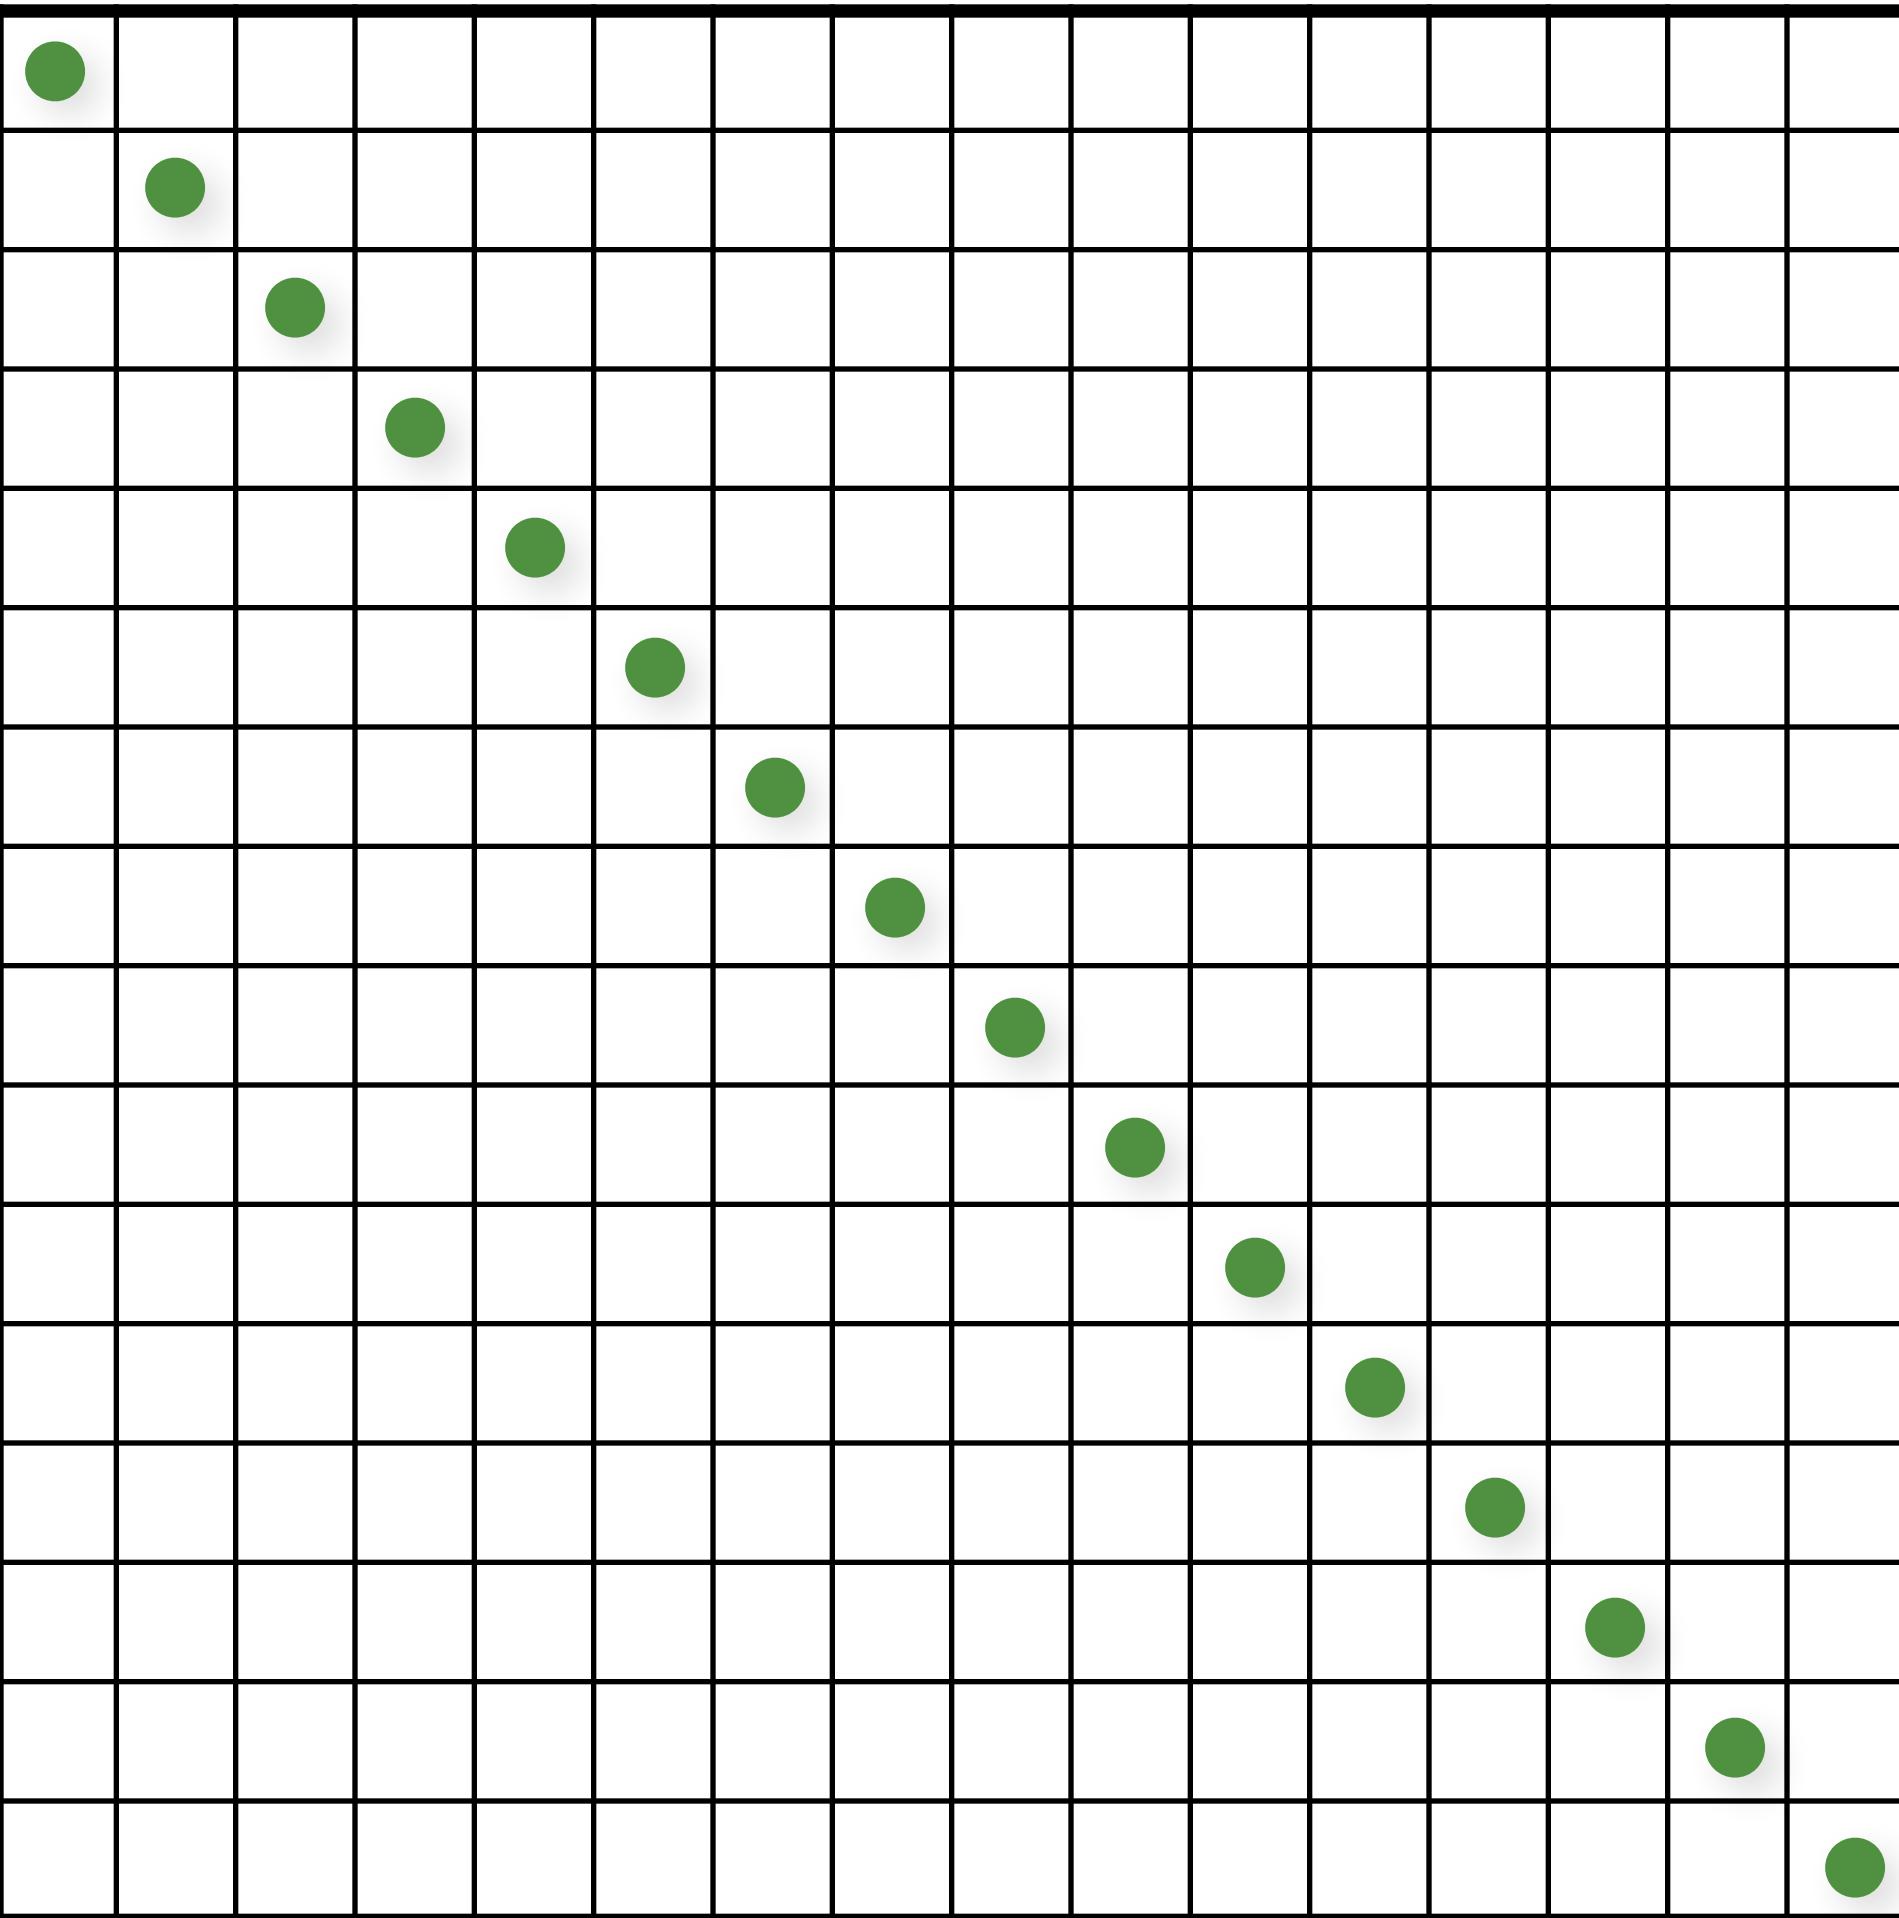
[Shirley 91]



Latin Hypercube (N-Rooks) Sampling

```
// initialize the diagonal
for (uint d = 0; d < numDimensions; d++)
    for (uint i = 0; i < numS; i++)
        samples(d,i) = (i + randf())/numS;
```

```
// shuffle each dimension independently
for (uint d = 0; d < numDimensions; d++)
    shuffle(samples(d,:));
```

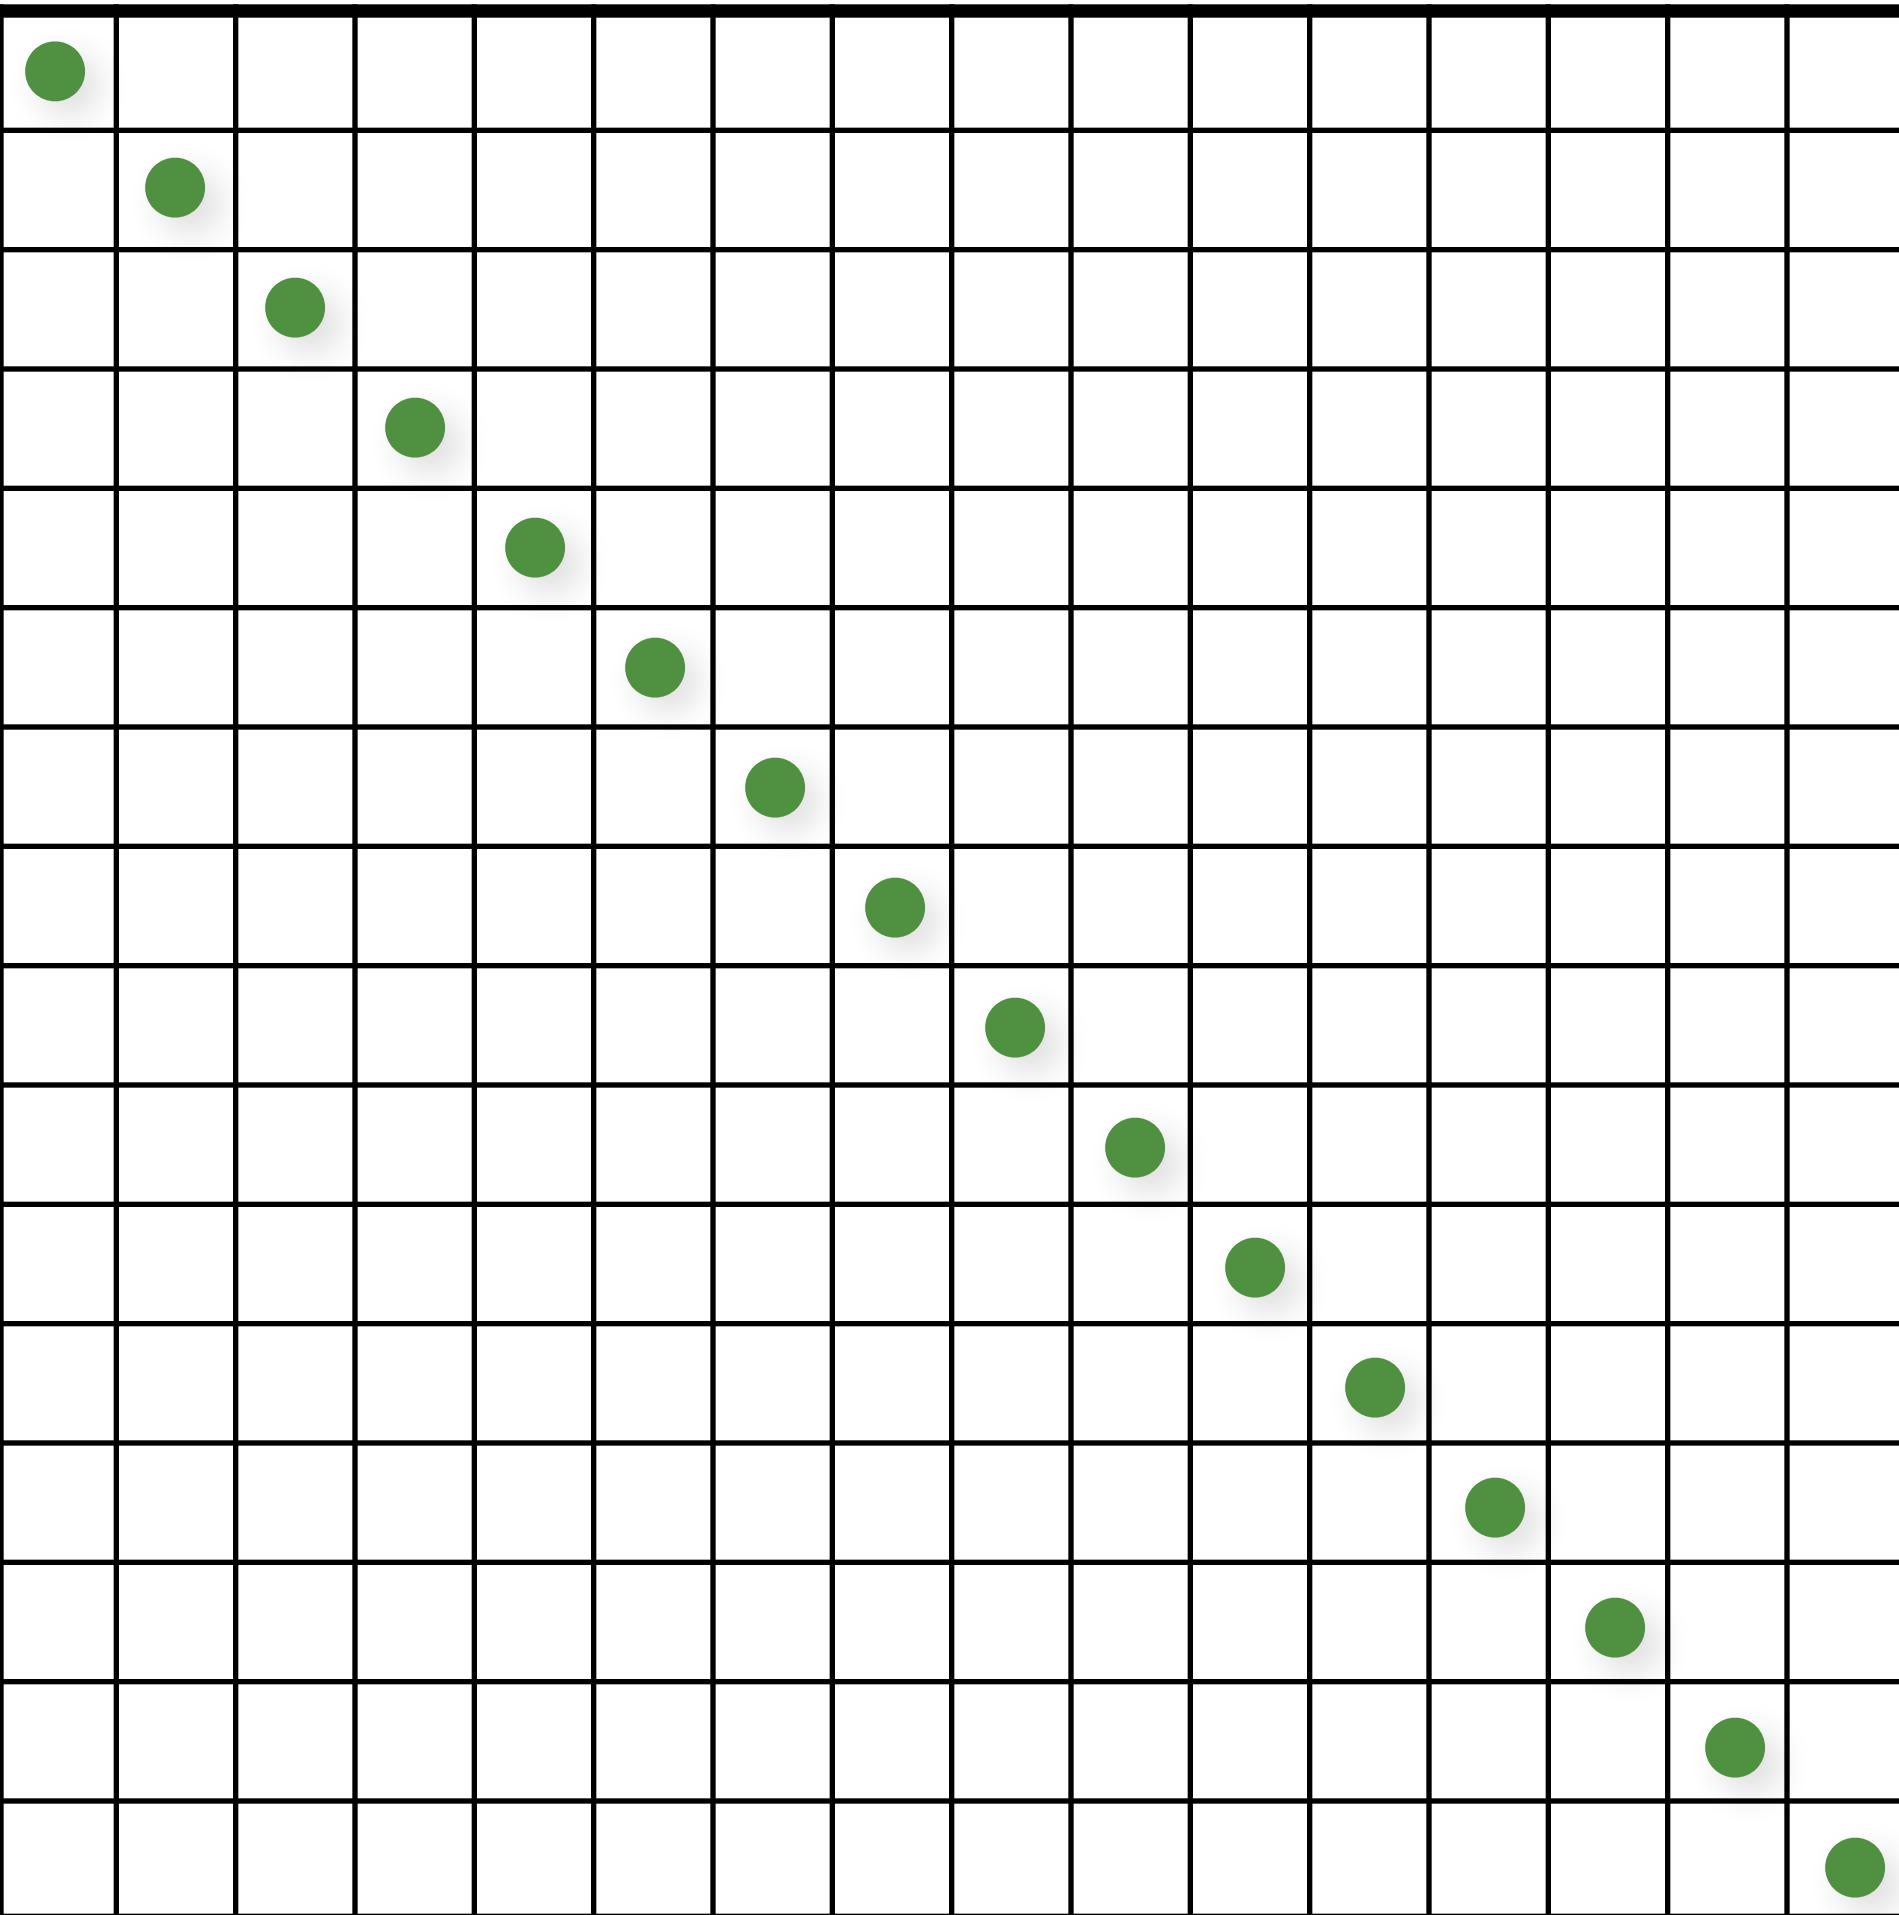


Initialize

Latin Hypercube (N-Rooks) Sampling

```
// initialize the diagonal
for (uint d = 0; d < numDimensions; d++)
    for (uint i = 0; i < numS; i++)
        samples(d,i) = (i + randf())/numS;

// shuffle each dimension independently
for (uint d = 0; d < numDimensions; d++)
    shuffle(samples(d,:));
```

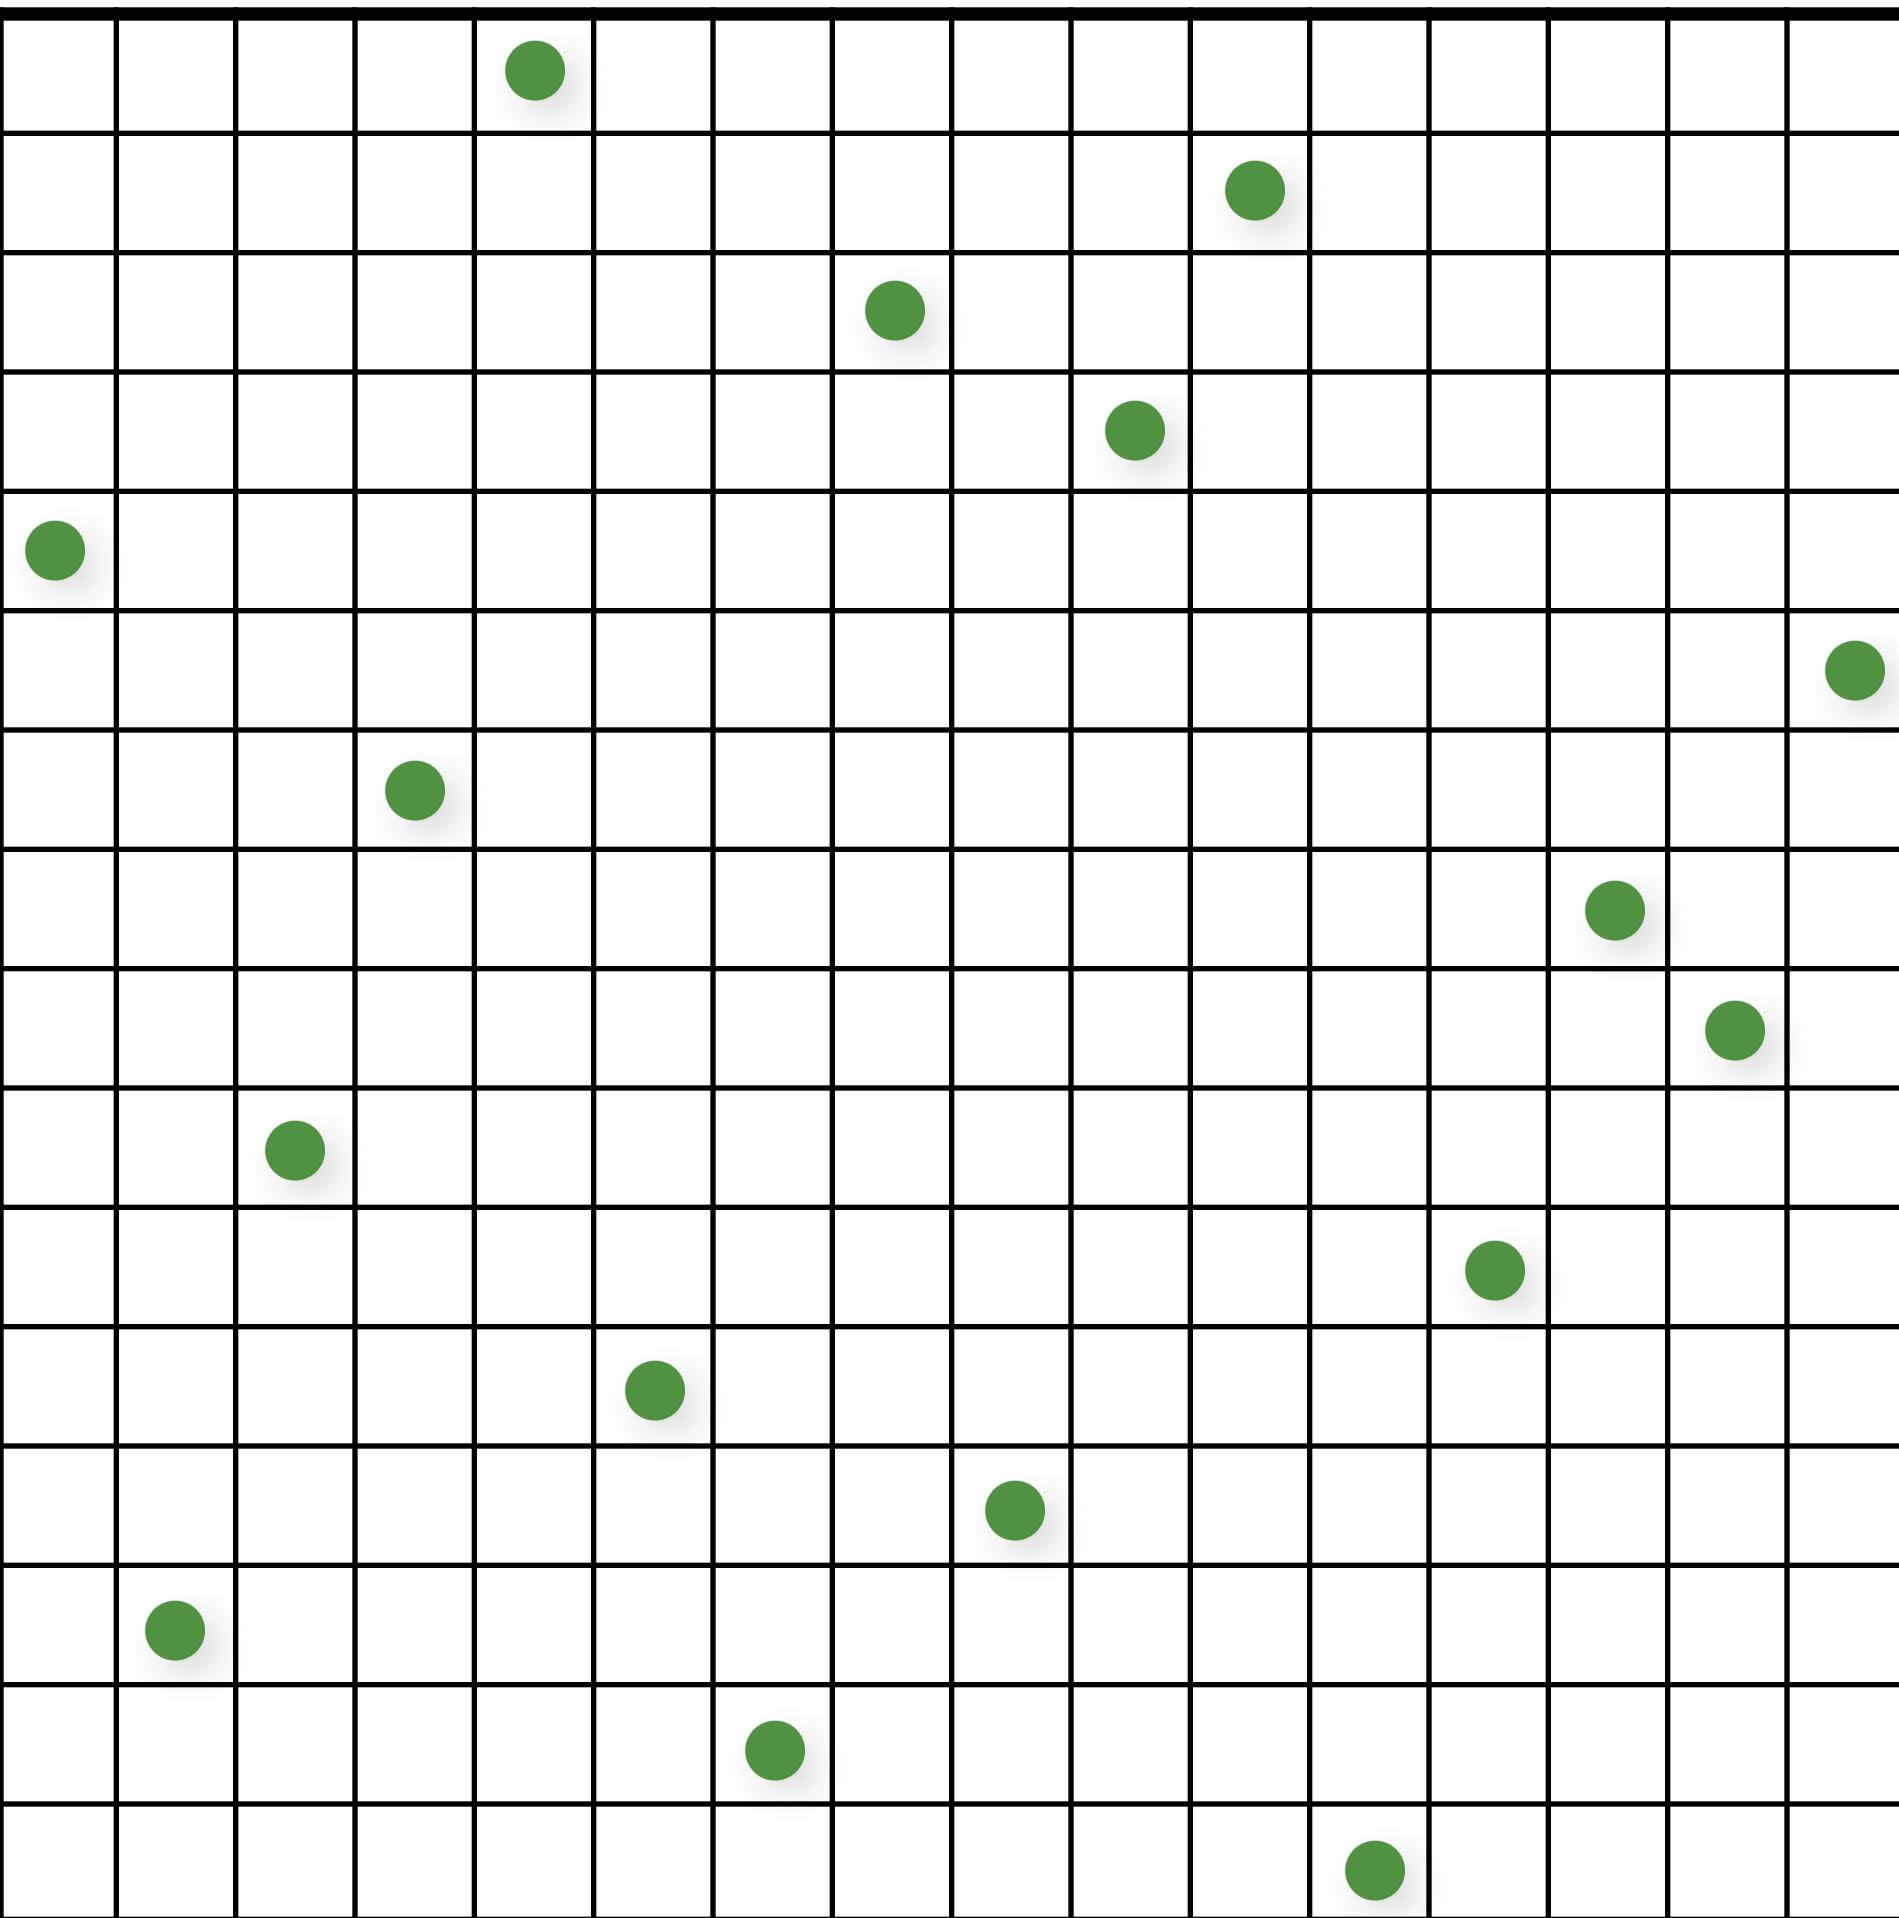


Shuffle rows

Latin Hypercube (N-Rooks) Sampling

```
// initialize the diagonal
for (uint d = 0; d < numDimensions; d++)
    for (uint i = 0; i < numS; i++)
        samples(d,i) = (i + randf())/numS;

// shuffle each dimension independently
for (uint d = 0; d < numDimensions; d++)
    shuffle(samples(d,:));
```

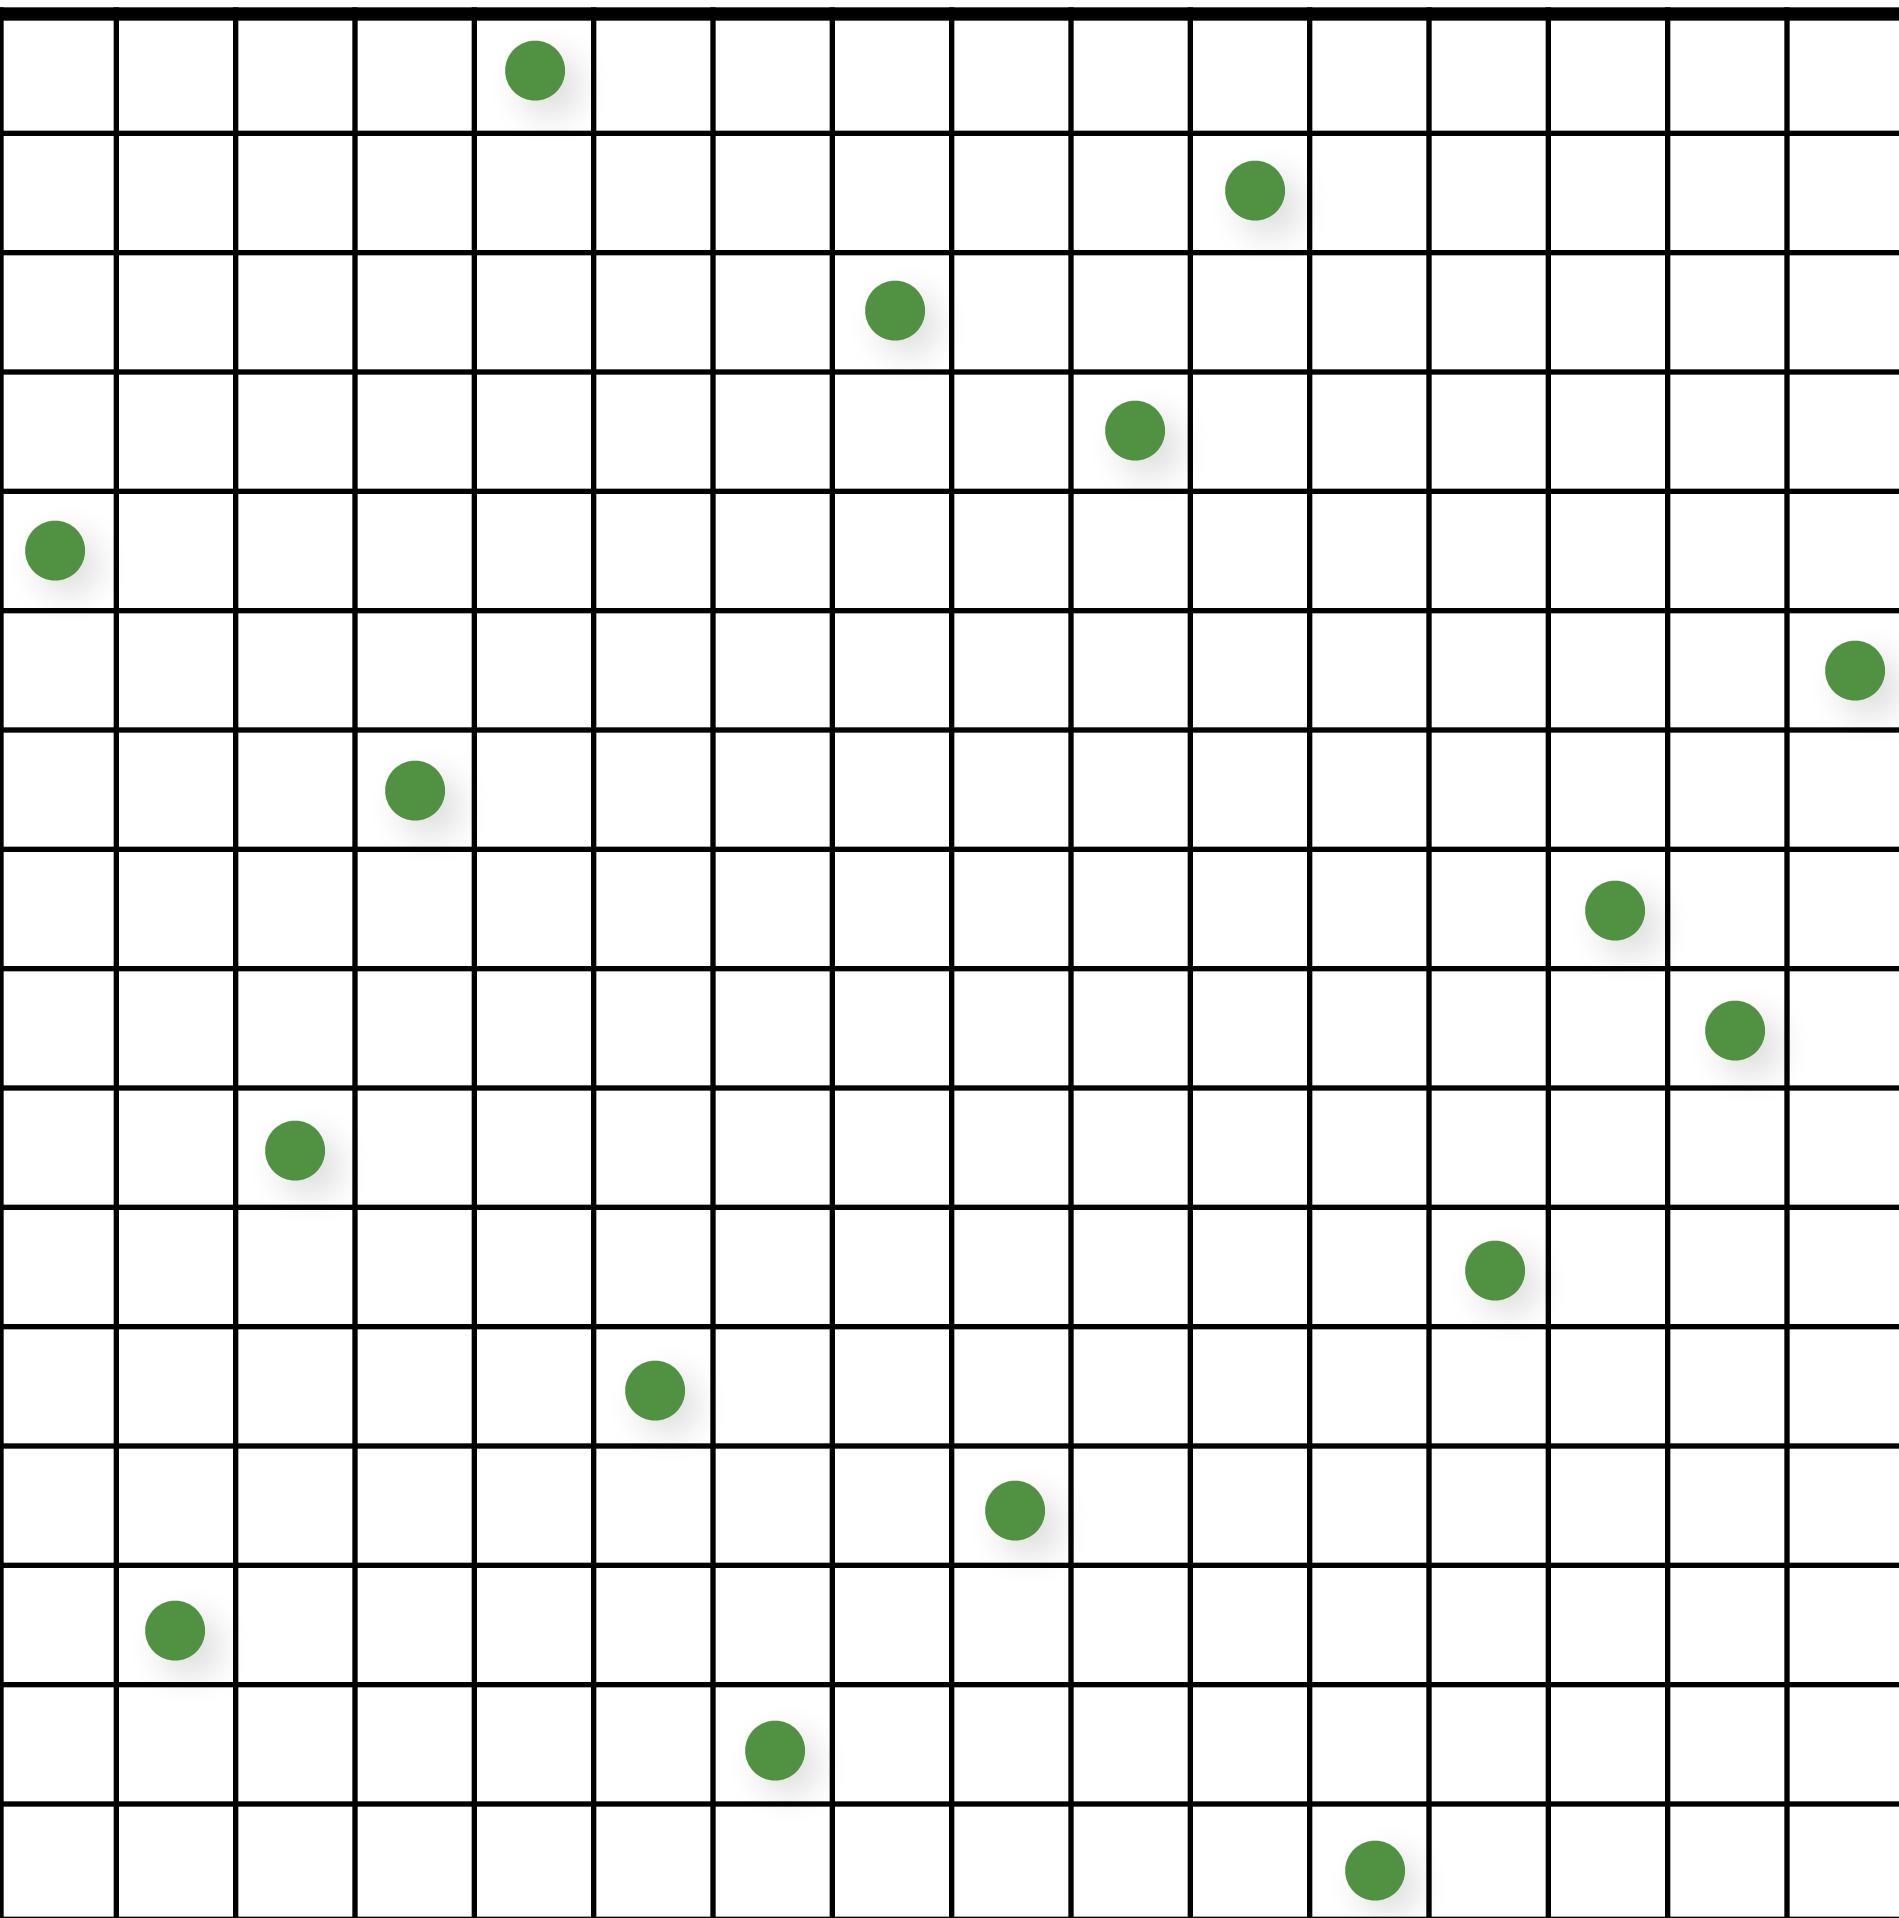


Shuffle rows

Latin Hypercube (N-Rooks) Sampling

```
// initialize the diagonal
for (uint d = 0; d < numDimensions; d++)
    for (uint i = 0; i < numS; i++)
        samples(d,i) = (i + randf())/numS;

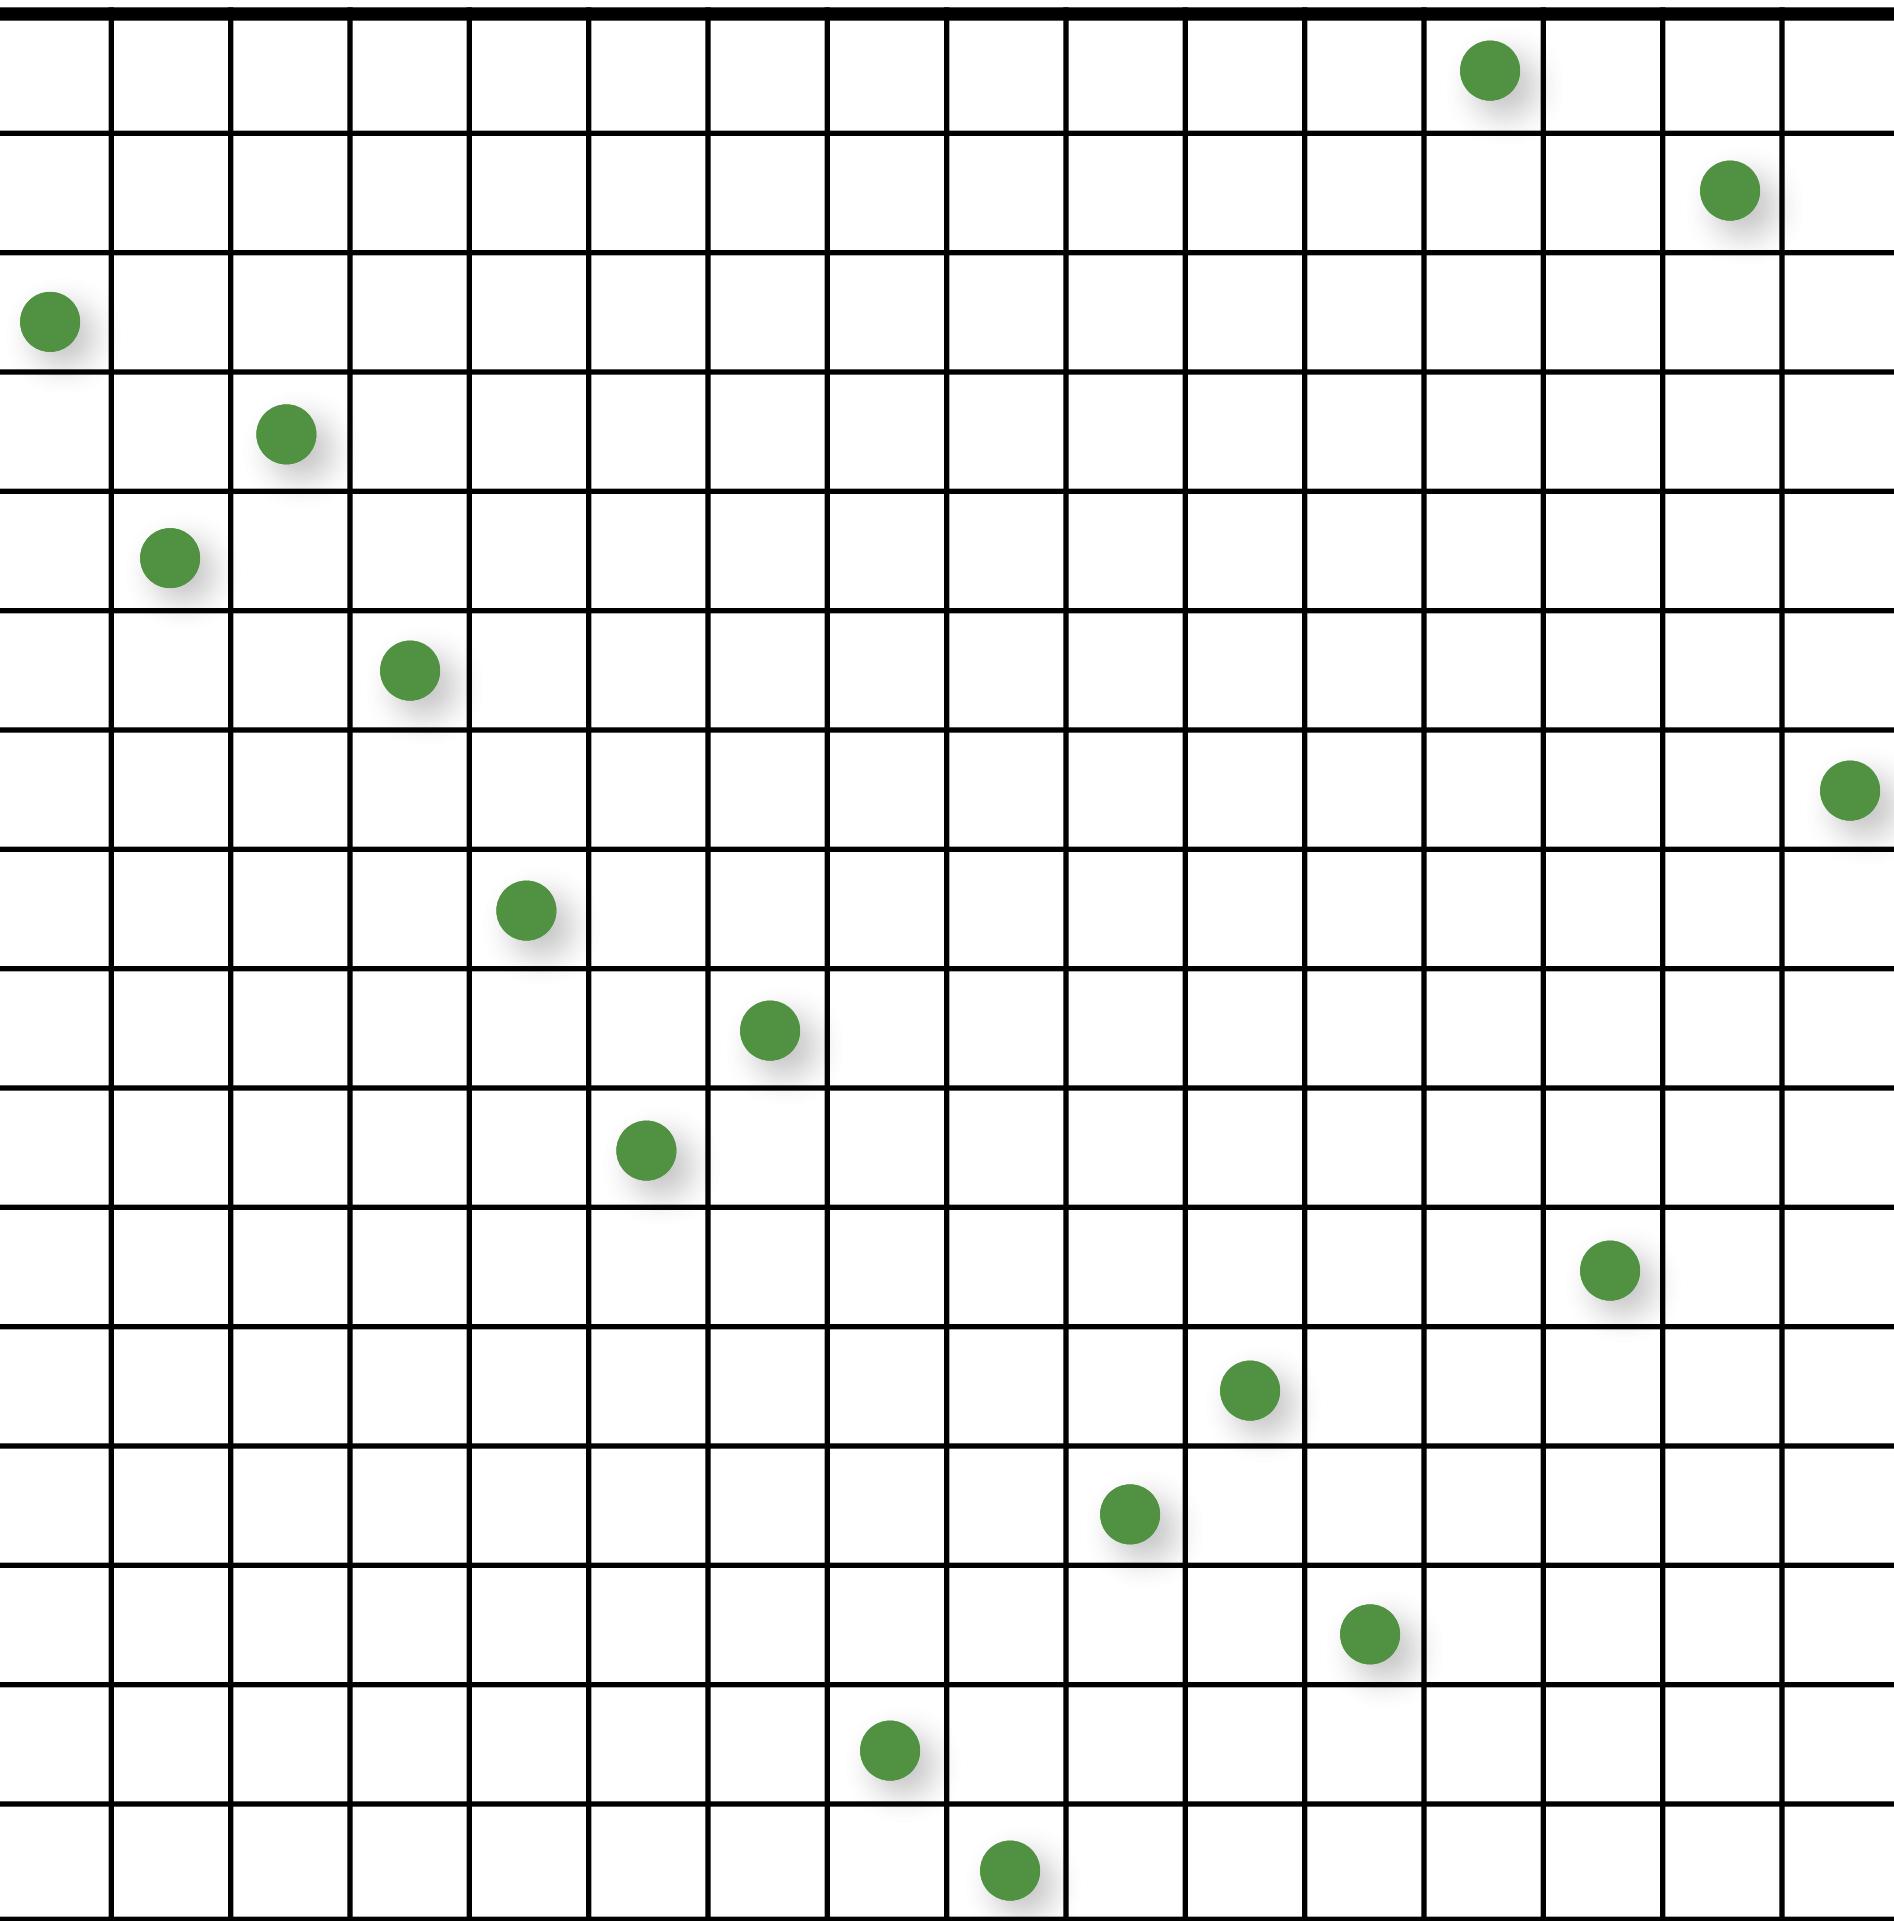
// shuffle each dimension independently
for (uint d = 0; d < numDimensions; d++)
    shuffle(samples(d,:));
```



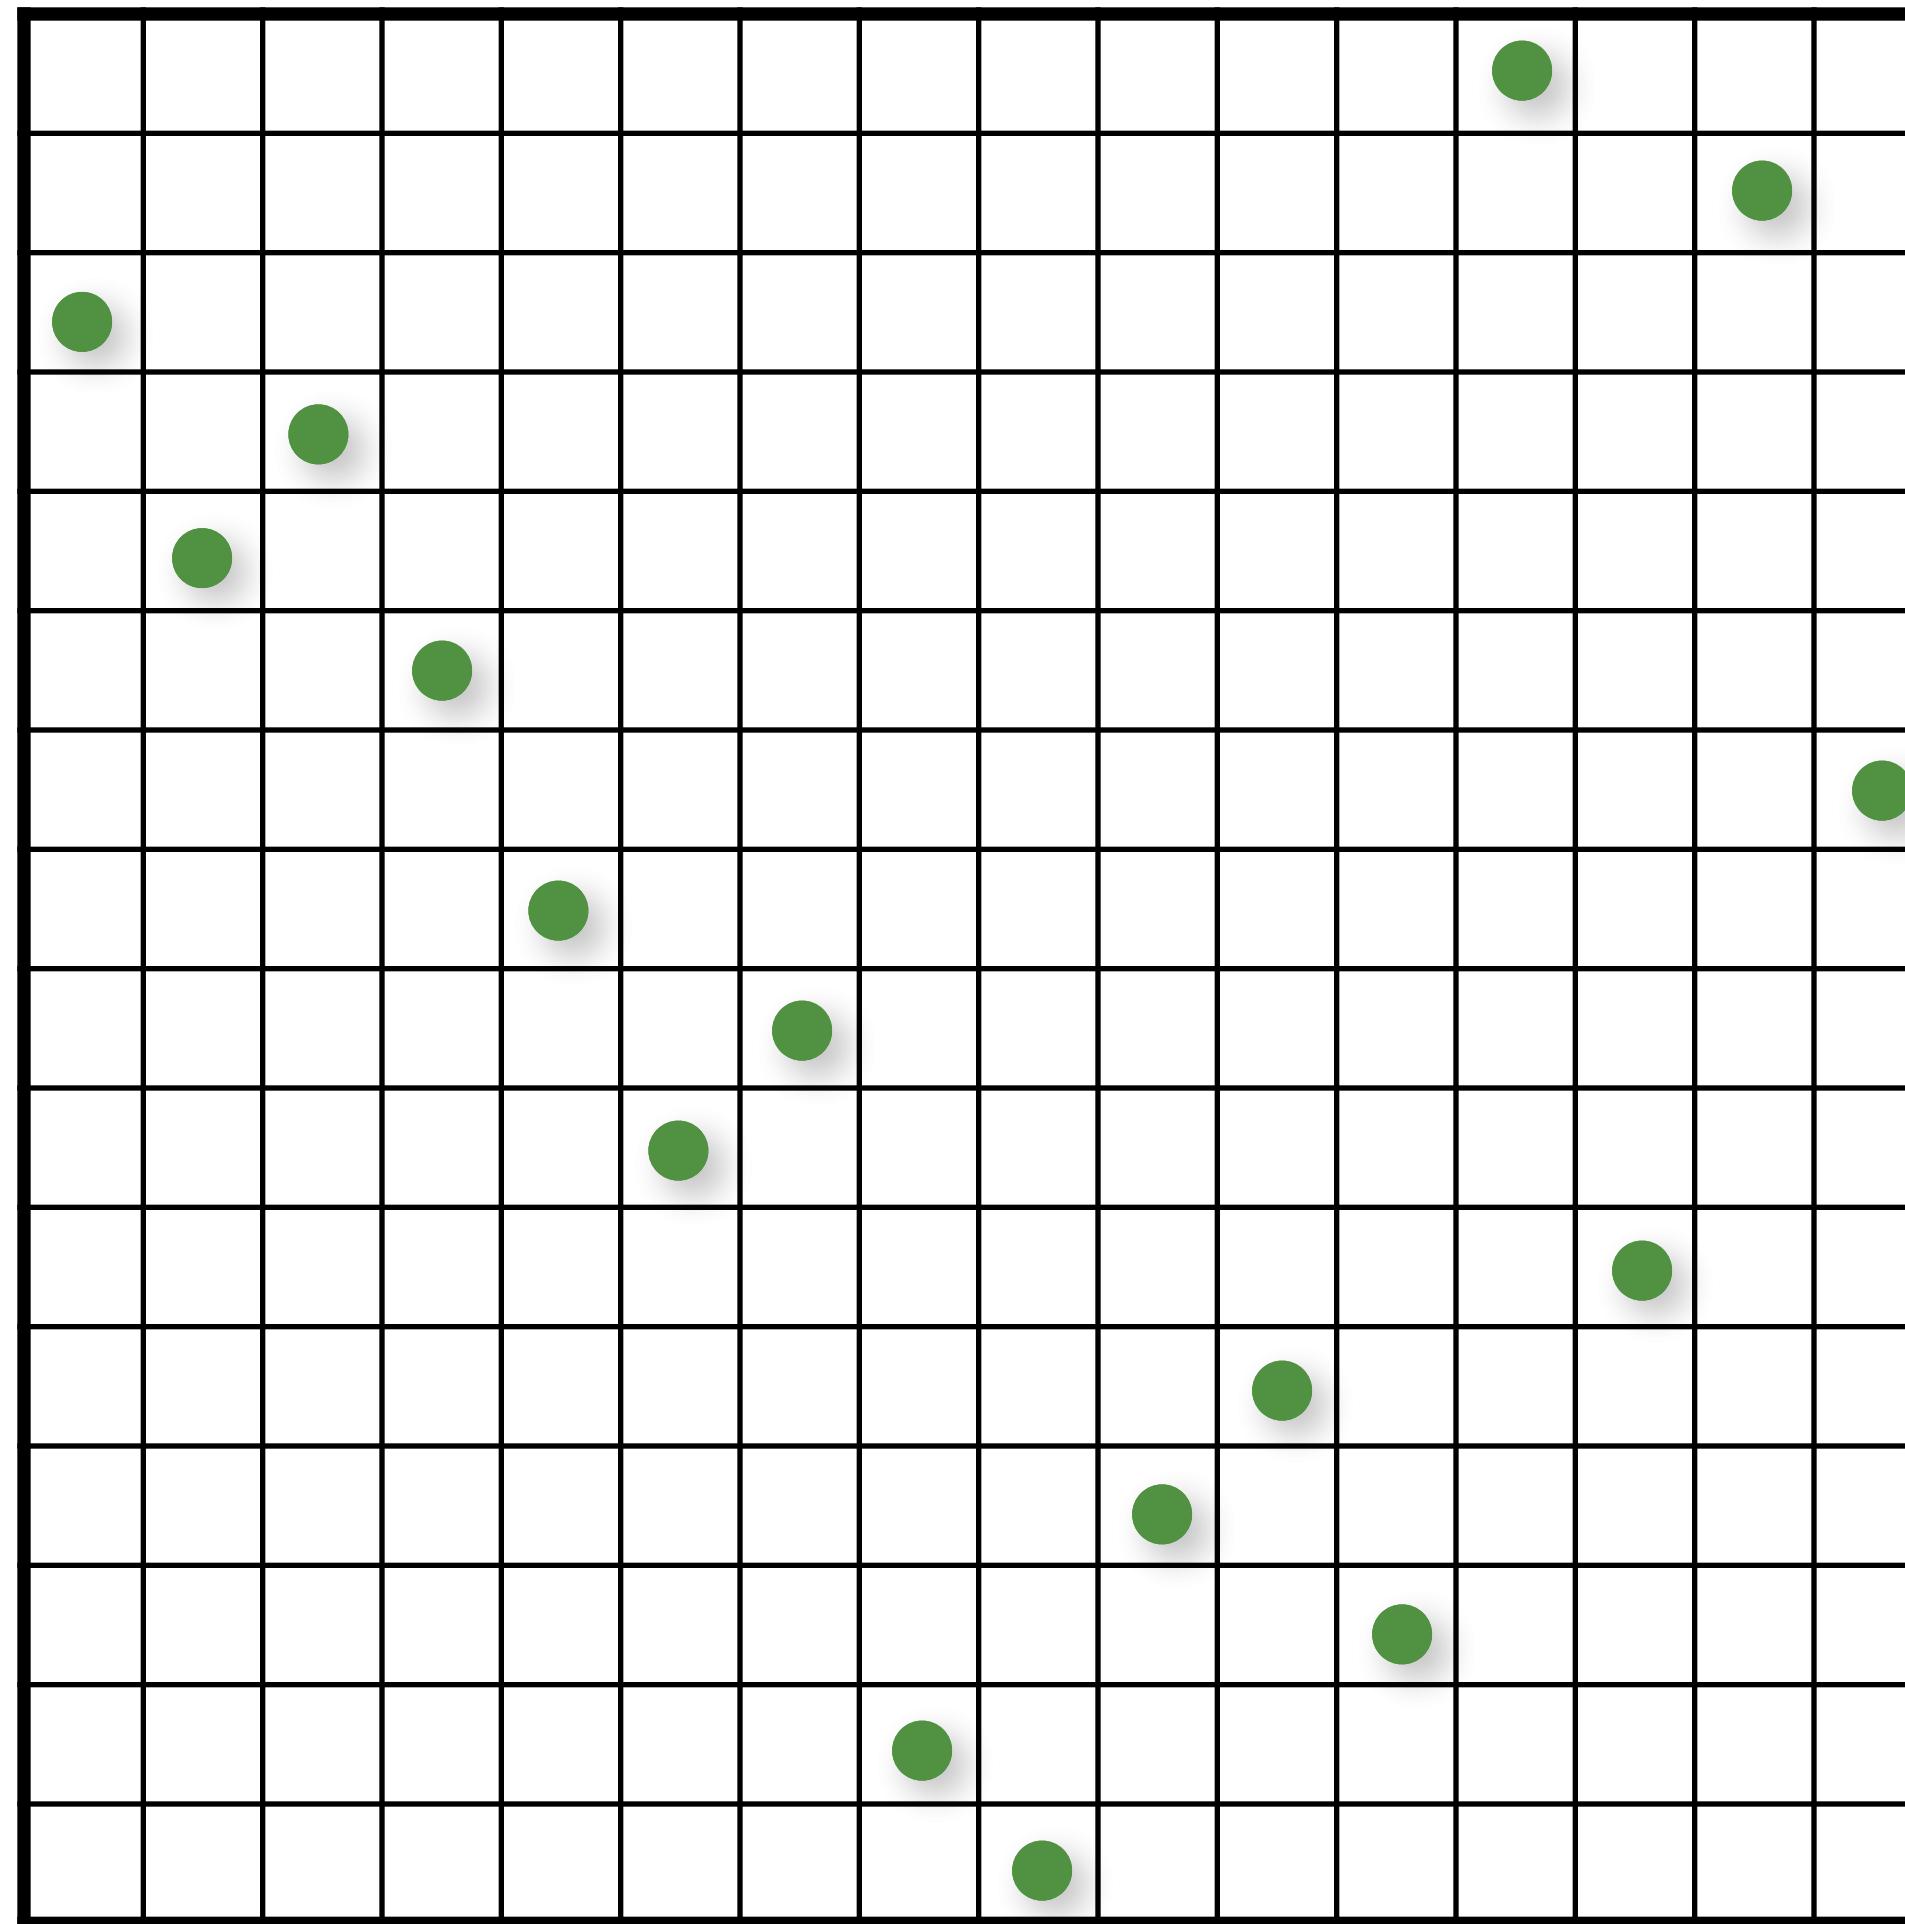
Latin Hypercube (N-Rooks) Sampling

```
// initialize the diagonal
for (uint d = 0; d < numDimensions; d++)
    for (uint i = 0; i < numS; i++)
        samples(d,i) = (i + randf())/numS;

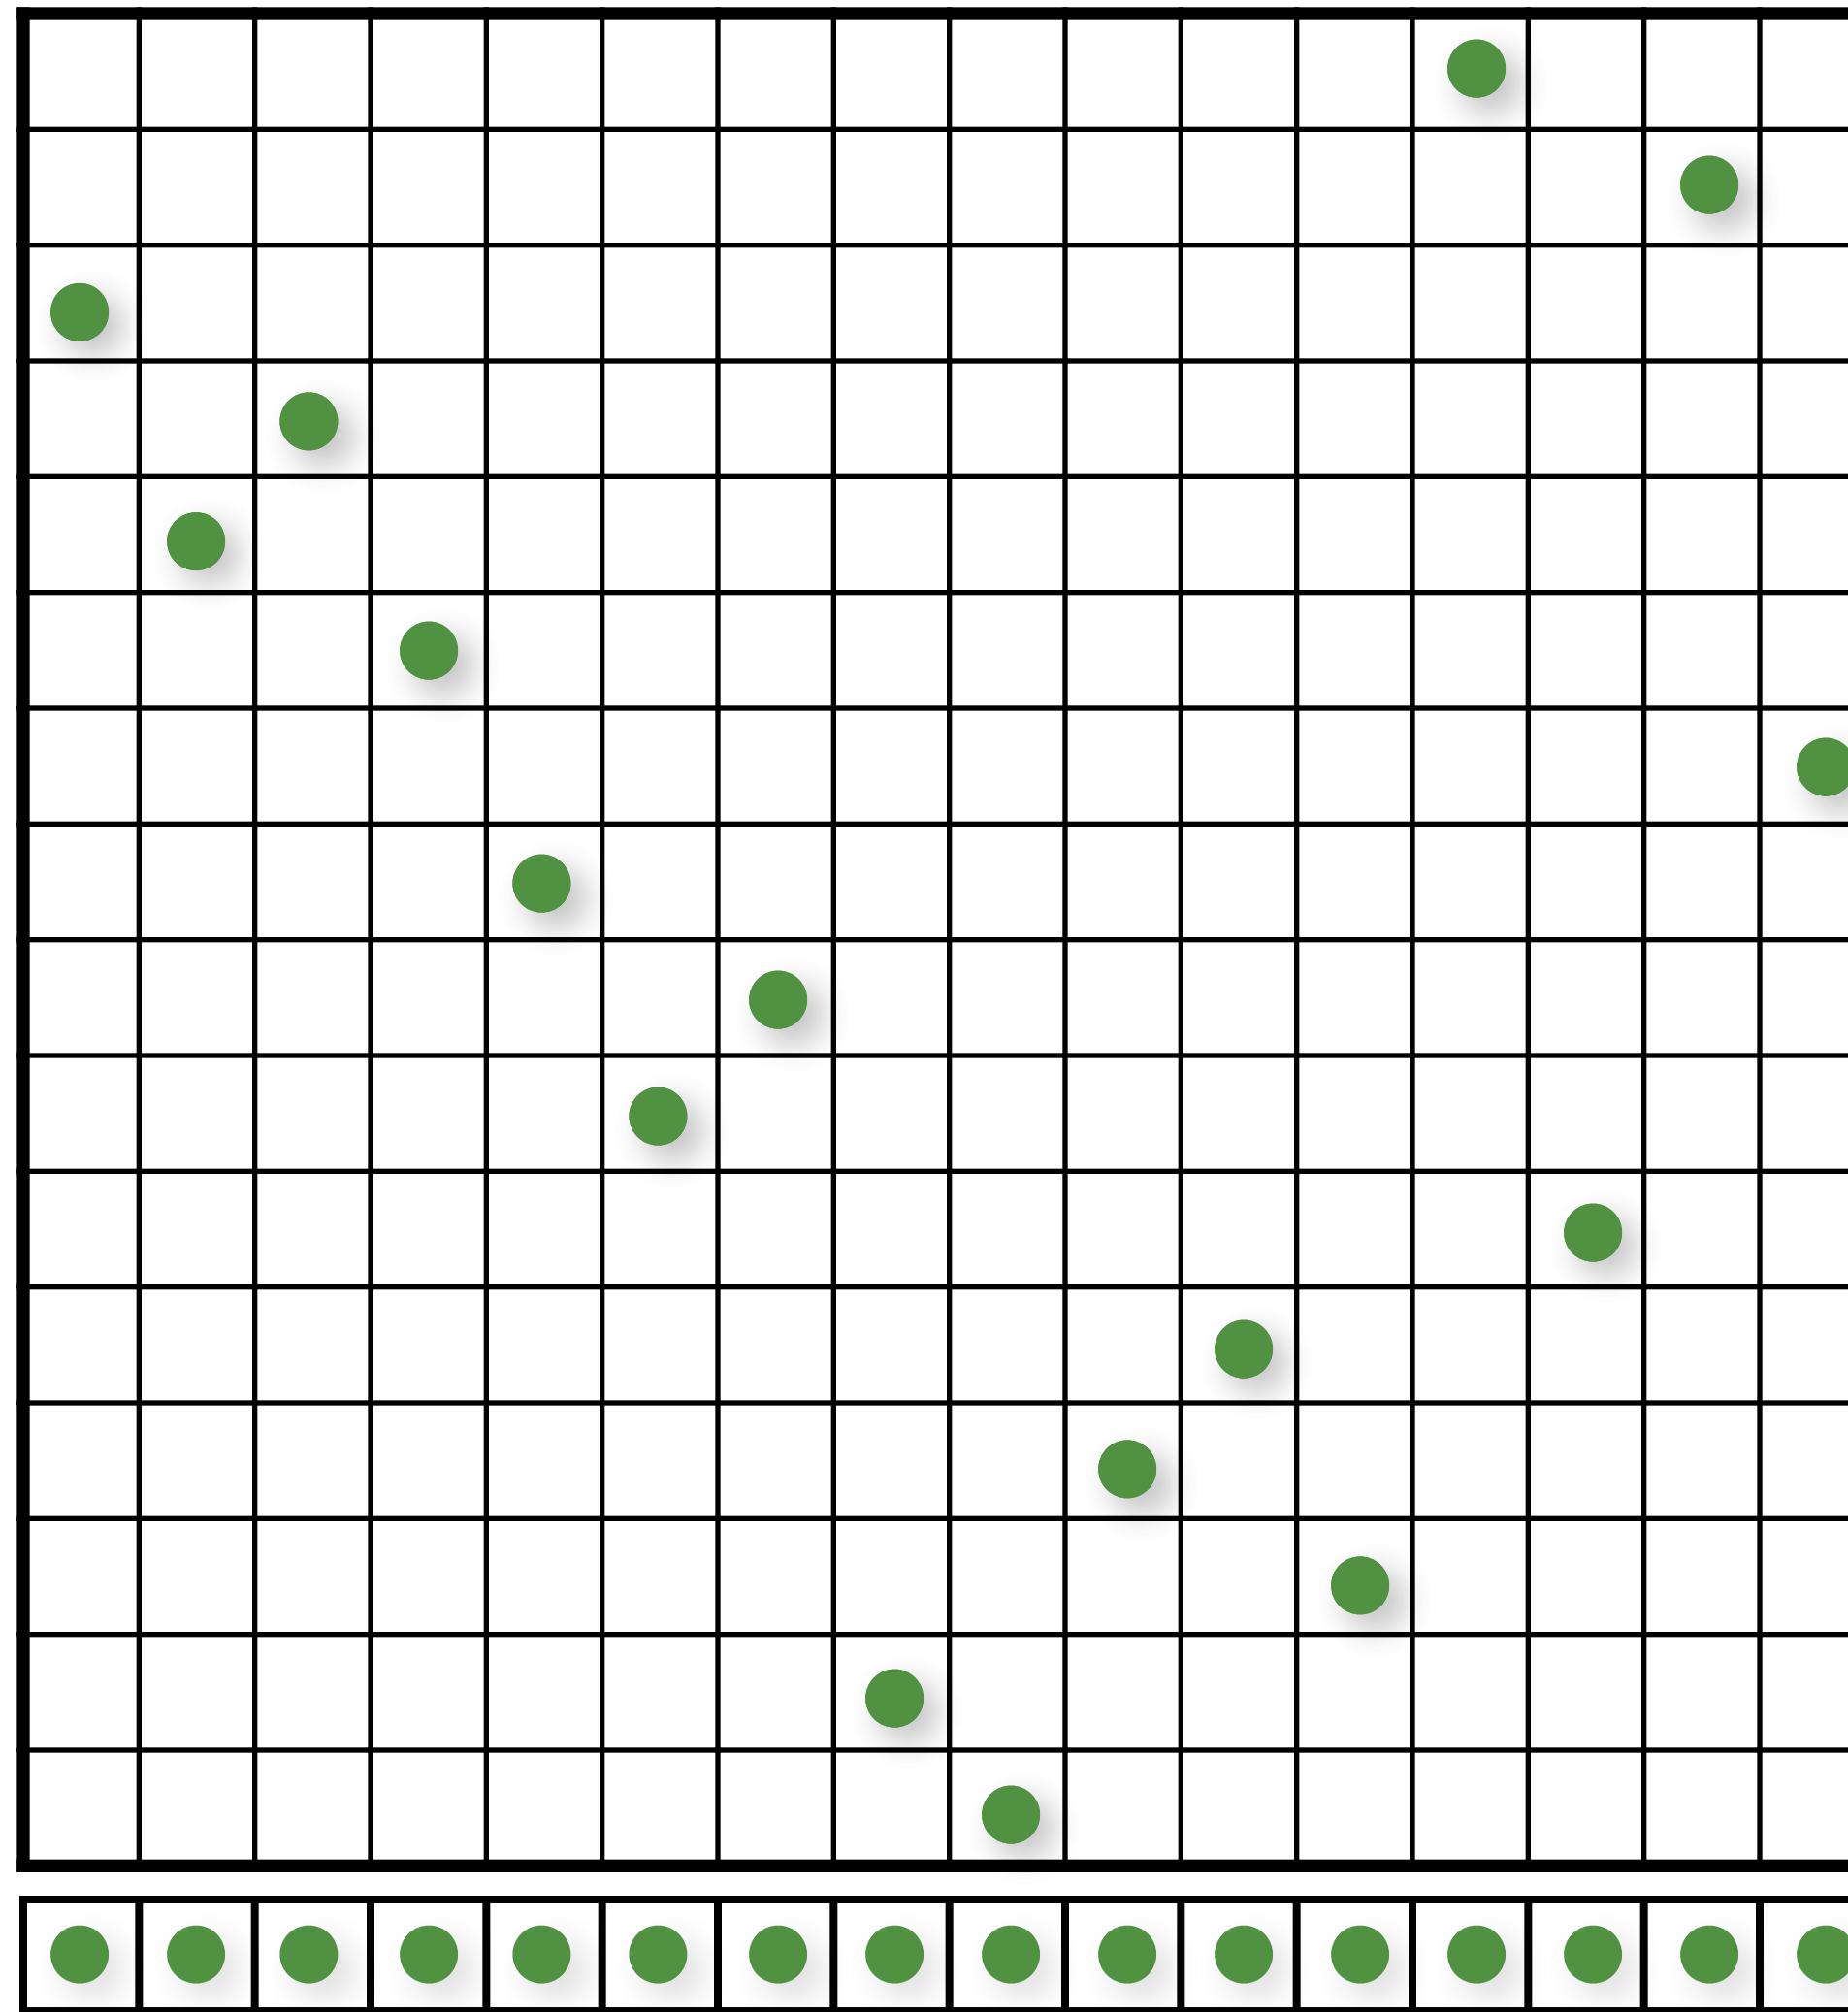
// shuffle each dimension independently
for (uint d = 0; d < numDimensions; d++)
    shuffle(samples(d,:));
```



Latin Hypercube (N-Rooks) Sampling



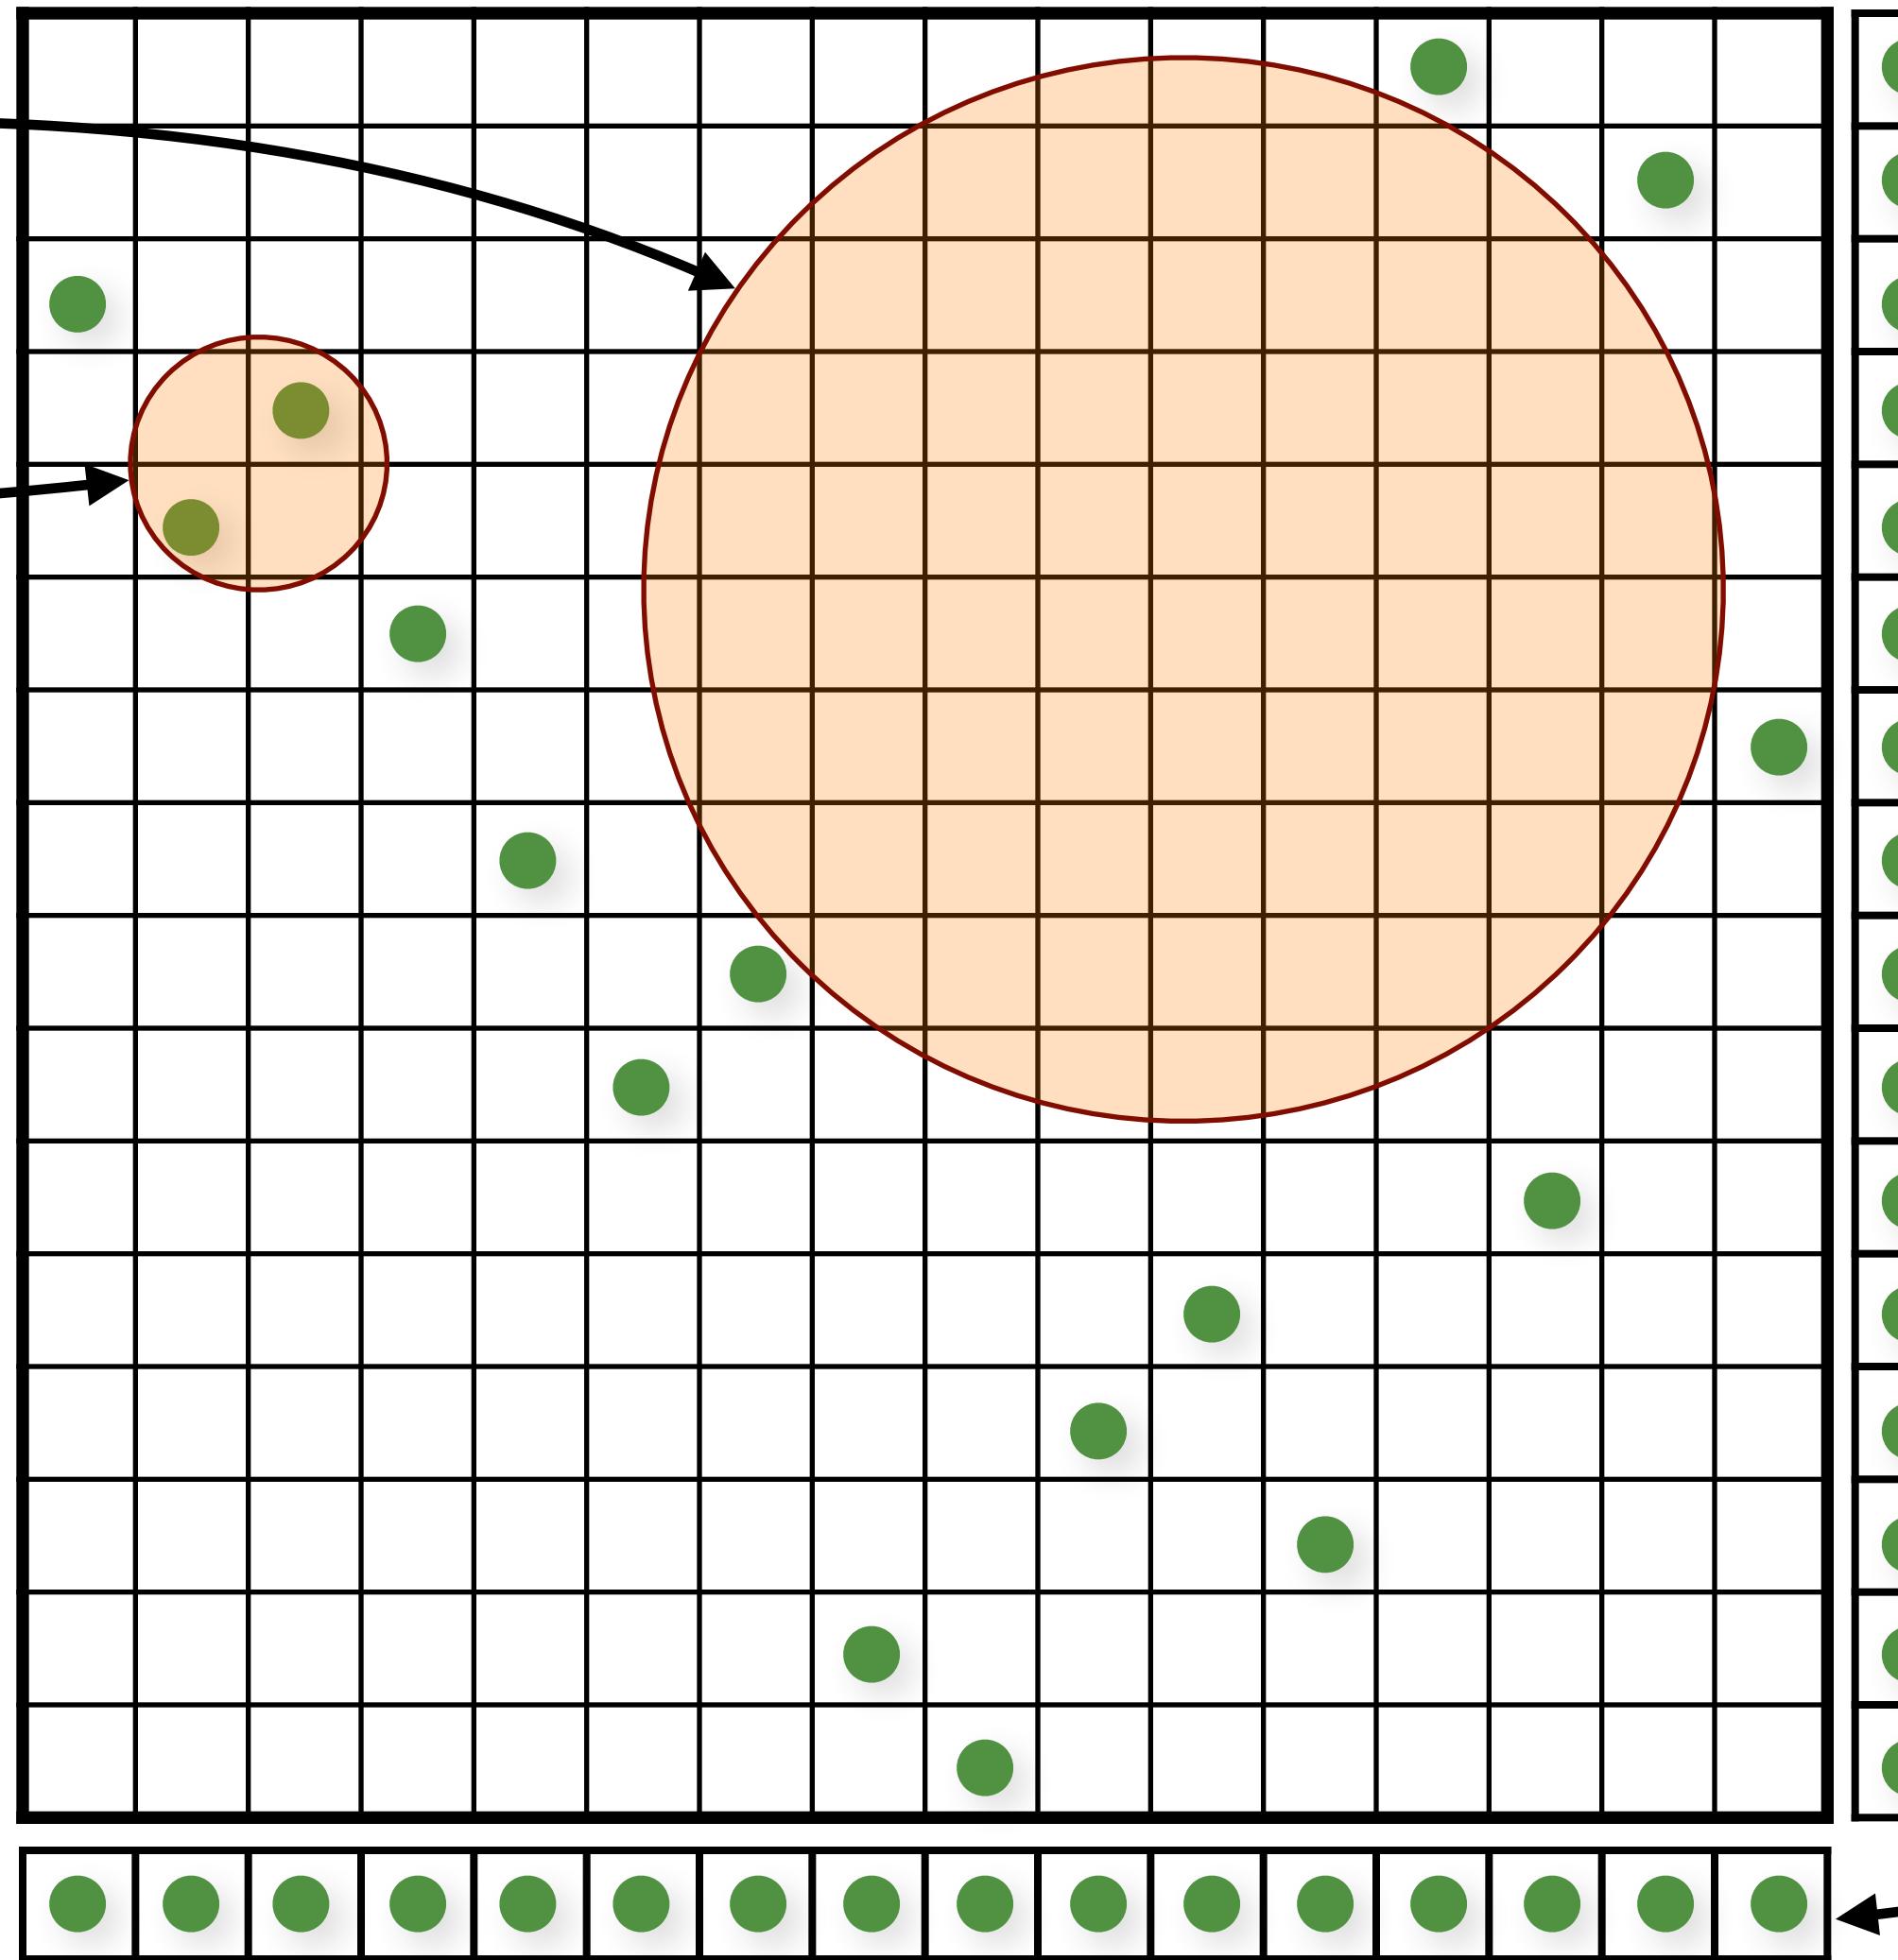
Latin Hypercube (N-Rooks) Sampling



Latin Hypercube (N-Rooks) Sampling

Unevenly distributed
in n-dimensions

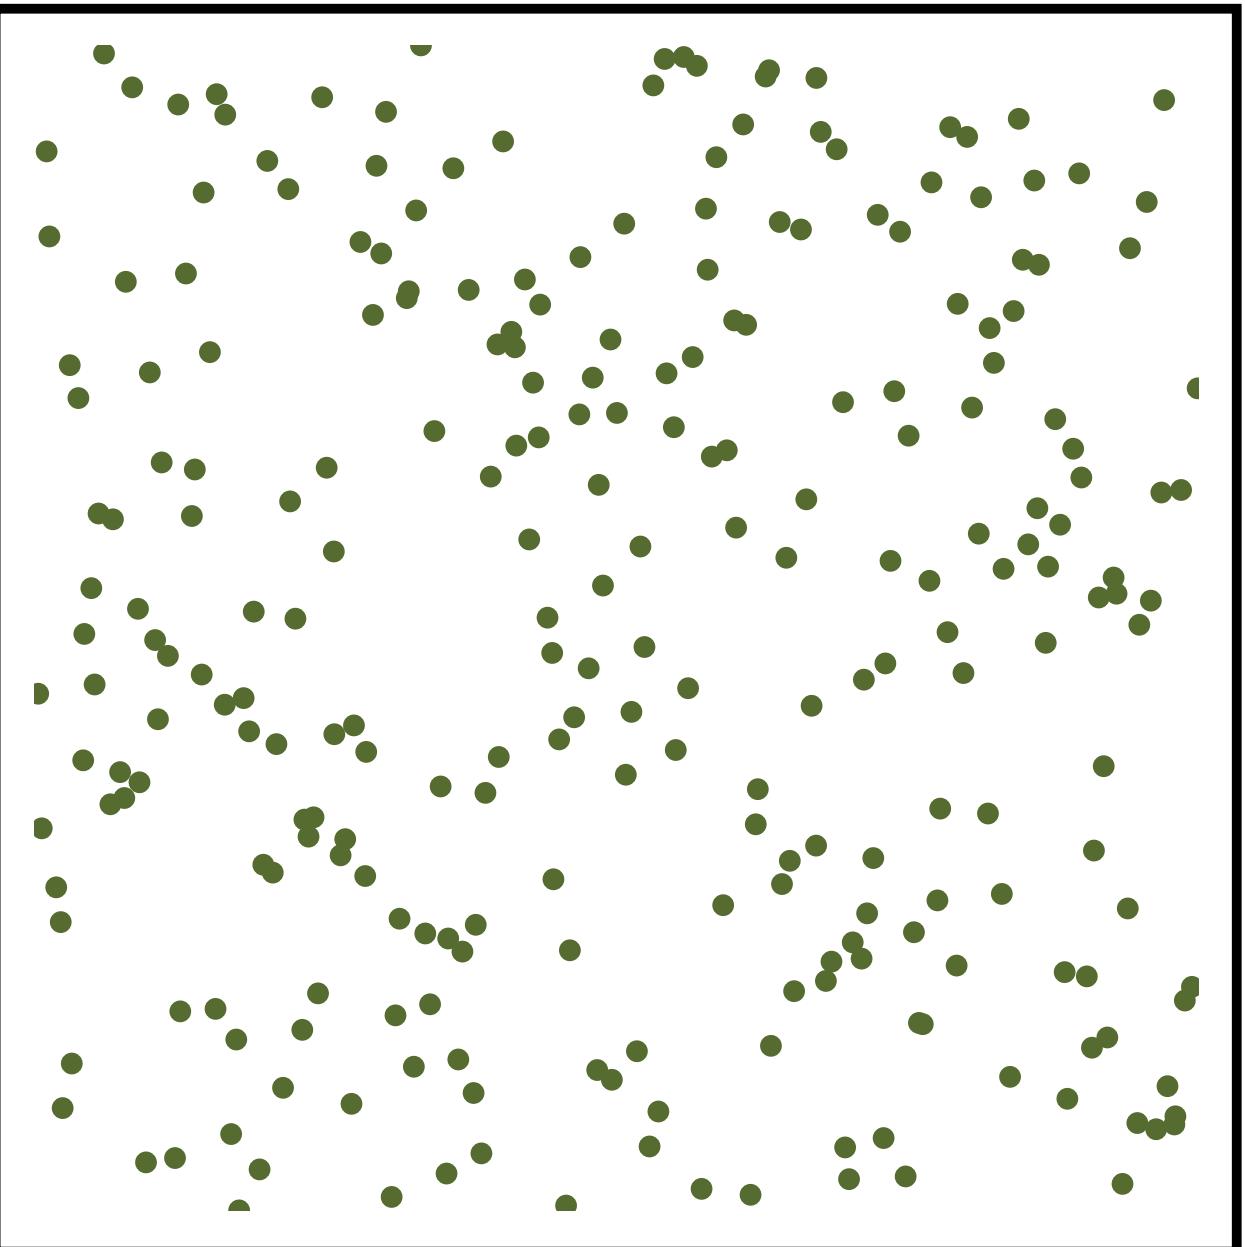
"even distributed"



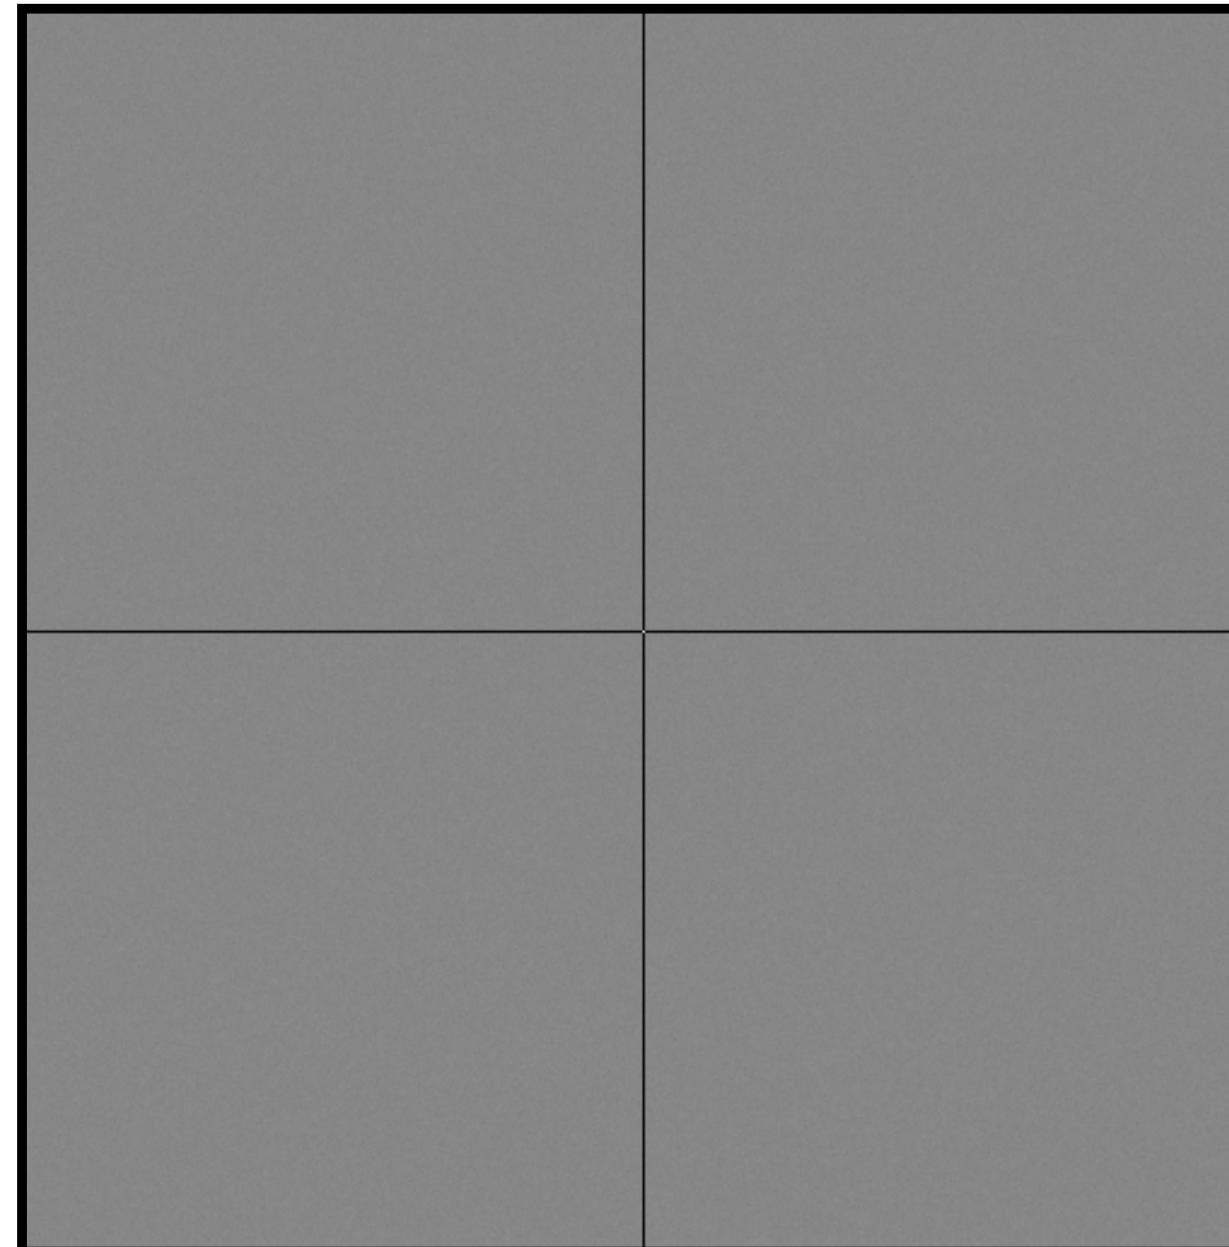
Evenly distributed in each
individual dimension

N-Rooks Sampling

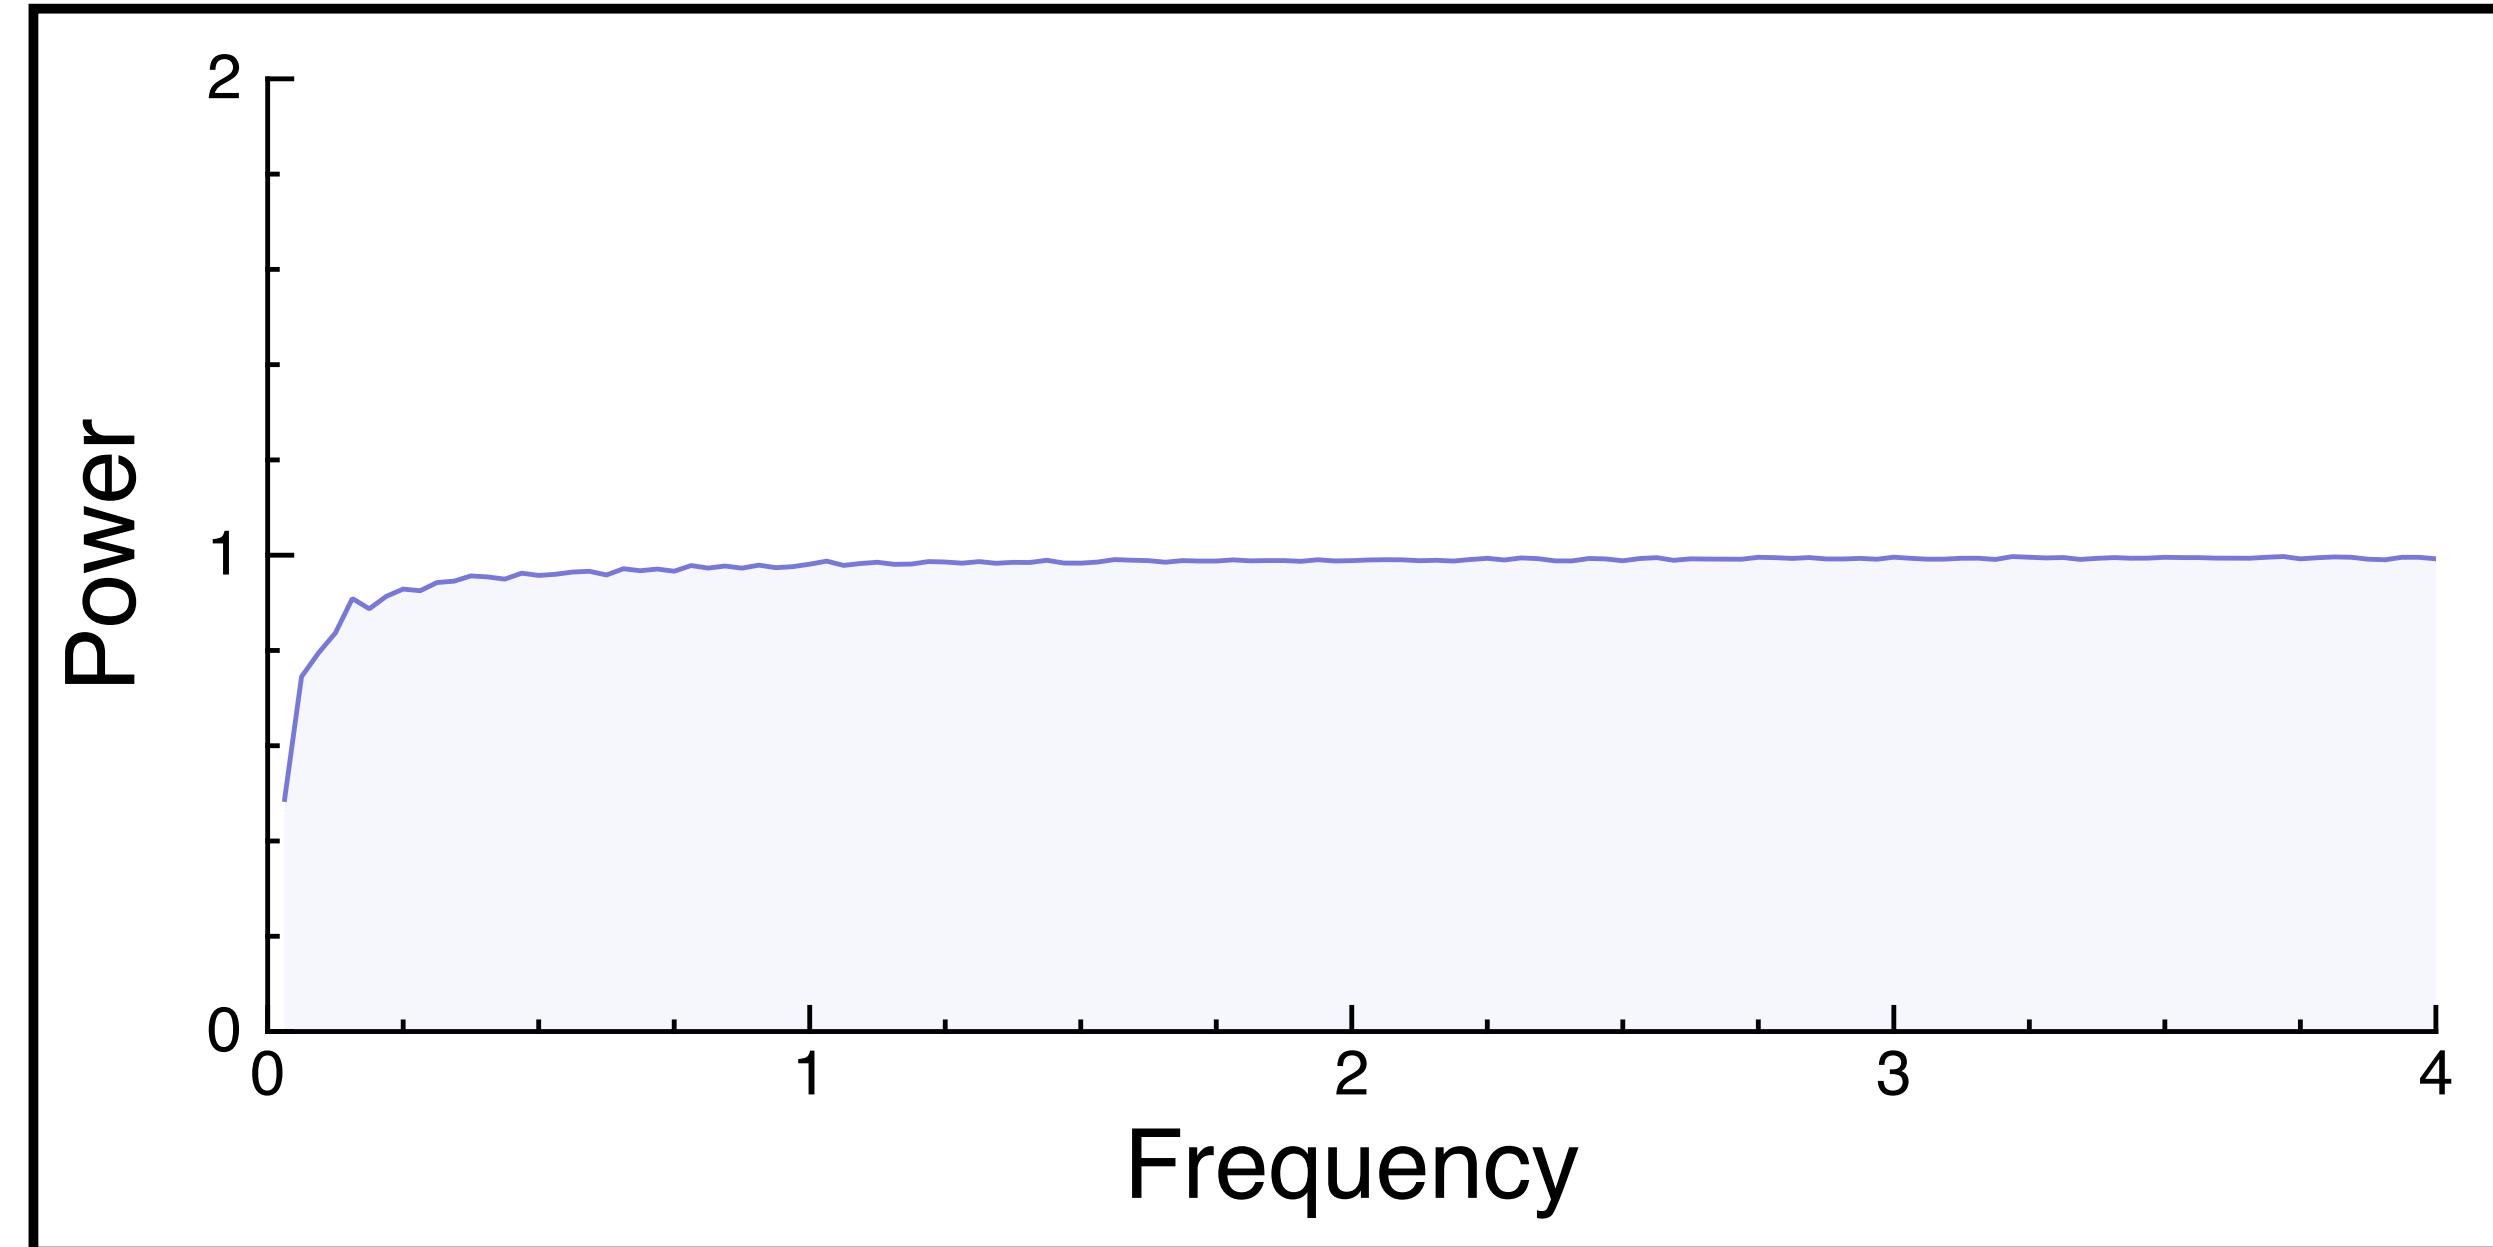
Samples



Power spectrum



Radial mean

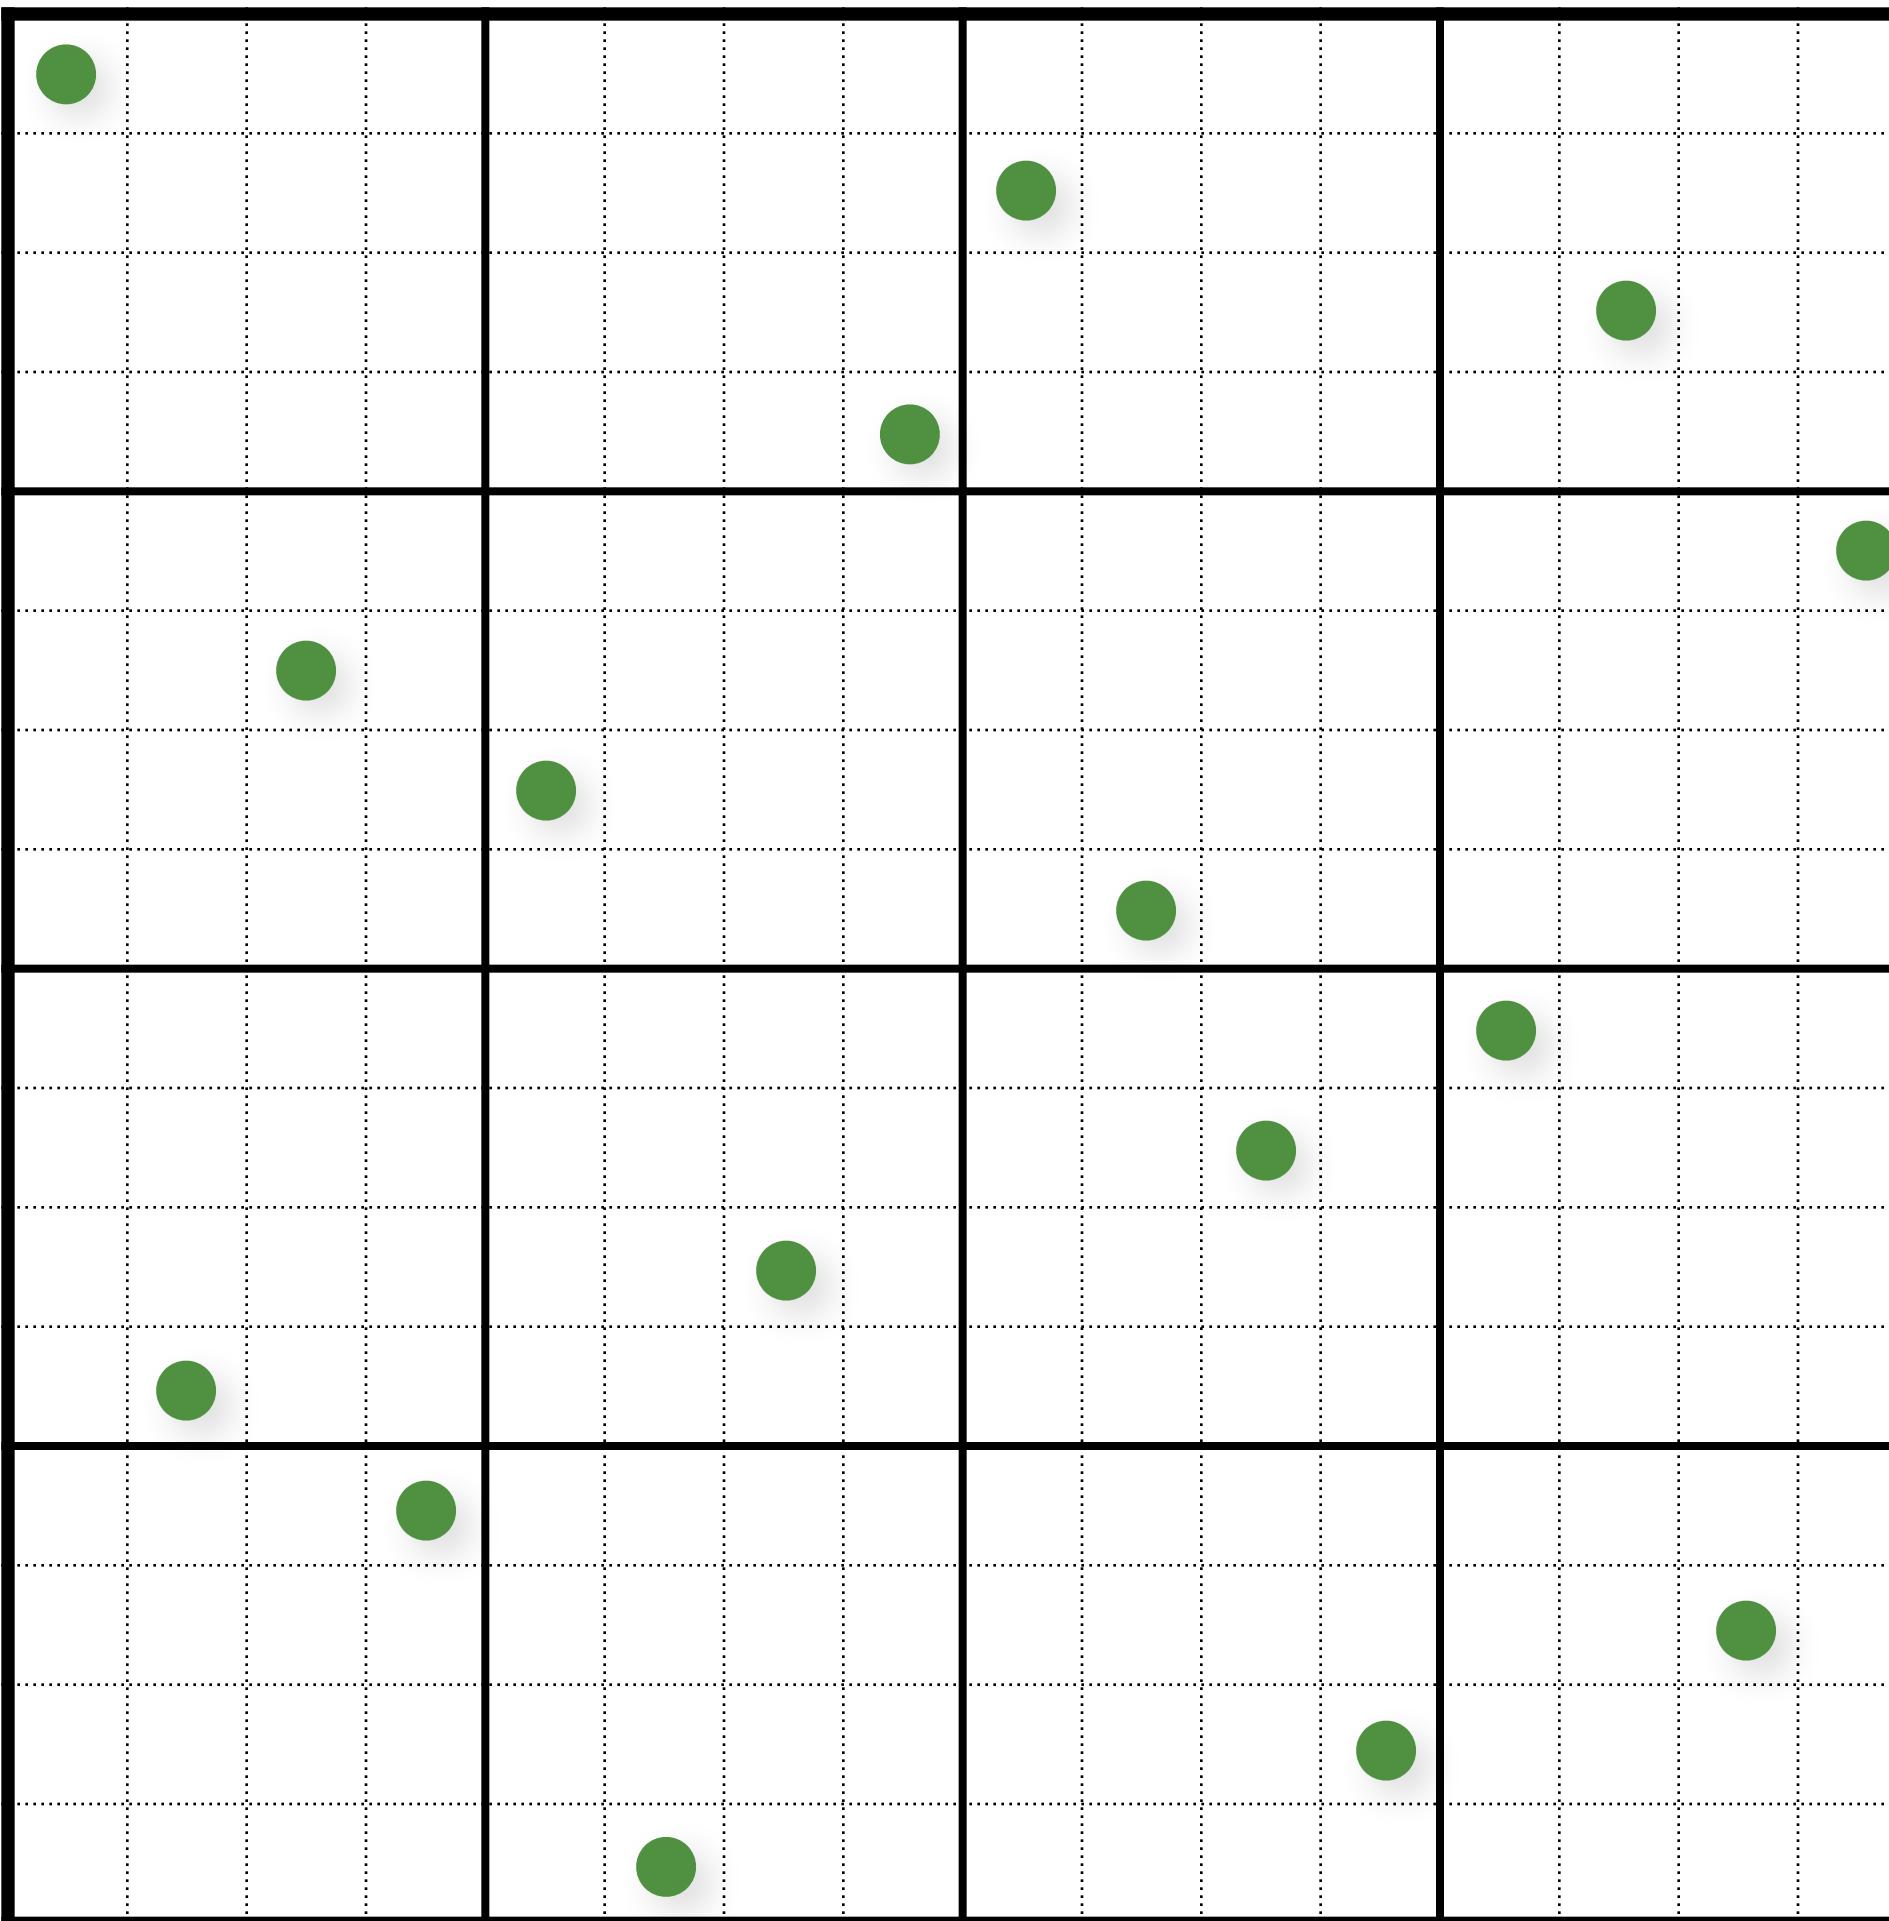


Multi-Jittered Sampling

Kenneth Chiu, Peter Shirley, and Changyaw Wang.
“Multi-jittered sampling.” In *Graphics Gems IV*, pp.
370–374. Academic Press, May 1994.

- combine N-Rooks and Jittered stratification constraints

Multi-Jittered Sampling



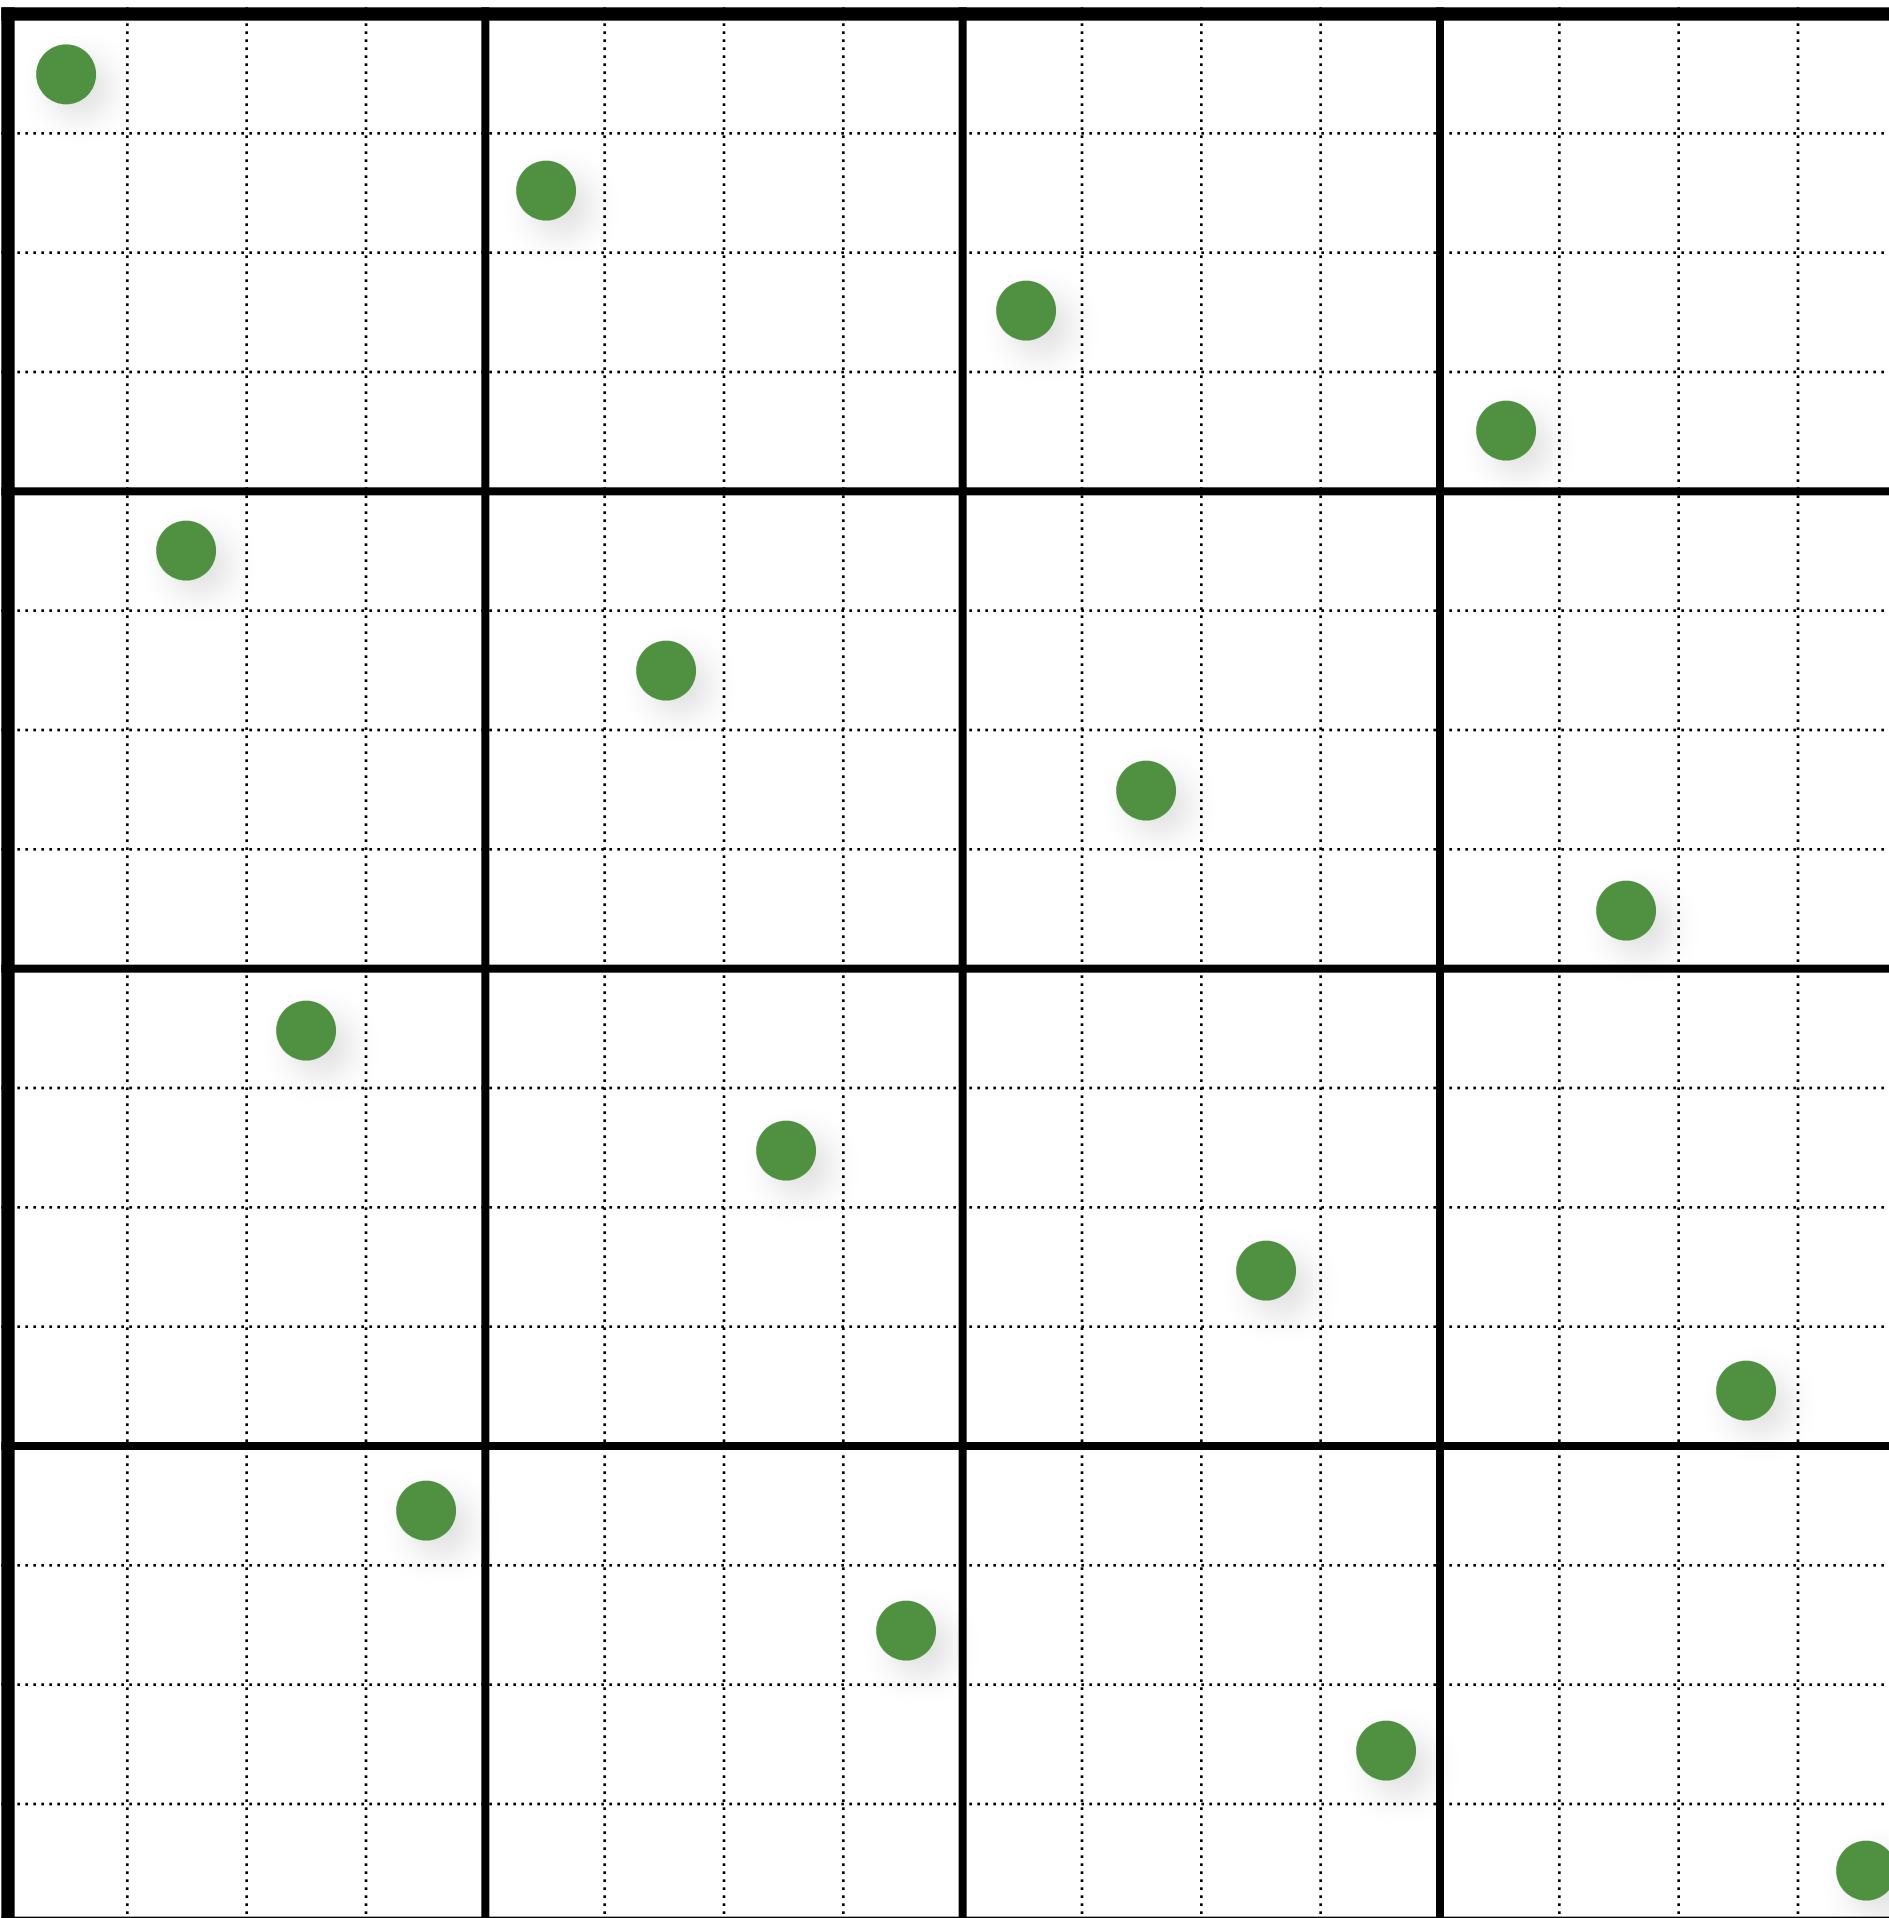
Multi-Jittered Sampling

```
// initialize
float cellsize = 1.0 / (resX*resY);
for (uint i = 0; i < resX; i++)
    for (uint j = 0; j < resY; j++)
    {
        samples(i,j).x = i/resX + (j+randf()) / (resX*resY);
        samples(i,j).y = j/resY + (i+randf()) / (resX*resY);
    }

// shuffle x coordinates within each column of cells
for (uint i = 0; i < resX; i++)
    for (uint j = resY-1; j >= 1; j--)
        swap(samples(i, j).x, samples(i, randi(0, j)).x);

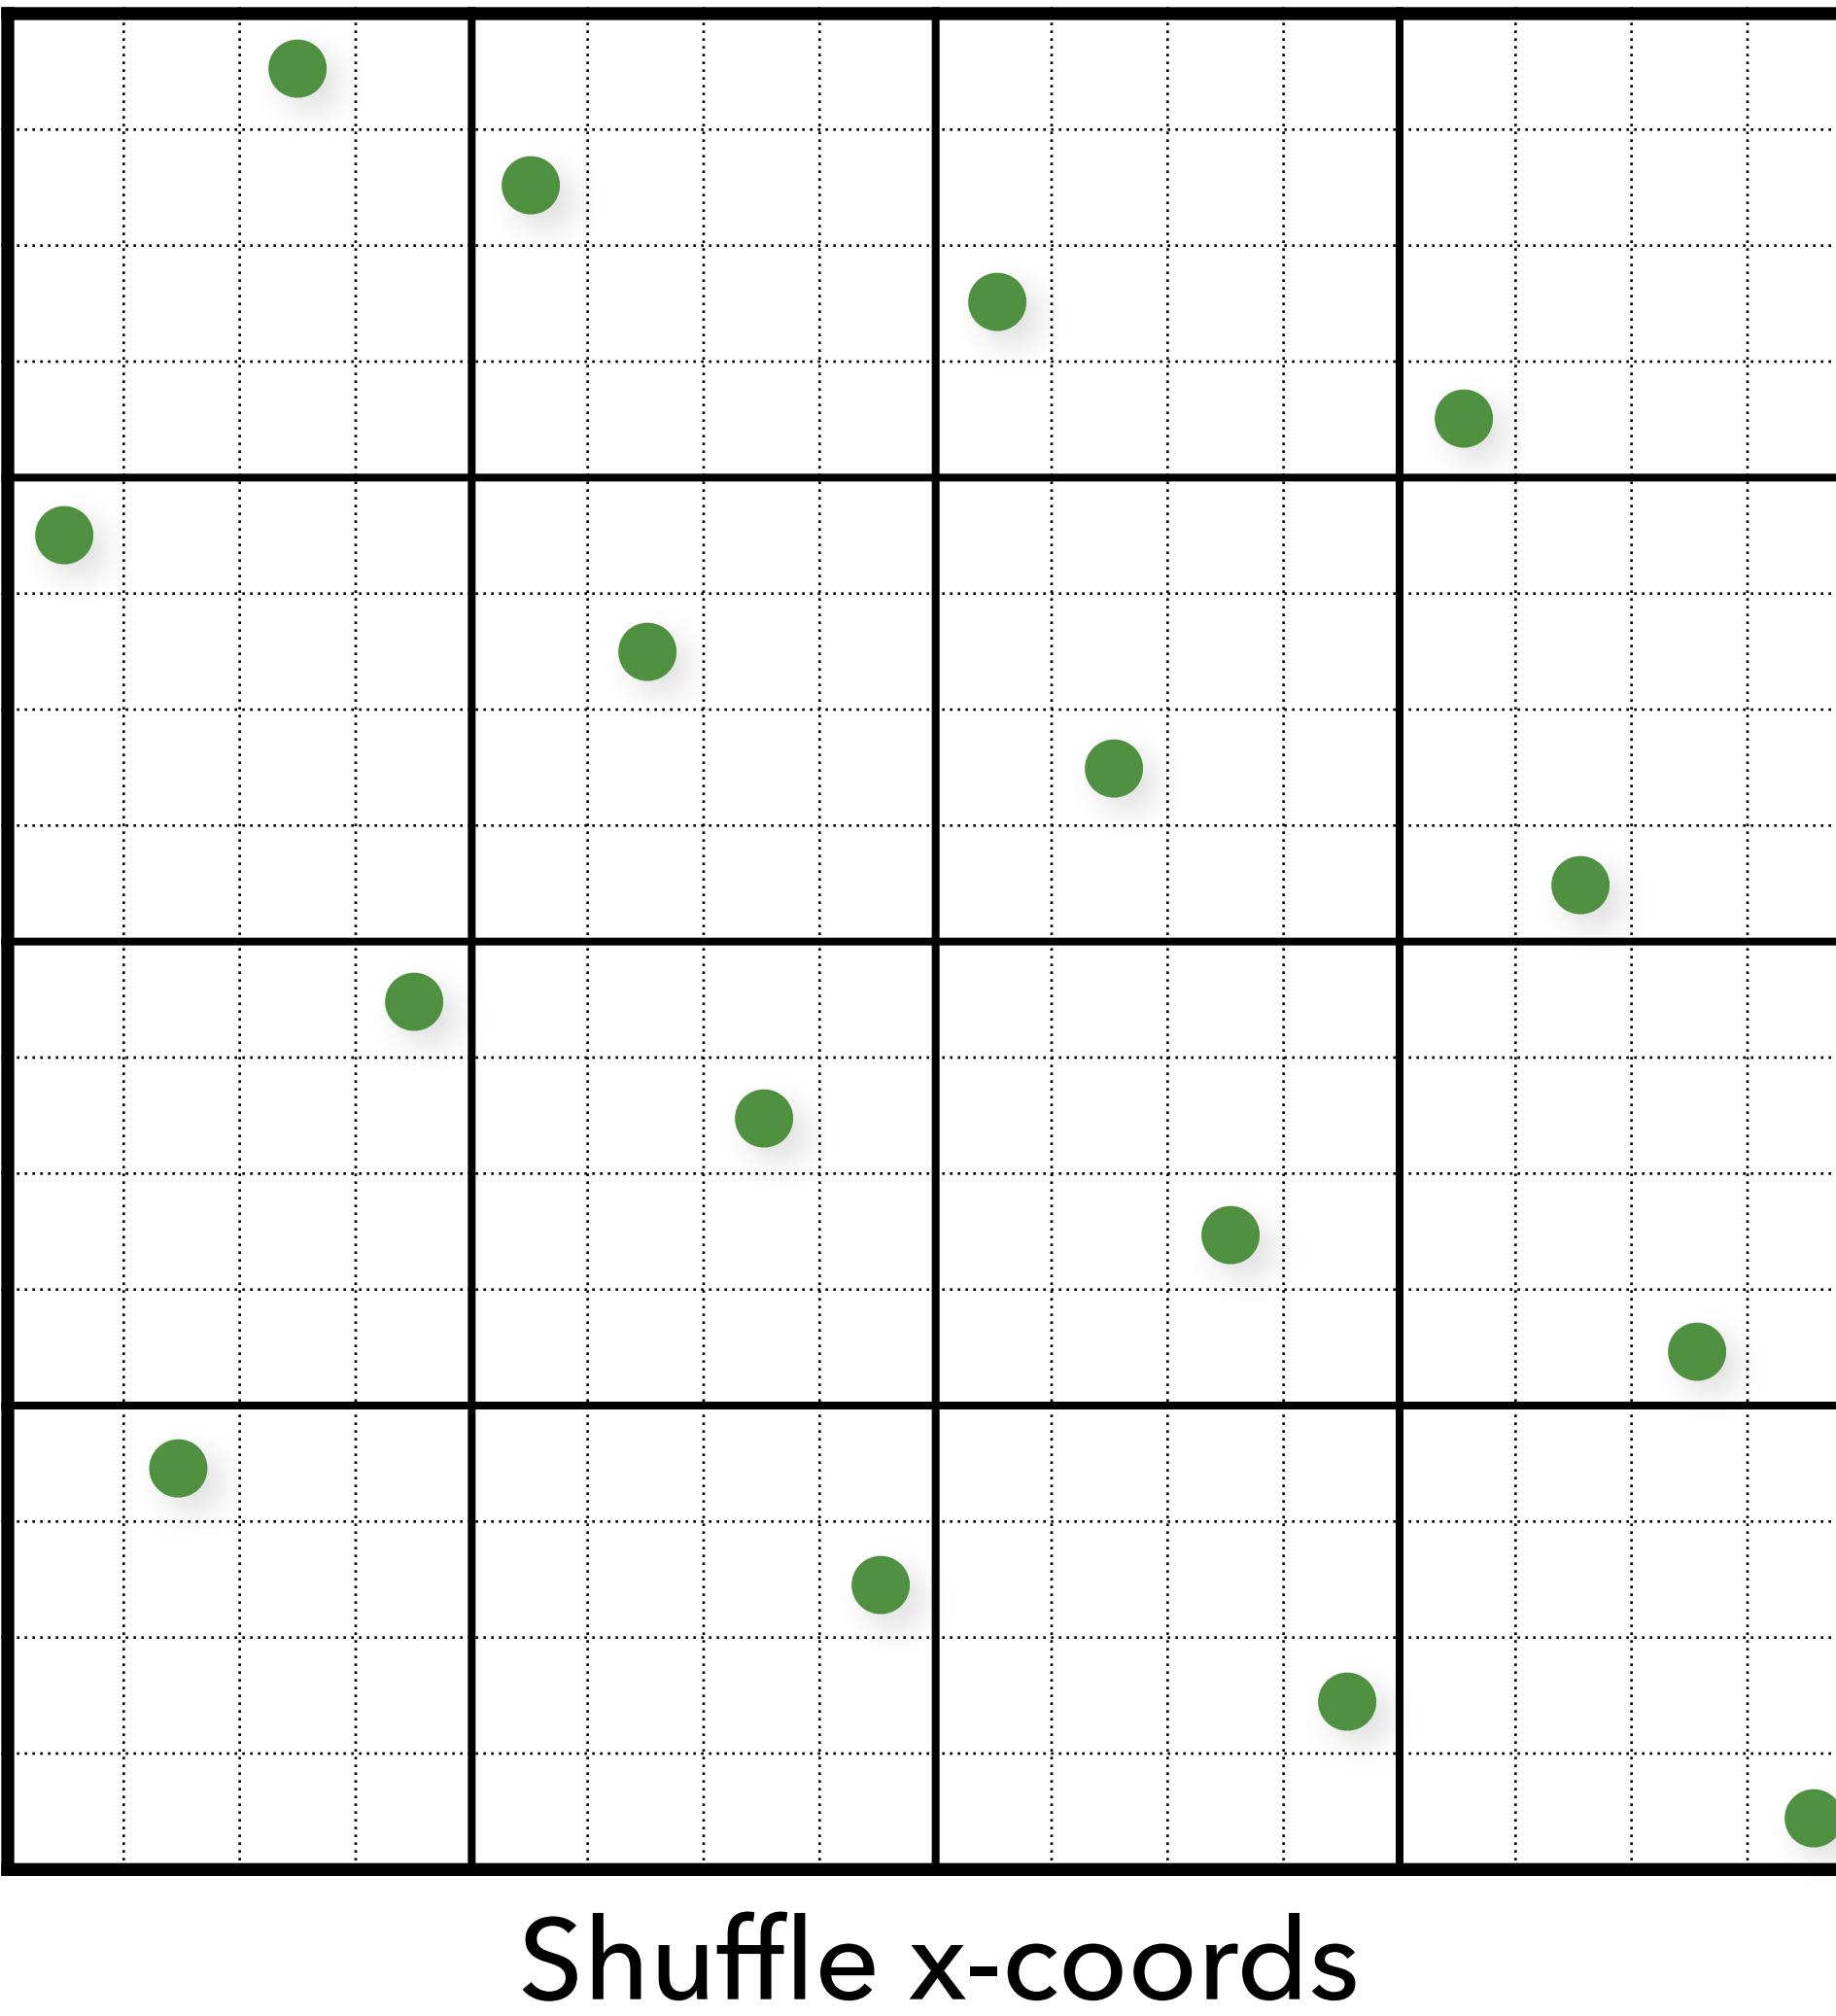
// shuffle y coordinates within each row of cells
for (unsigned j = 0; j < resY; j++)
    for (unsigned i = resX-1; i >= 1; i--)
        swap(samples(i, j).y, samples(randi(0, i), j).y);
```

Multi-Jittered Sampling

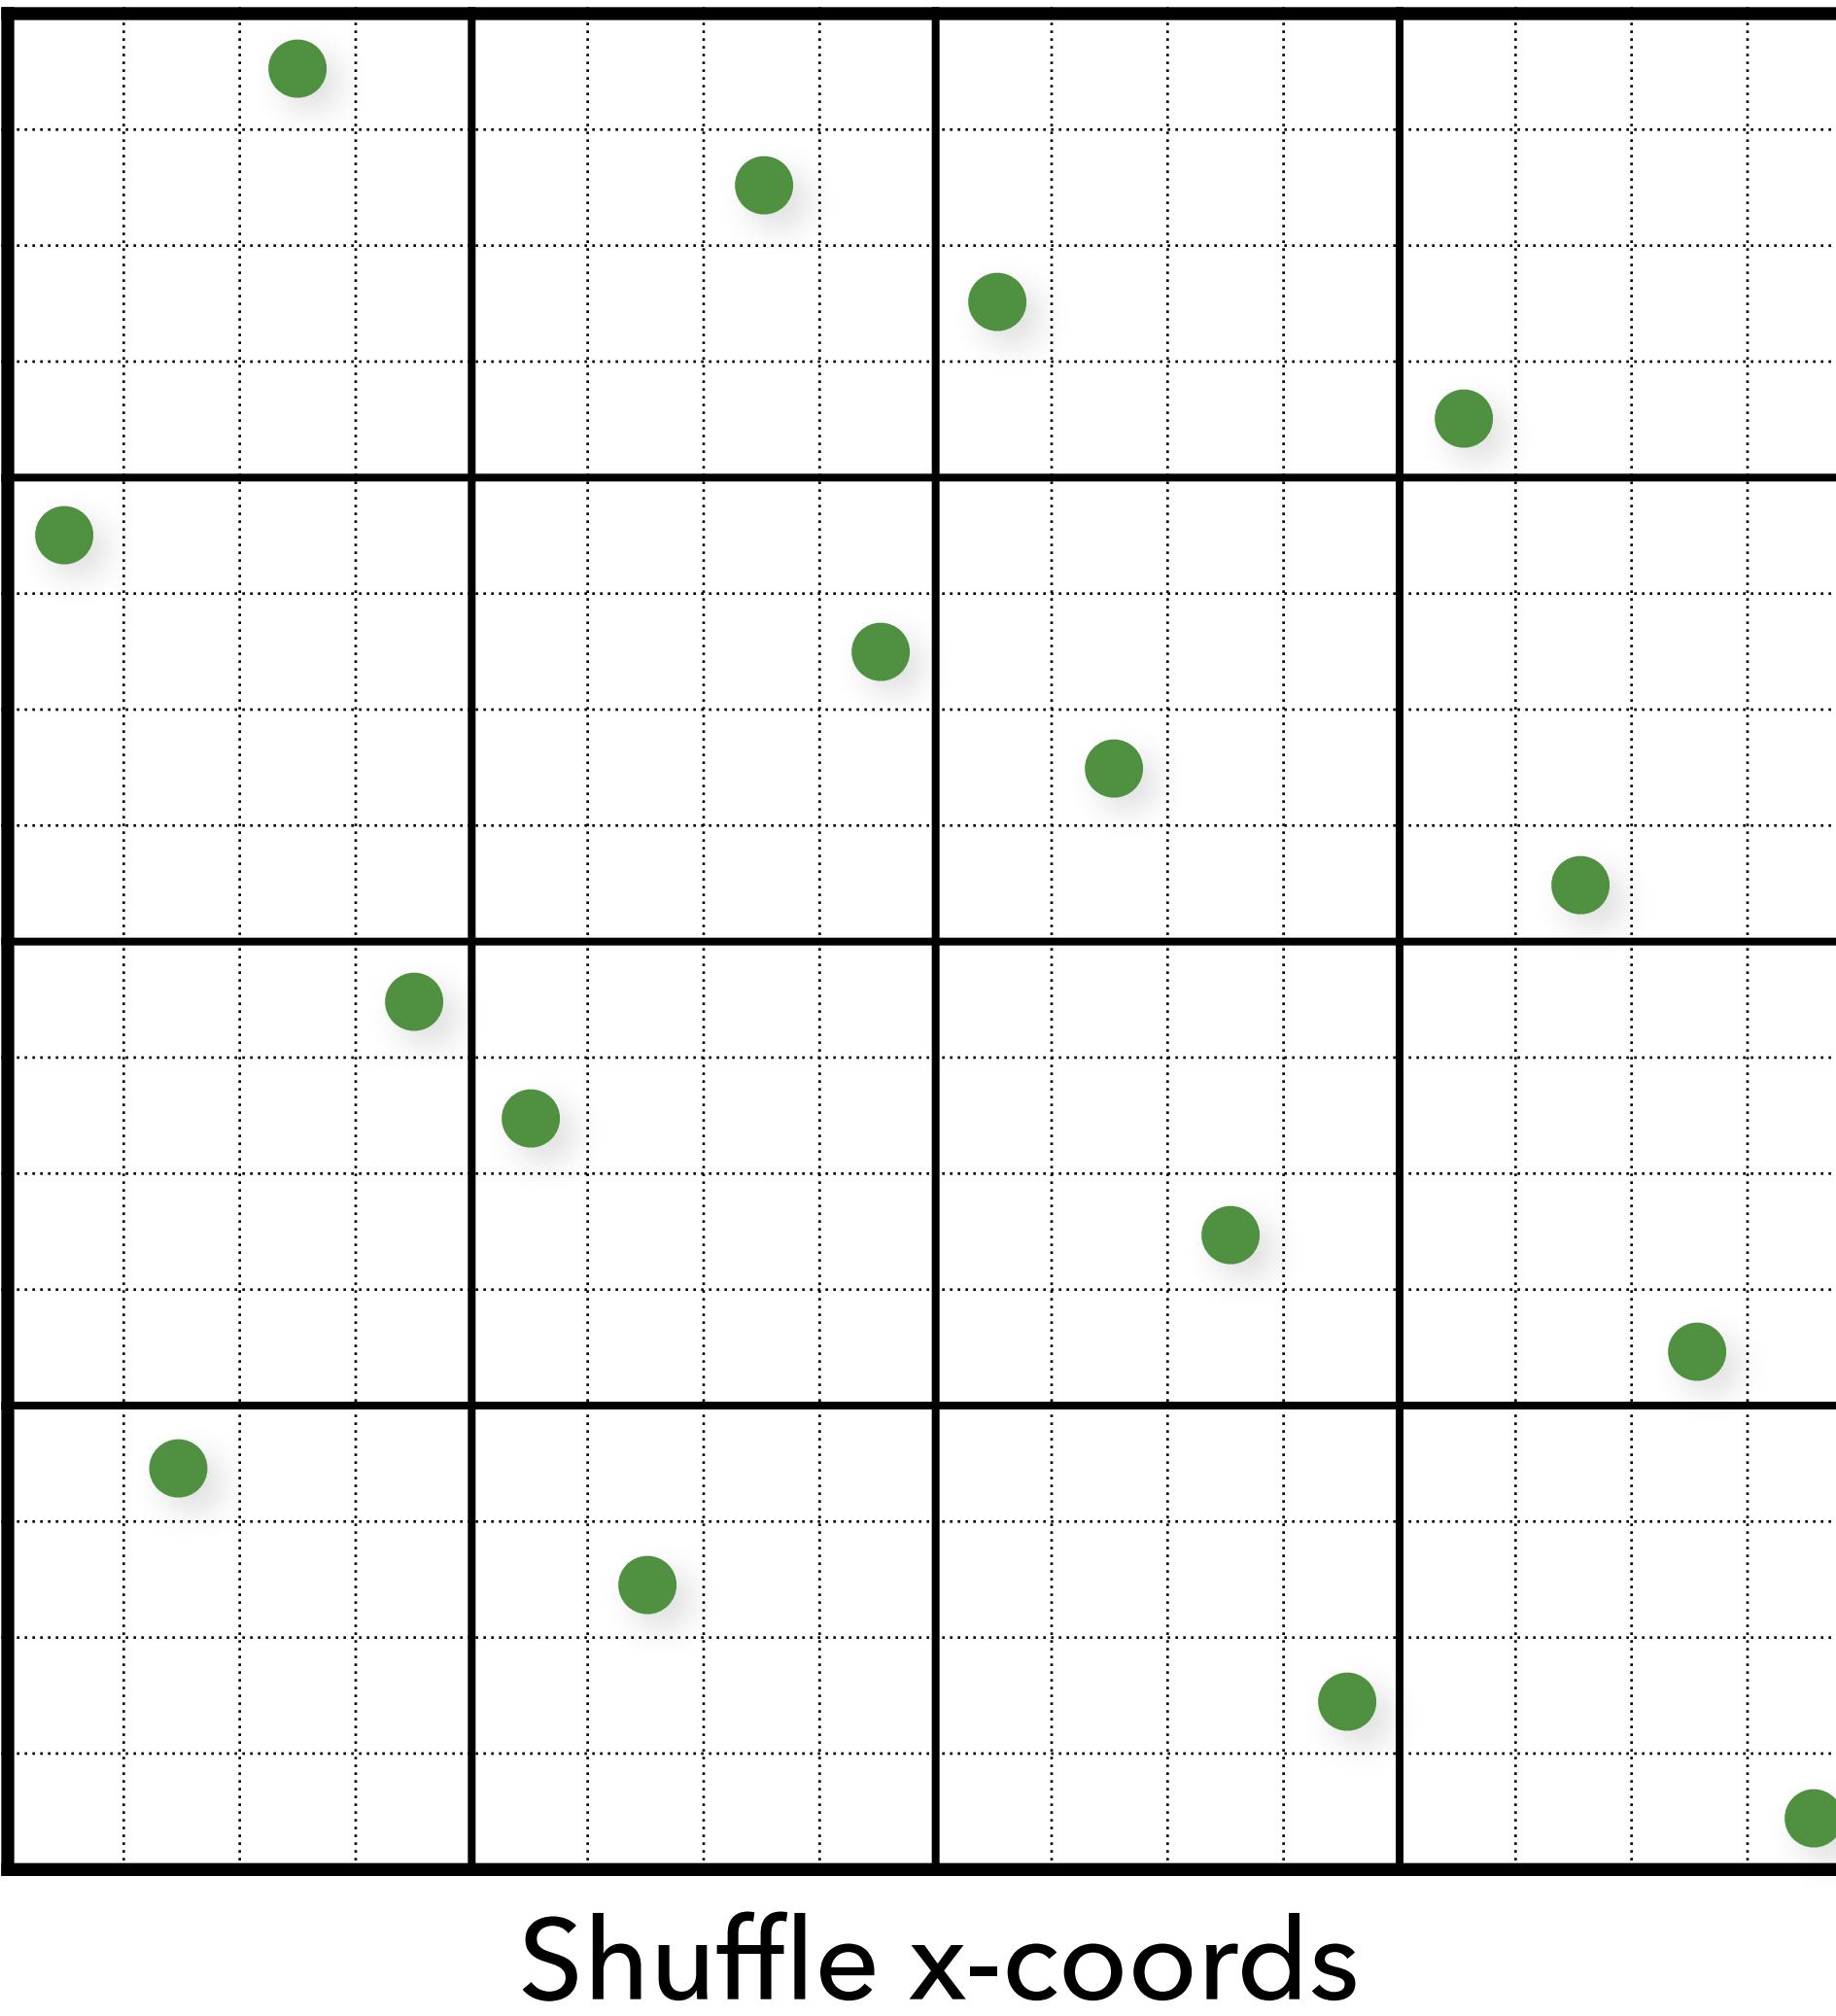


Shuffled coordinates

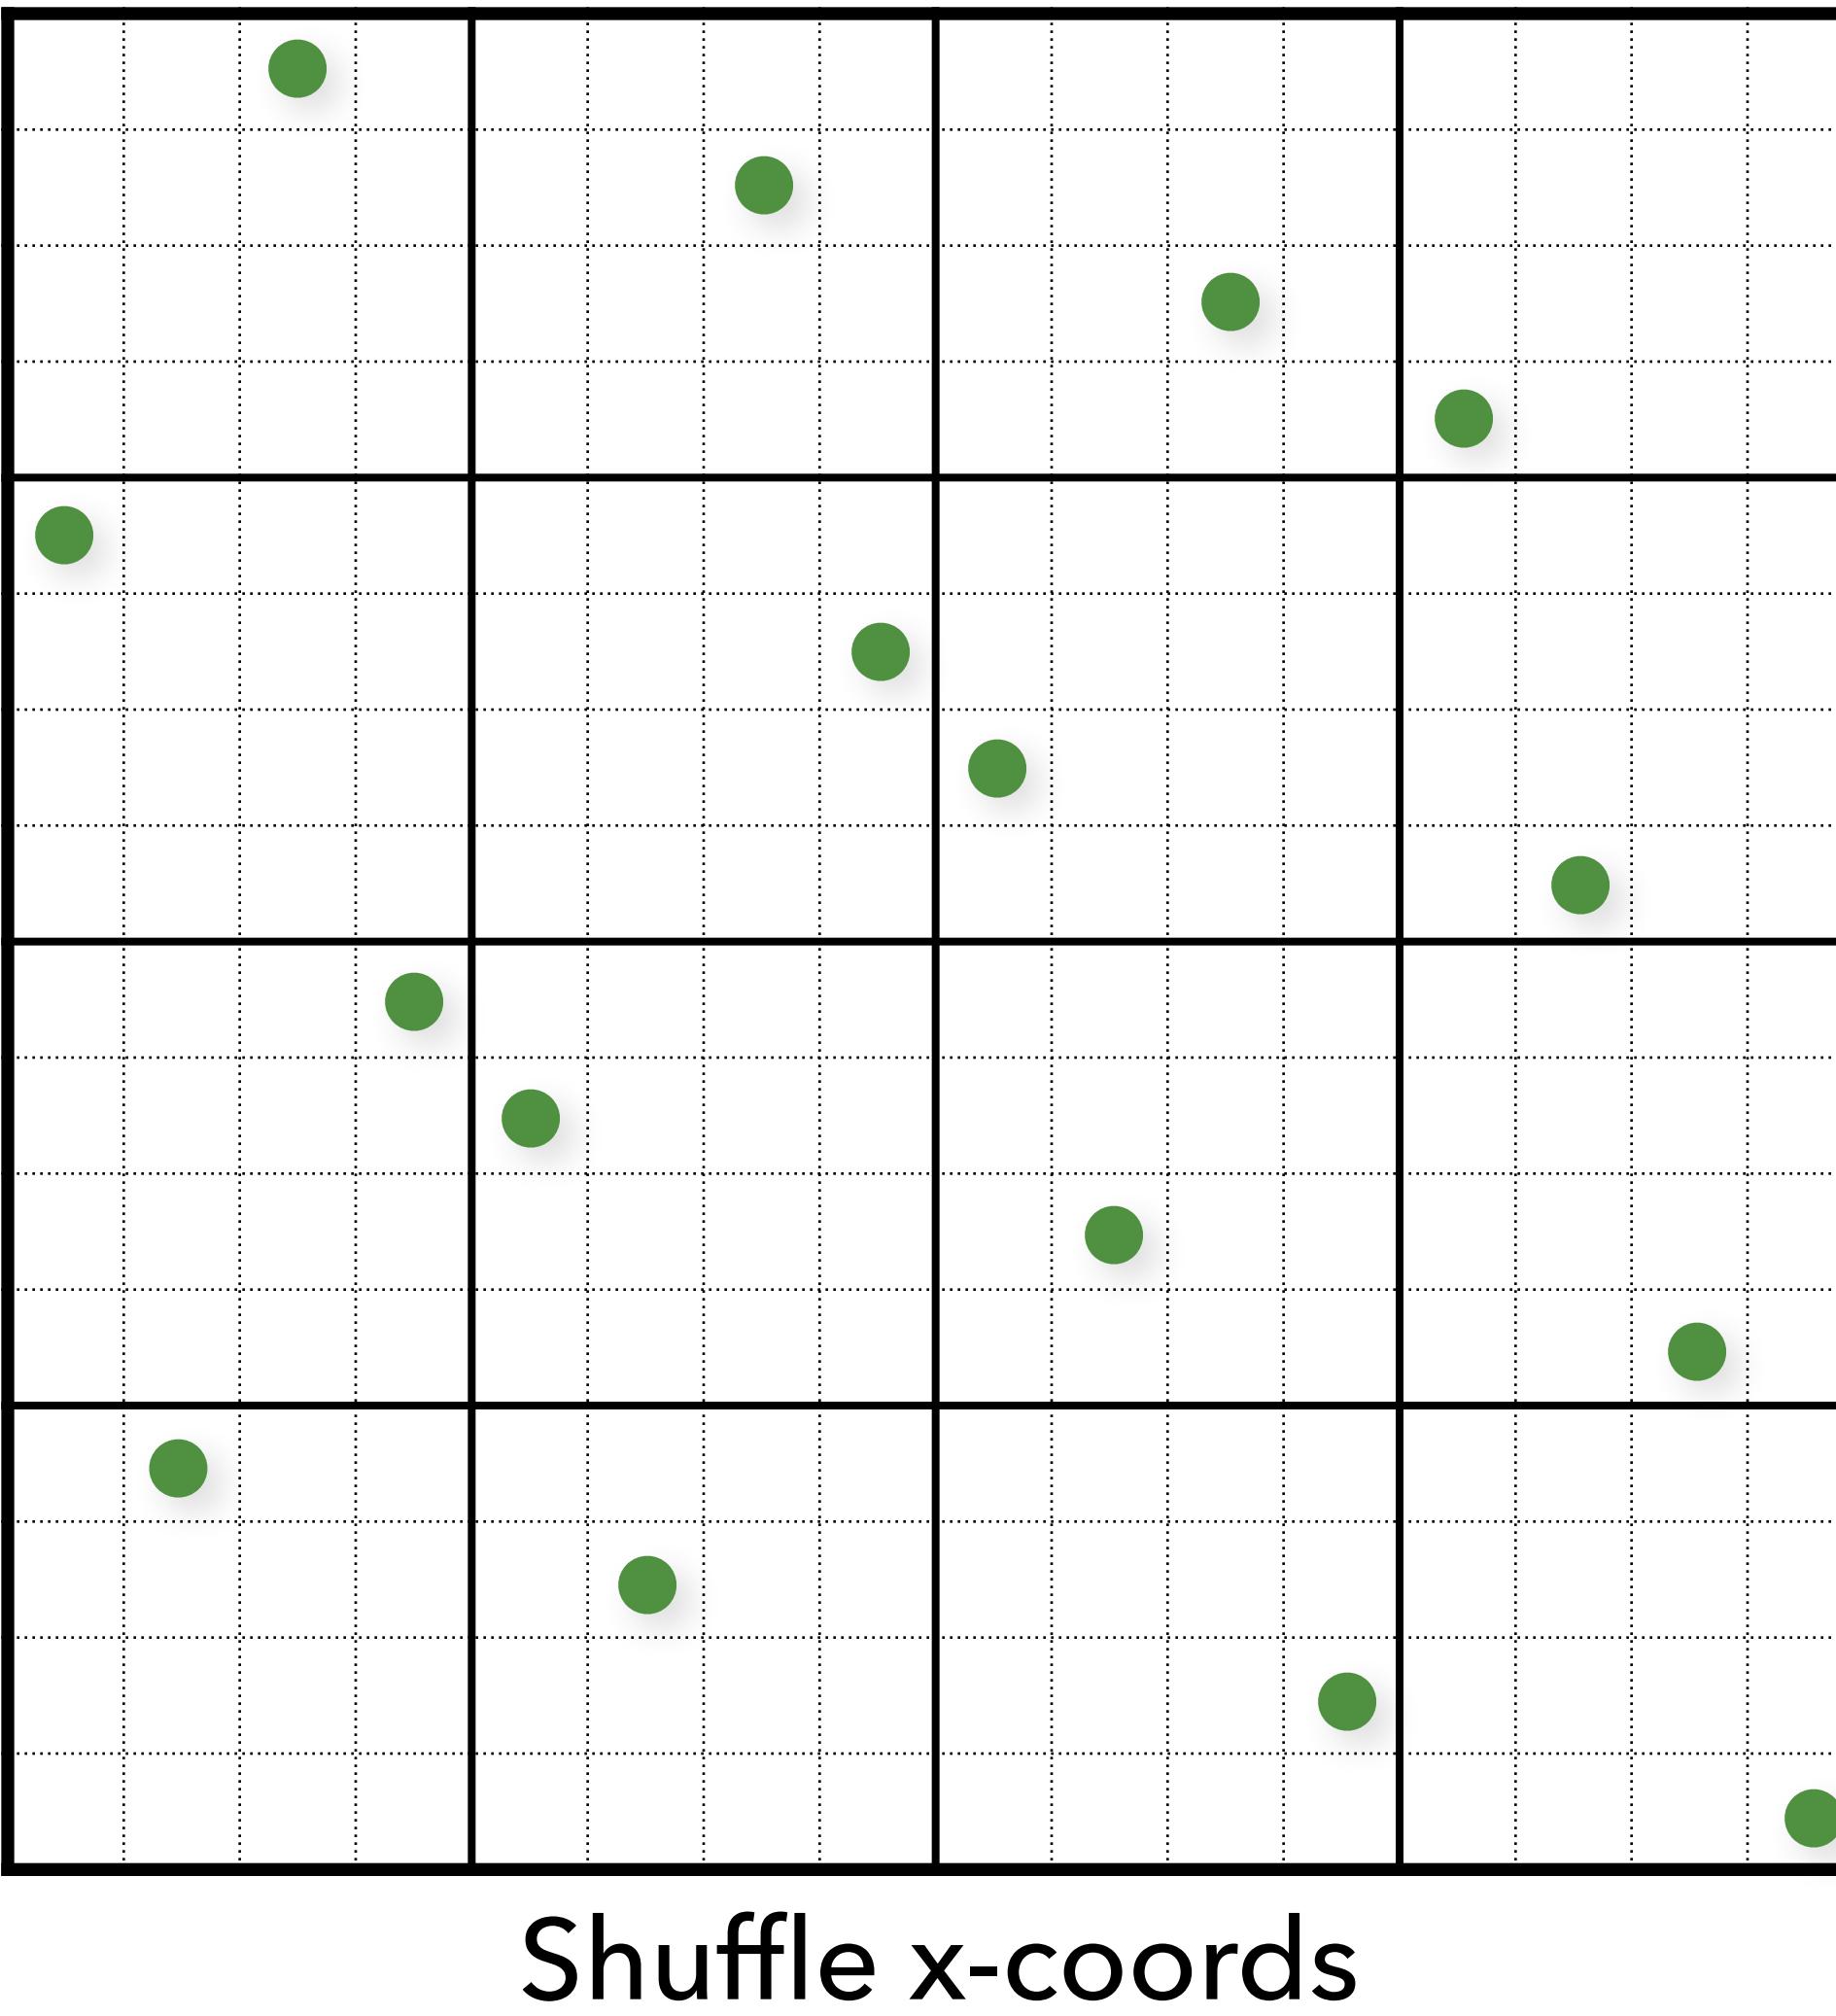
Multi-Jittered Sampling



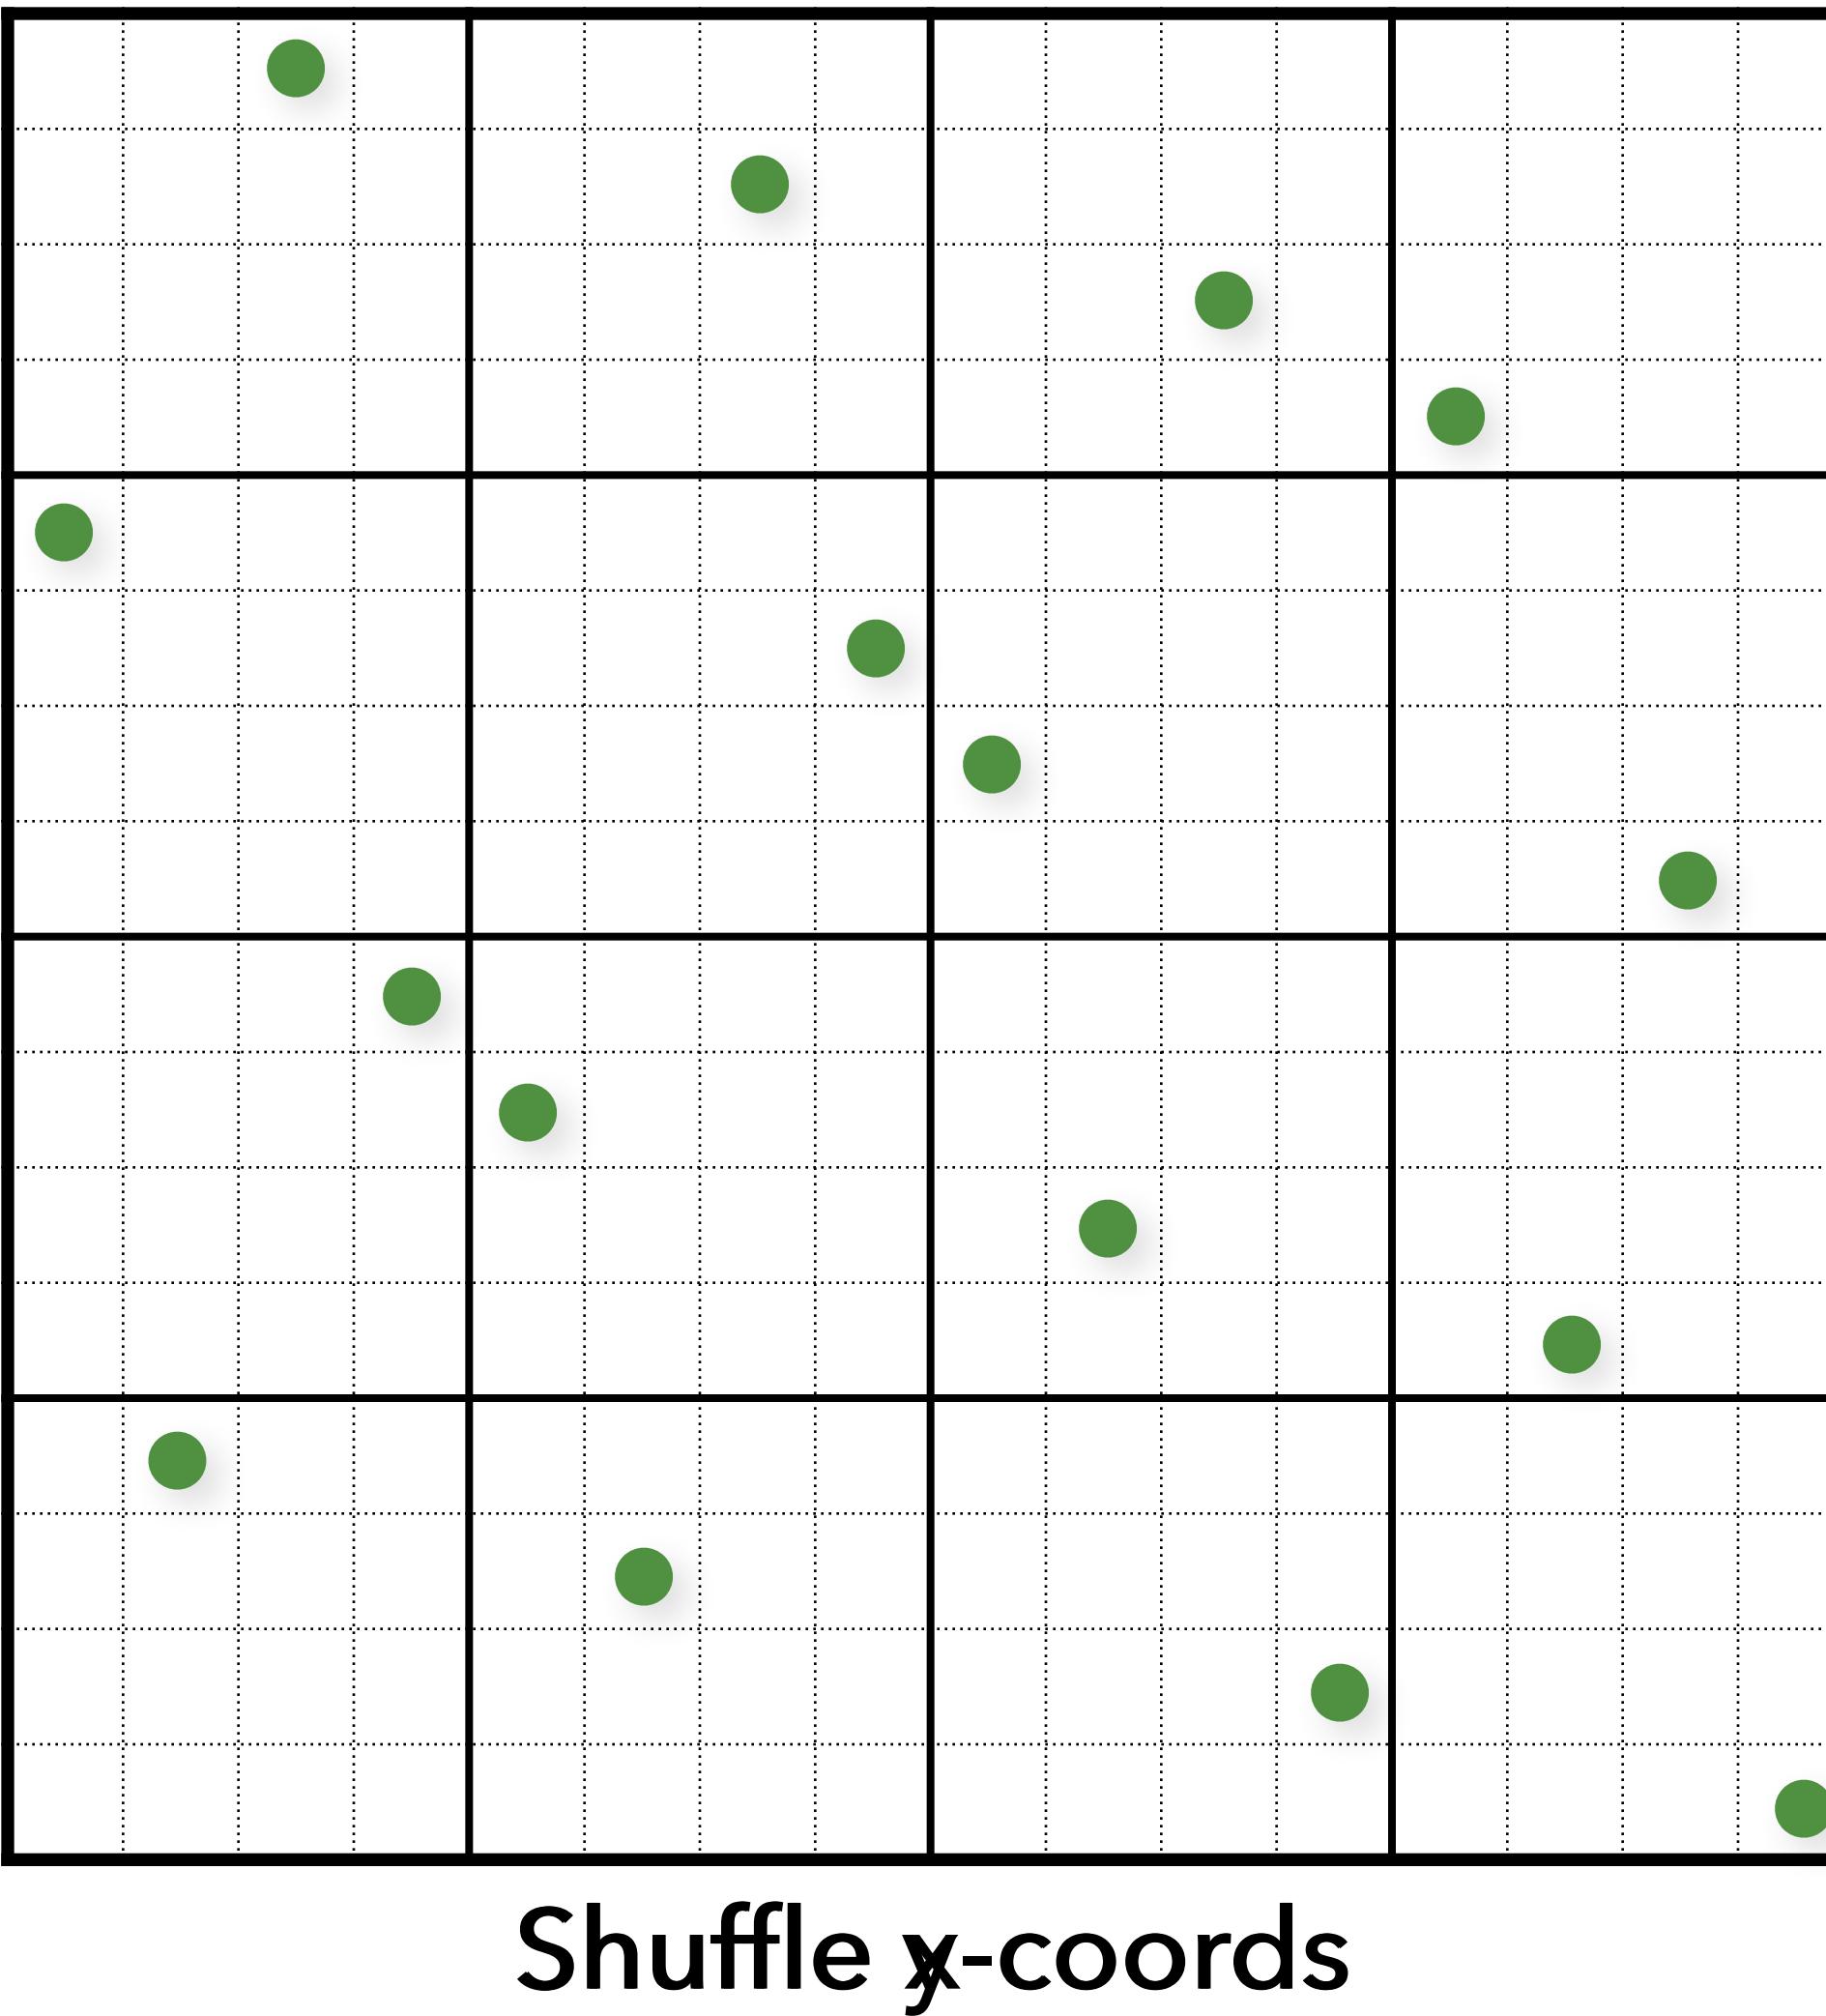
Multi-Jittered Sampling



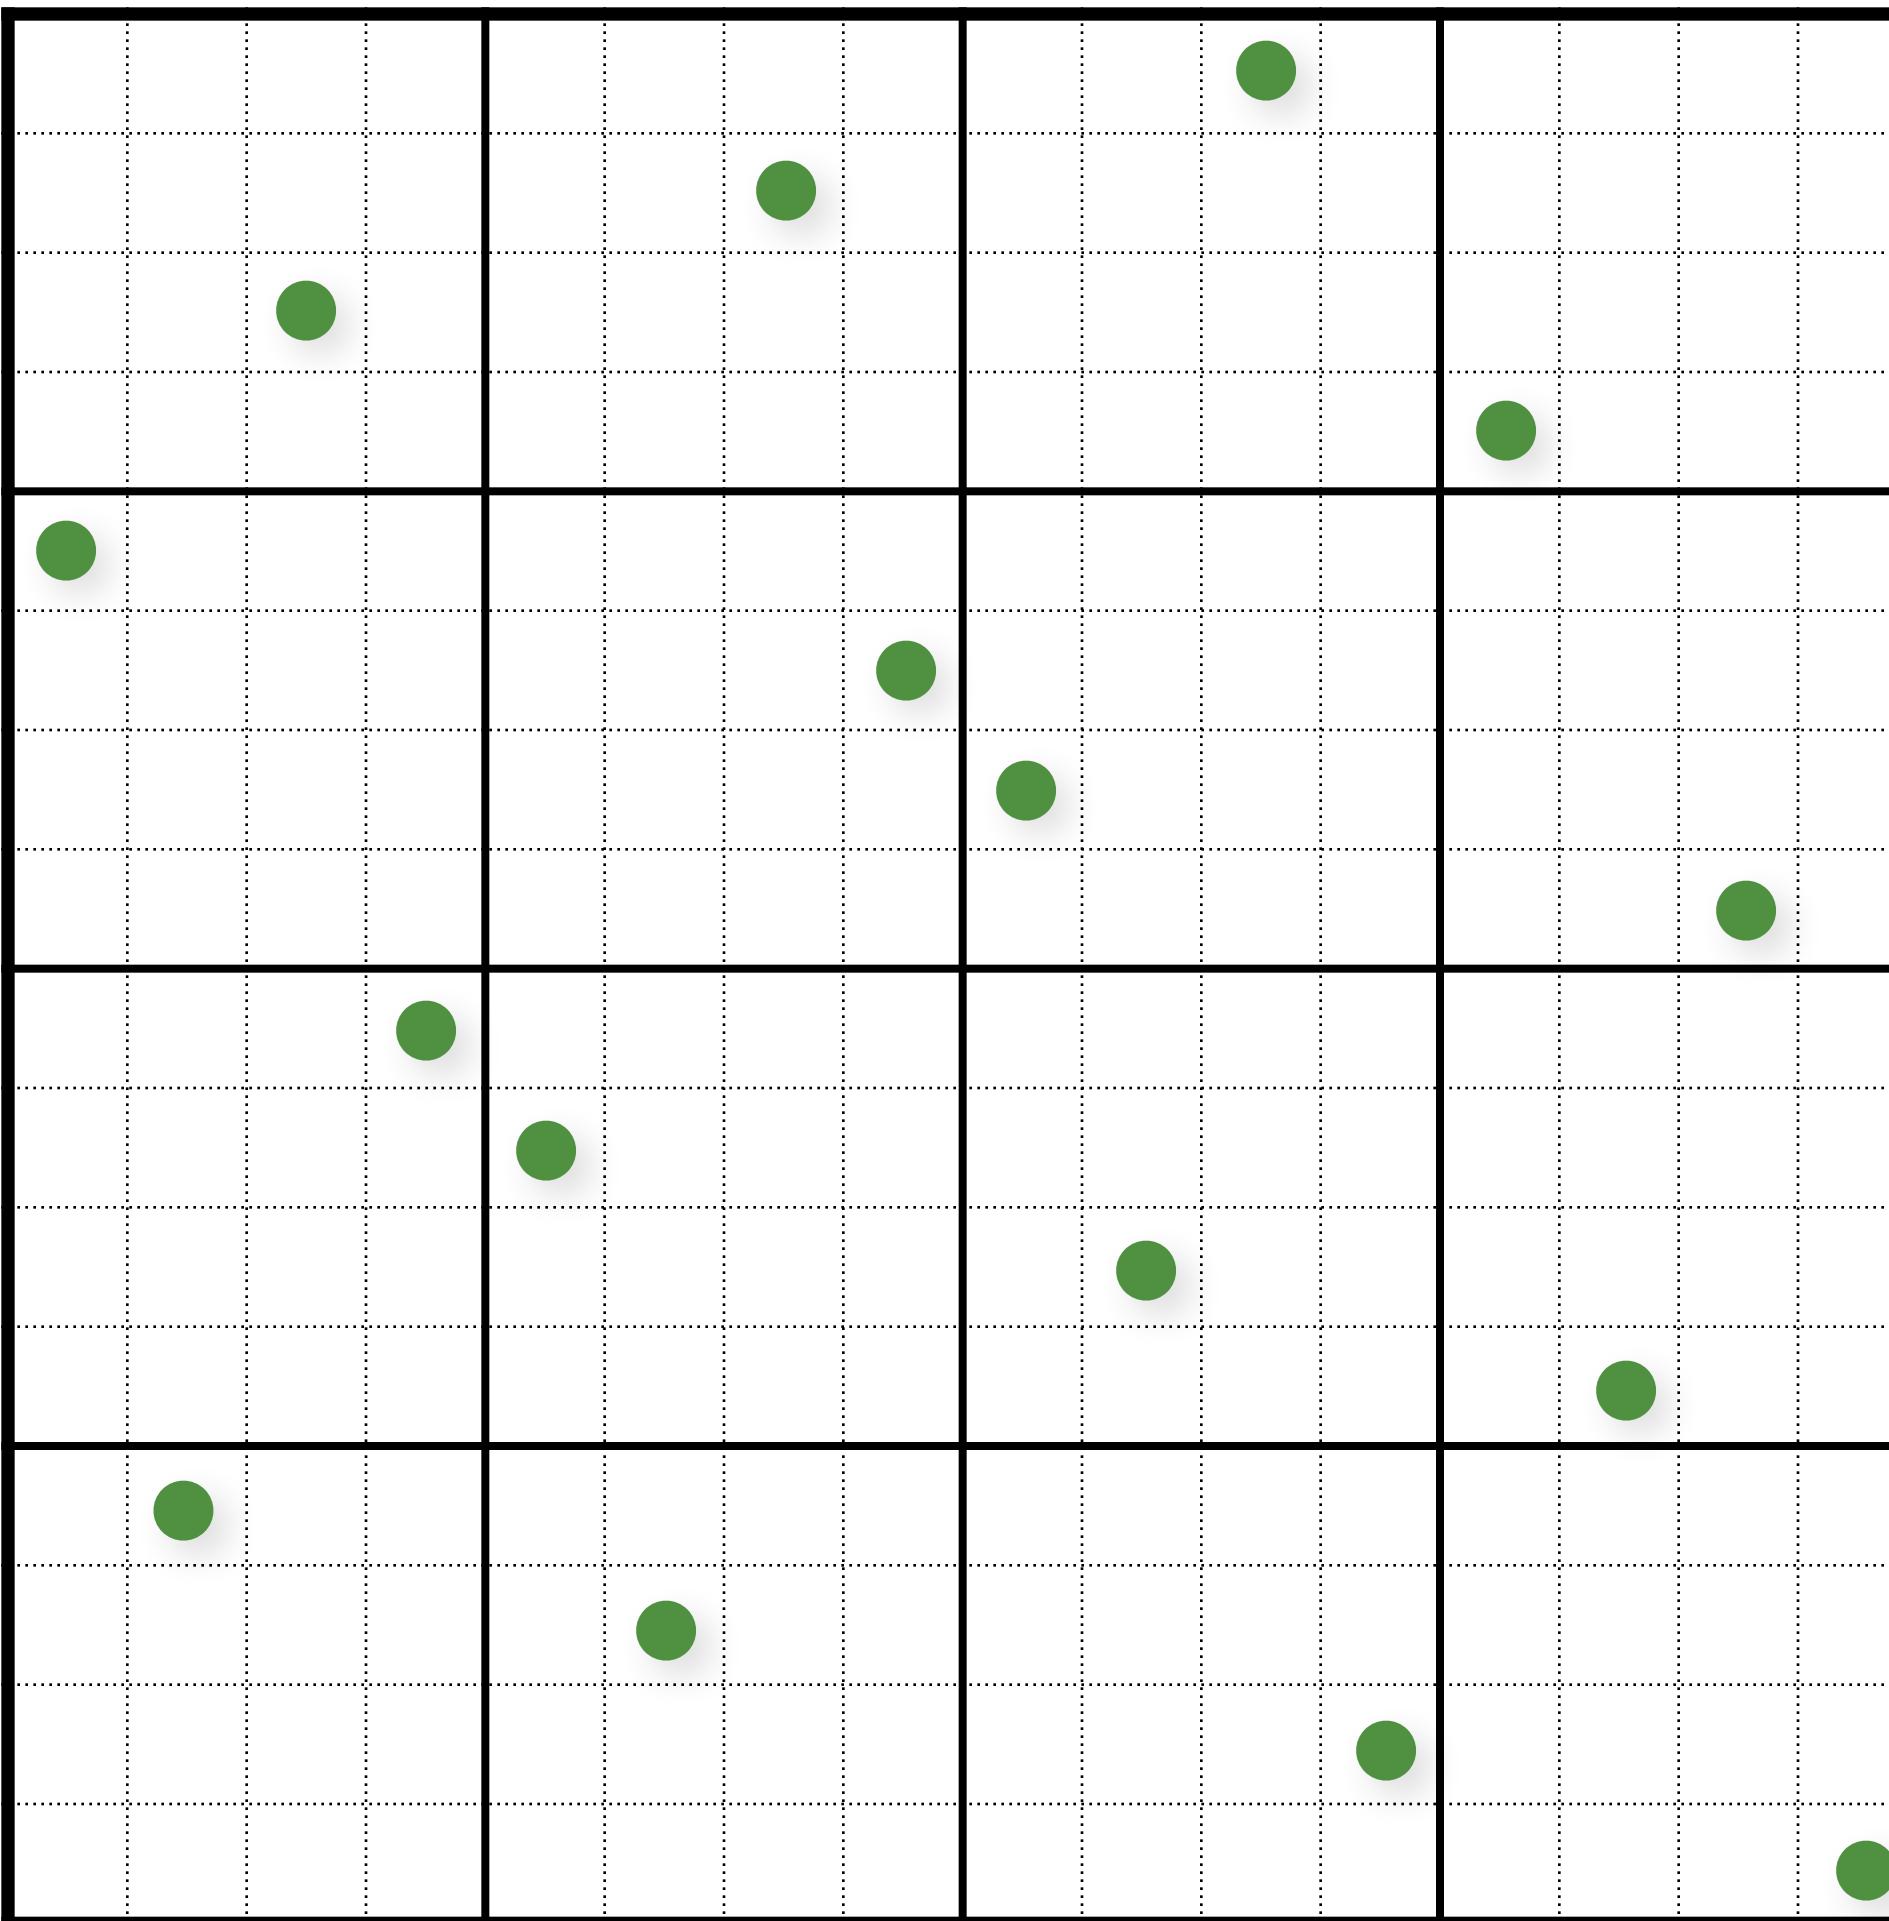
Multi-Jittered Sampling



Multi-Jittered Sampling

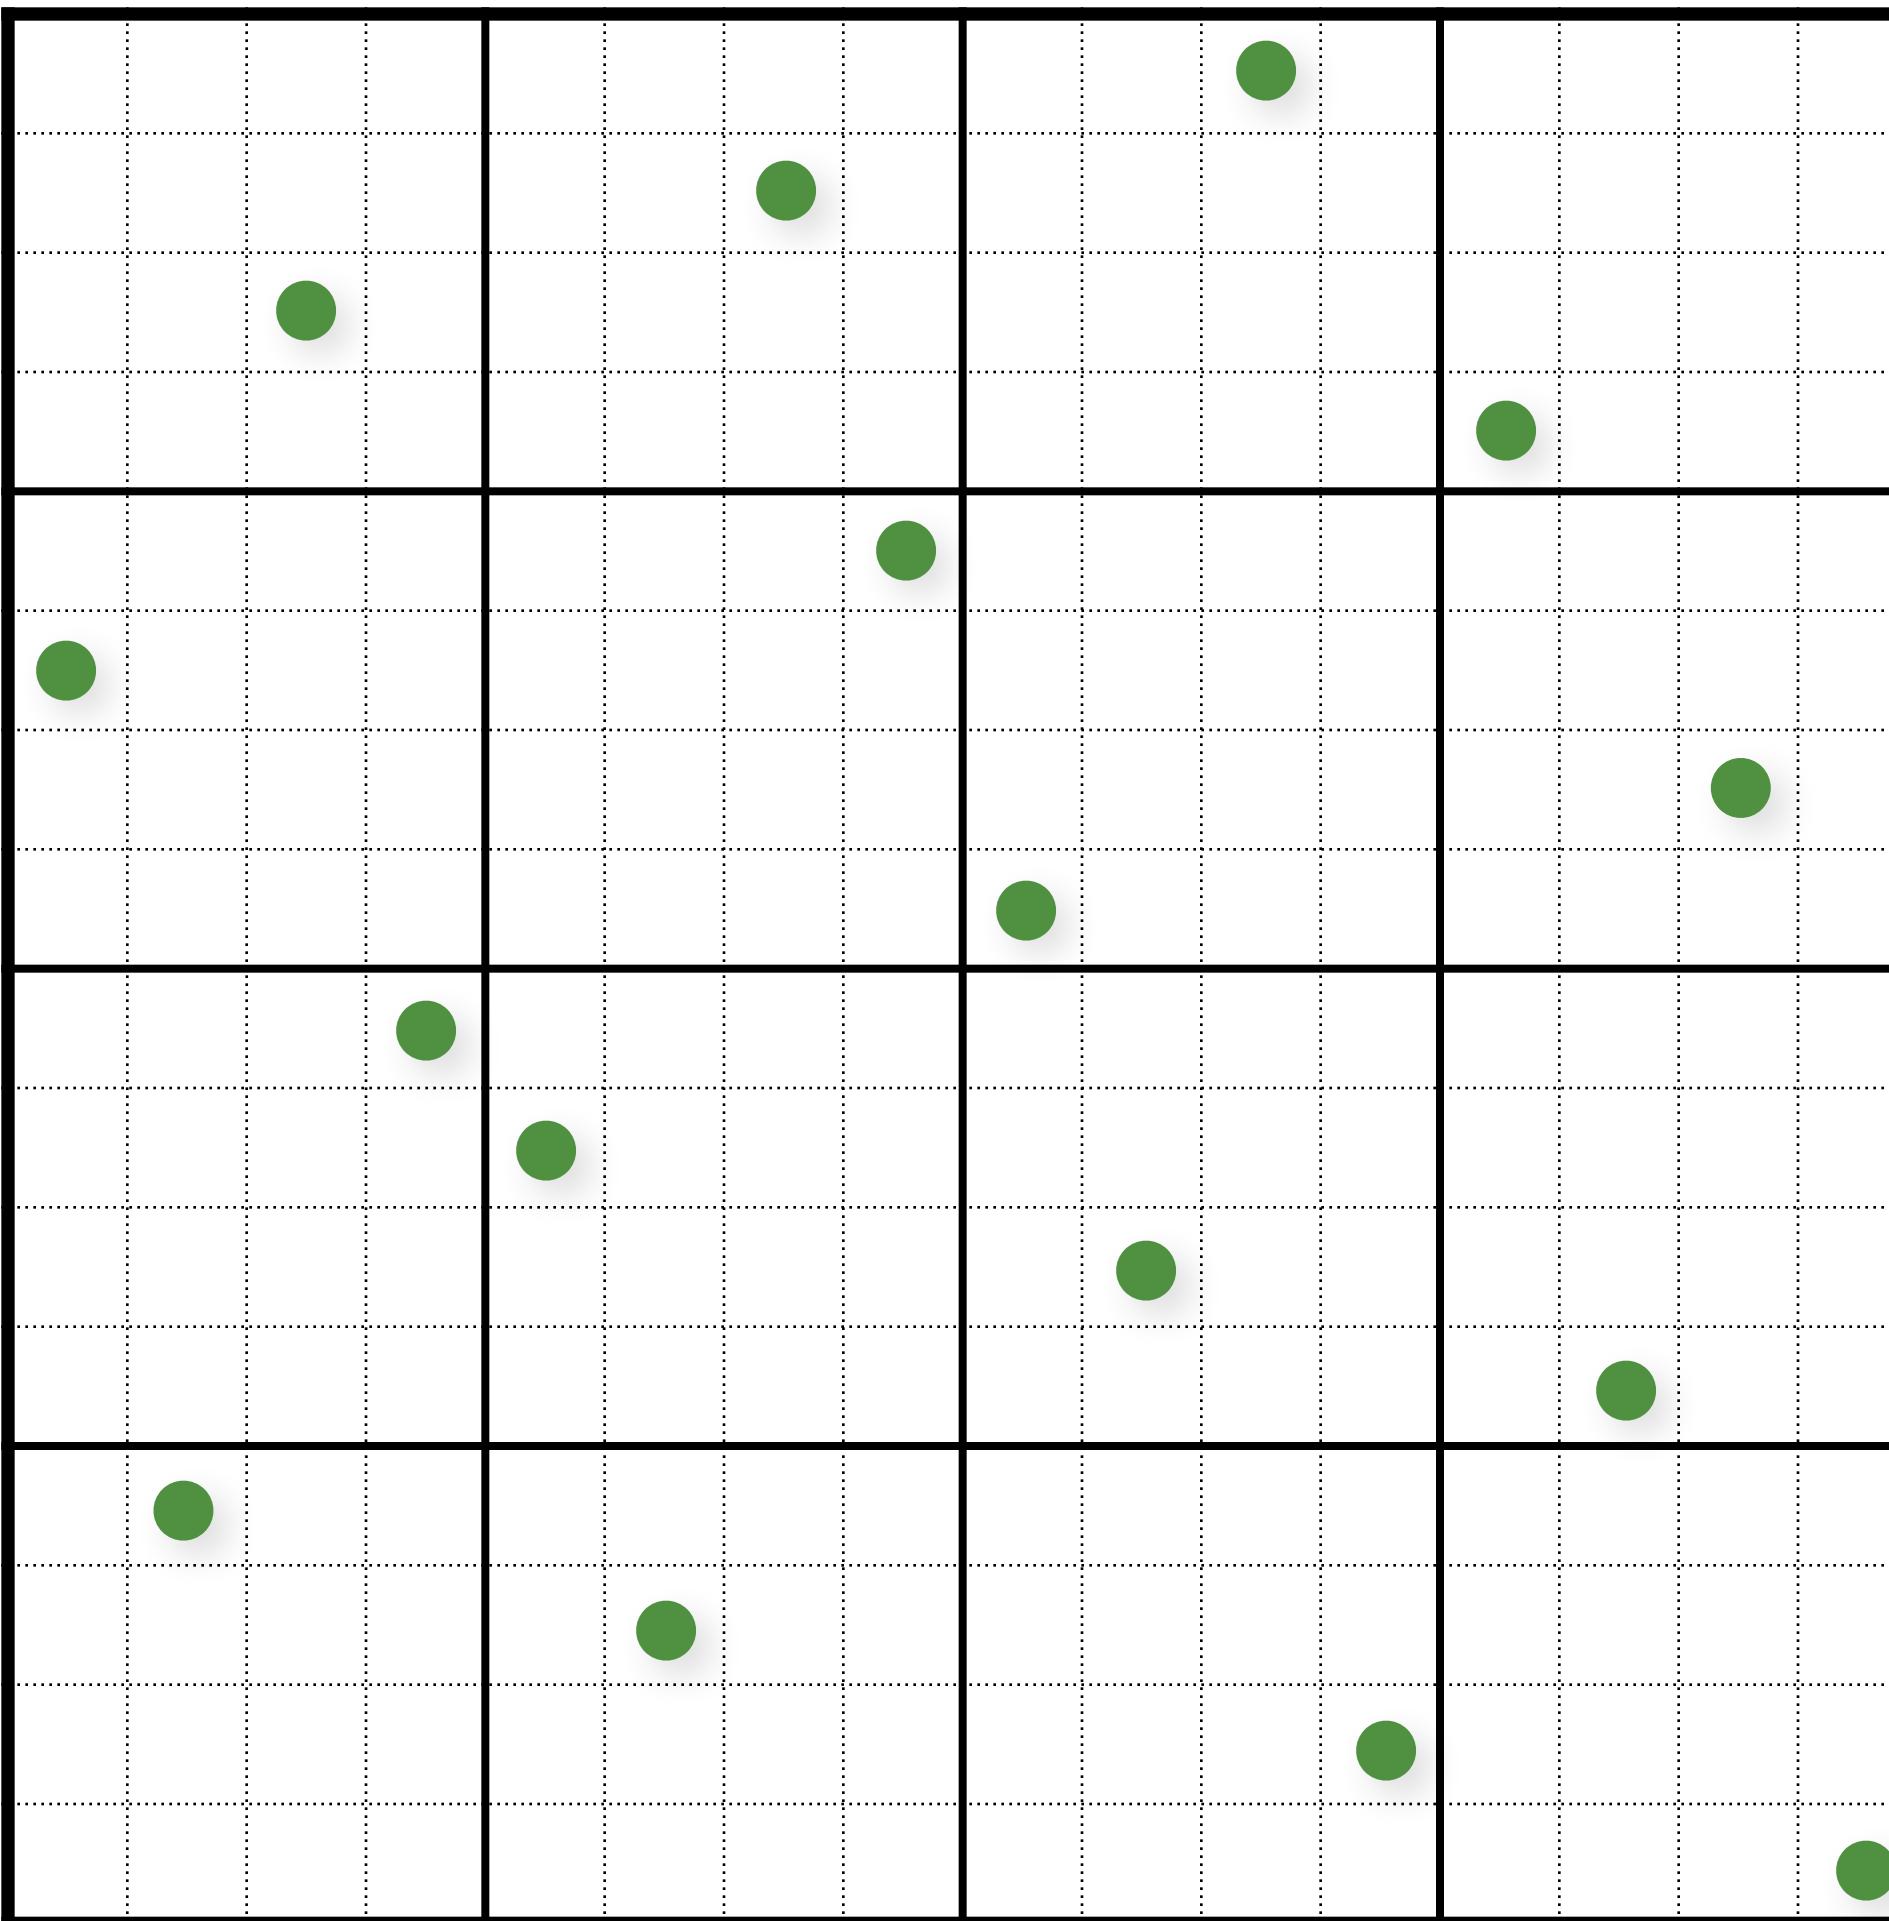


Multi-Jittered Sampling



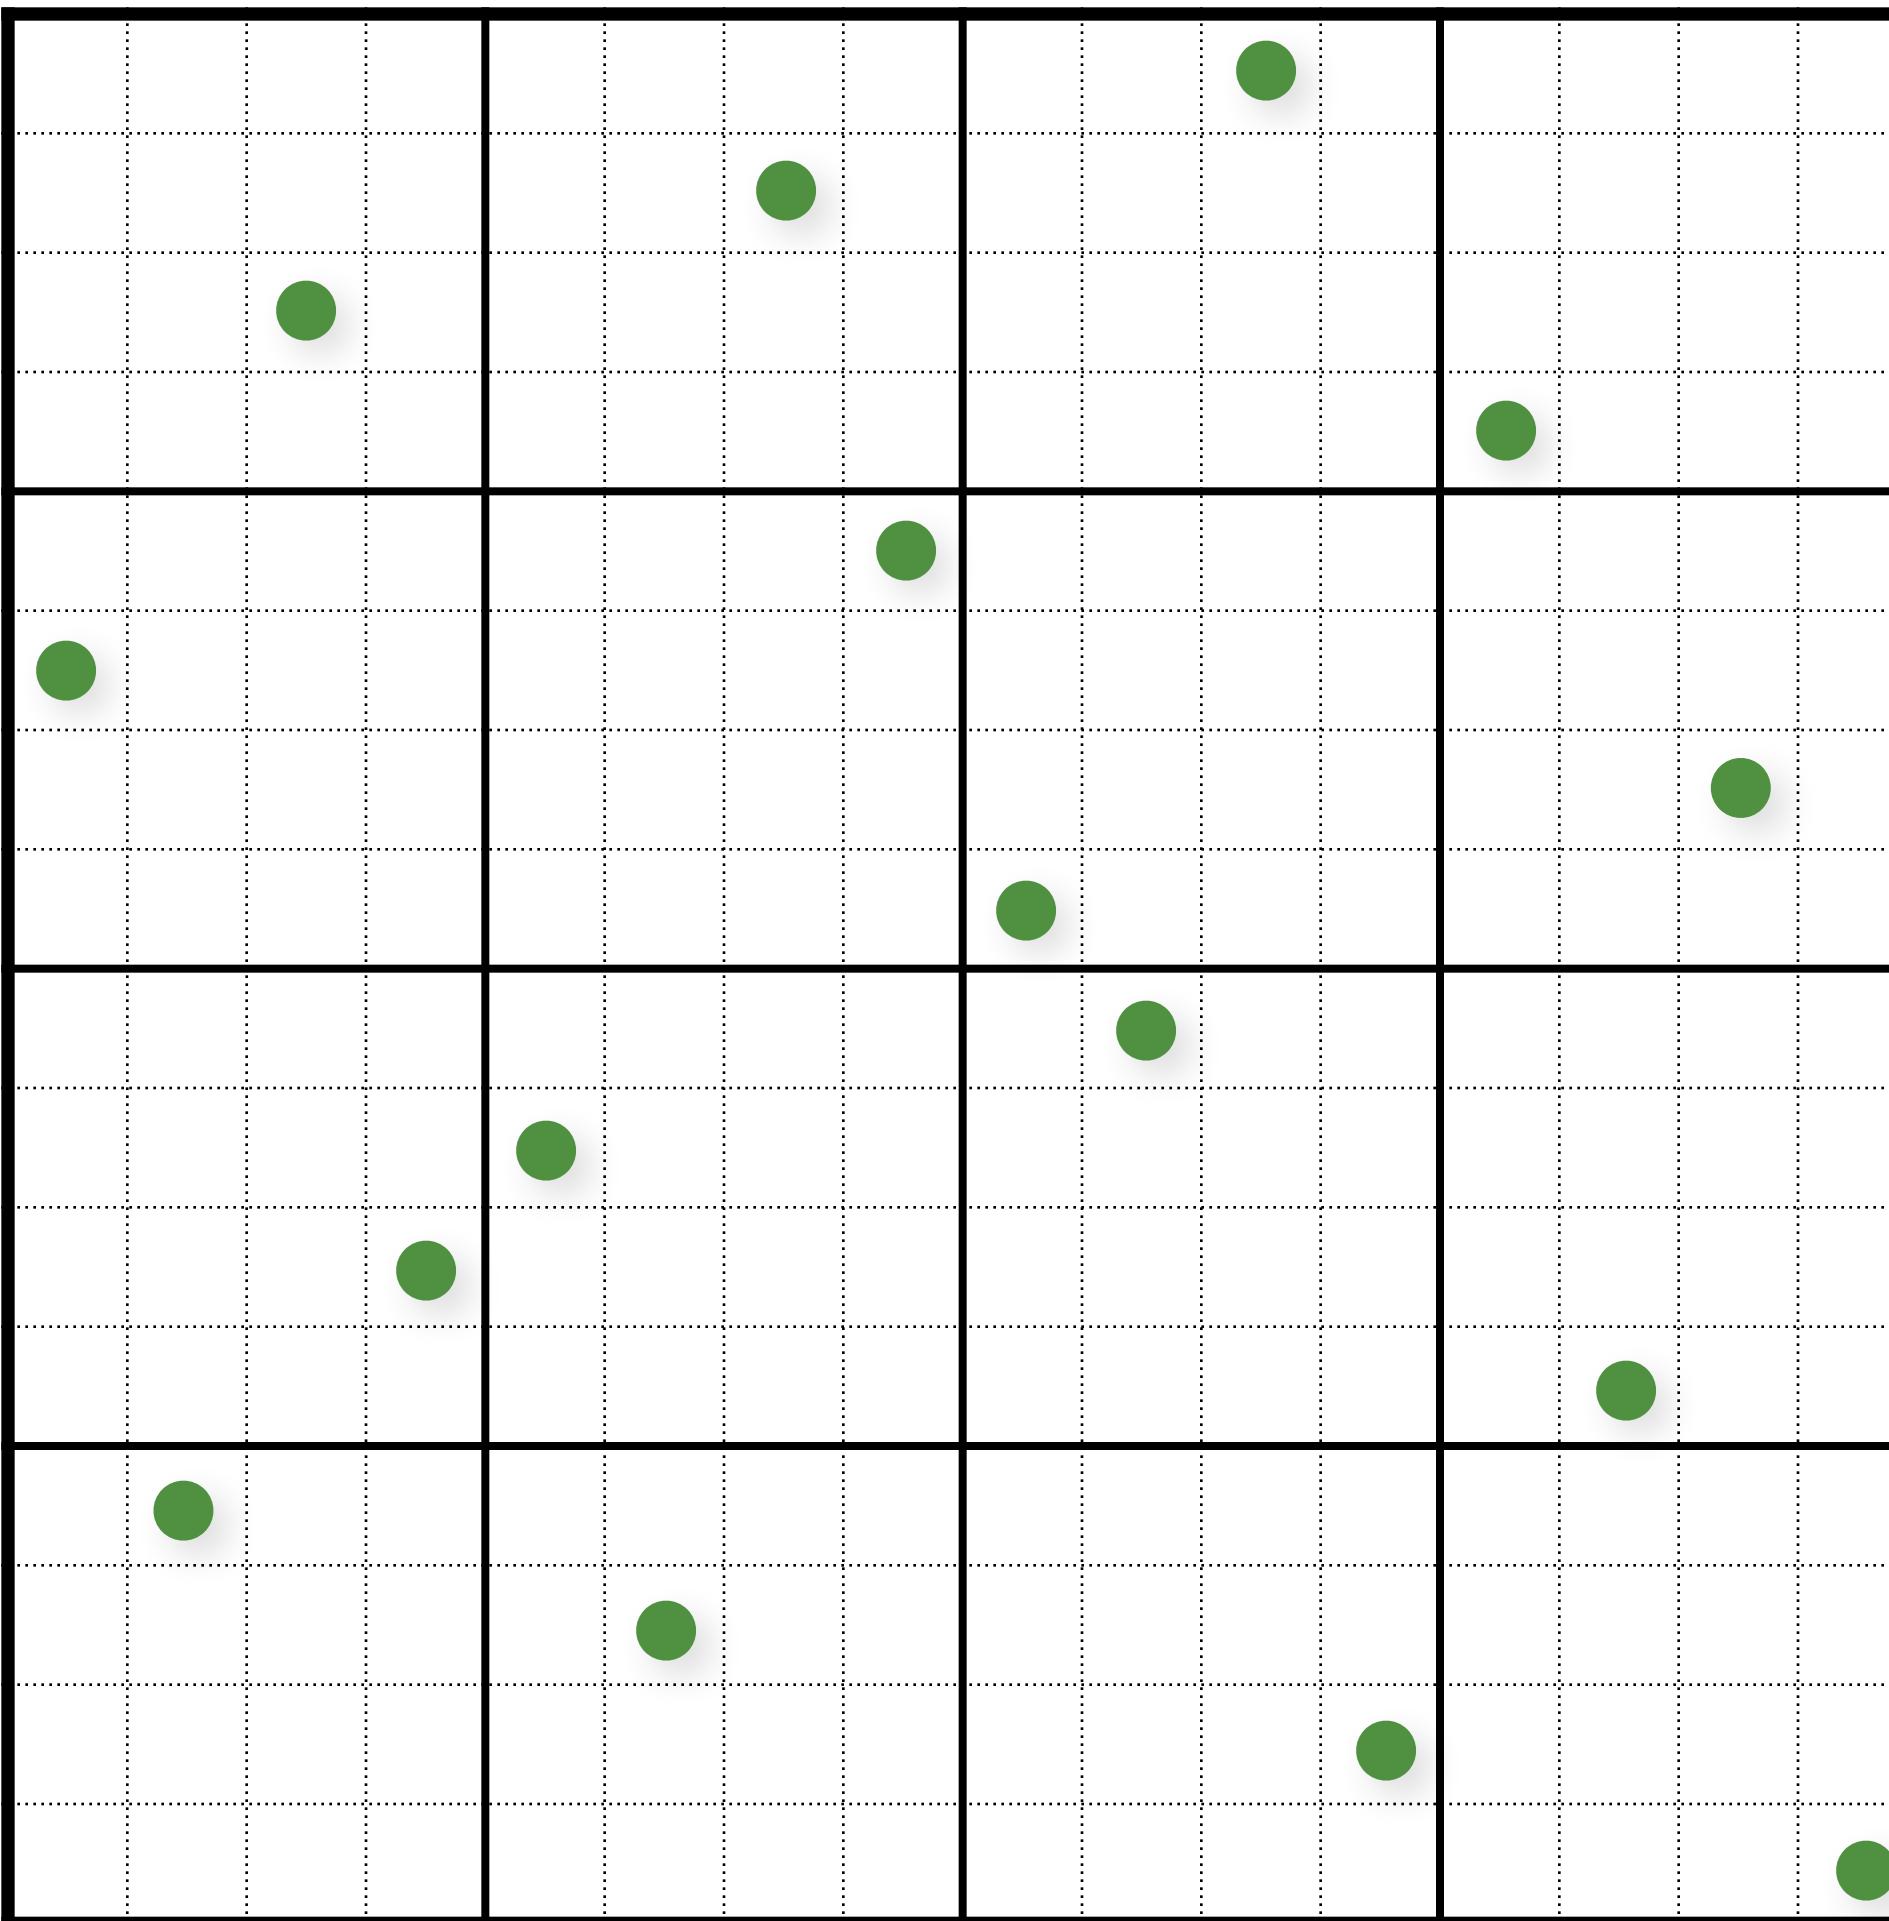
Shuffle y-coords

Multi-Jittered Sampling



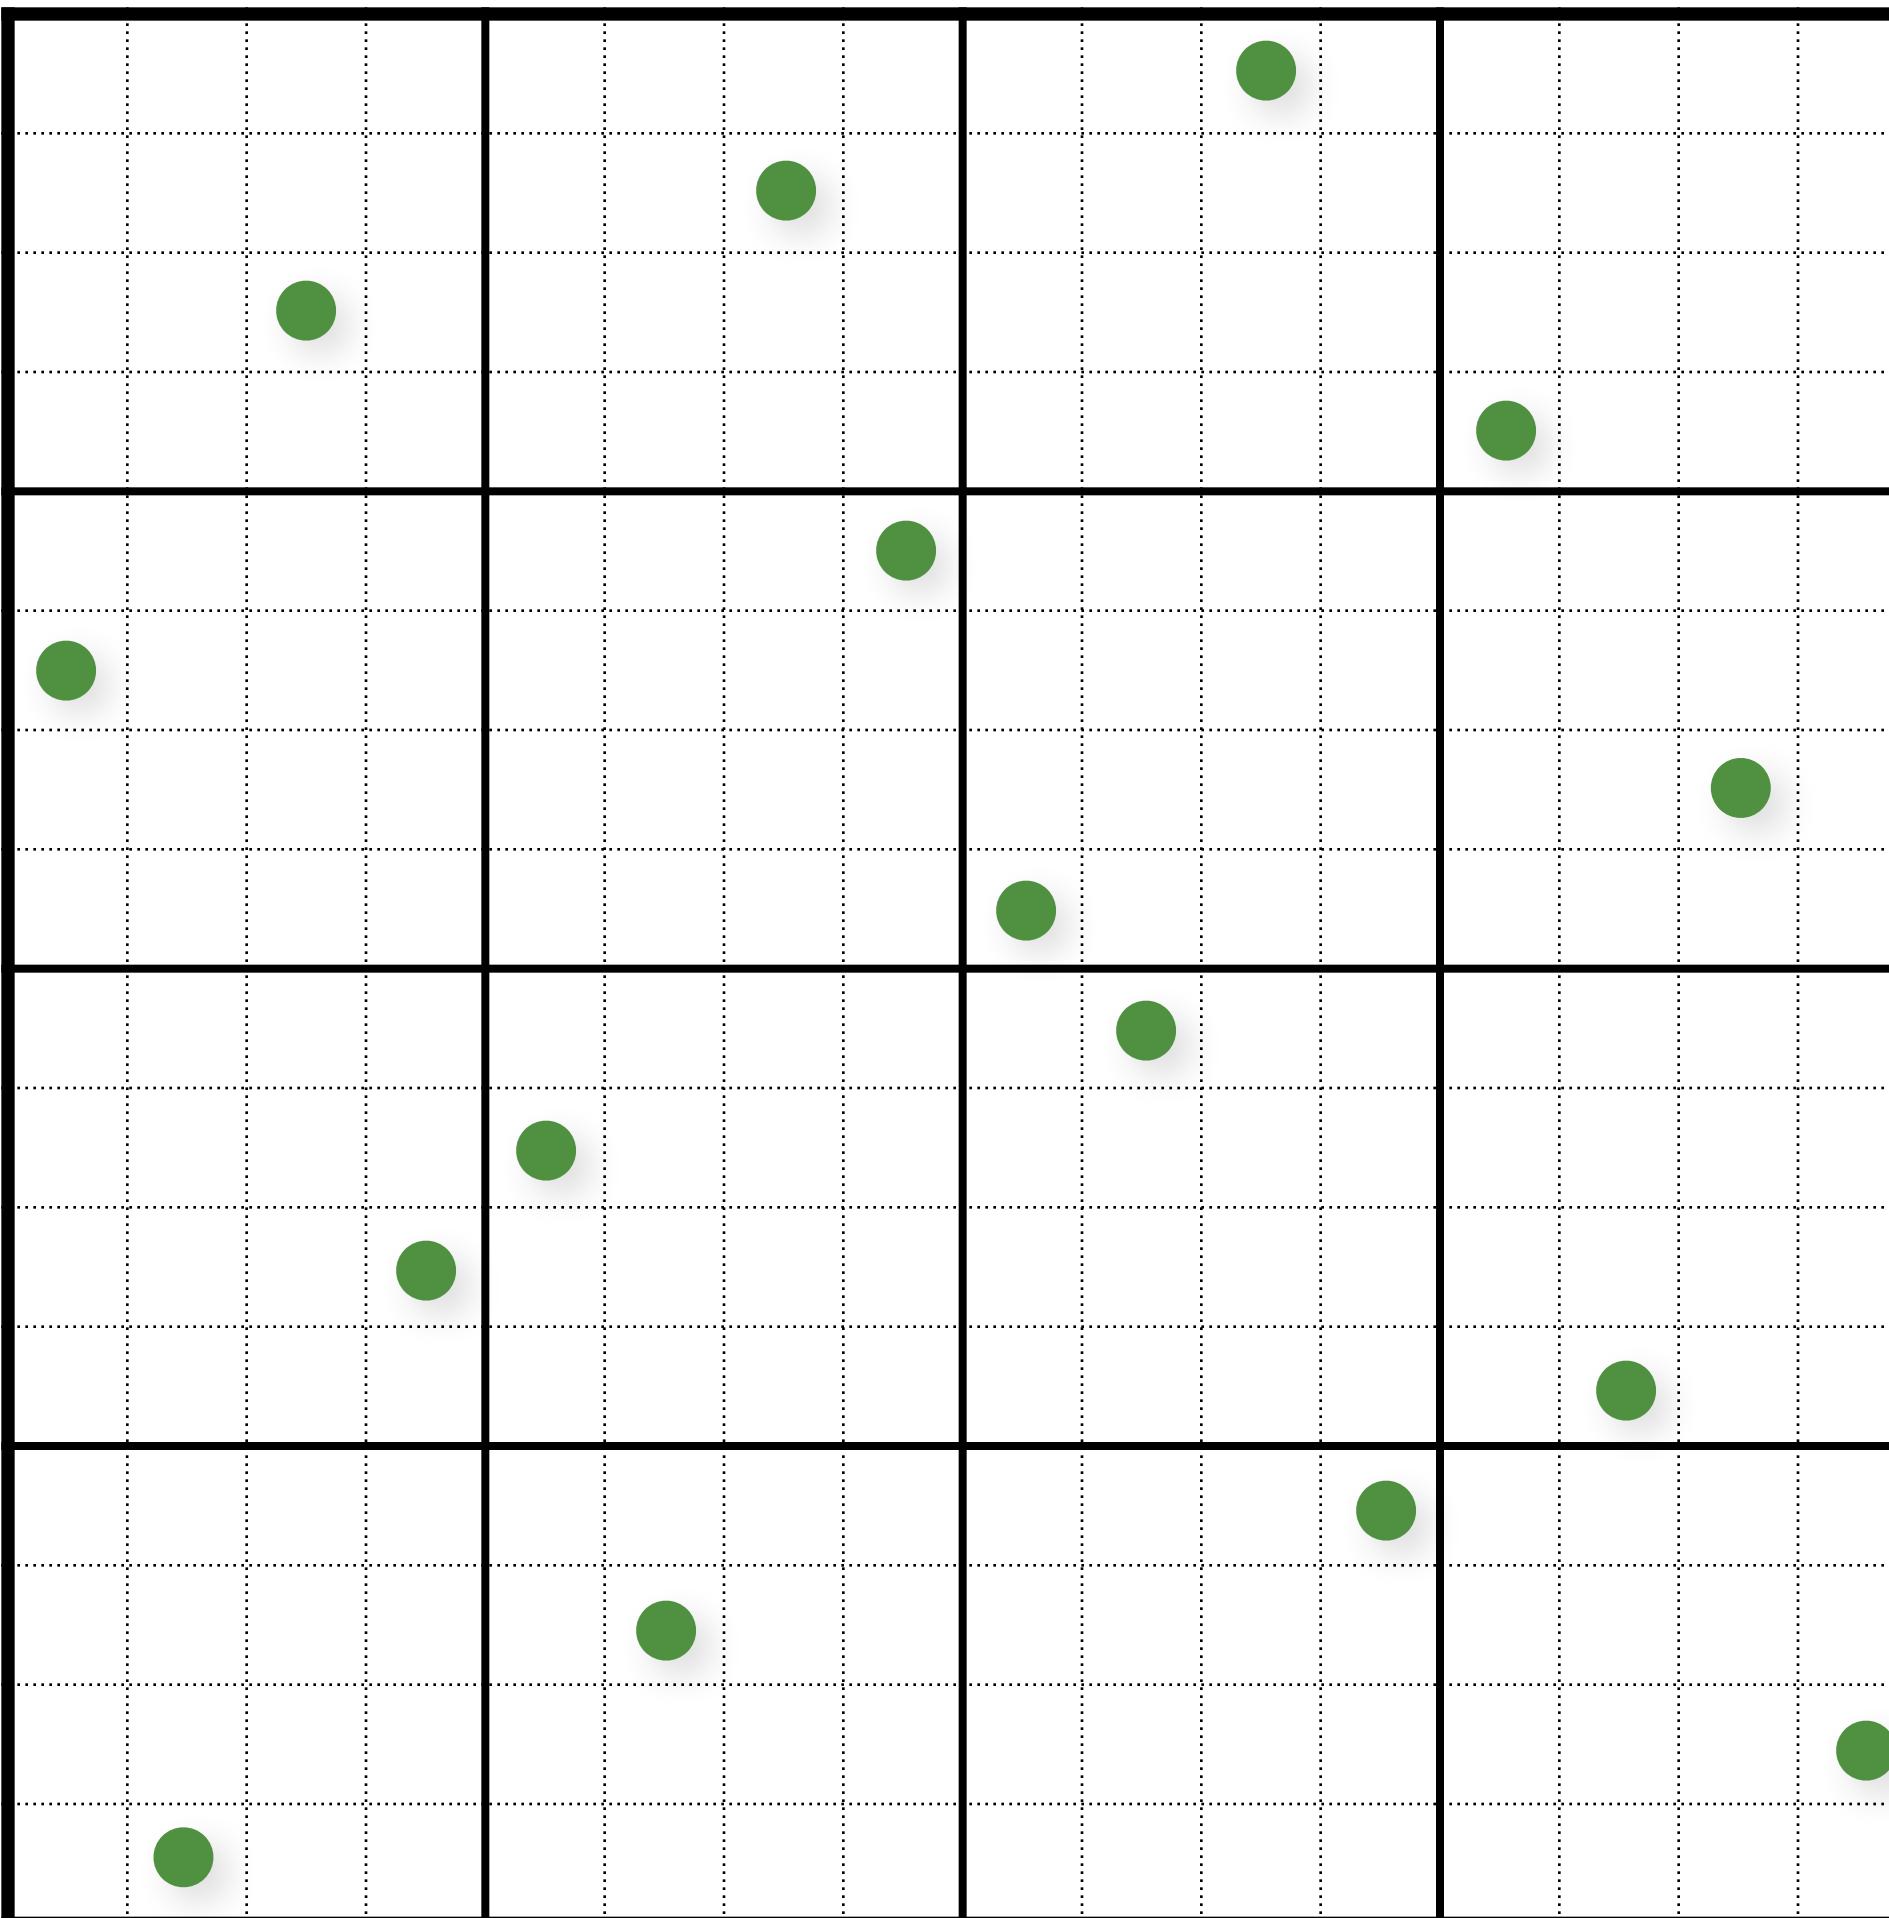
Shuffle y-coords

Multi-Jittered Sampling

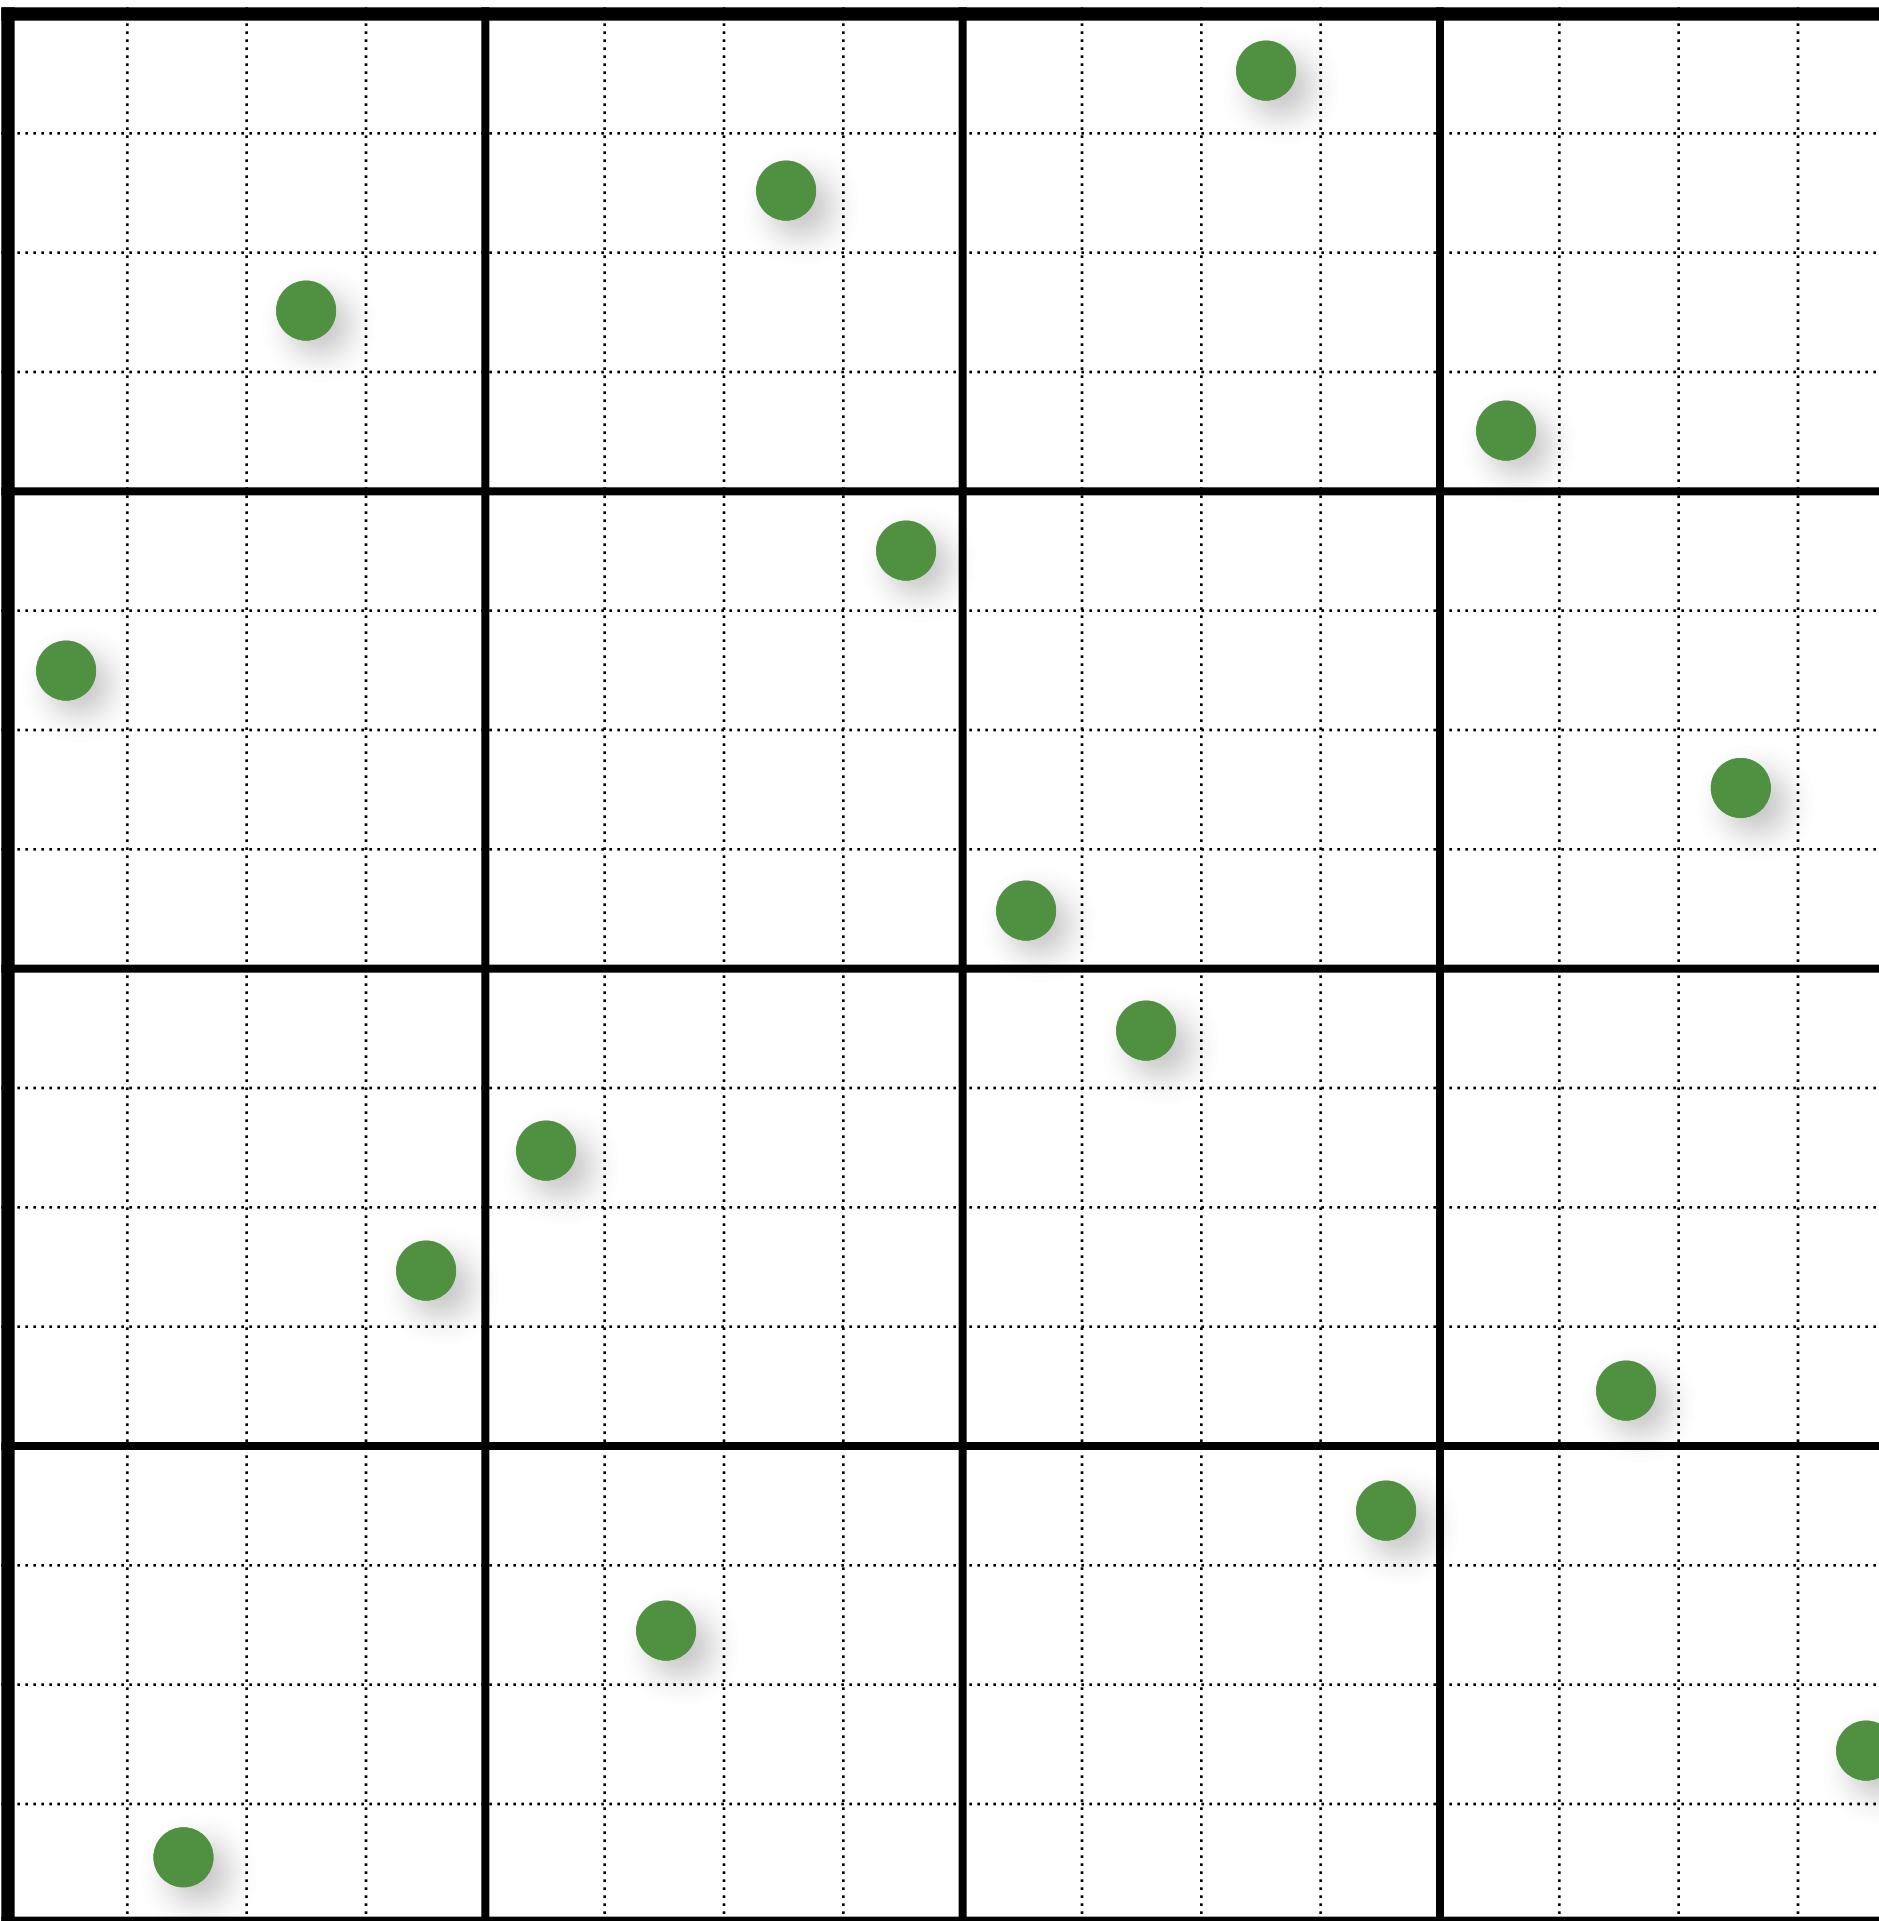


Shuffle y-coords

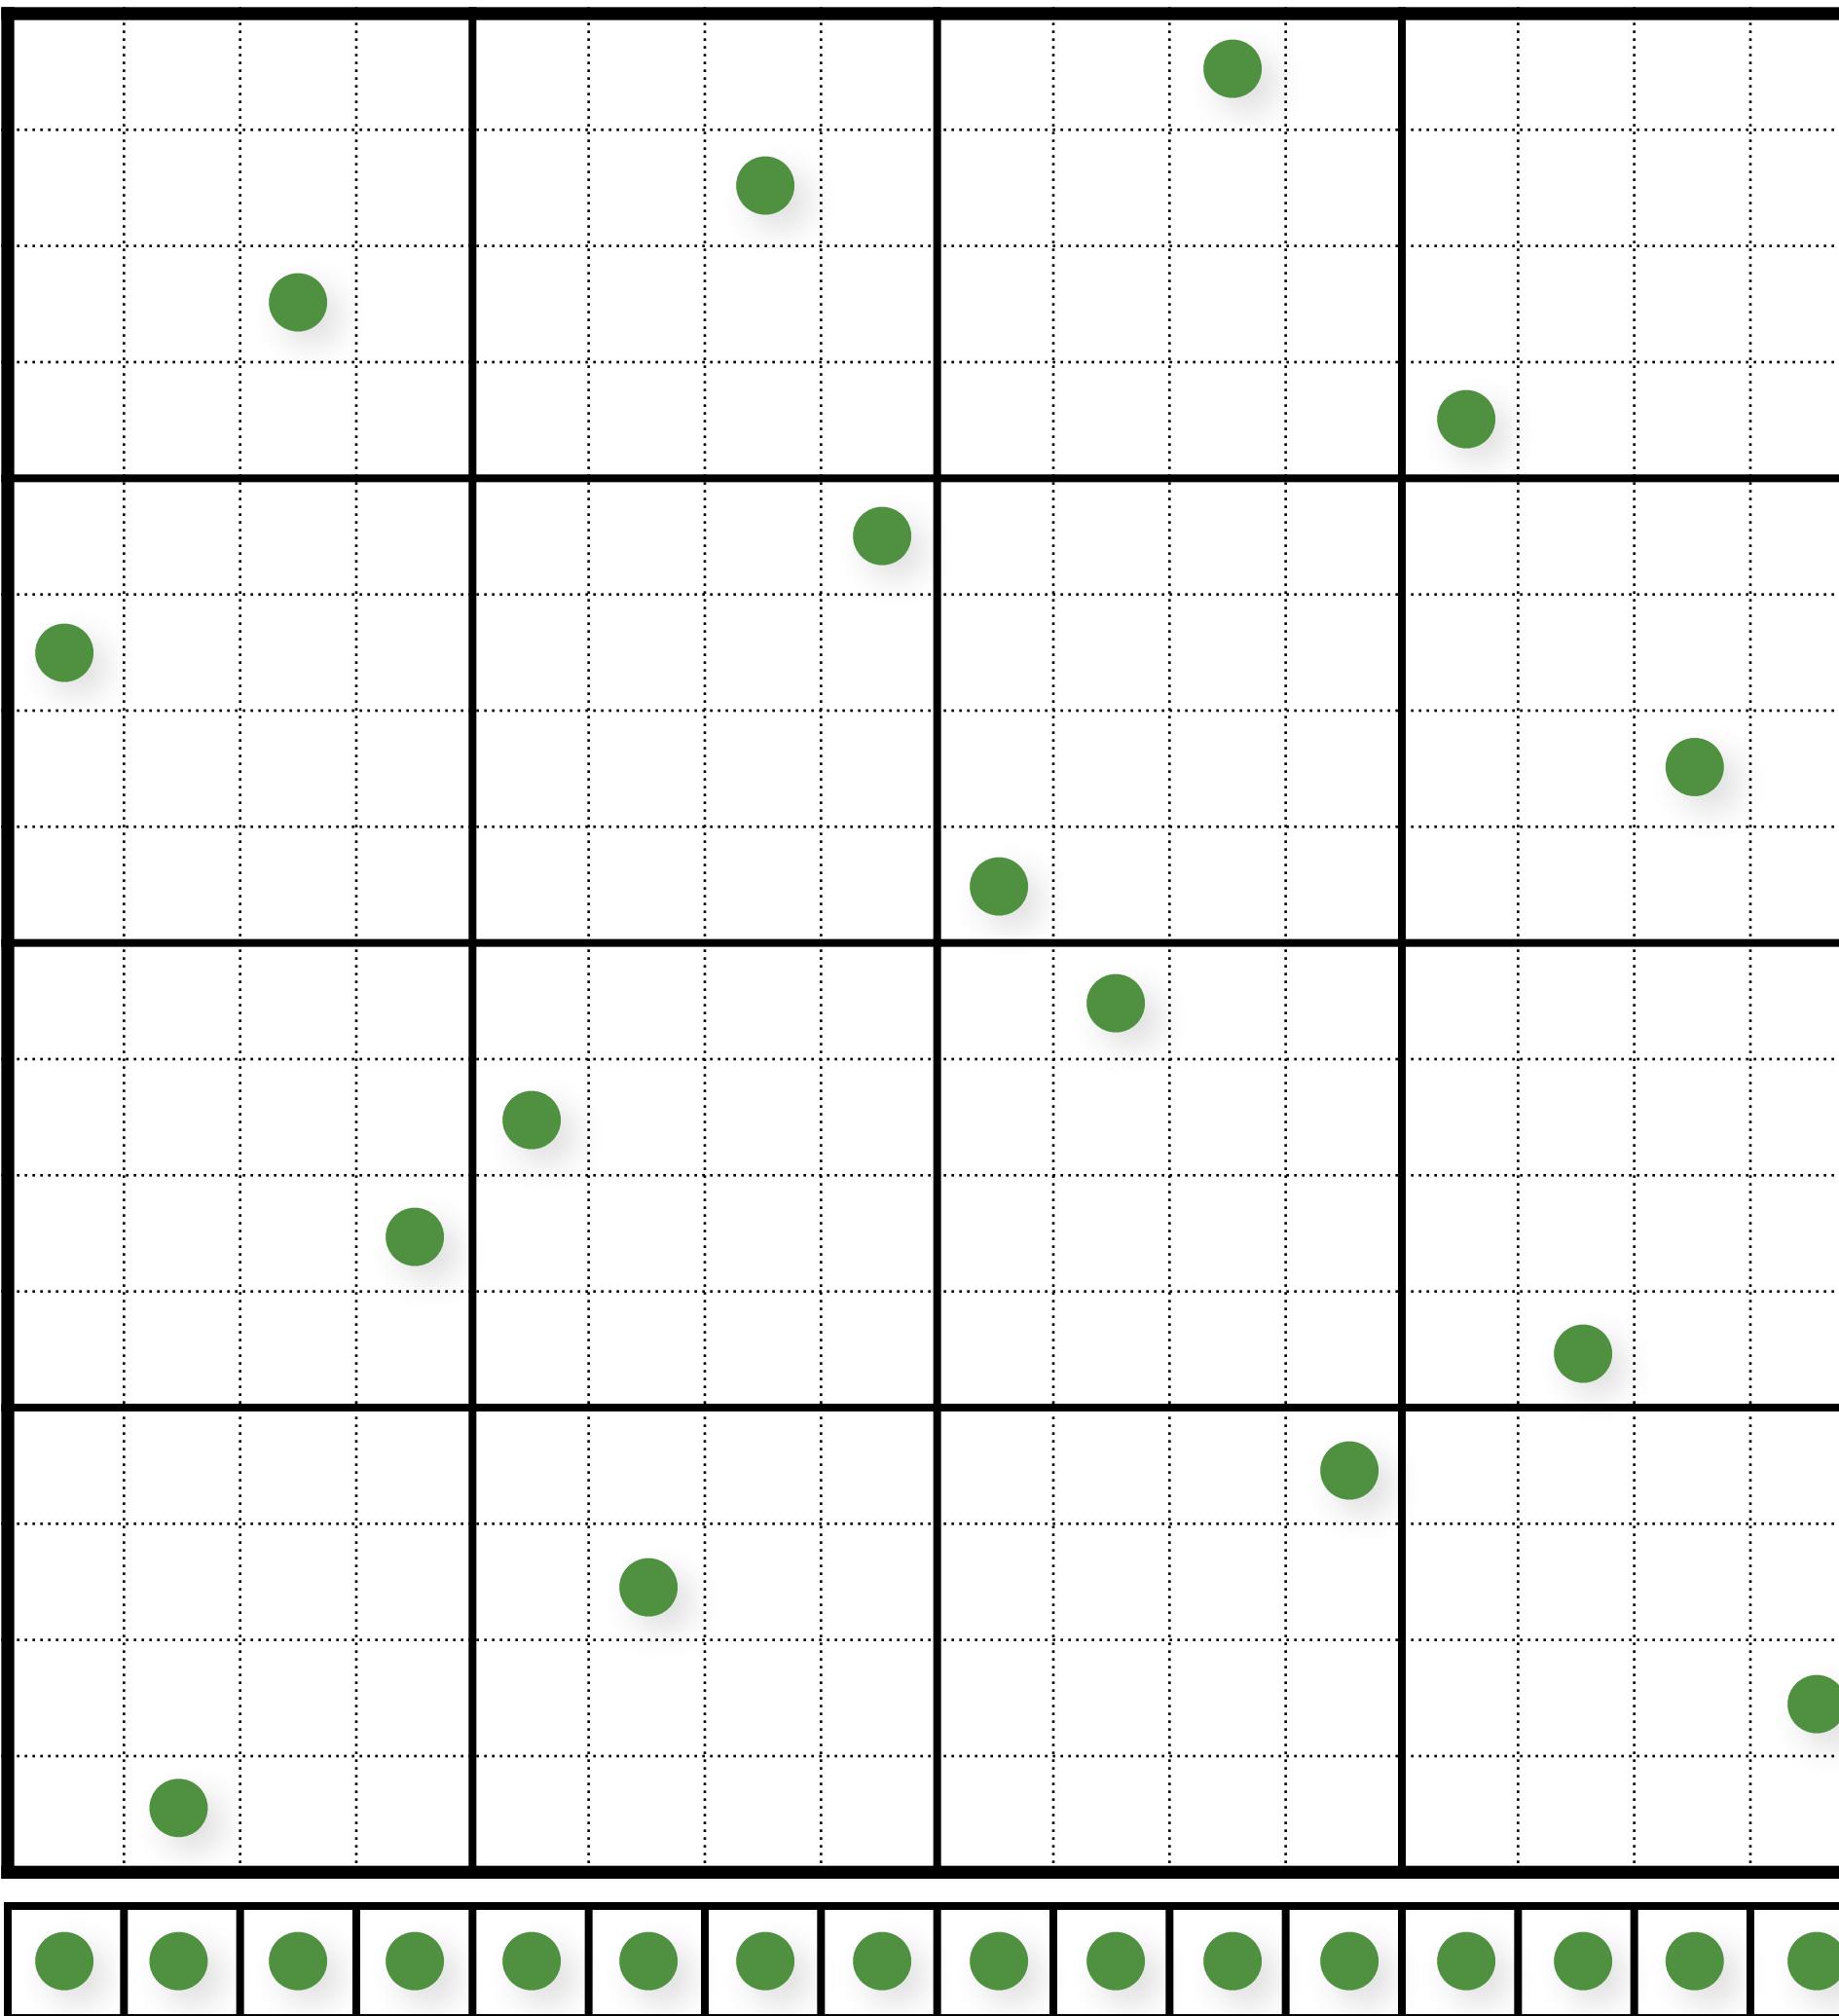
Multi-Jittered Sampling



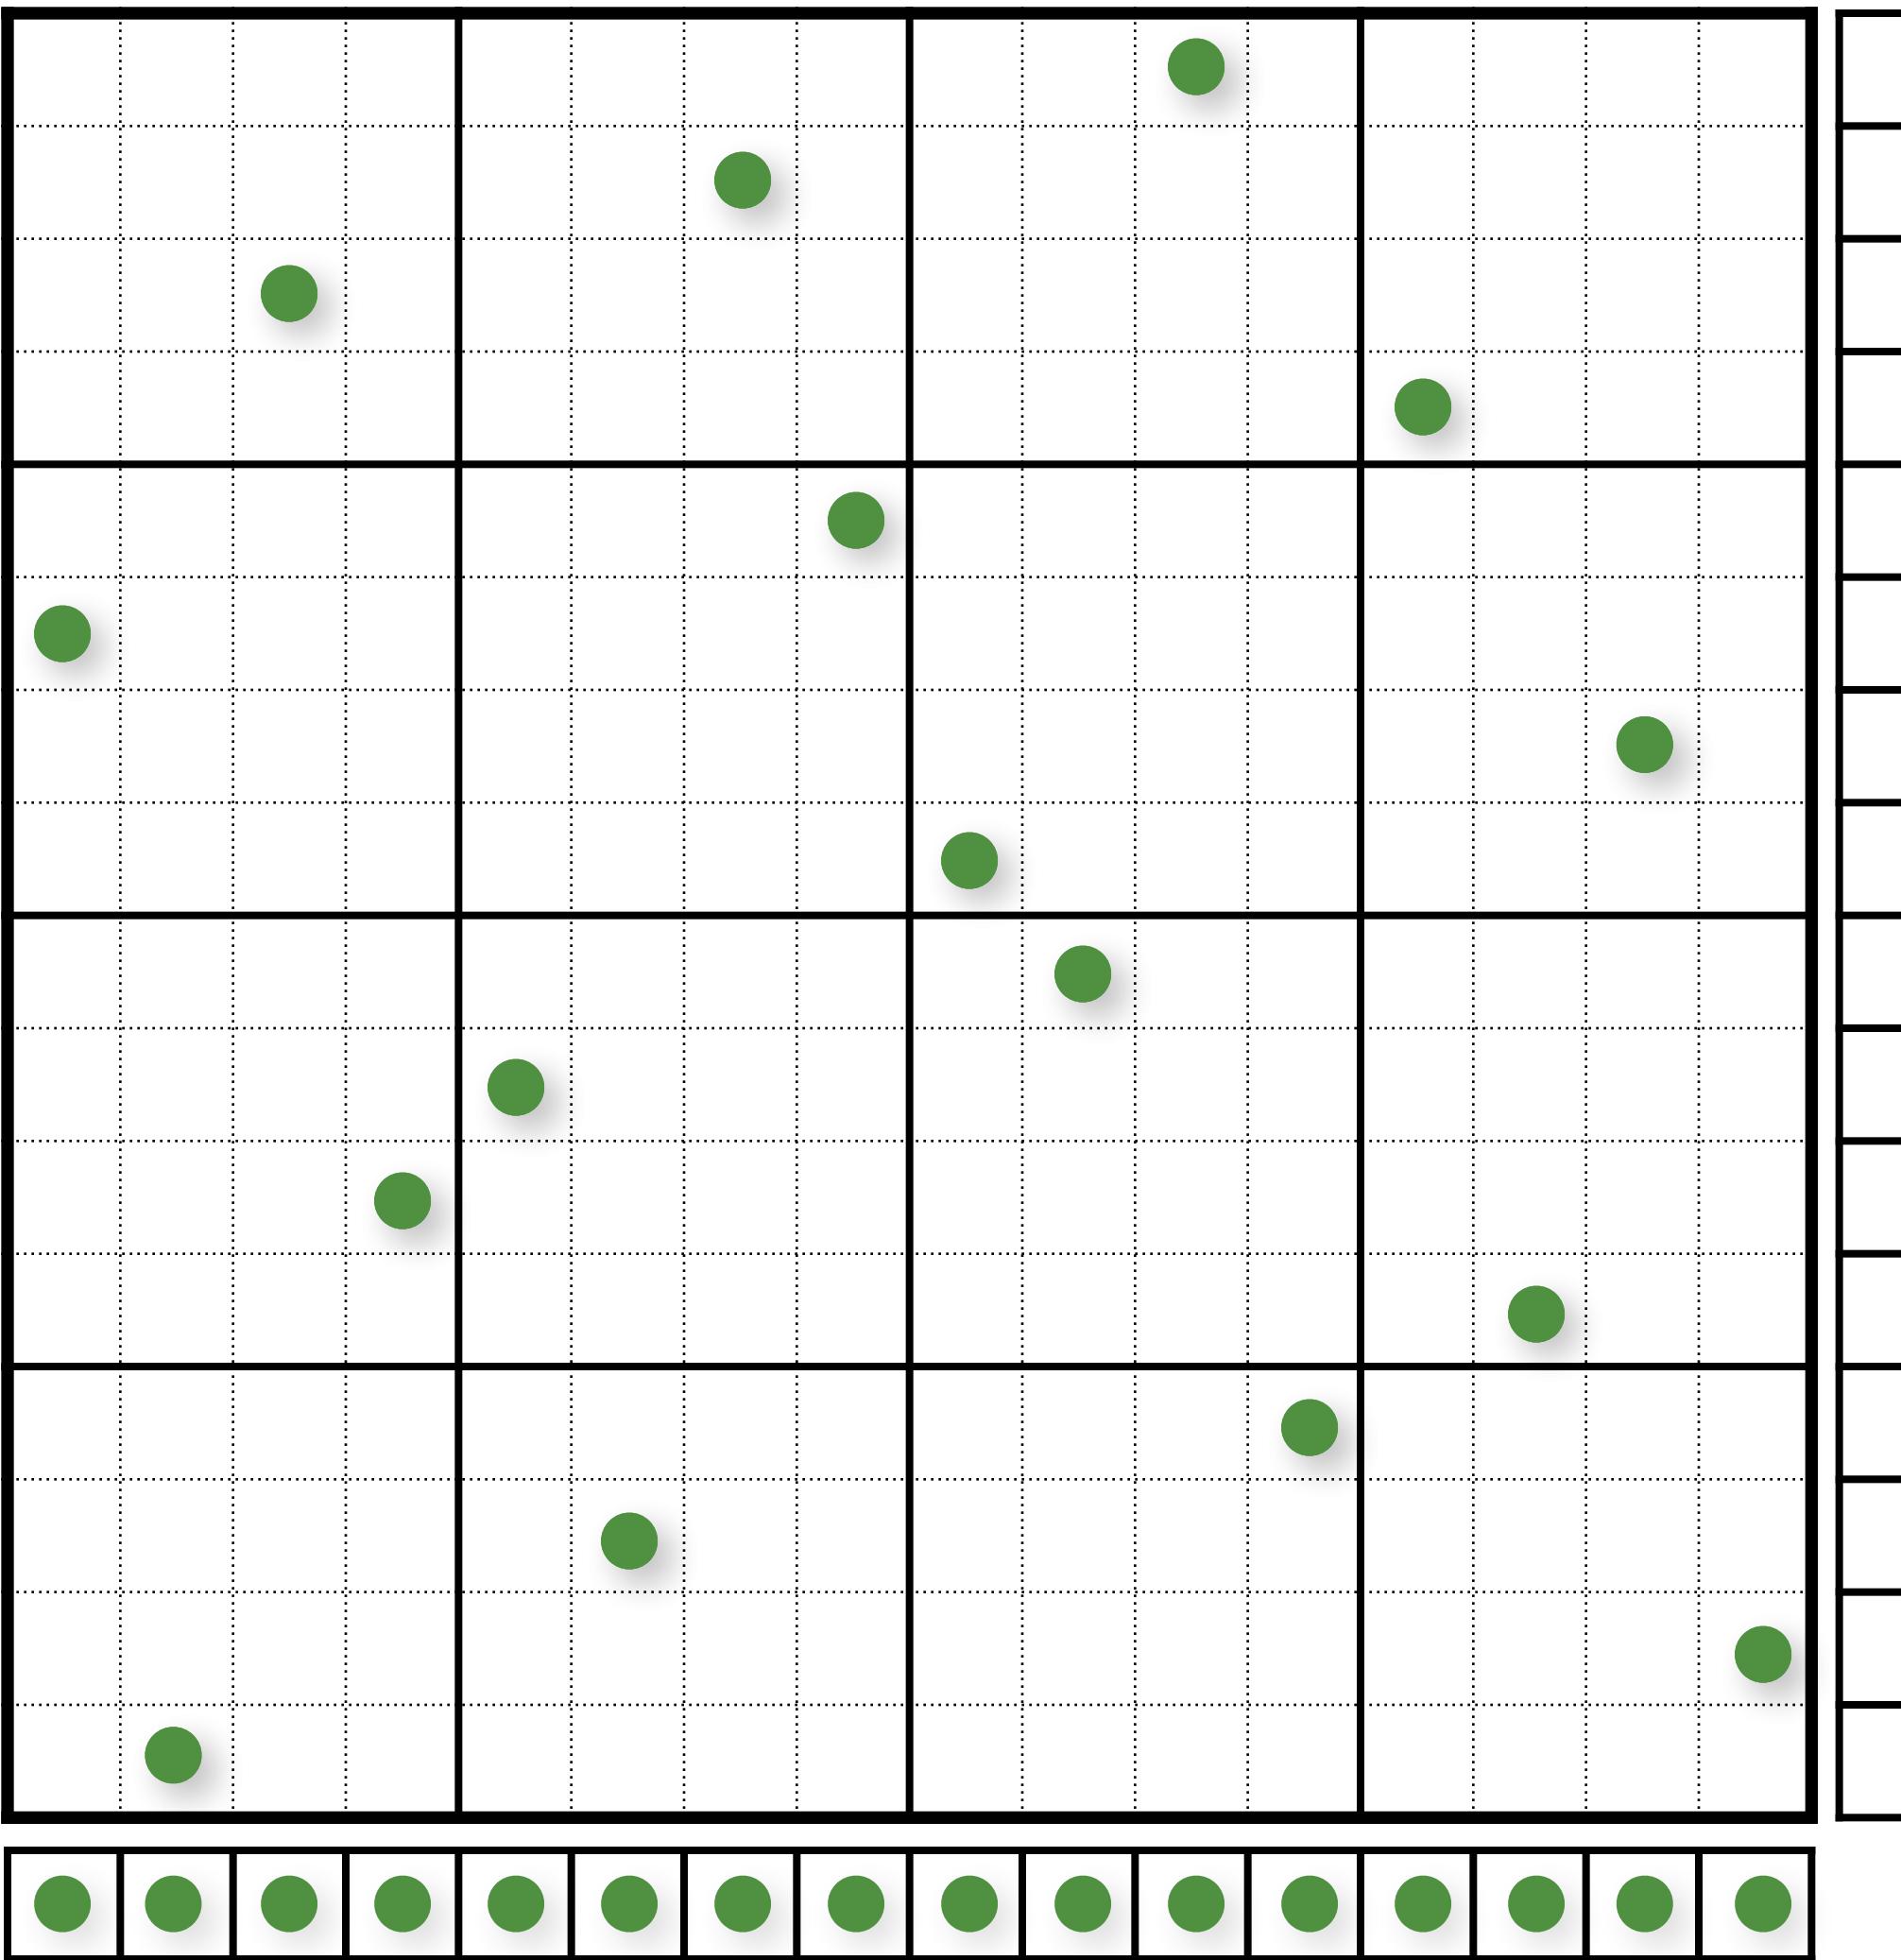
Multi-Jittered Sampling (Projections)



Multi-Jittered Sampling (Projections)

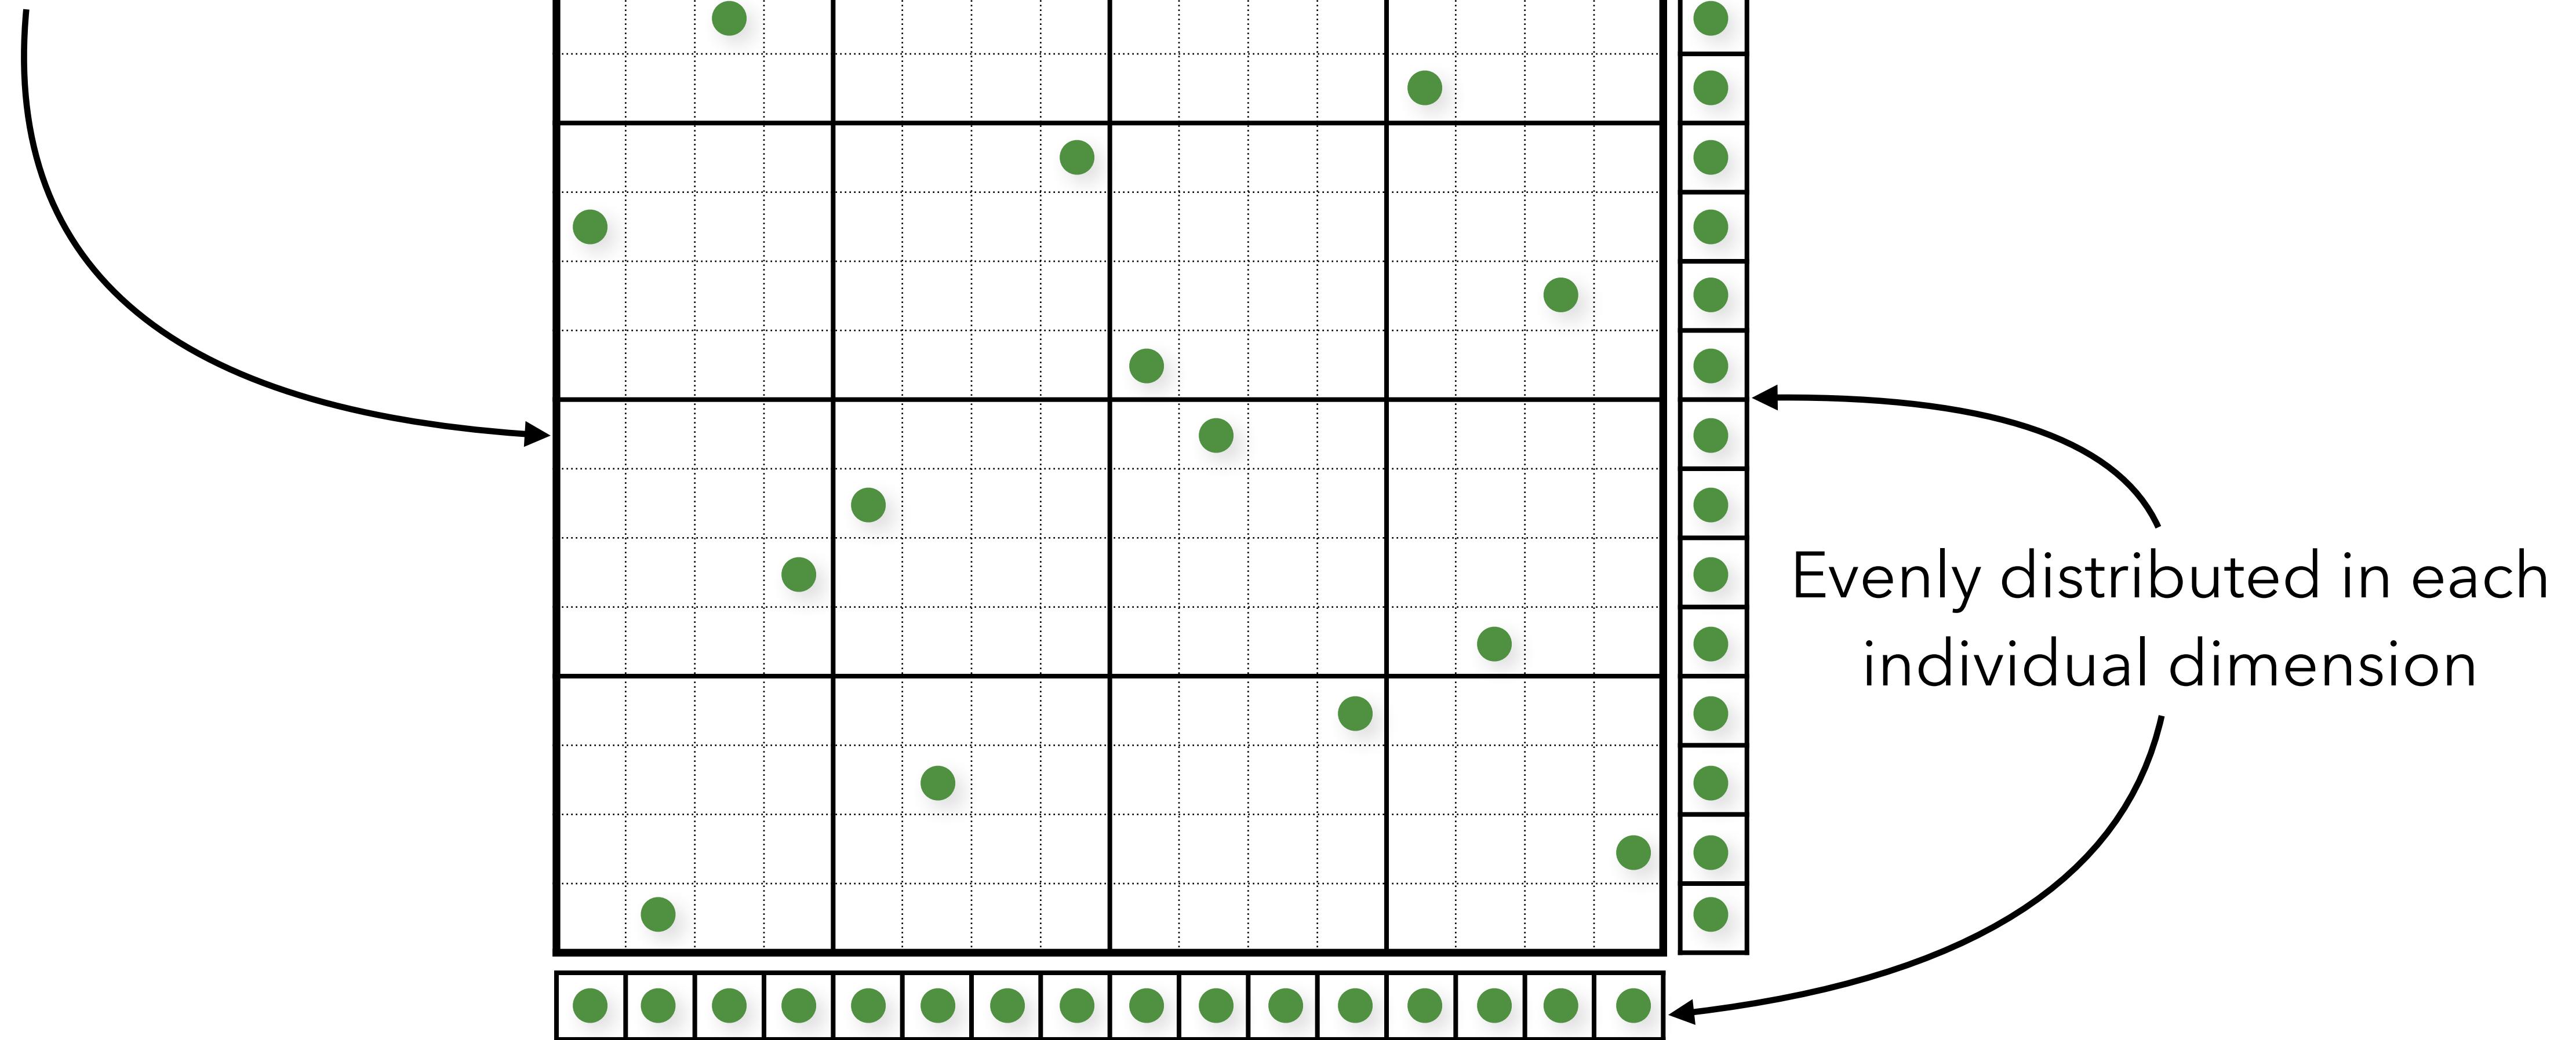


Multi-Jittered Sampling (Projections)



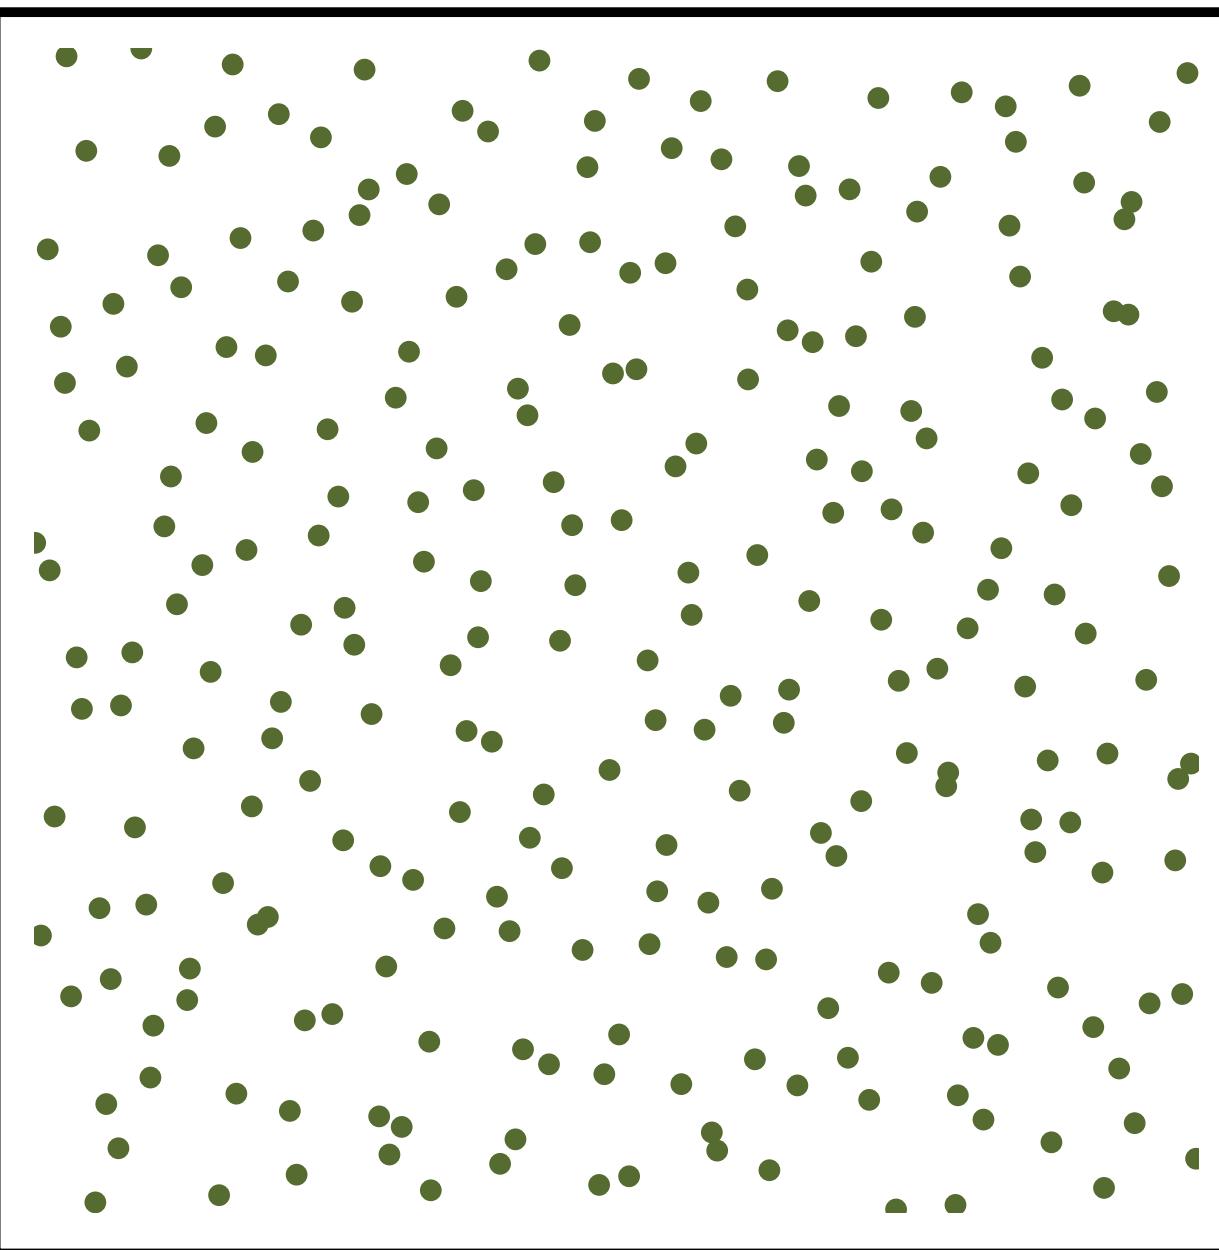
Multi-Jittered Sampling (Projections)

Evenly distributed in 2D!

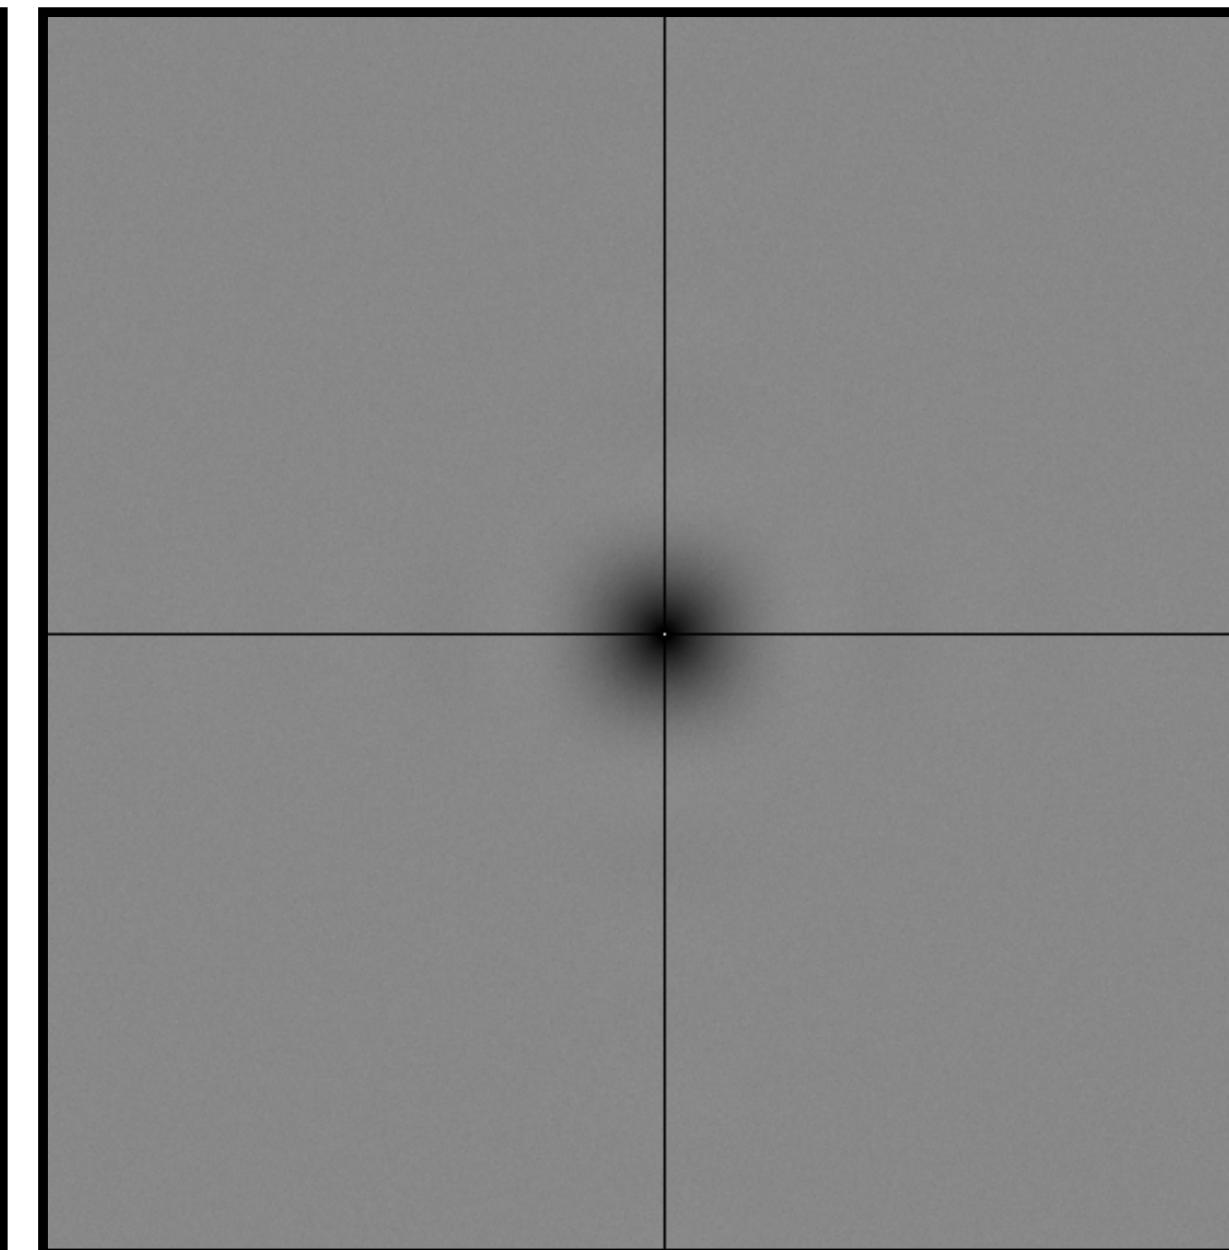


Multi-Jittered Sampling

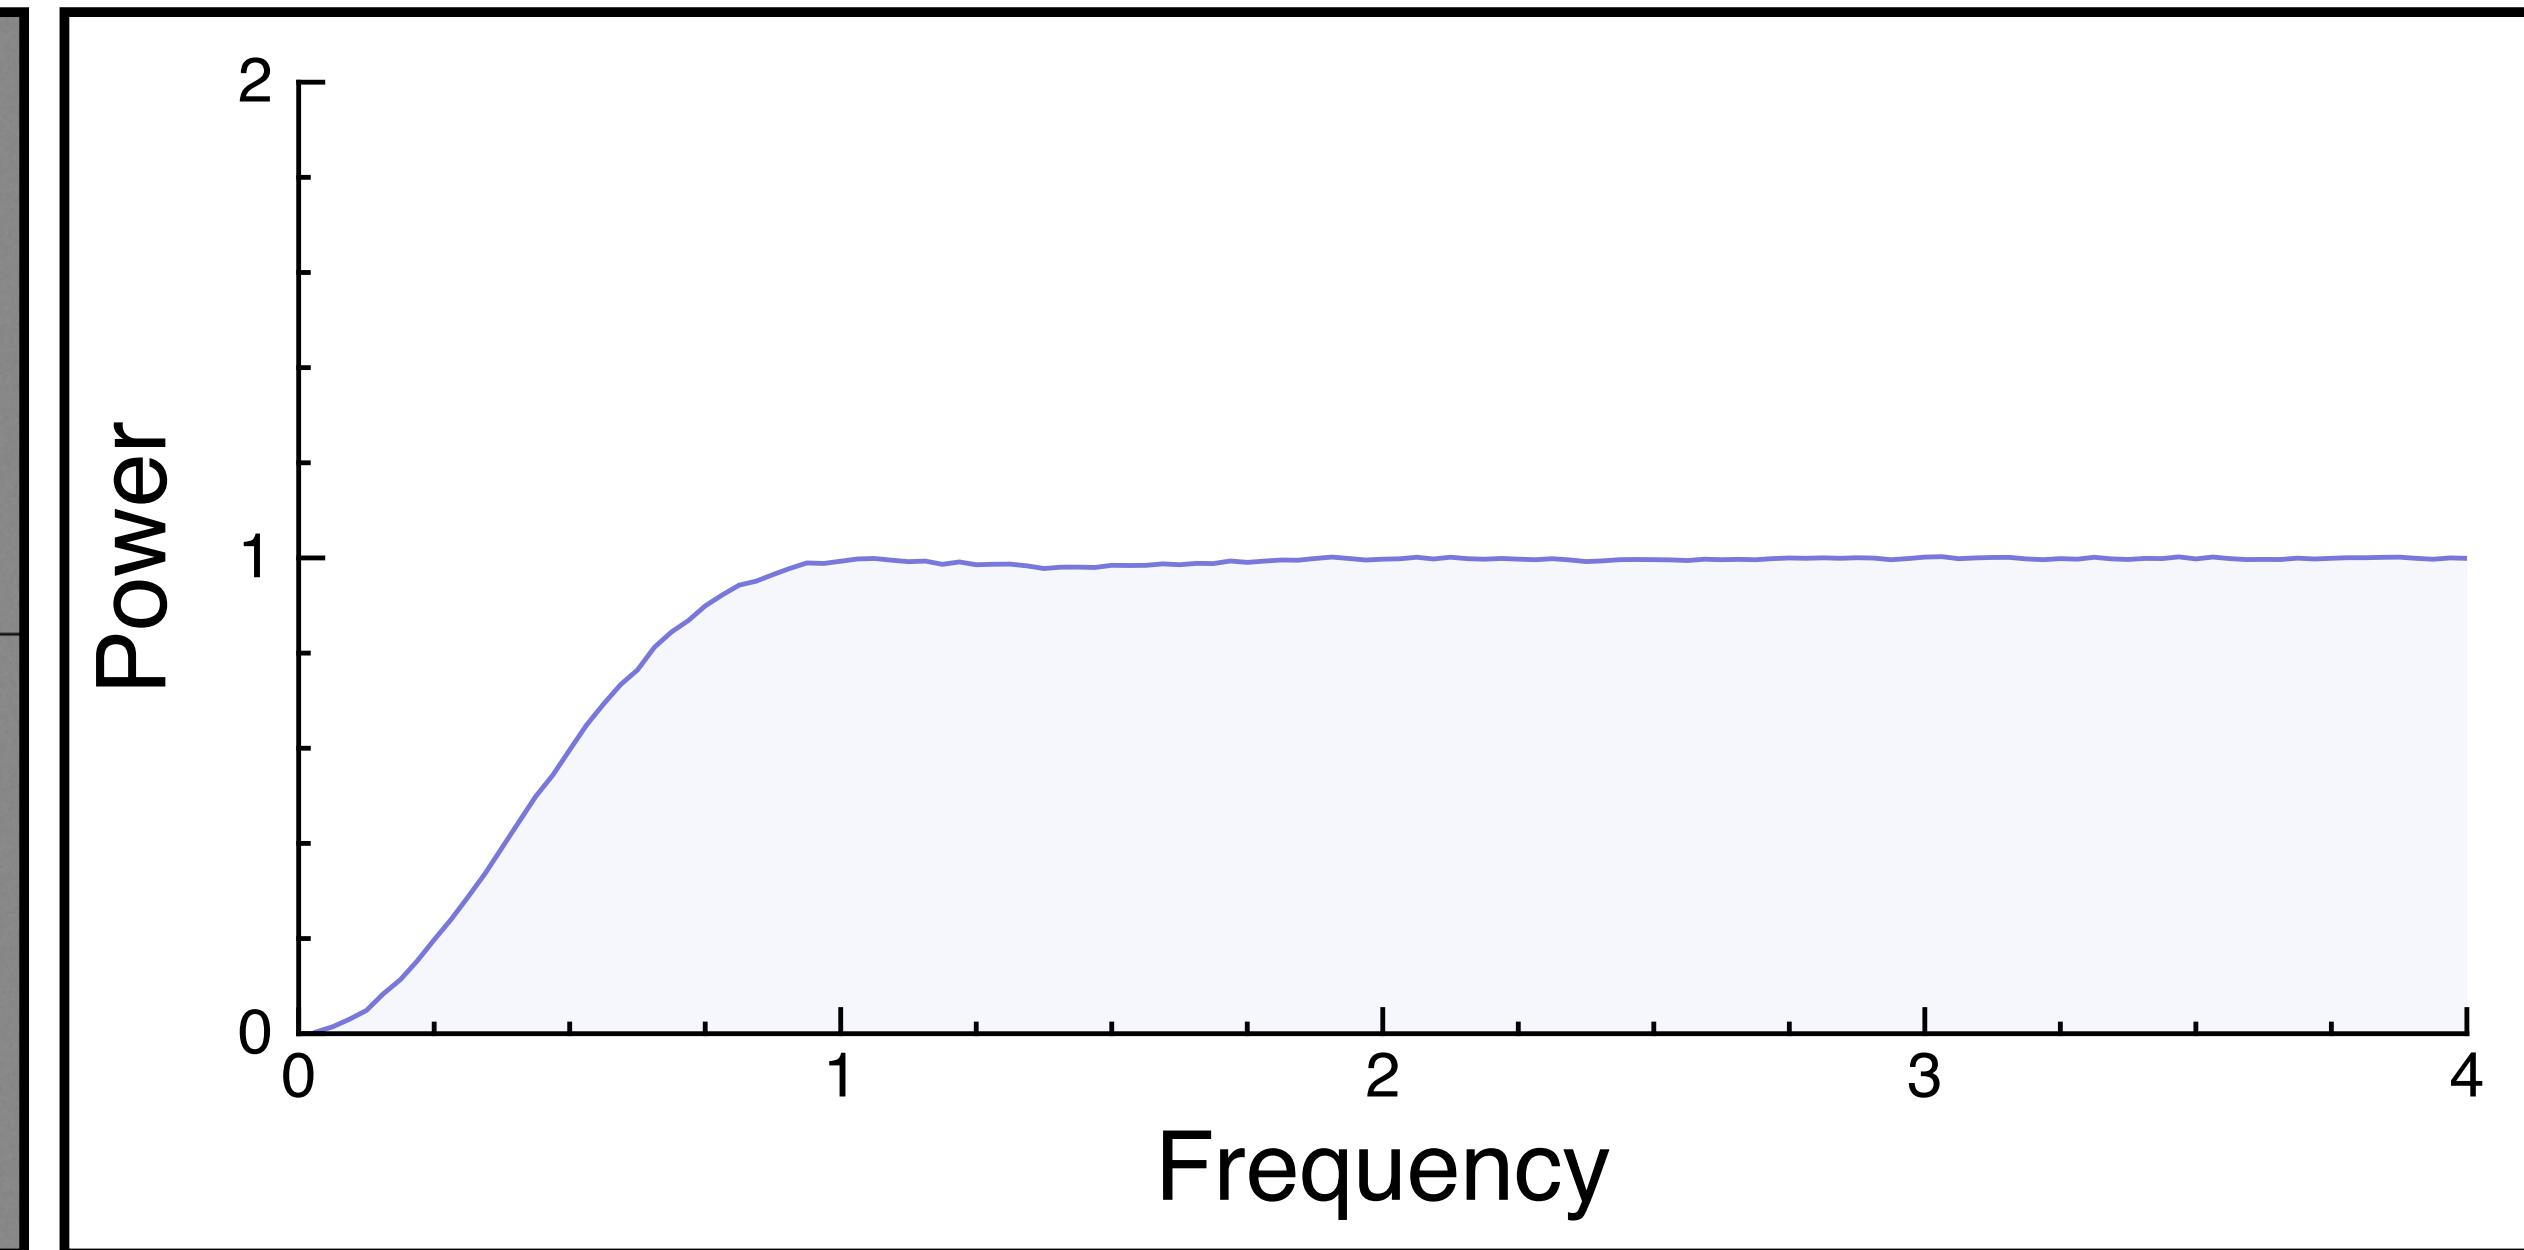
Samples



Power spectrum

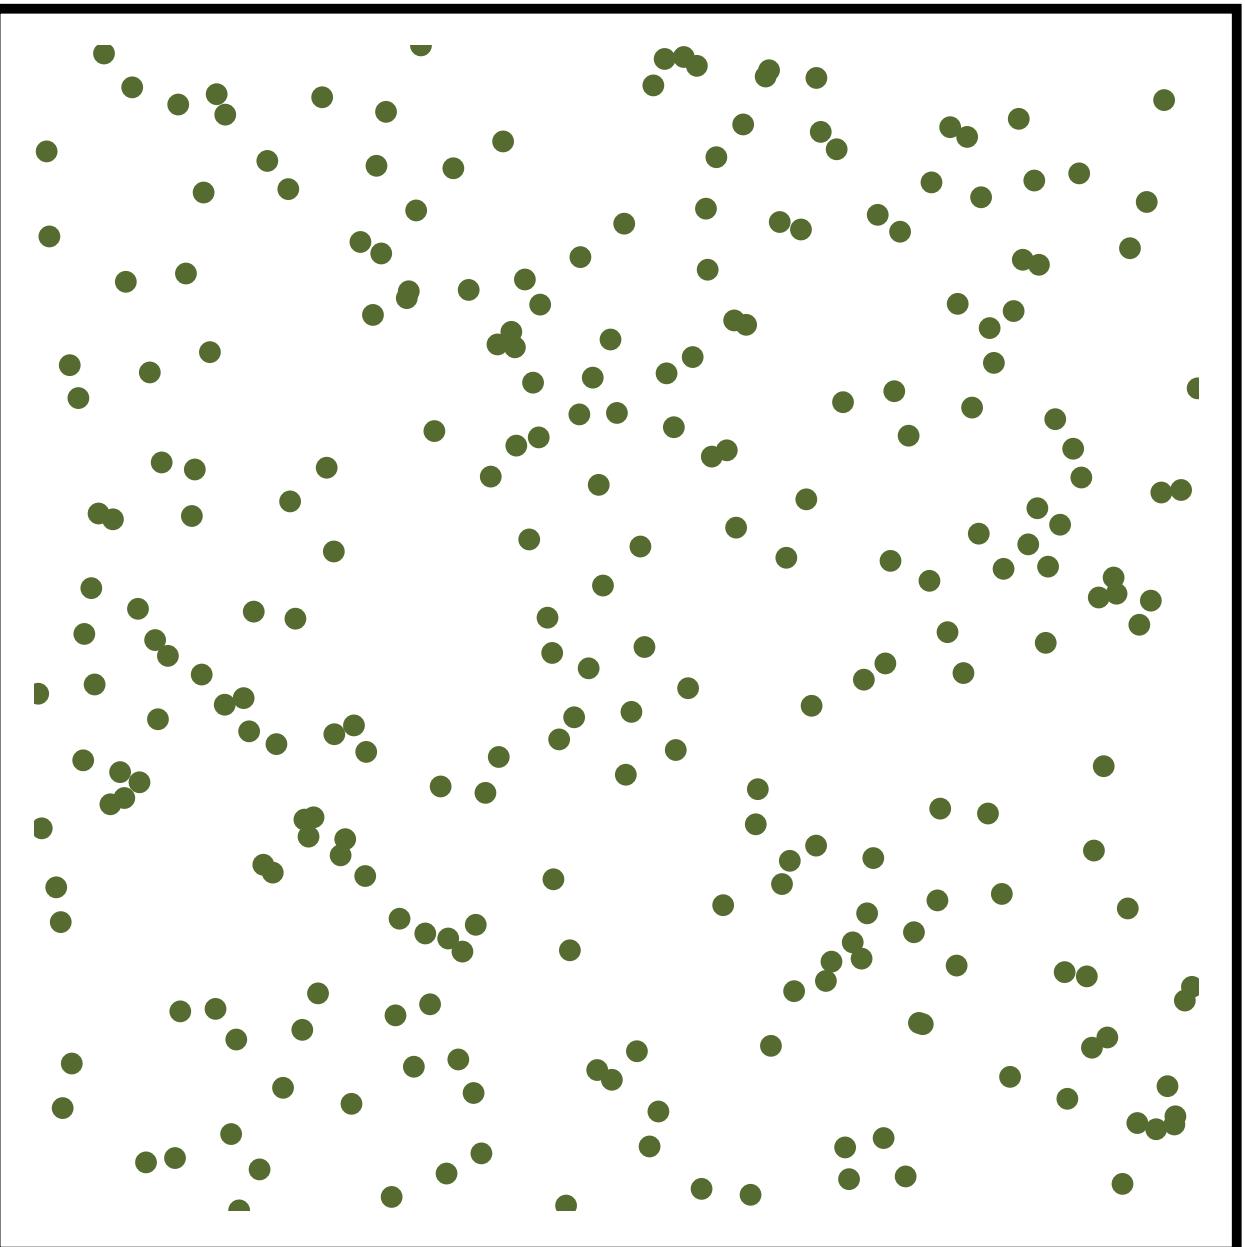


Radial mean

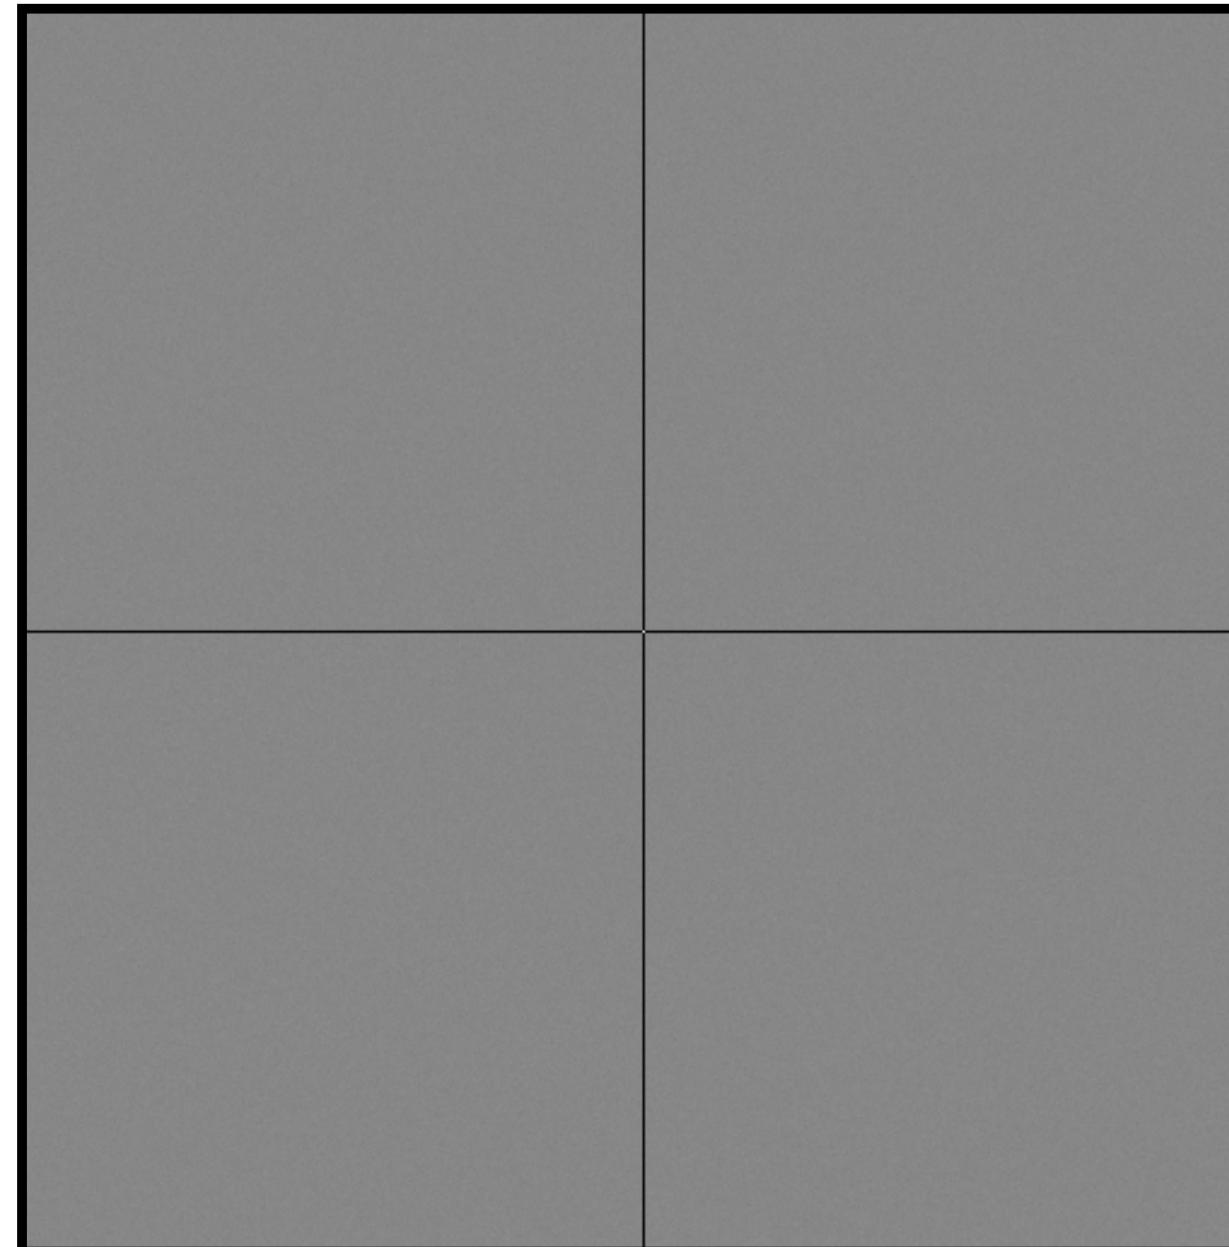


N-Rooks Sampling

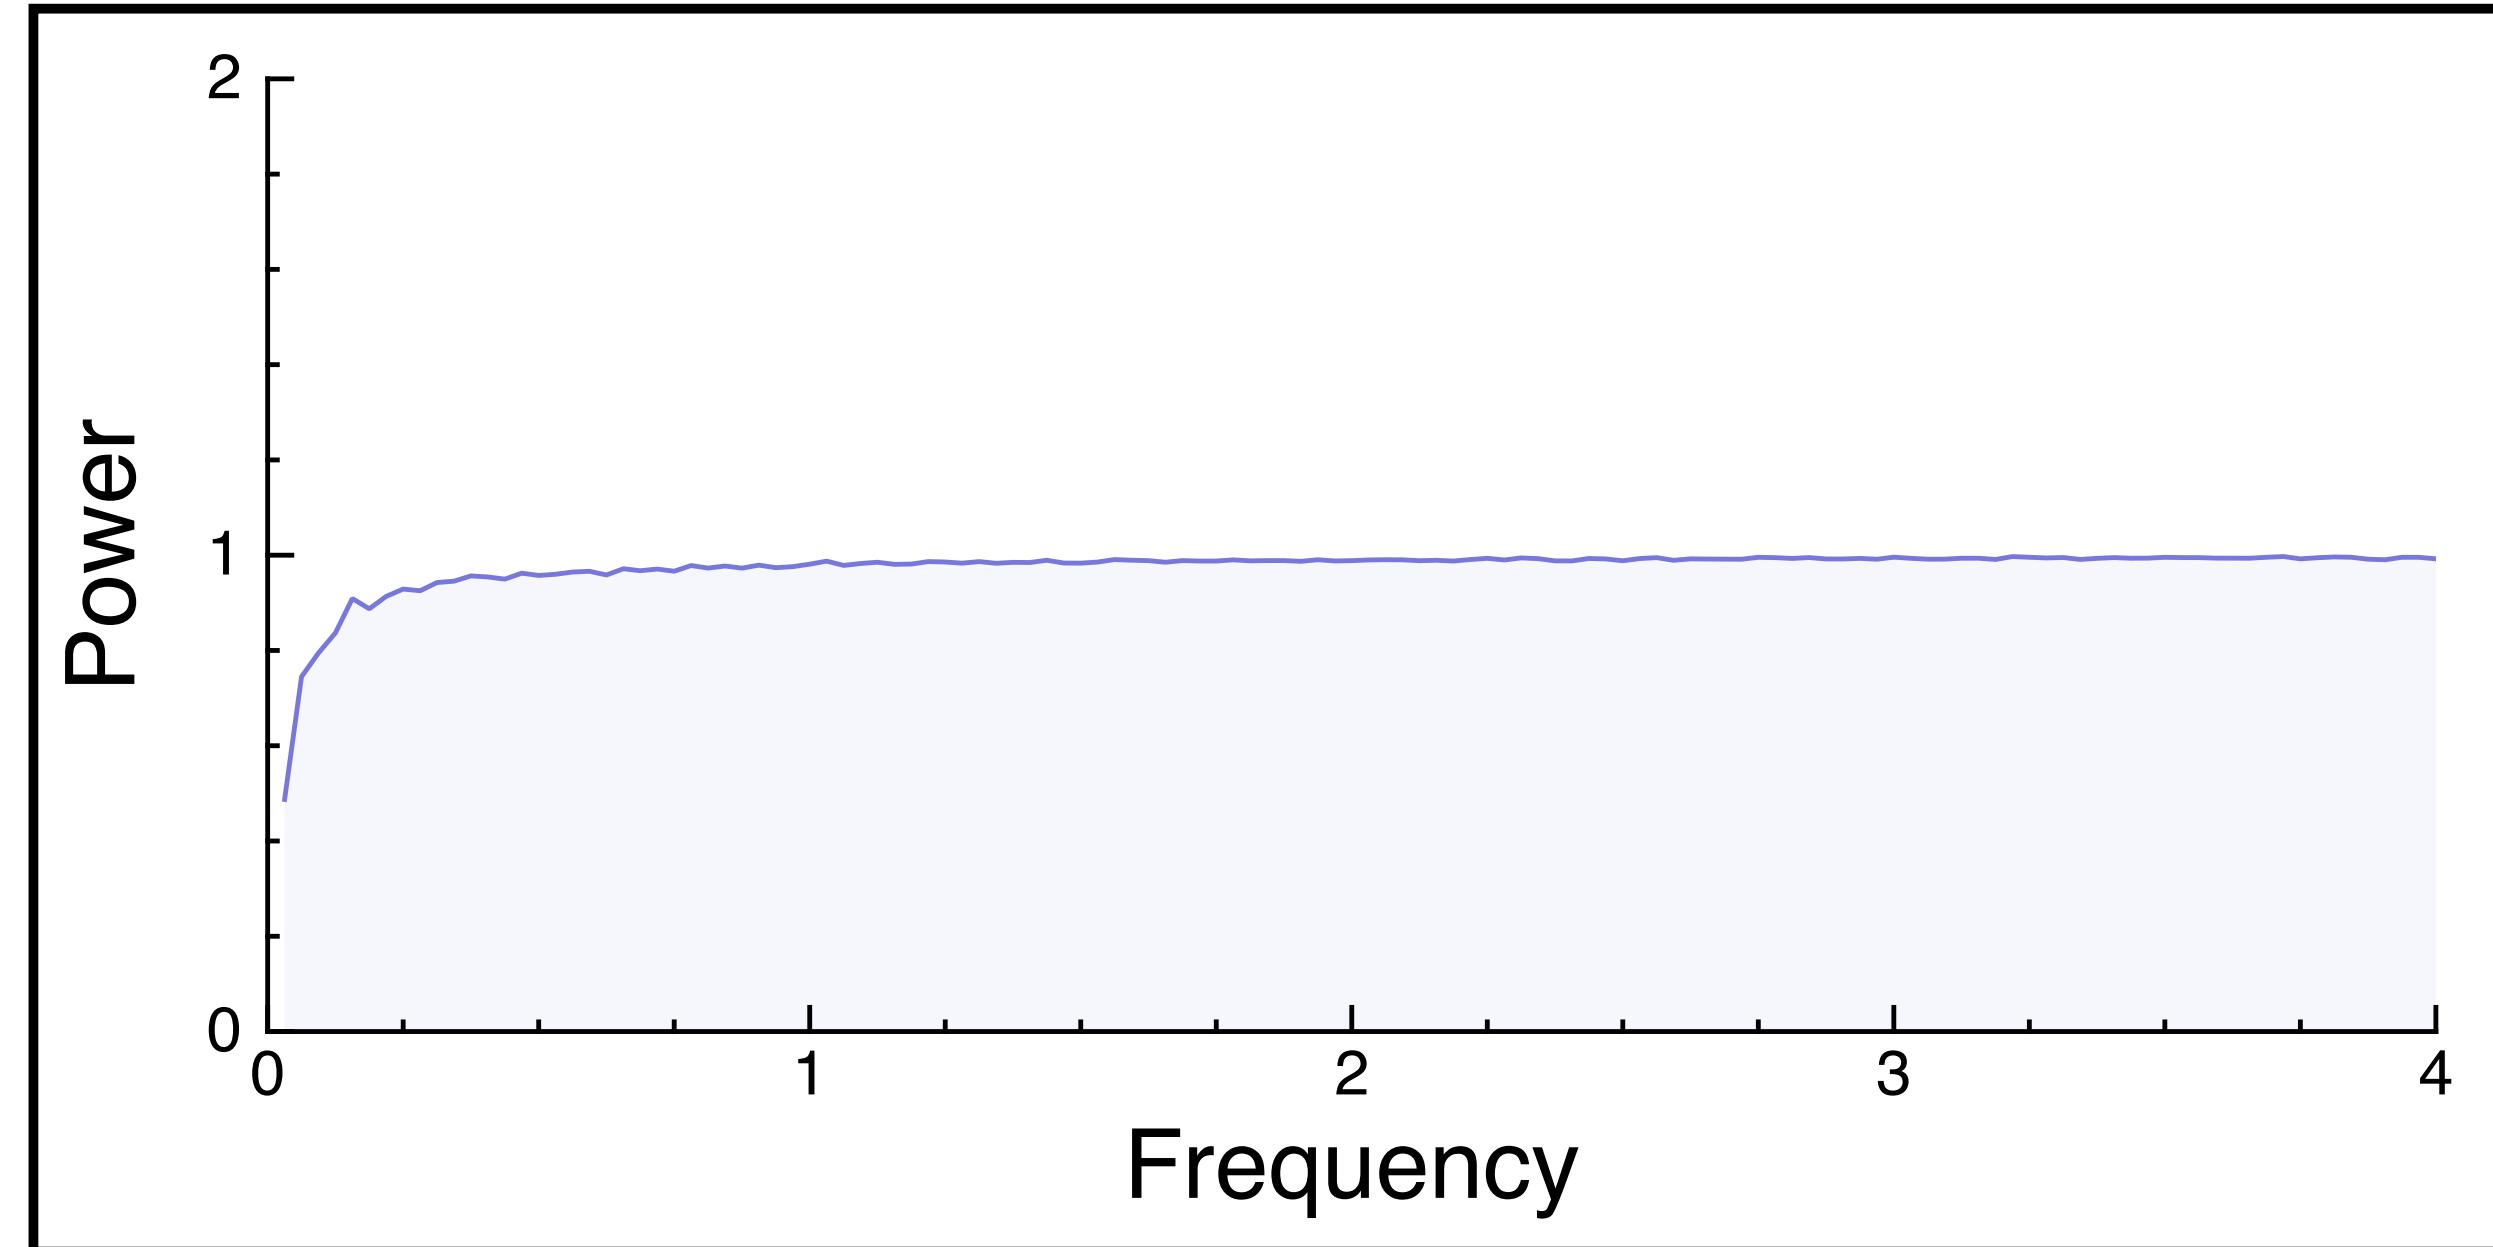
Samples



Power spectrum

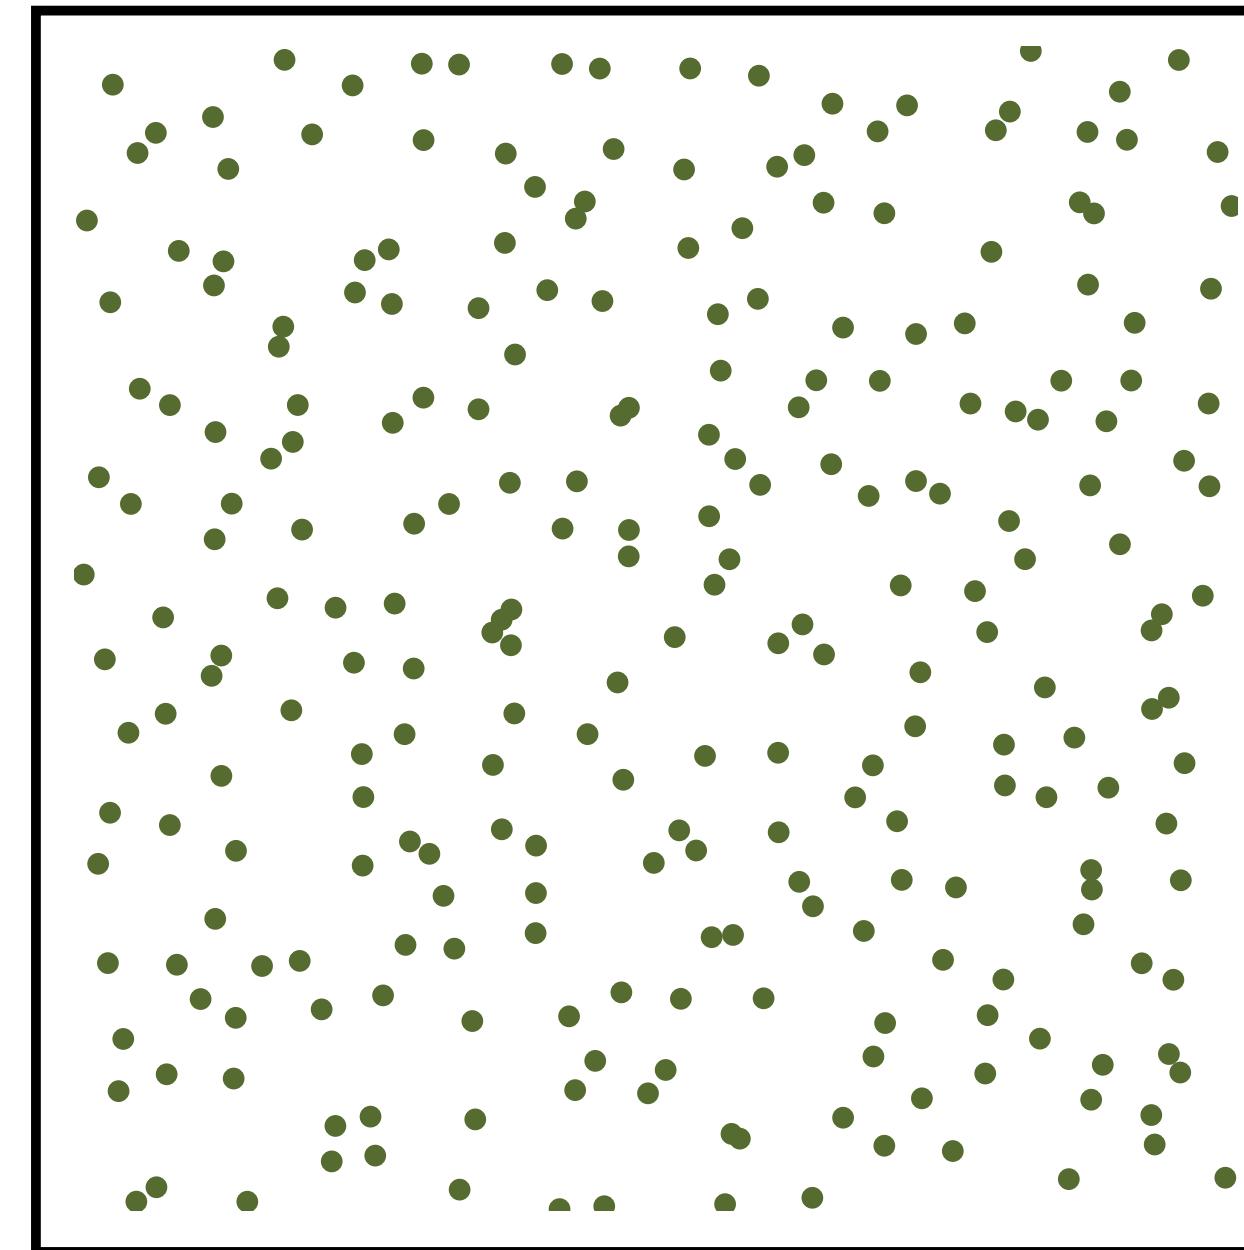


Radial mean

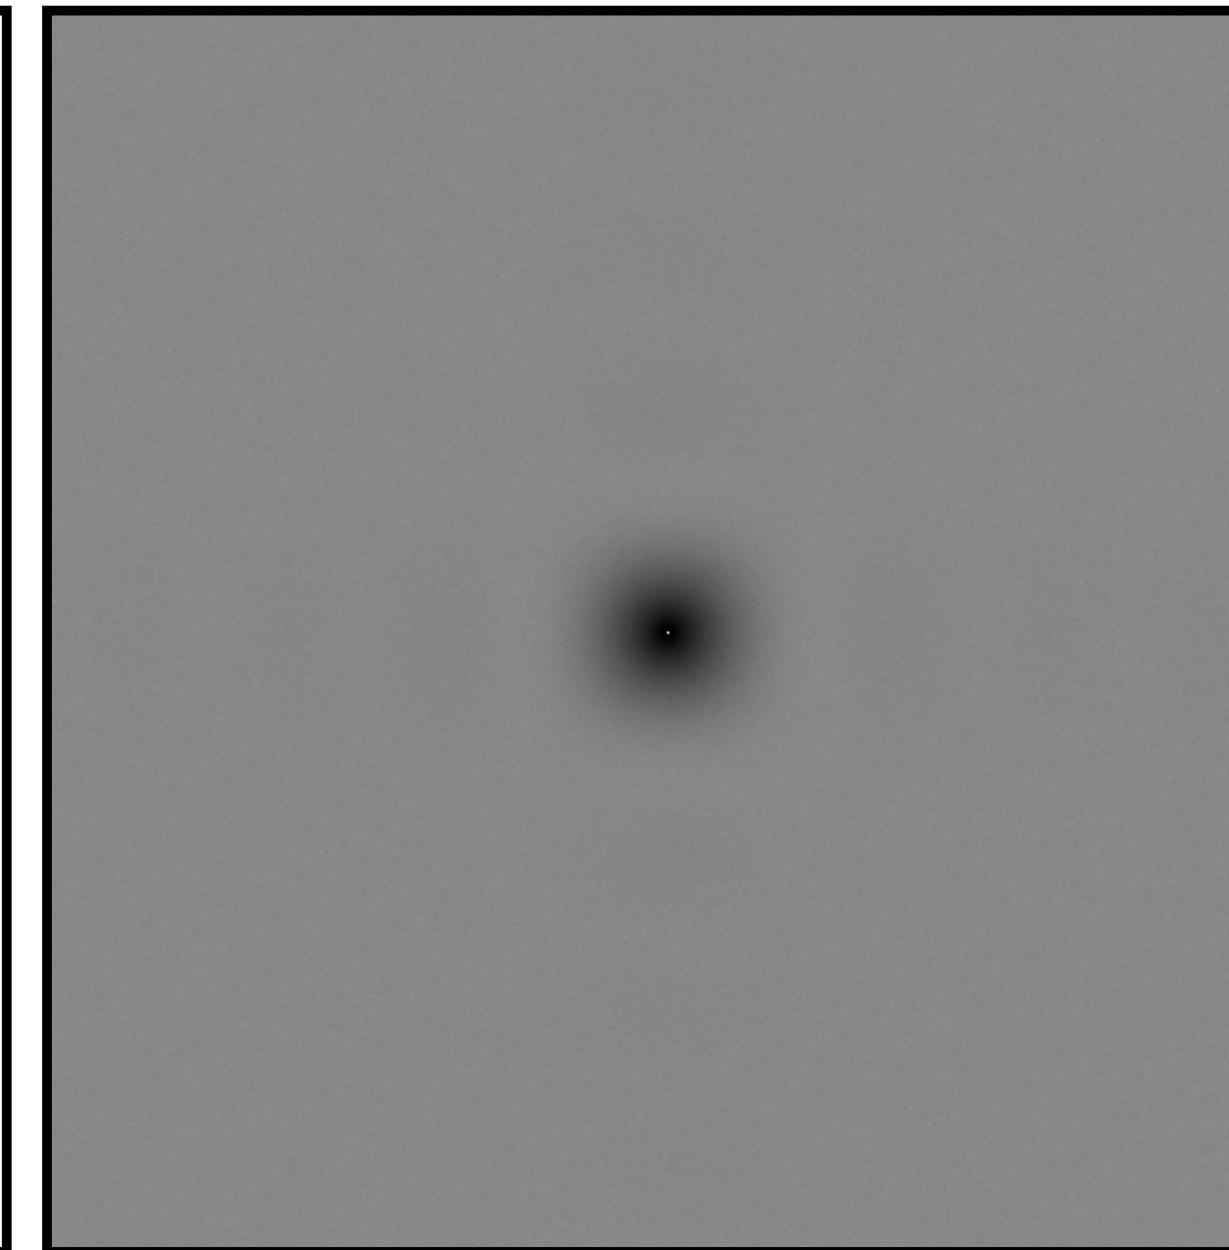


Jittered Sampling

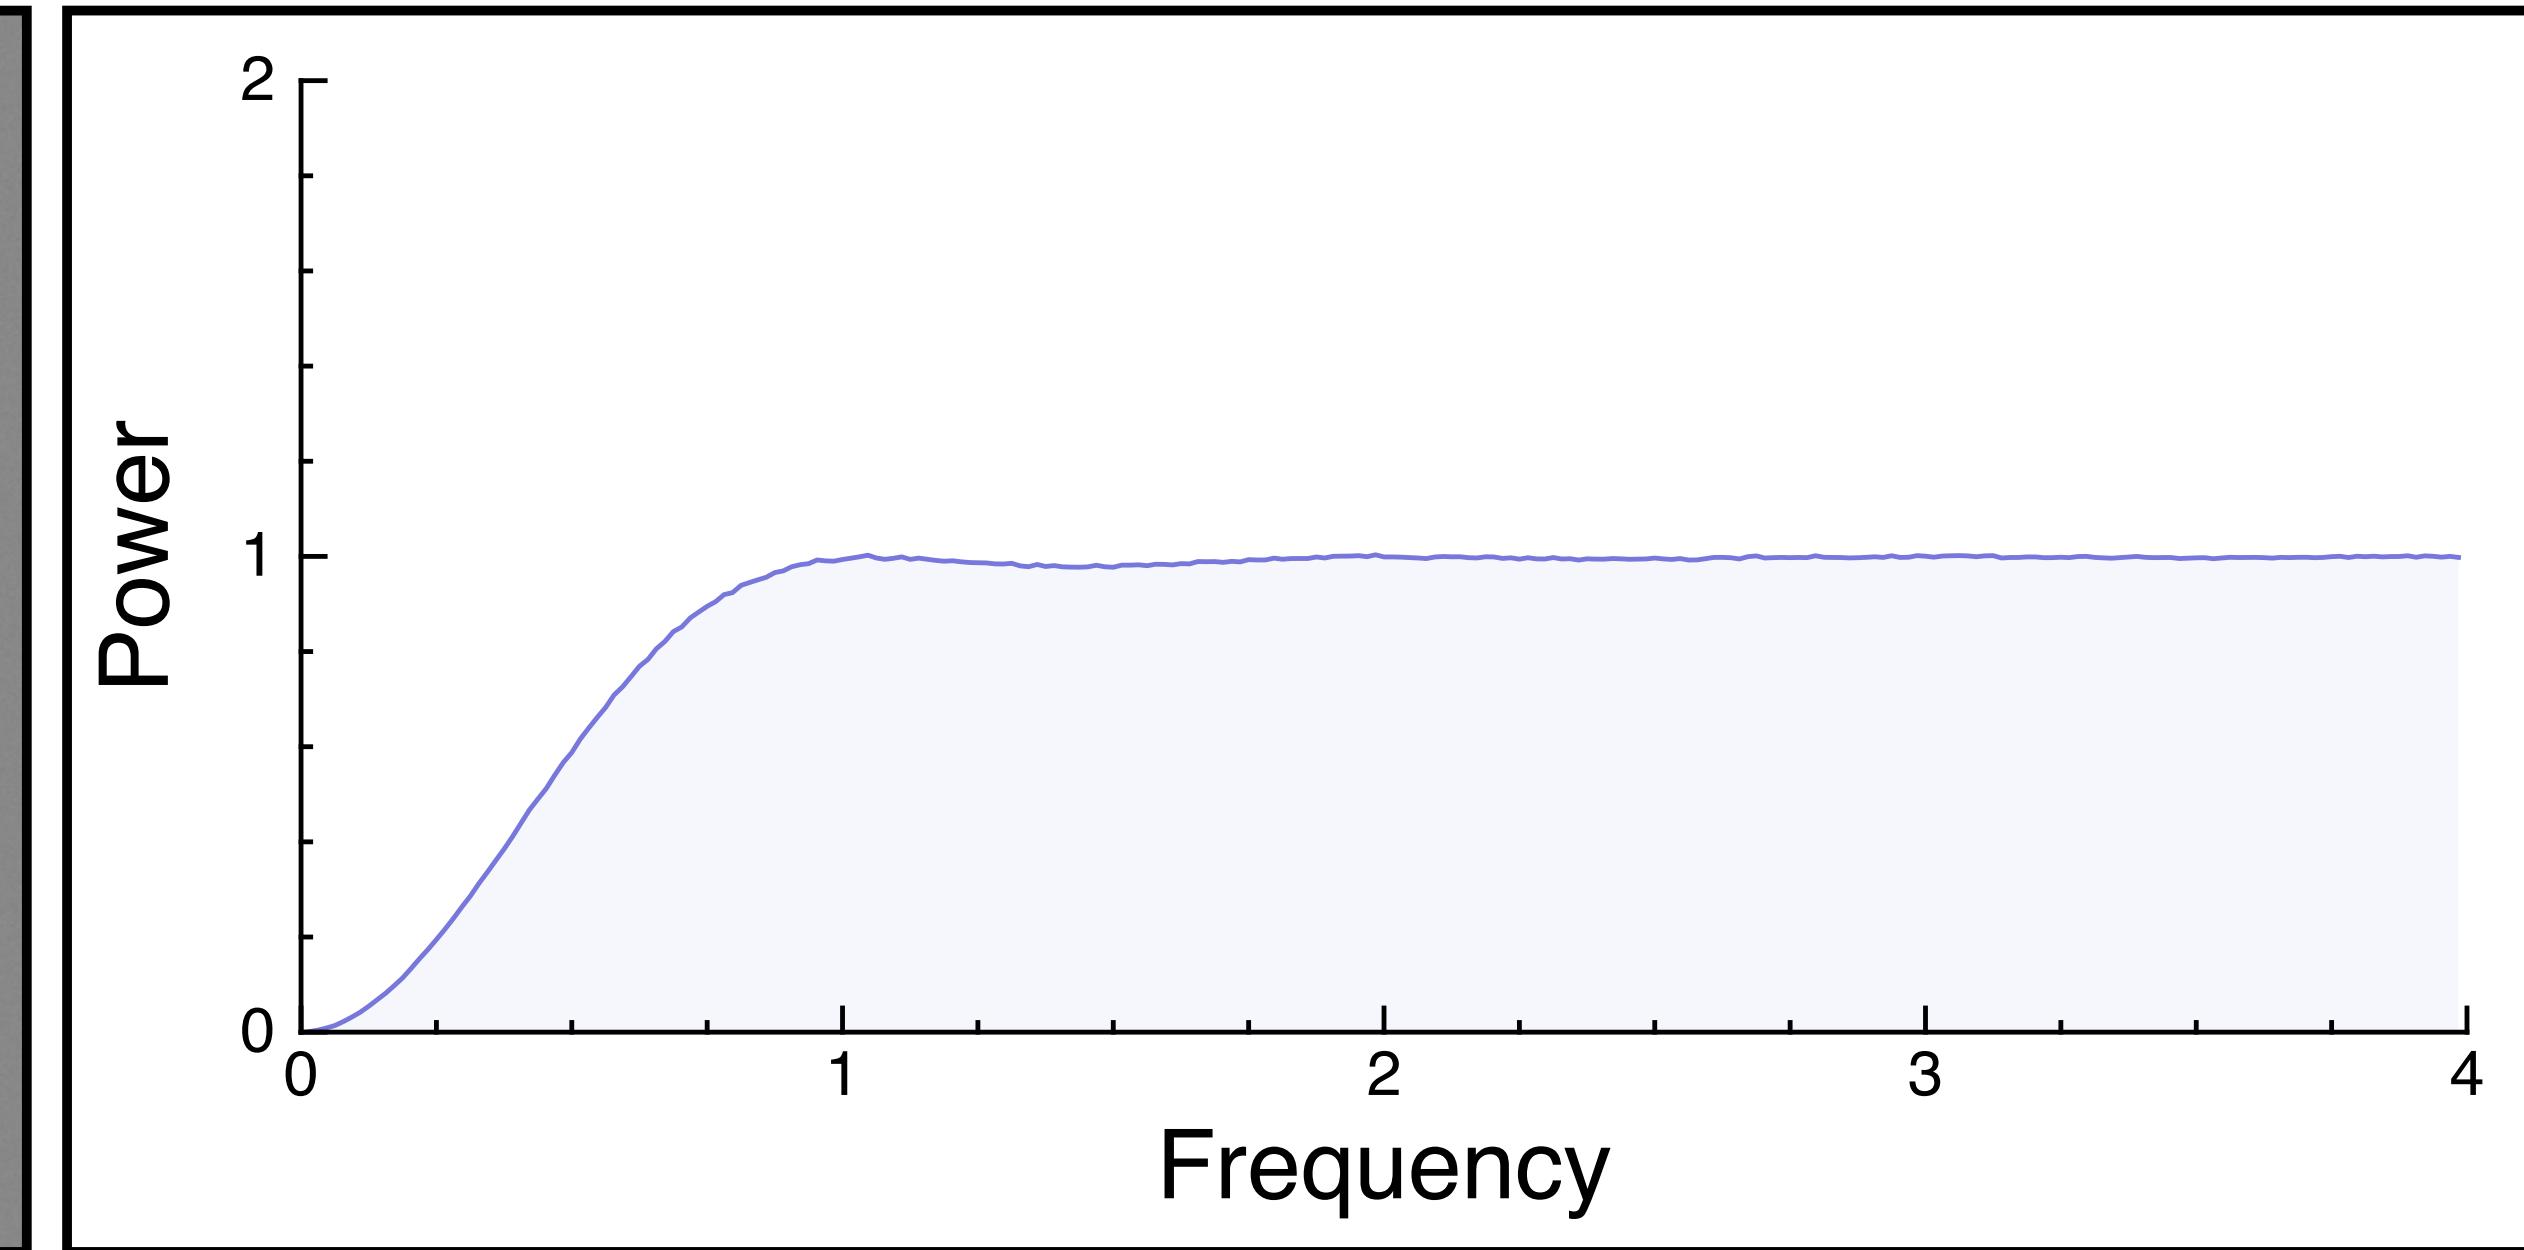
Samples



Power spectrum



Radial mean



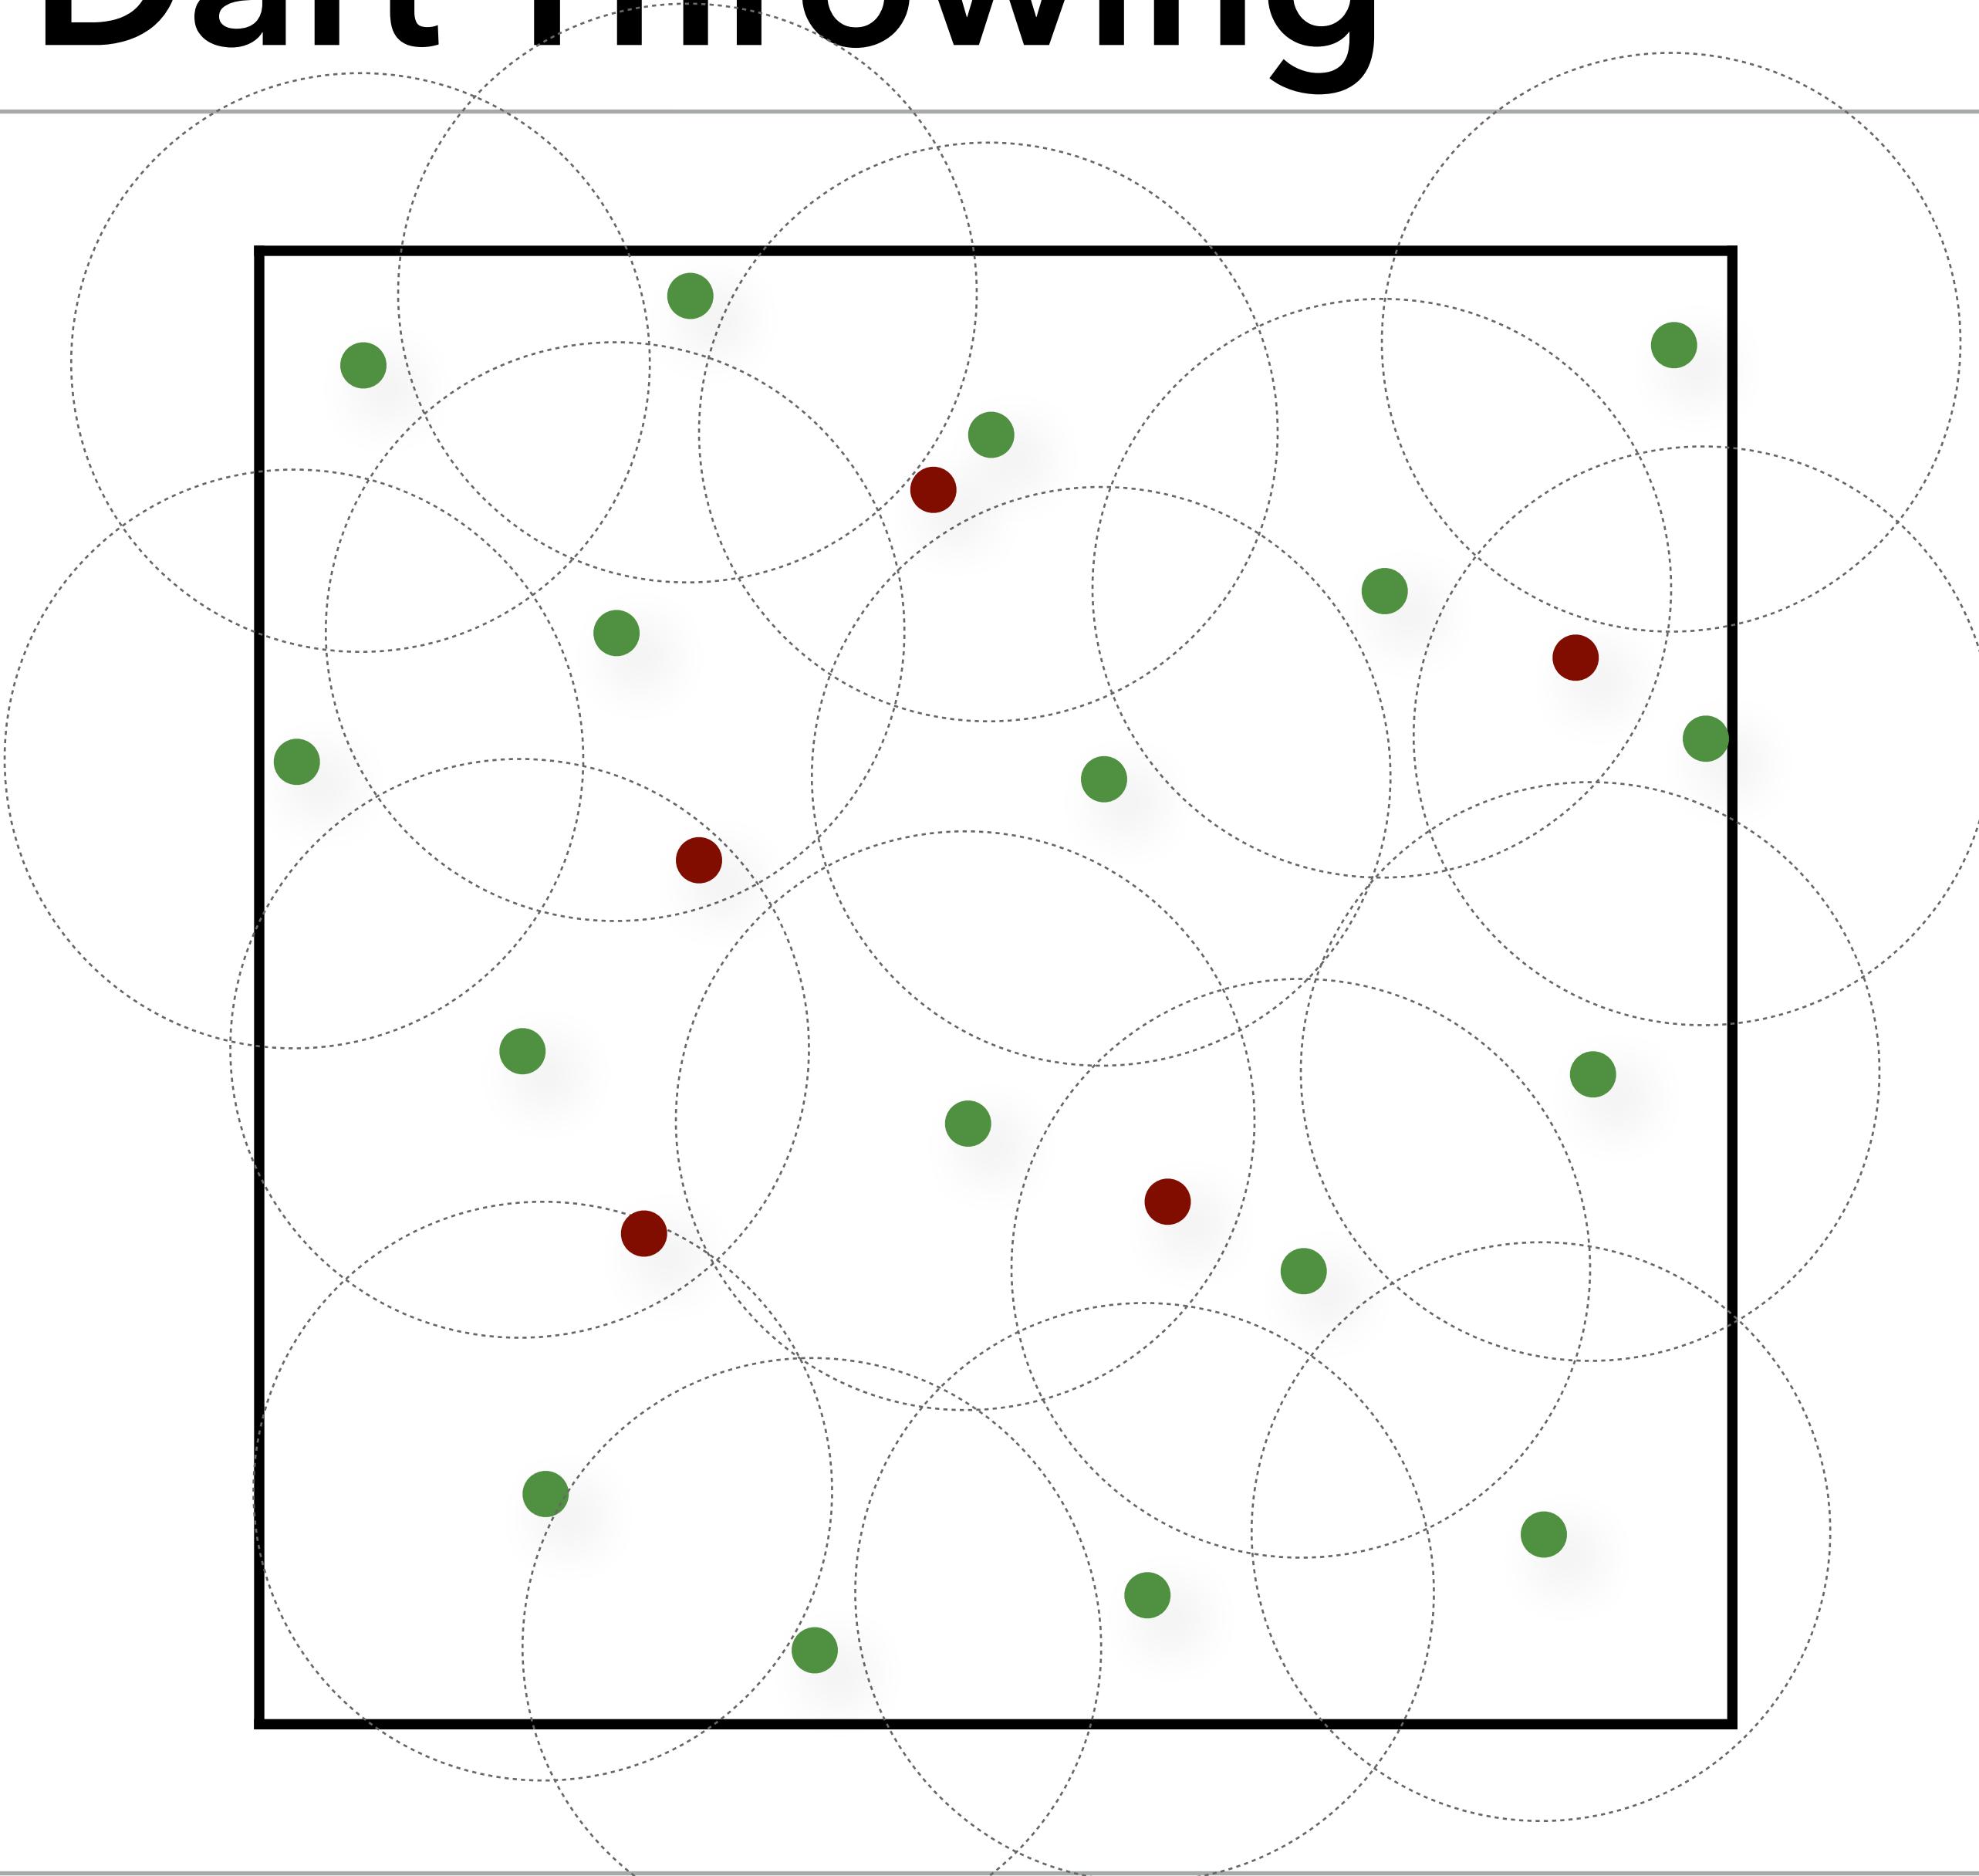
Poisson-Disk/Blue-Noise Sampling

Enforce a minimum distance between points

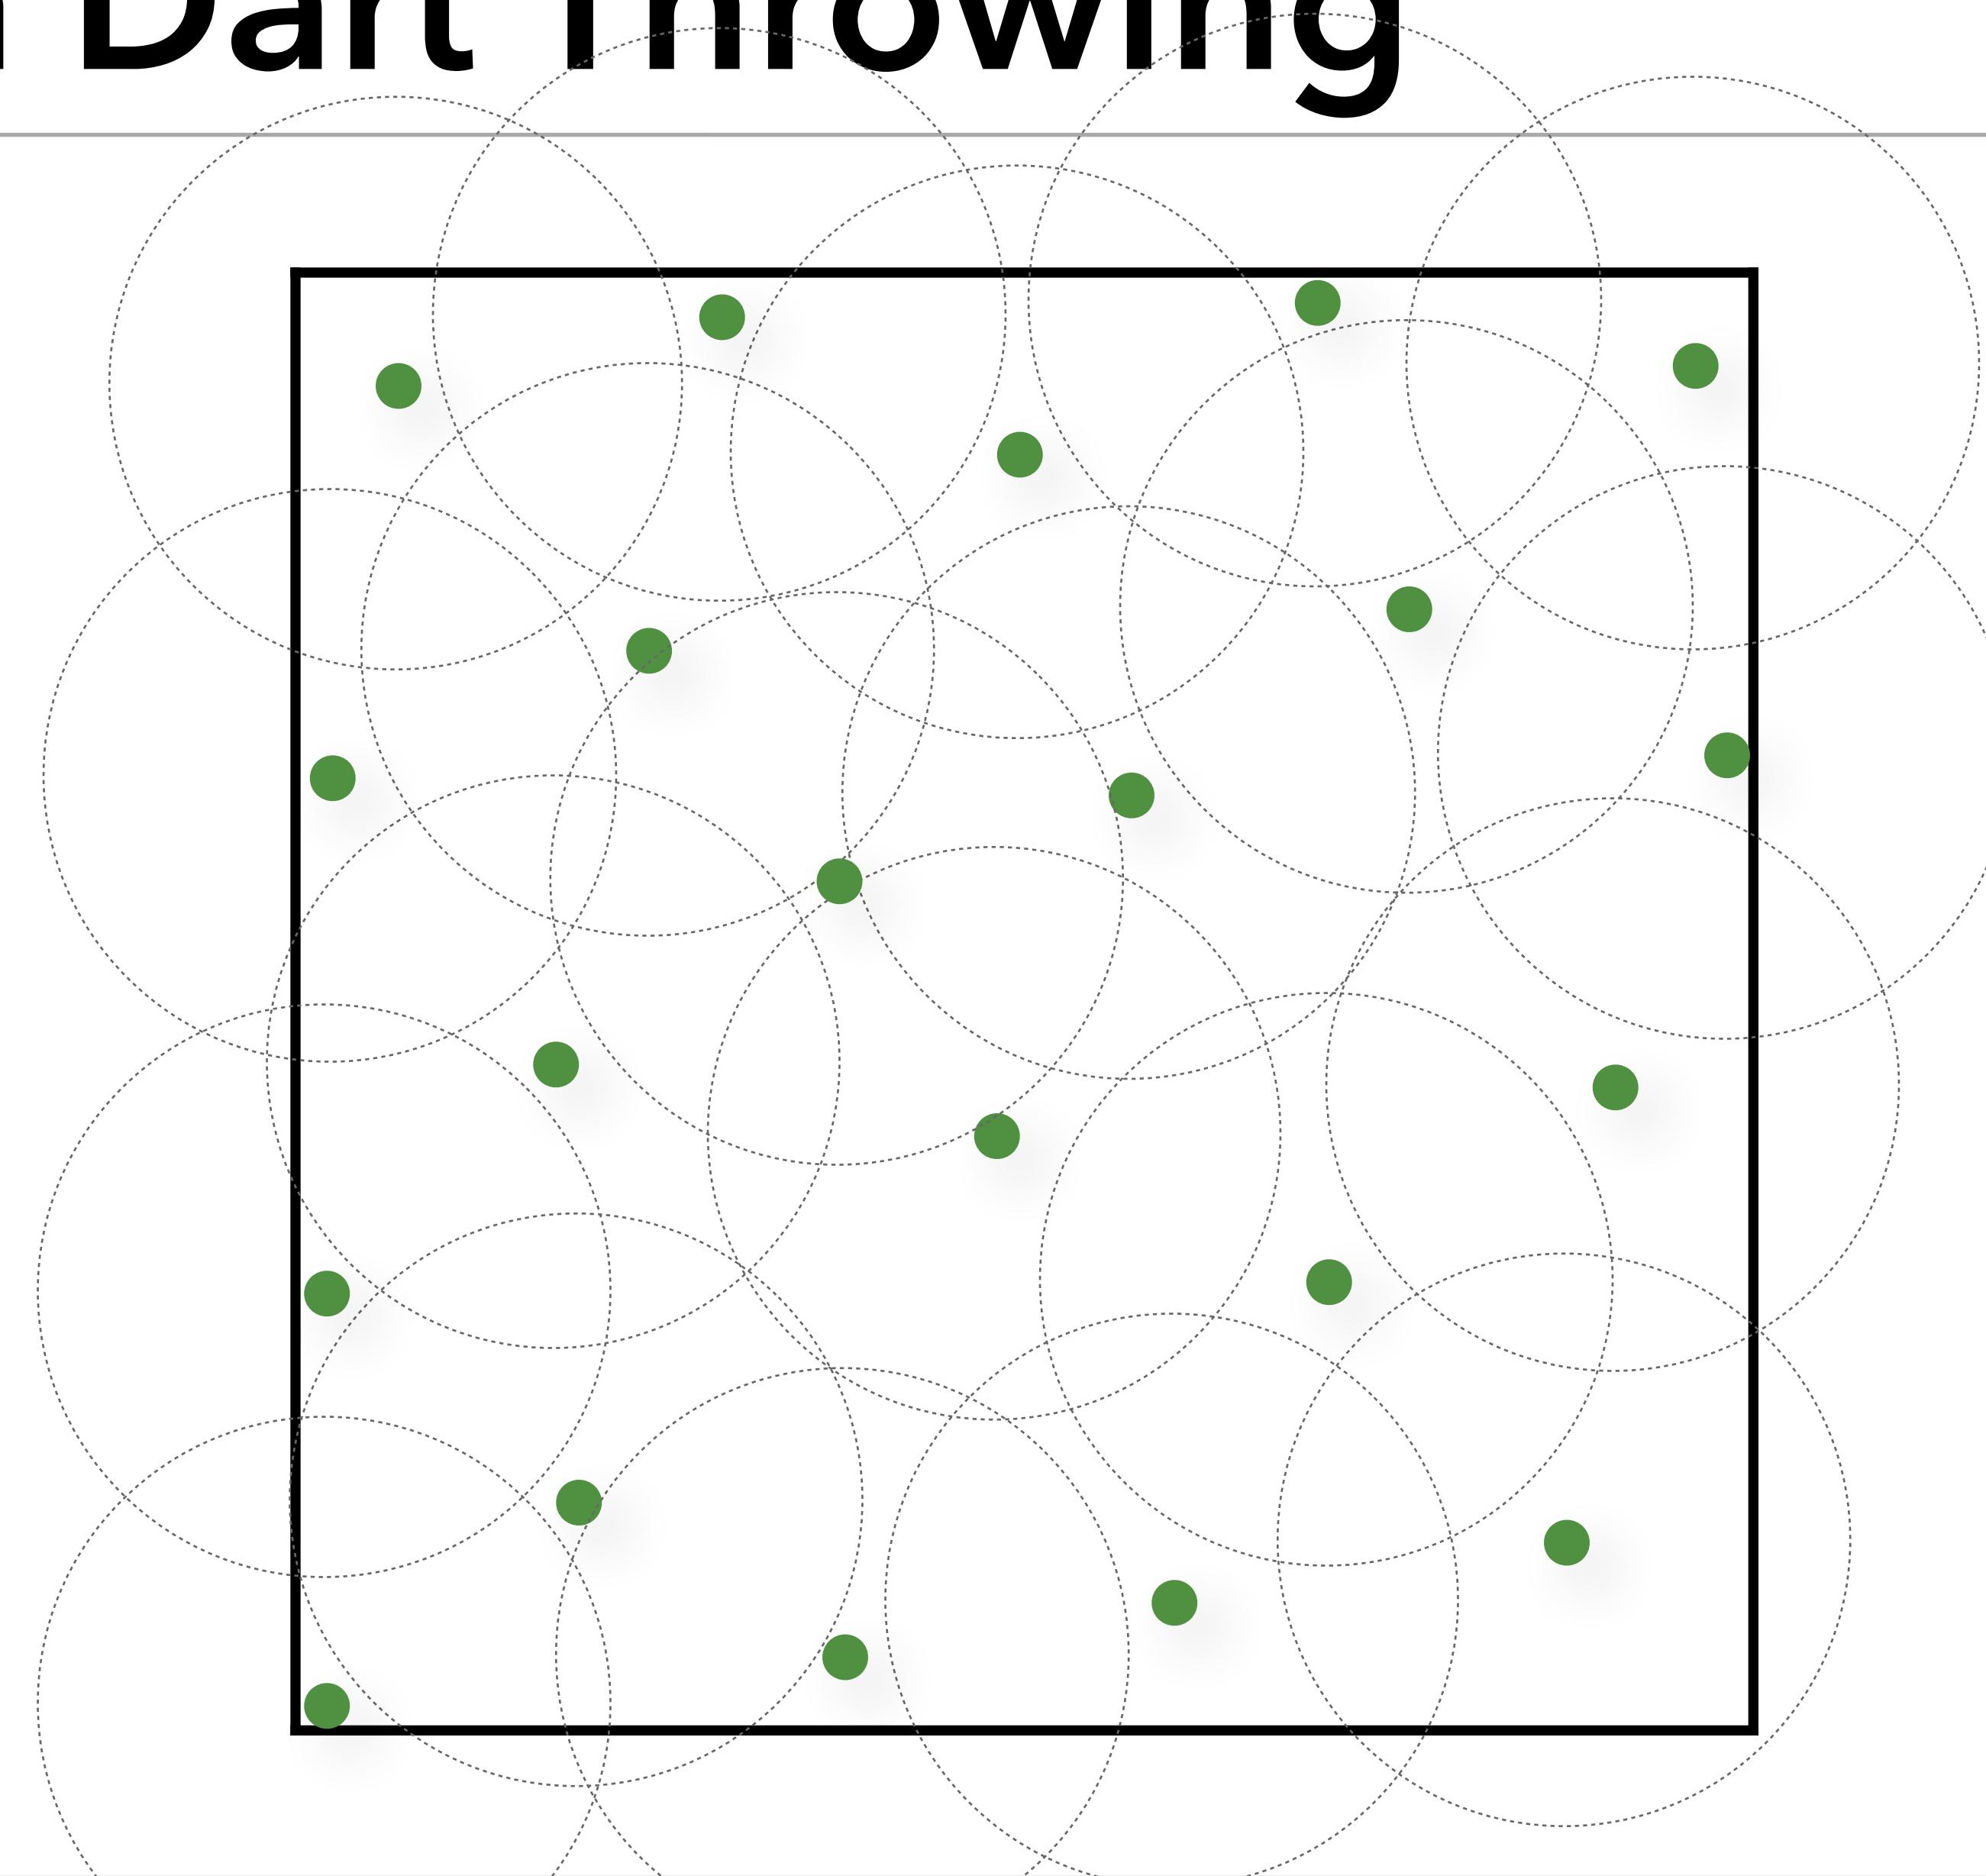
Poisson-Disk Sampling:

- Mark A. Z. Dippé and Erling Henry Wold. "Antialiasing through stochastic sampling." *ACM SIGGRAPH*, 1985.
- Robert L. Cook. "Stochastic sampling in computer graphics." *ACM Transactions on Graphics*, 1986.
- Ares Lagae and Philip Dutré. "A comparison of methods for generating Poisson disk distributions." *Computer Graphics Forum*, 2008.

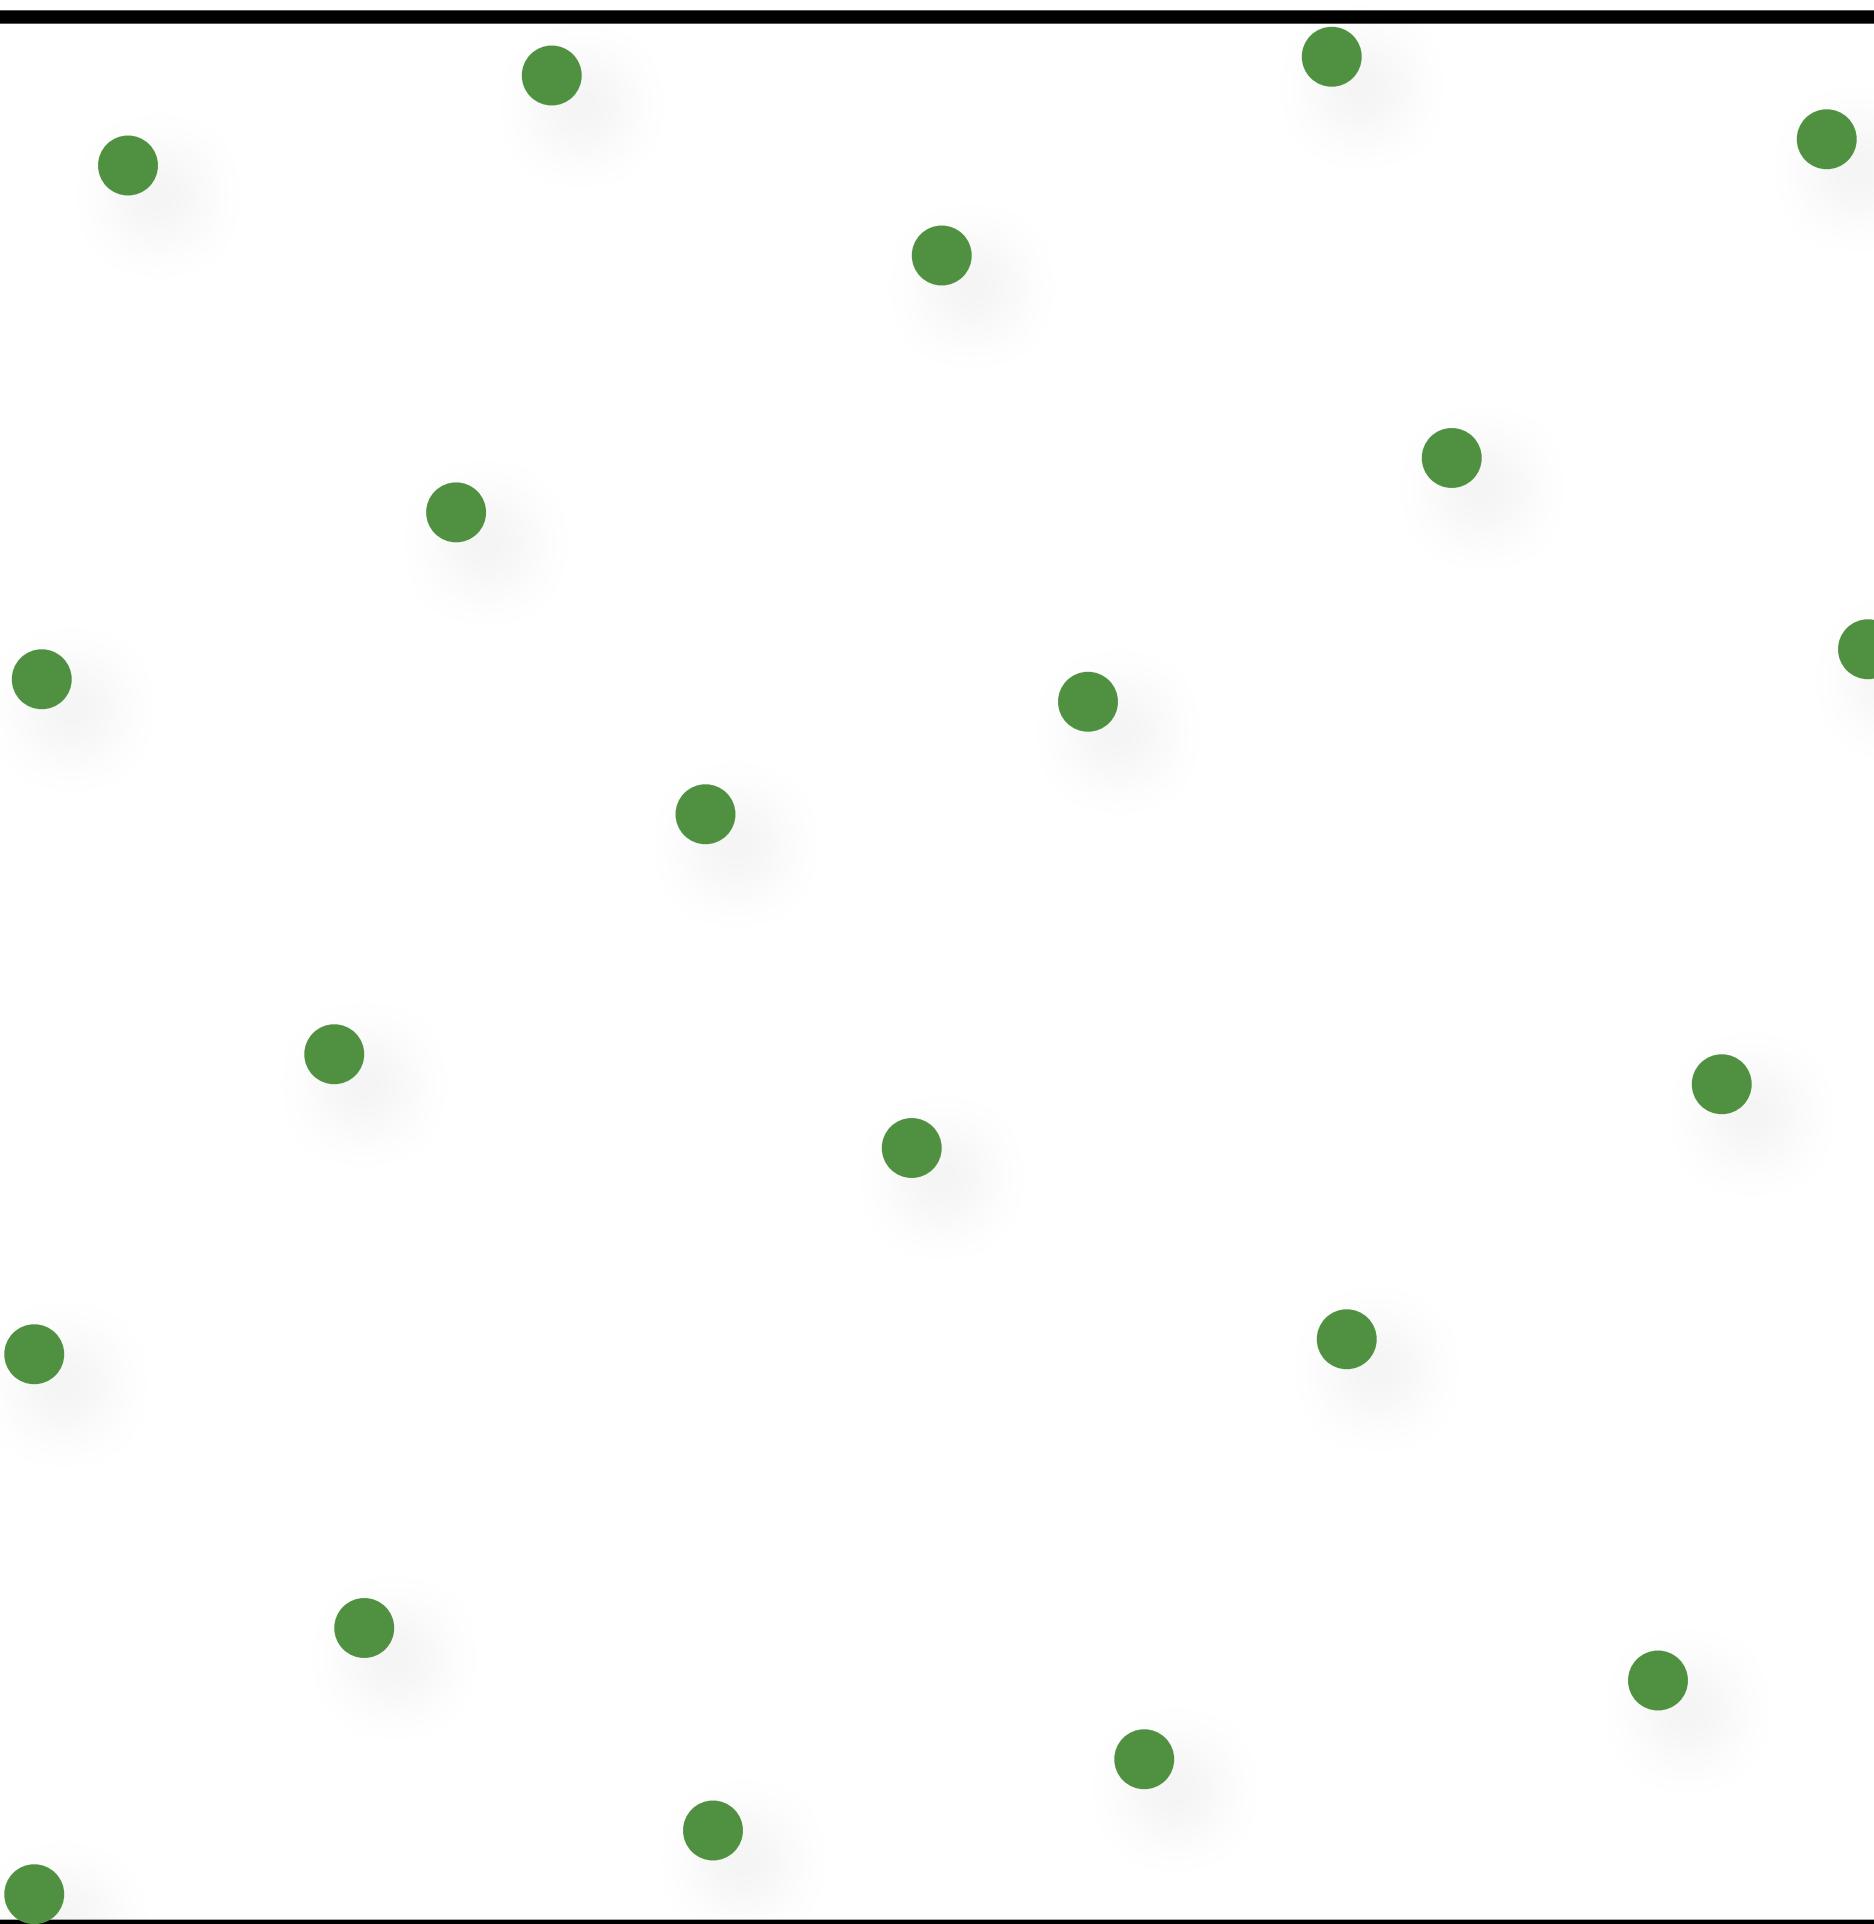
Random Dart Throwing



Random Dart Throwing

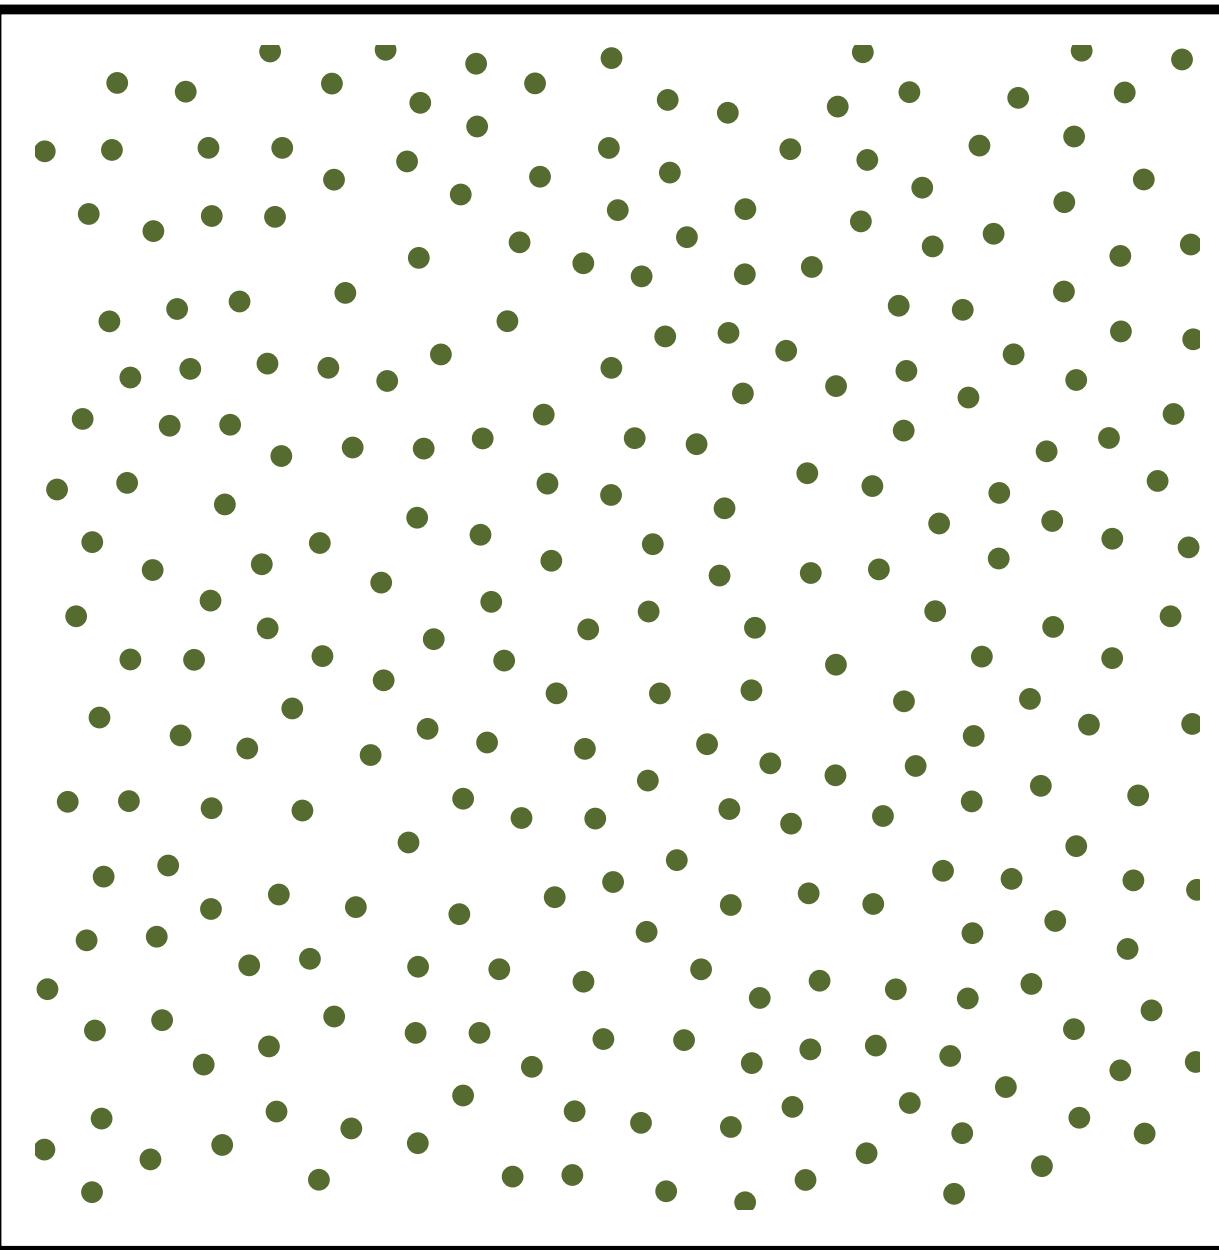


Random Dart Throwing

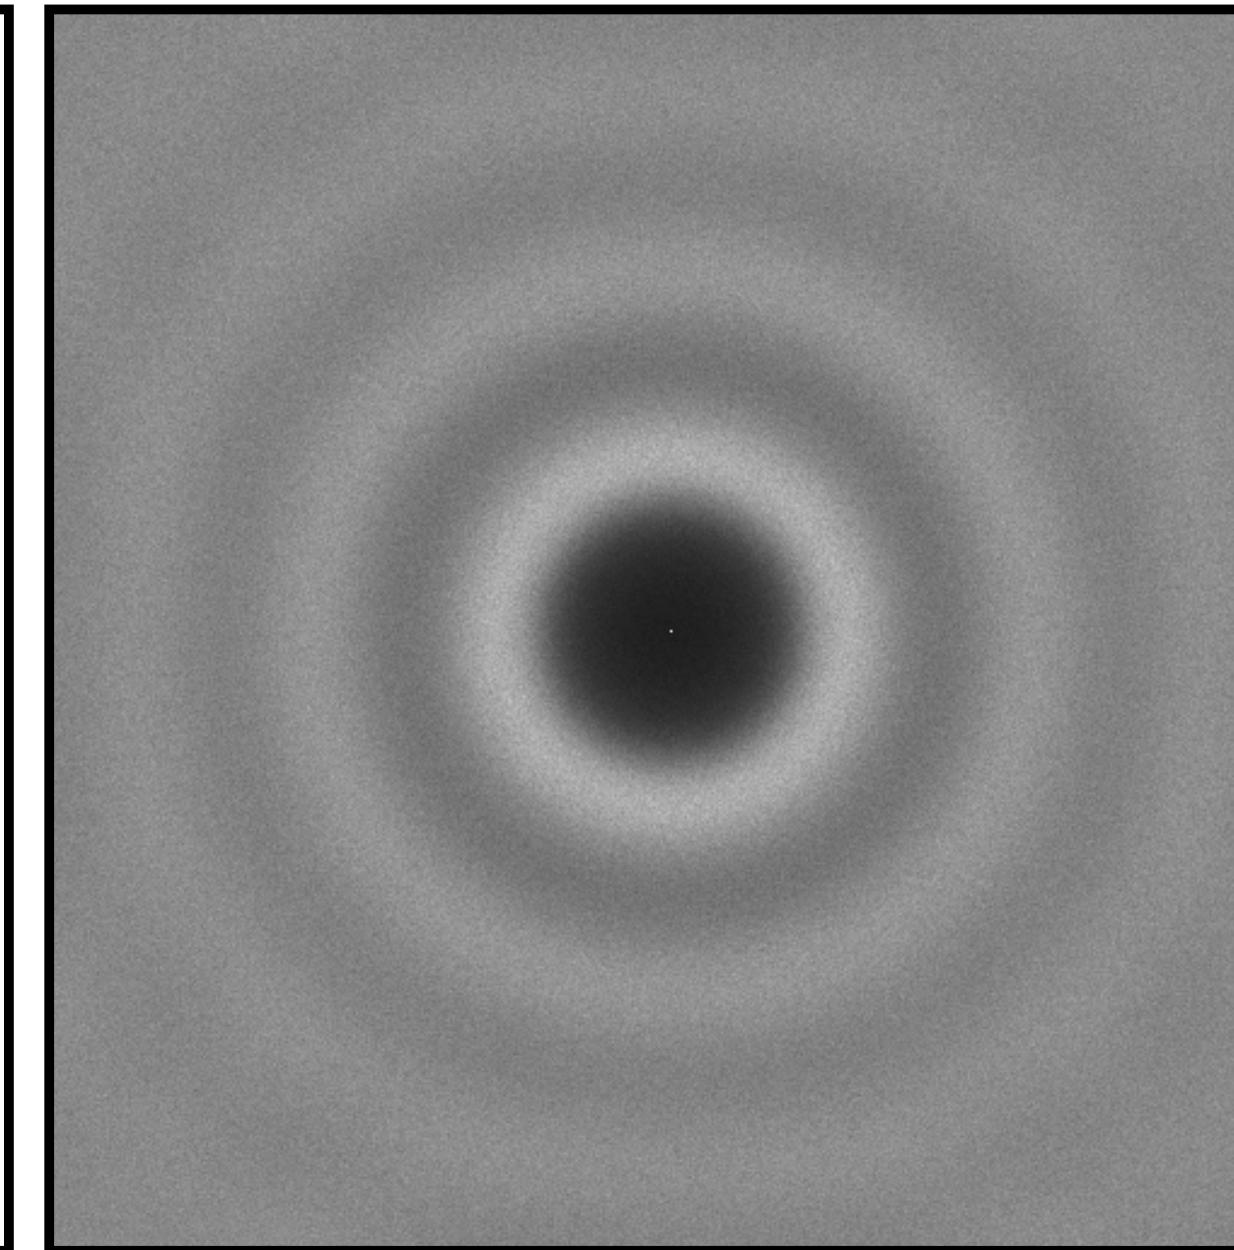


Poisson Disk Sampling

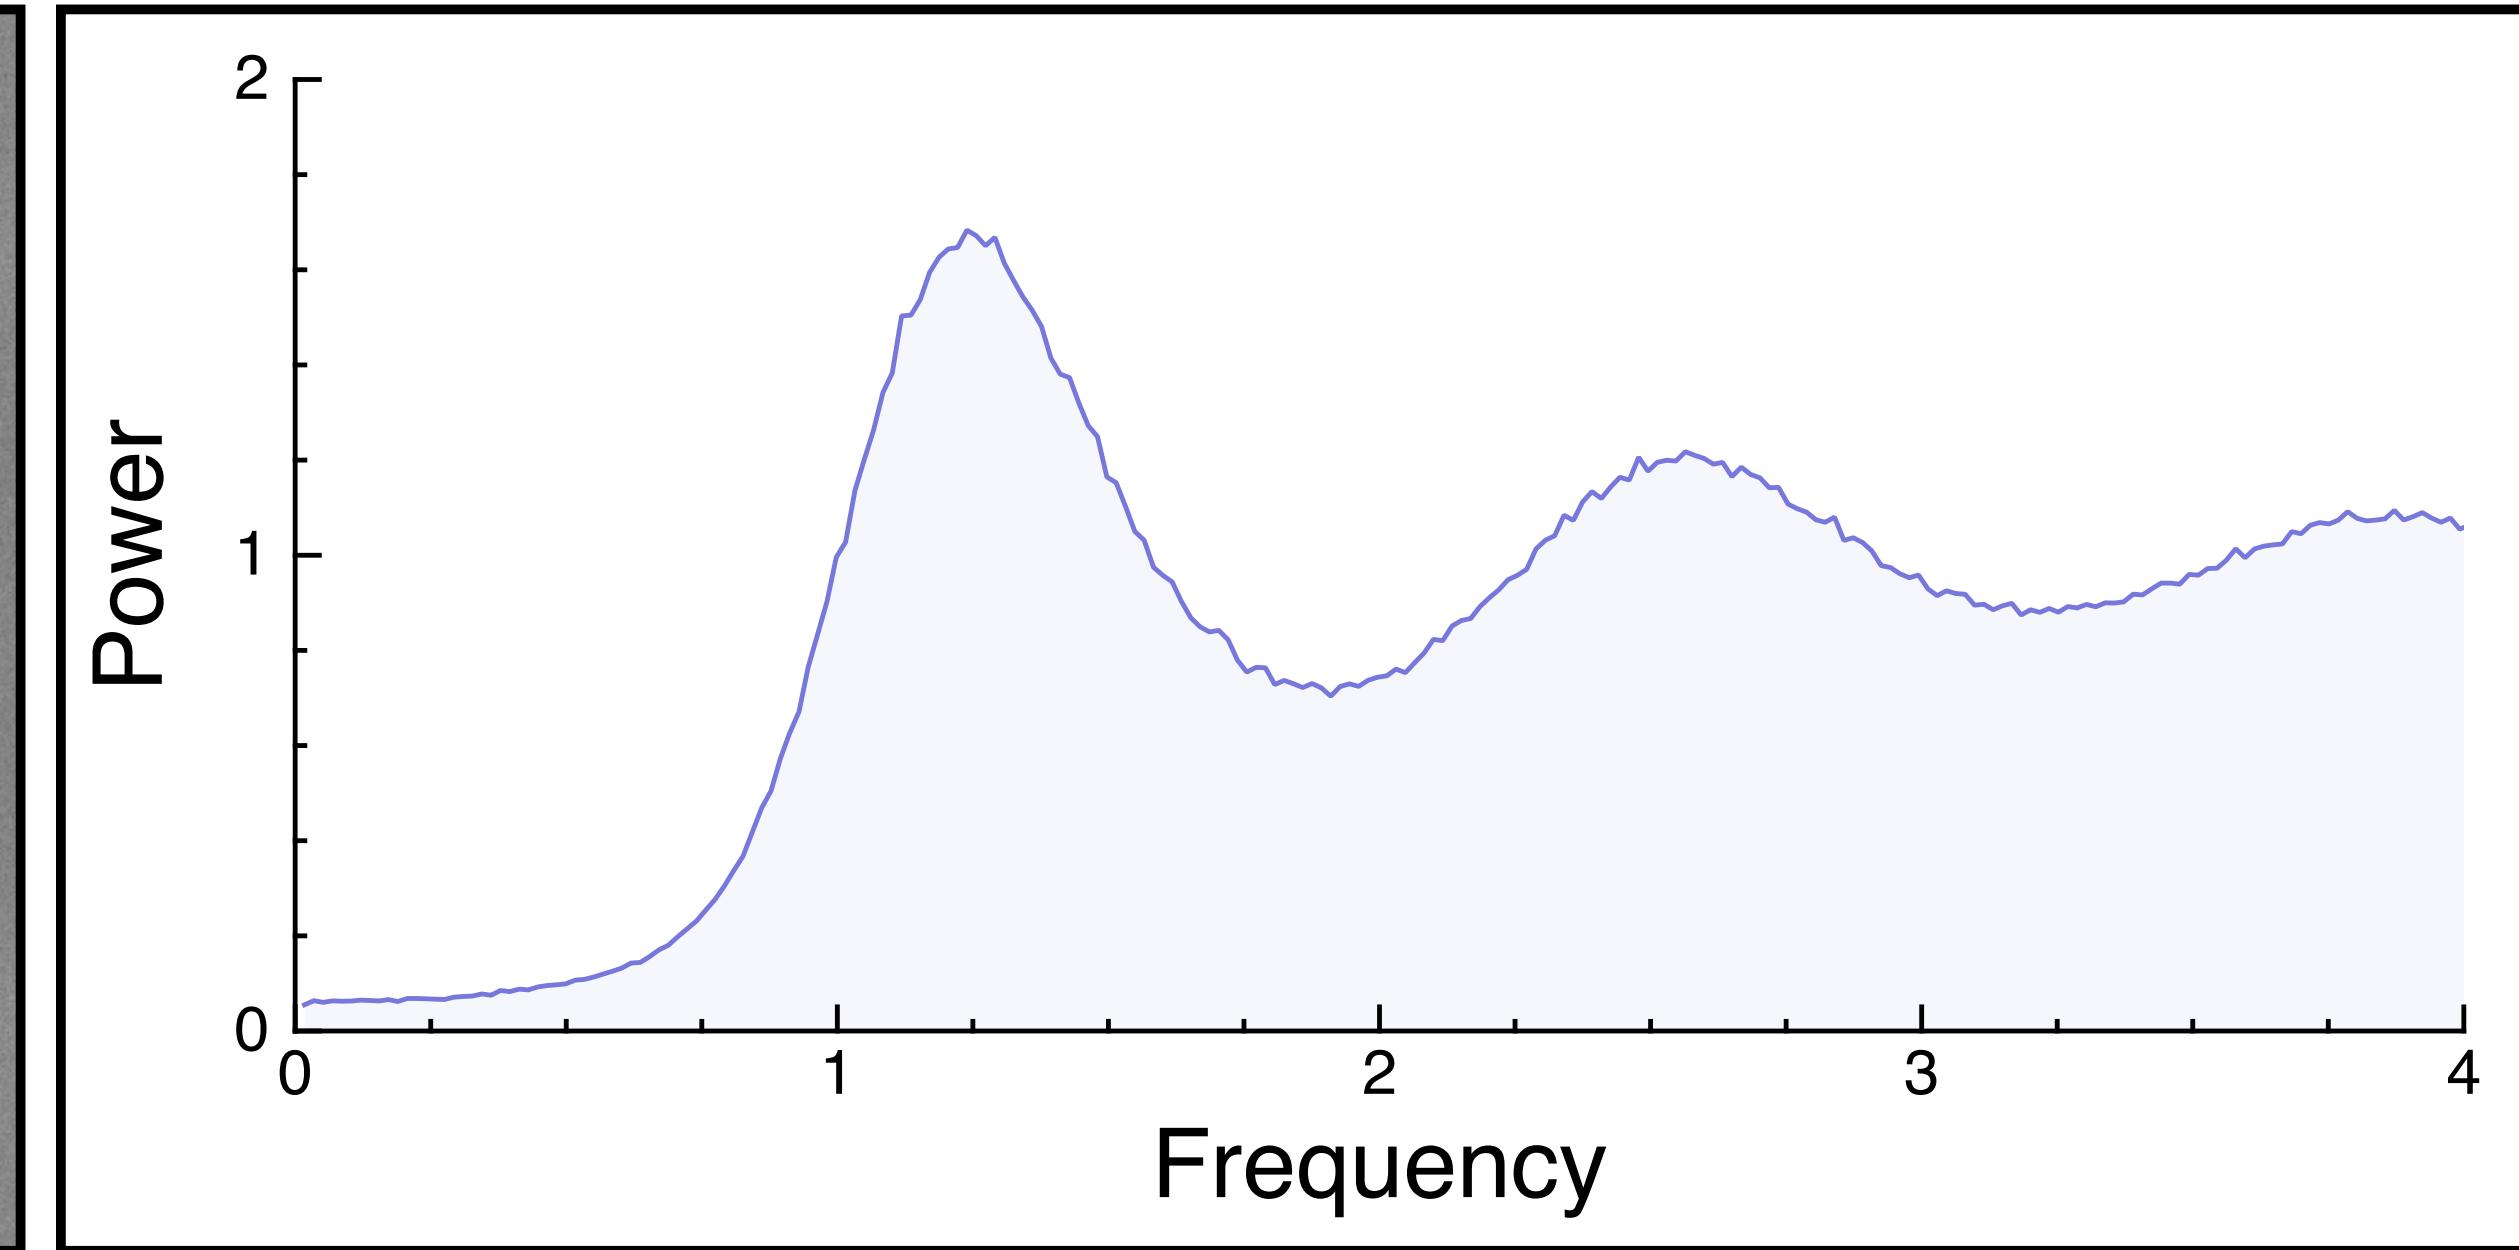
Samples



Power spectrum



Radial mean

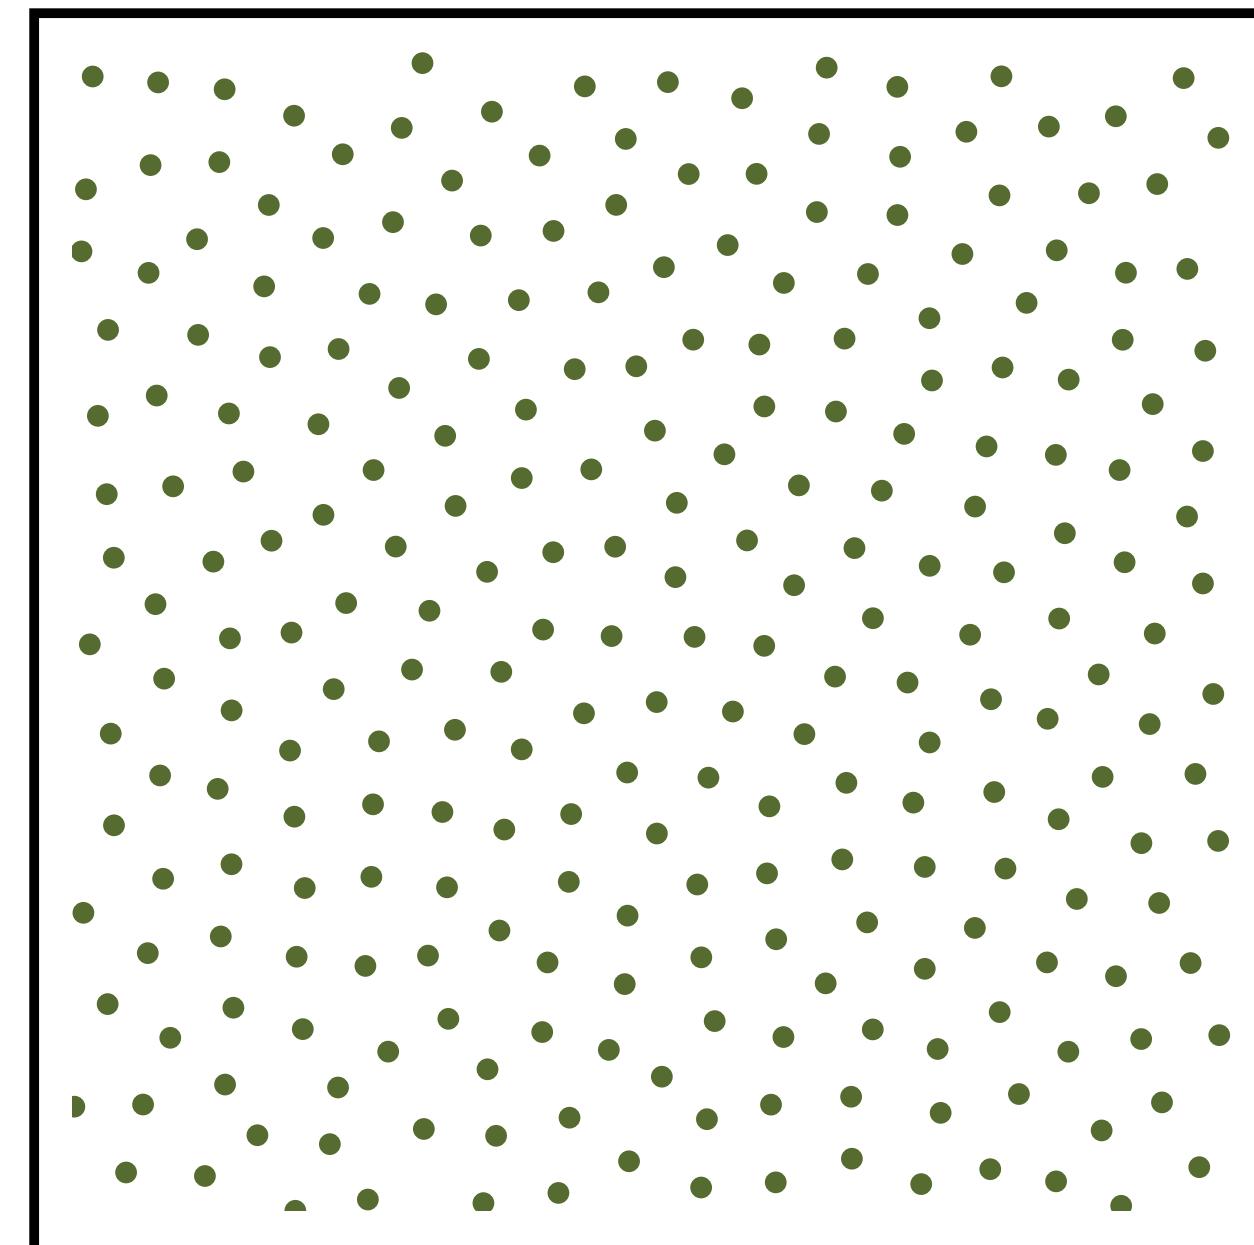


Blue-Noise Sampling (Relaxation-based)

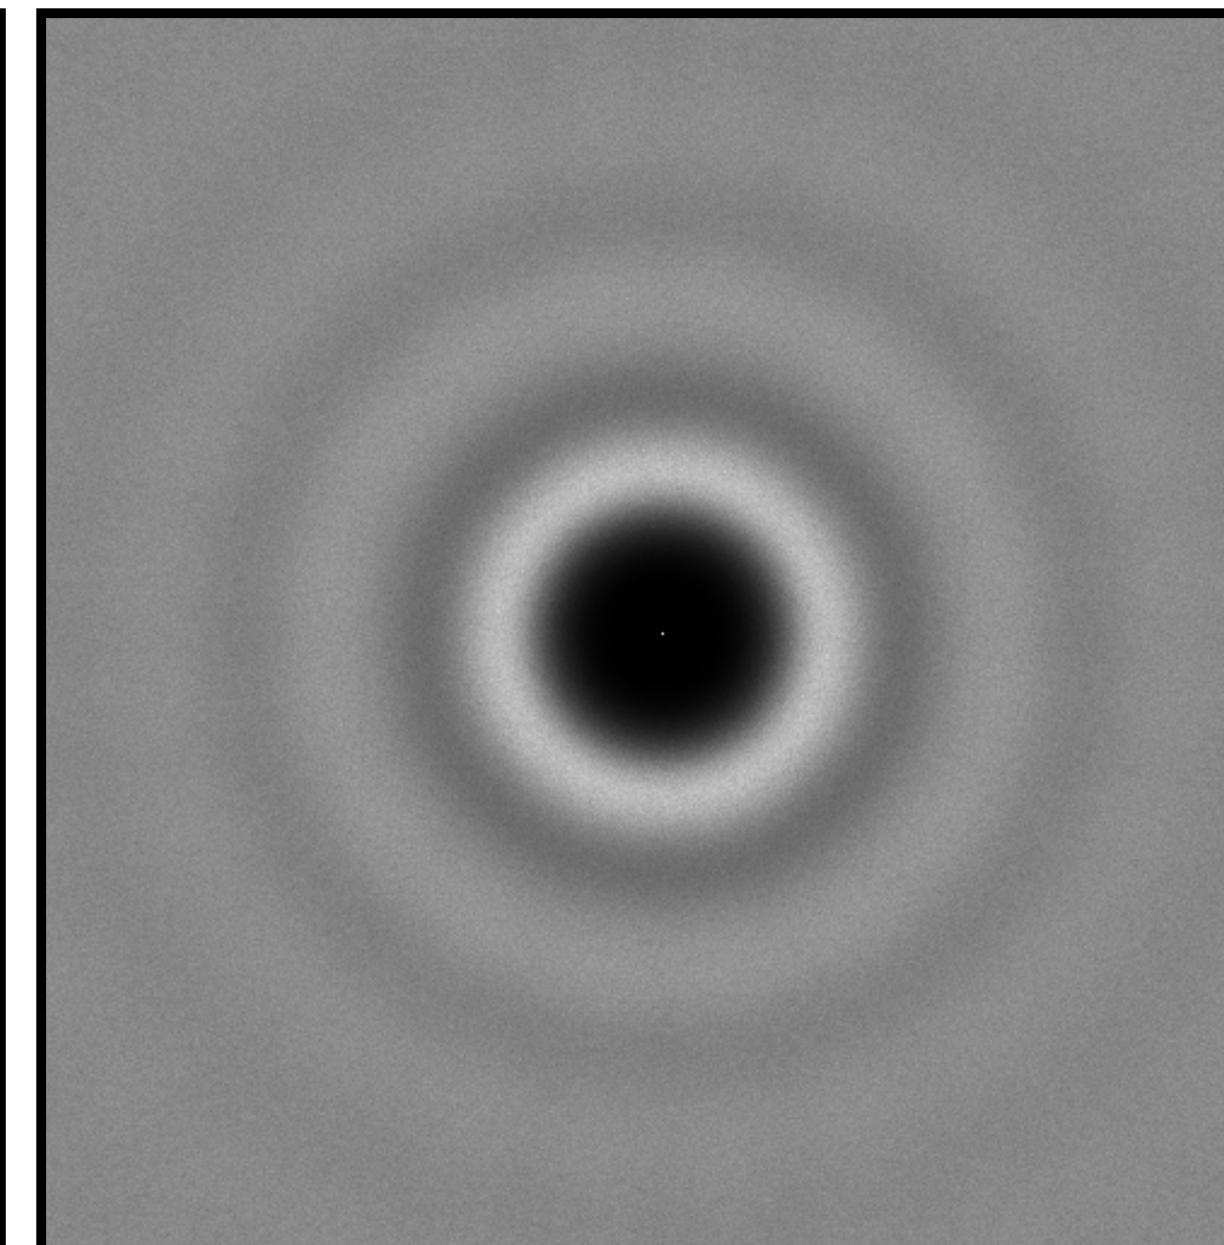
1. Initialize sample positions (e.g. random)
2. Use an iterative relaxation to move samples away from each other.

CCVT Sampling [Balzer et al. 2009]

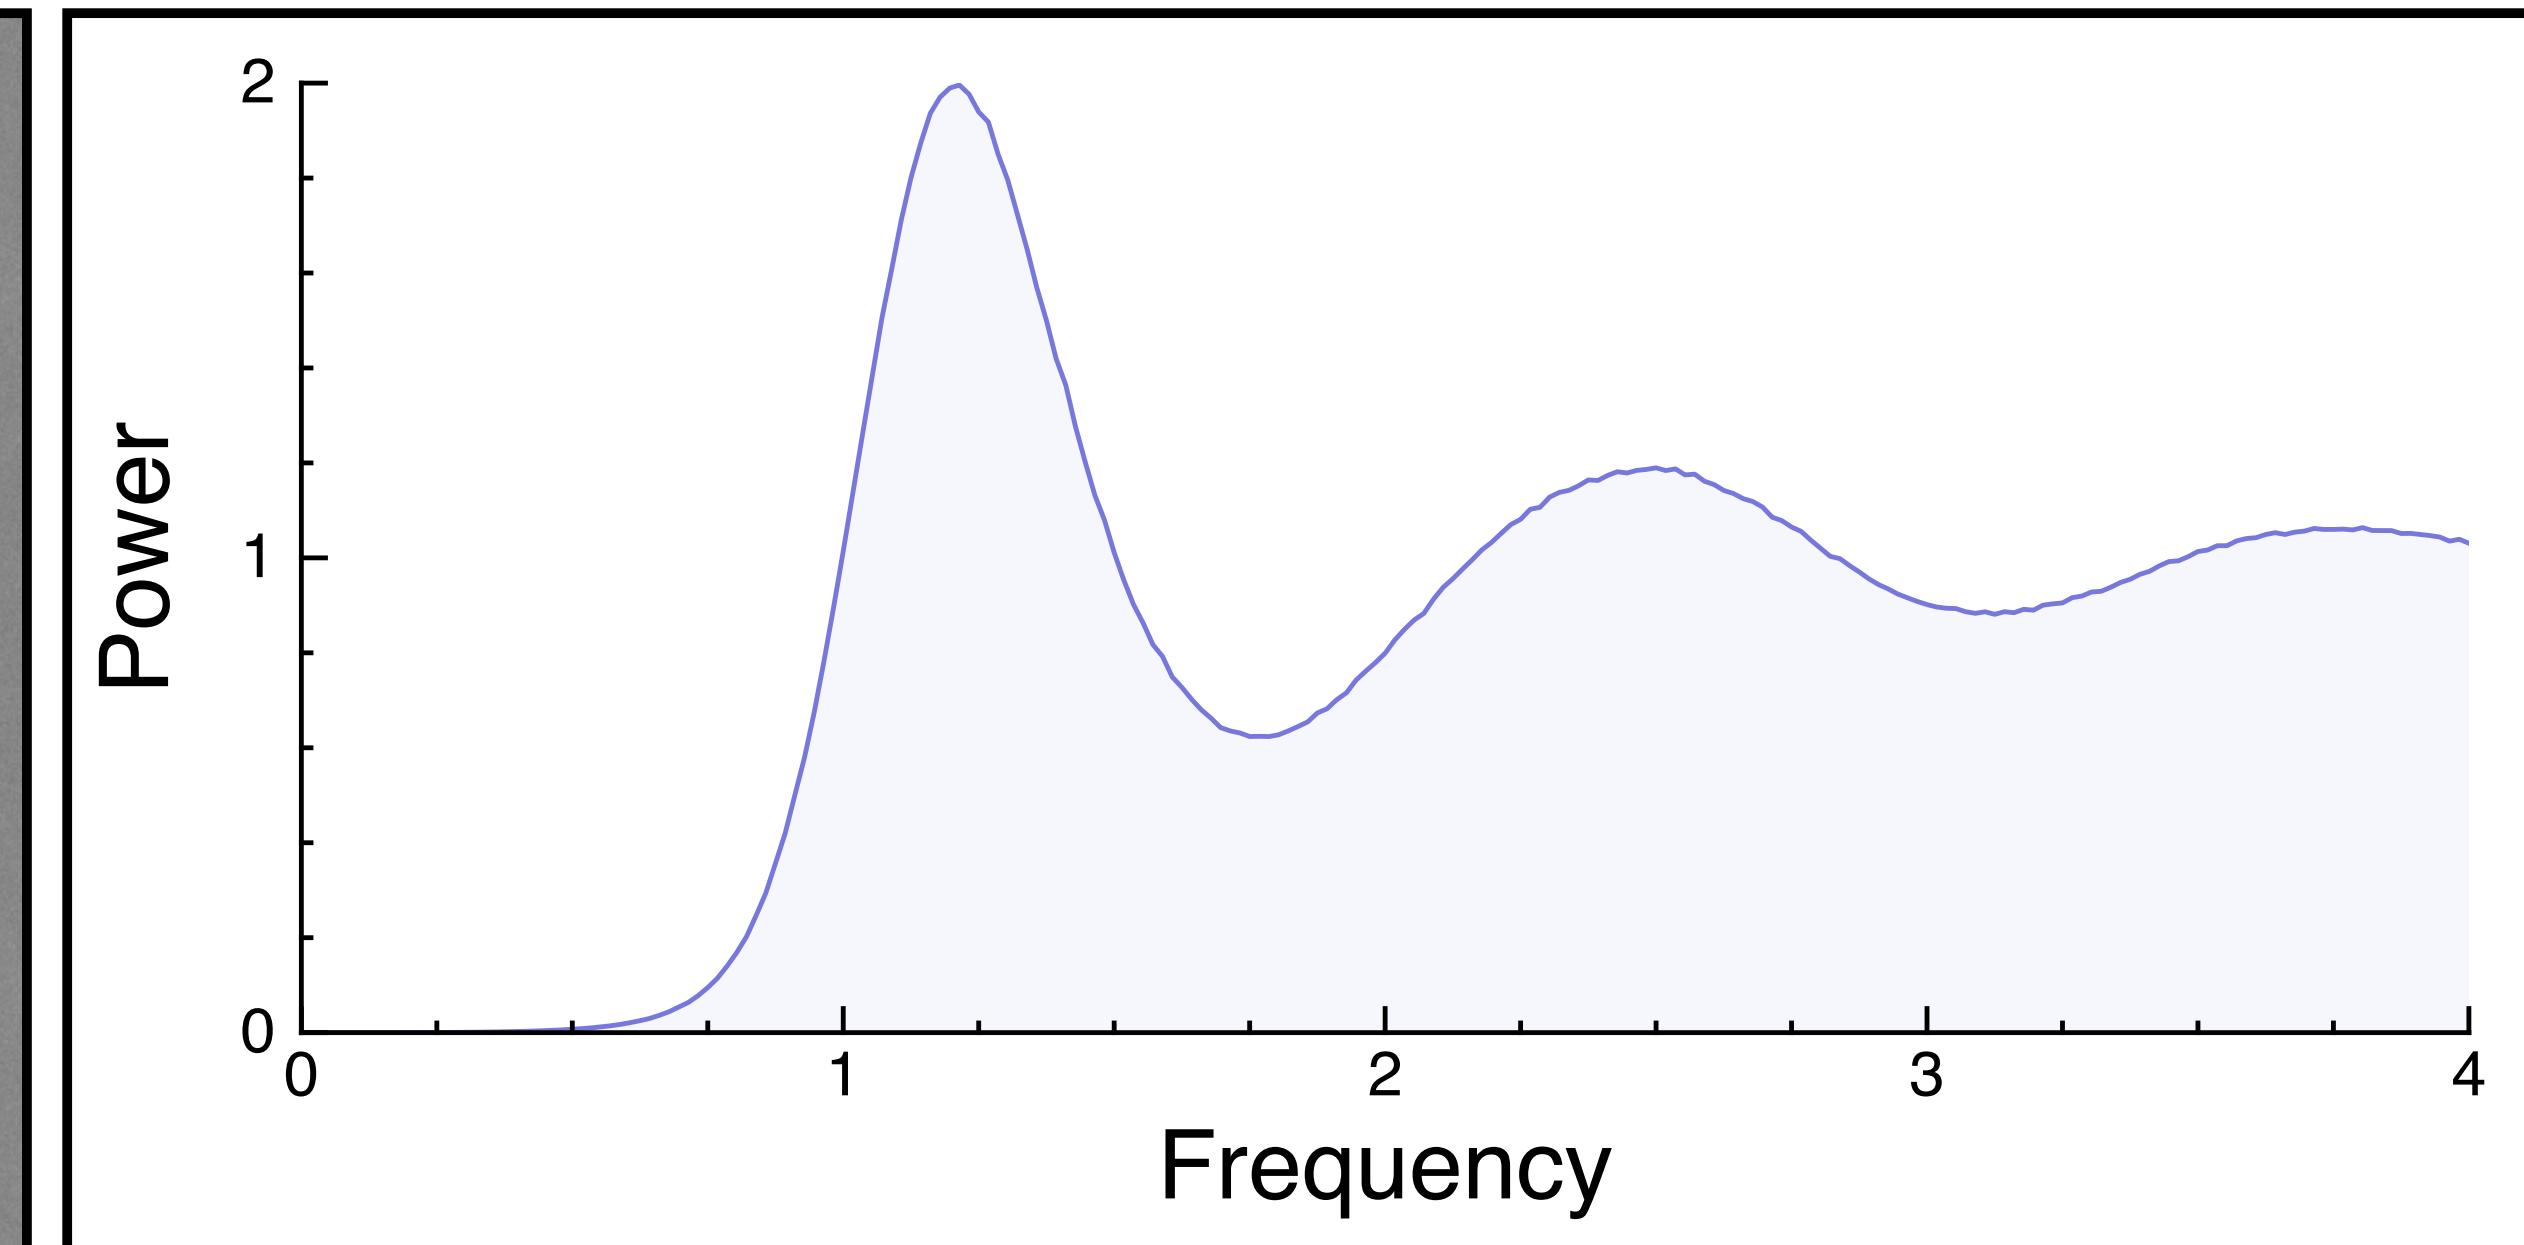
Samples



Power spectrum

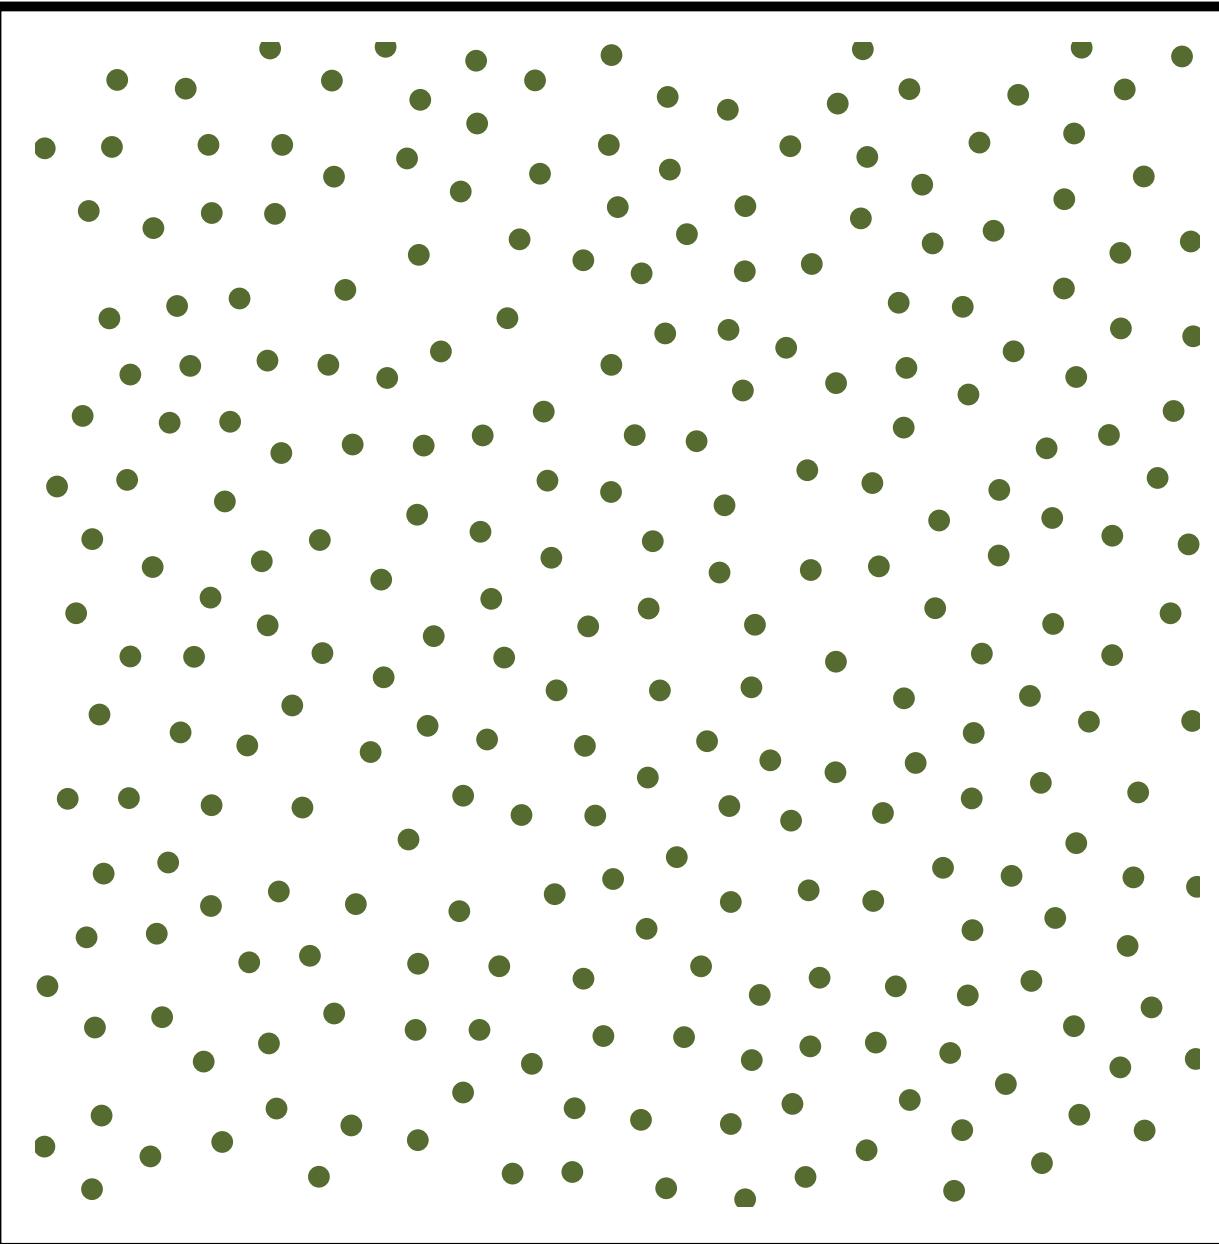


Radial mean

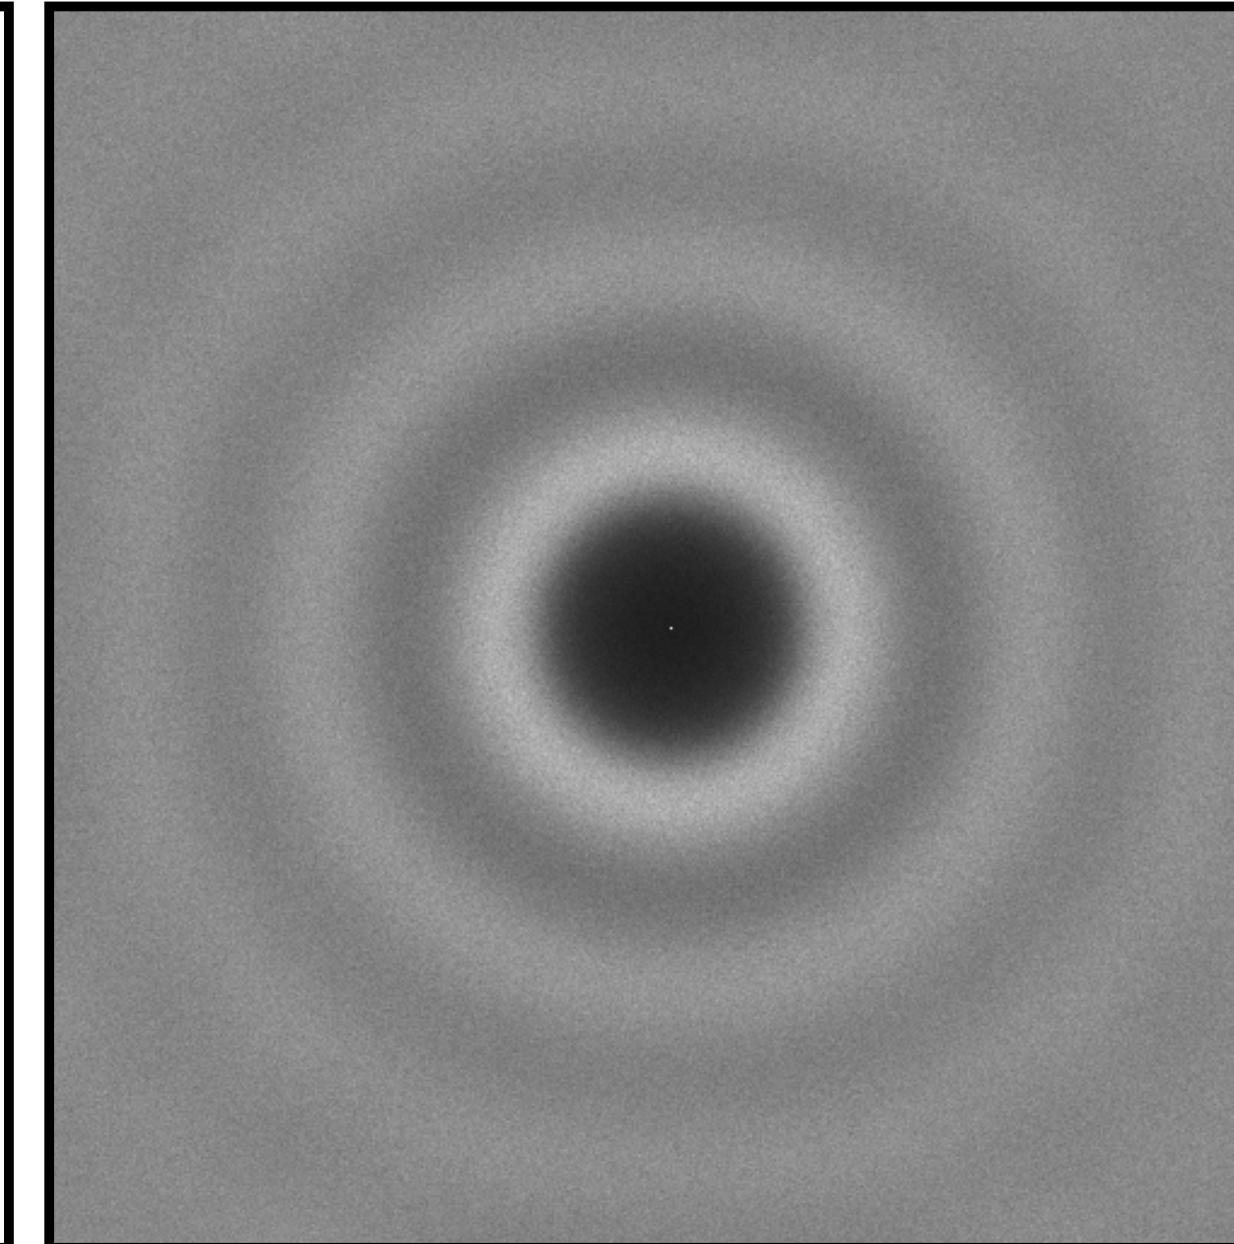


Poisson Disk Sampling

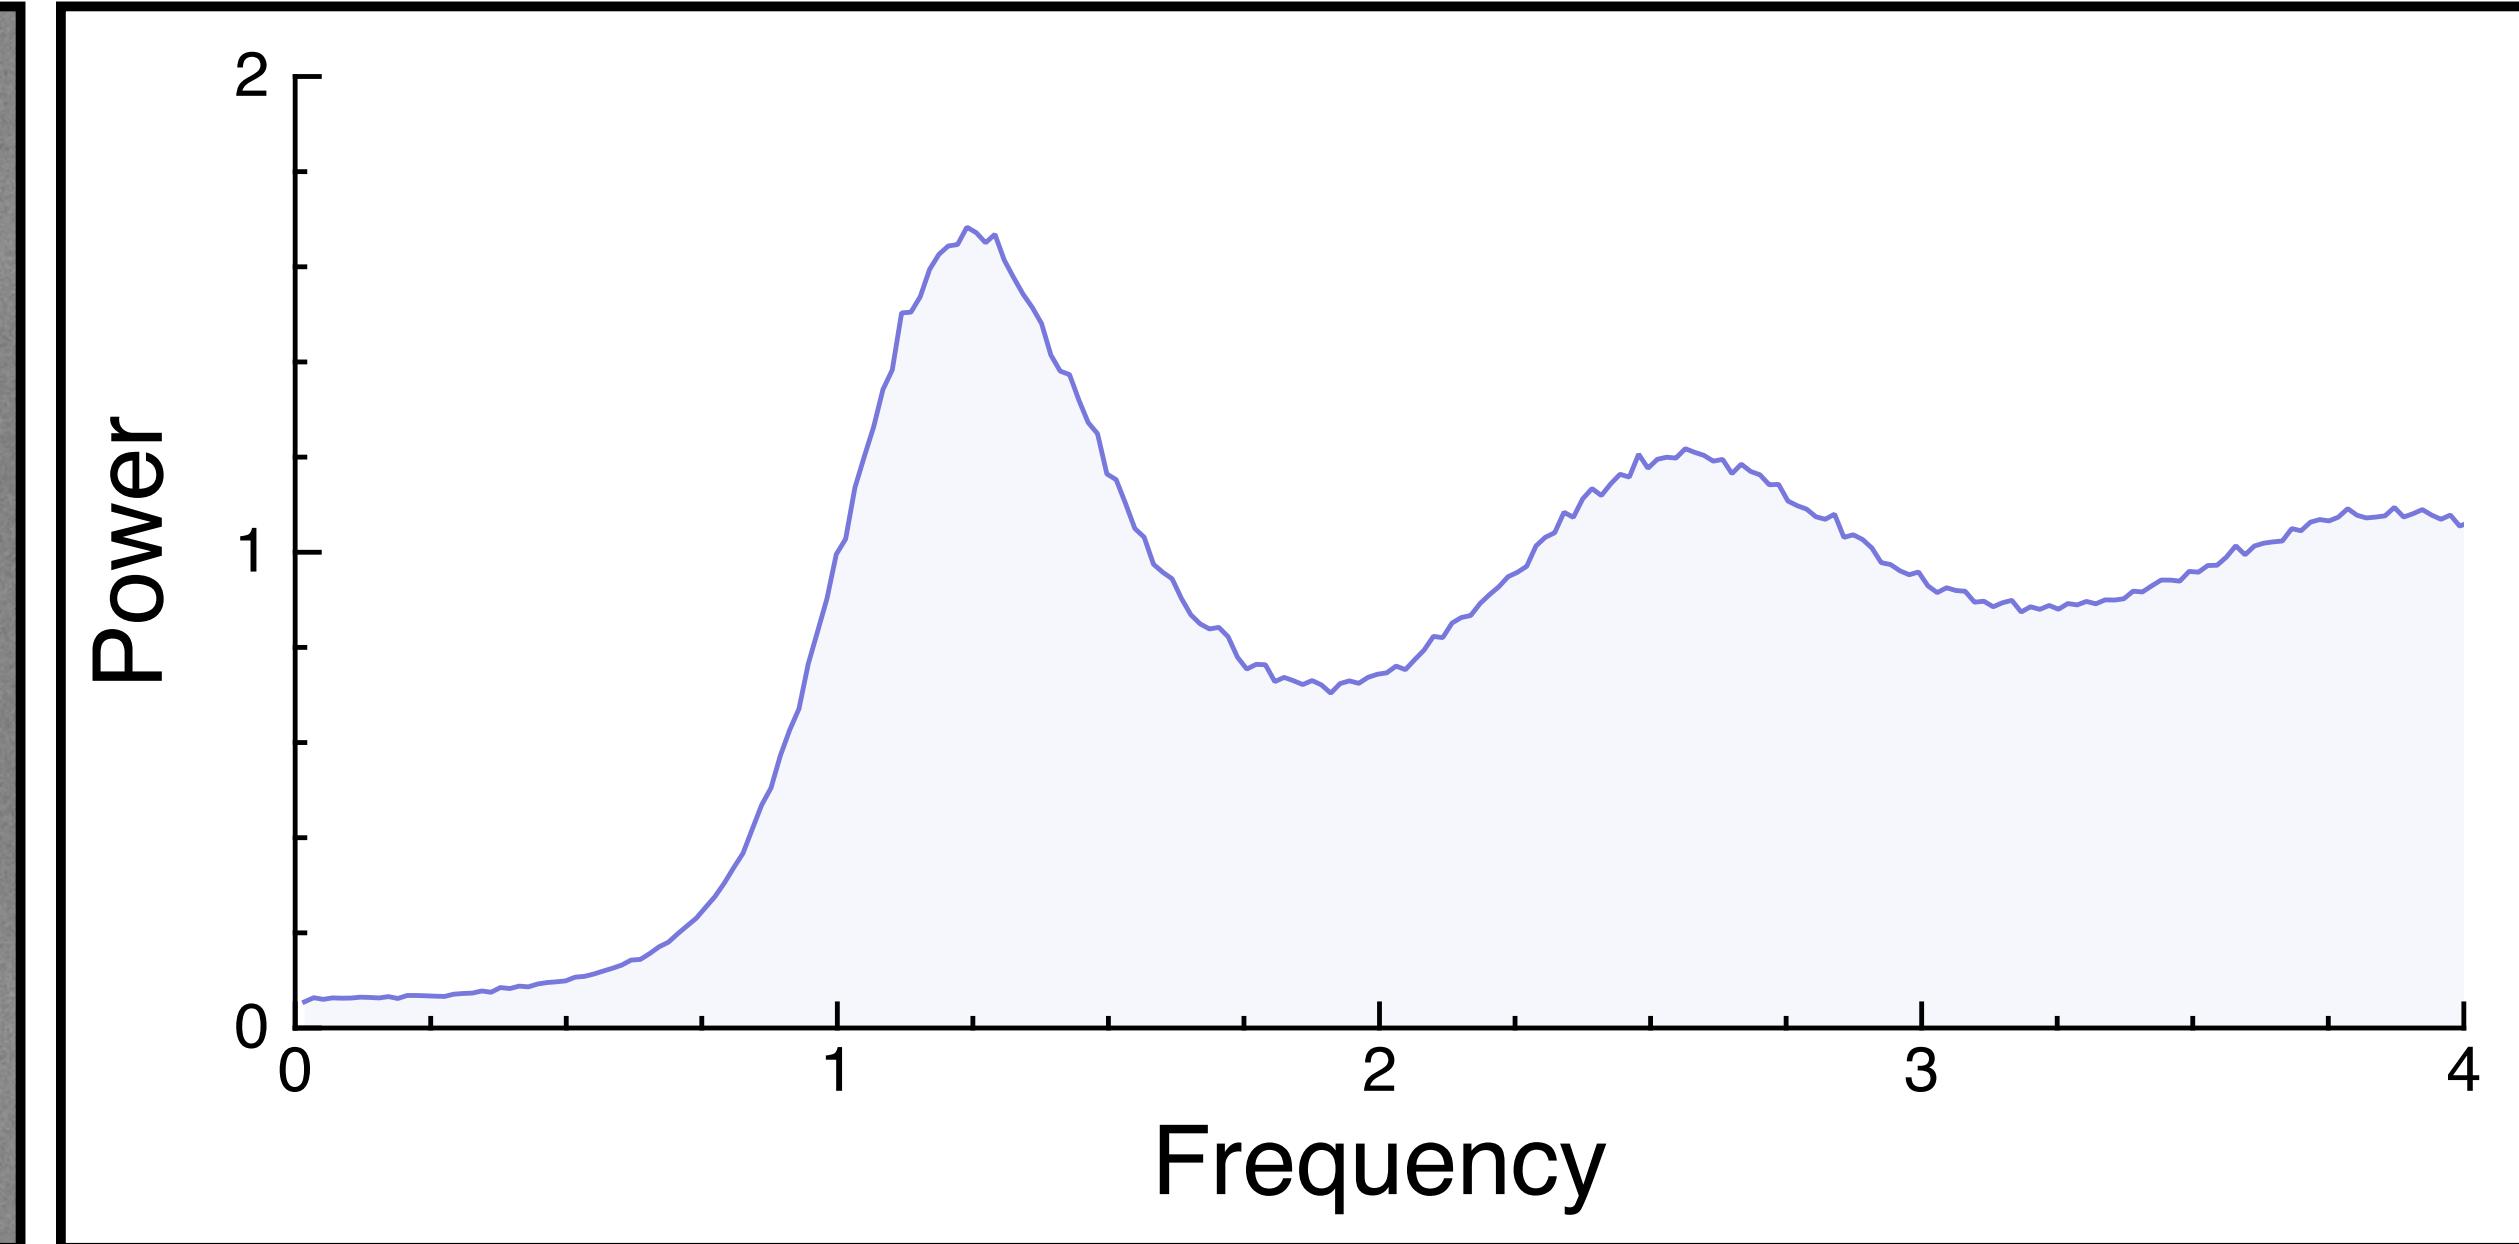
Samples



Power spectrum



Radial mean



Low-Discrepancy Sampling

Deterministic sets of points specially crafted to be evenly distributed (have low discrepancy).

Entire field of study called Quasi-Monte Carlo (QMC)

The Van der Corput Sequence

Radical Inverse Φ_b in base 2

Subsequent points “fall into
biggest holes”

n	Base 2	Φ_b
1	1	.1 = 1/2
2	10	.01 = 1/4
3	11	.11 = 3/4
4	100	.001 = 1/8
5	101	.101 = 5/8
6	110	.011 = 3/8
7	111	.111 = 7/8
...		



Halton and Hammersley Points

Halton: Radical inverse with different base for each dimension:

$$x_i = (\Phi_2(i), \Phi_3(i), \Phi_5(i), \dots, \Phi_{p_n}(i))$$

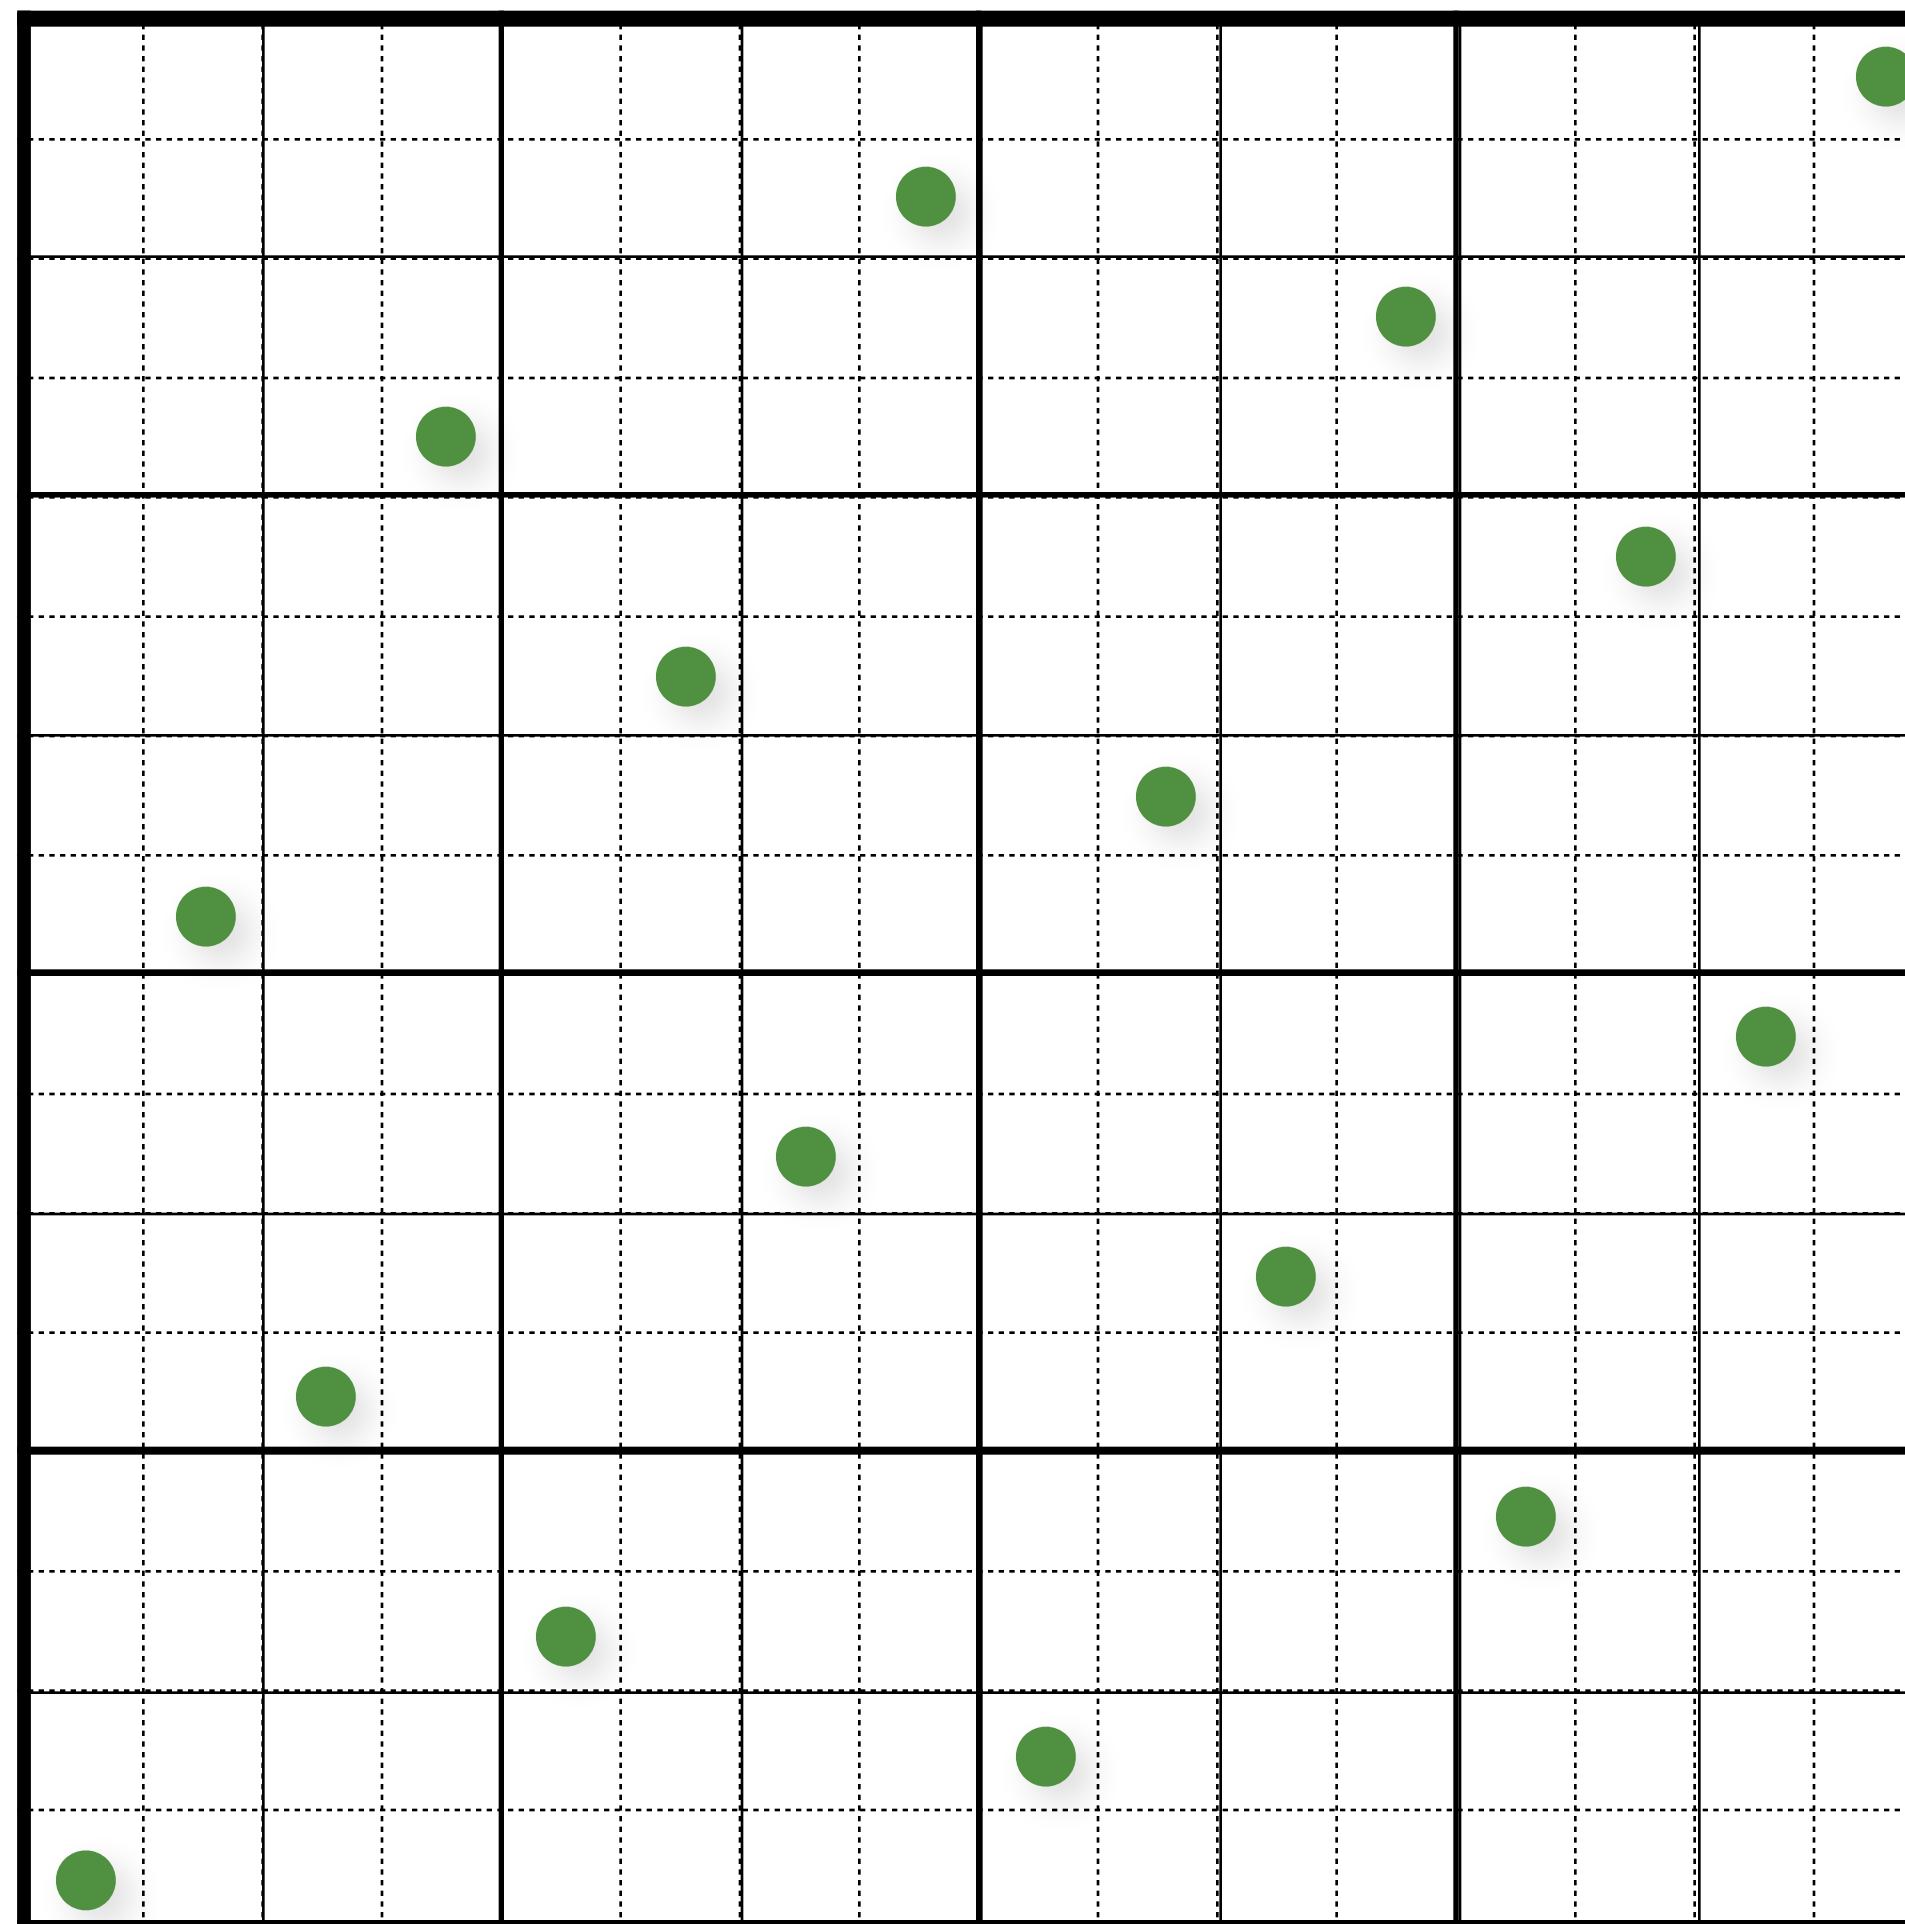
- The bases should all be relatively prime.
- Incremental/progressive generation of samples

Hammersley: Same as Halton, but first dimension is i/N :

$$x_i = \left(\frac{i}{N}, \Phi_2(i), \Phi_3(i), \dots, \Phi_{p_n}(i) \right)$$

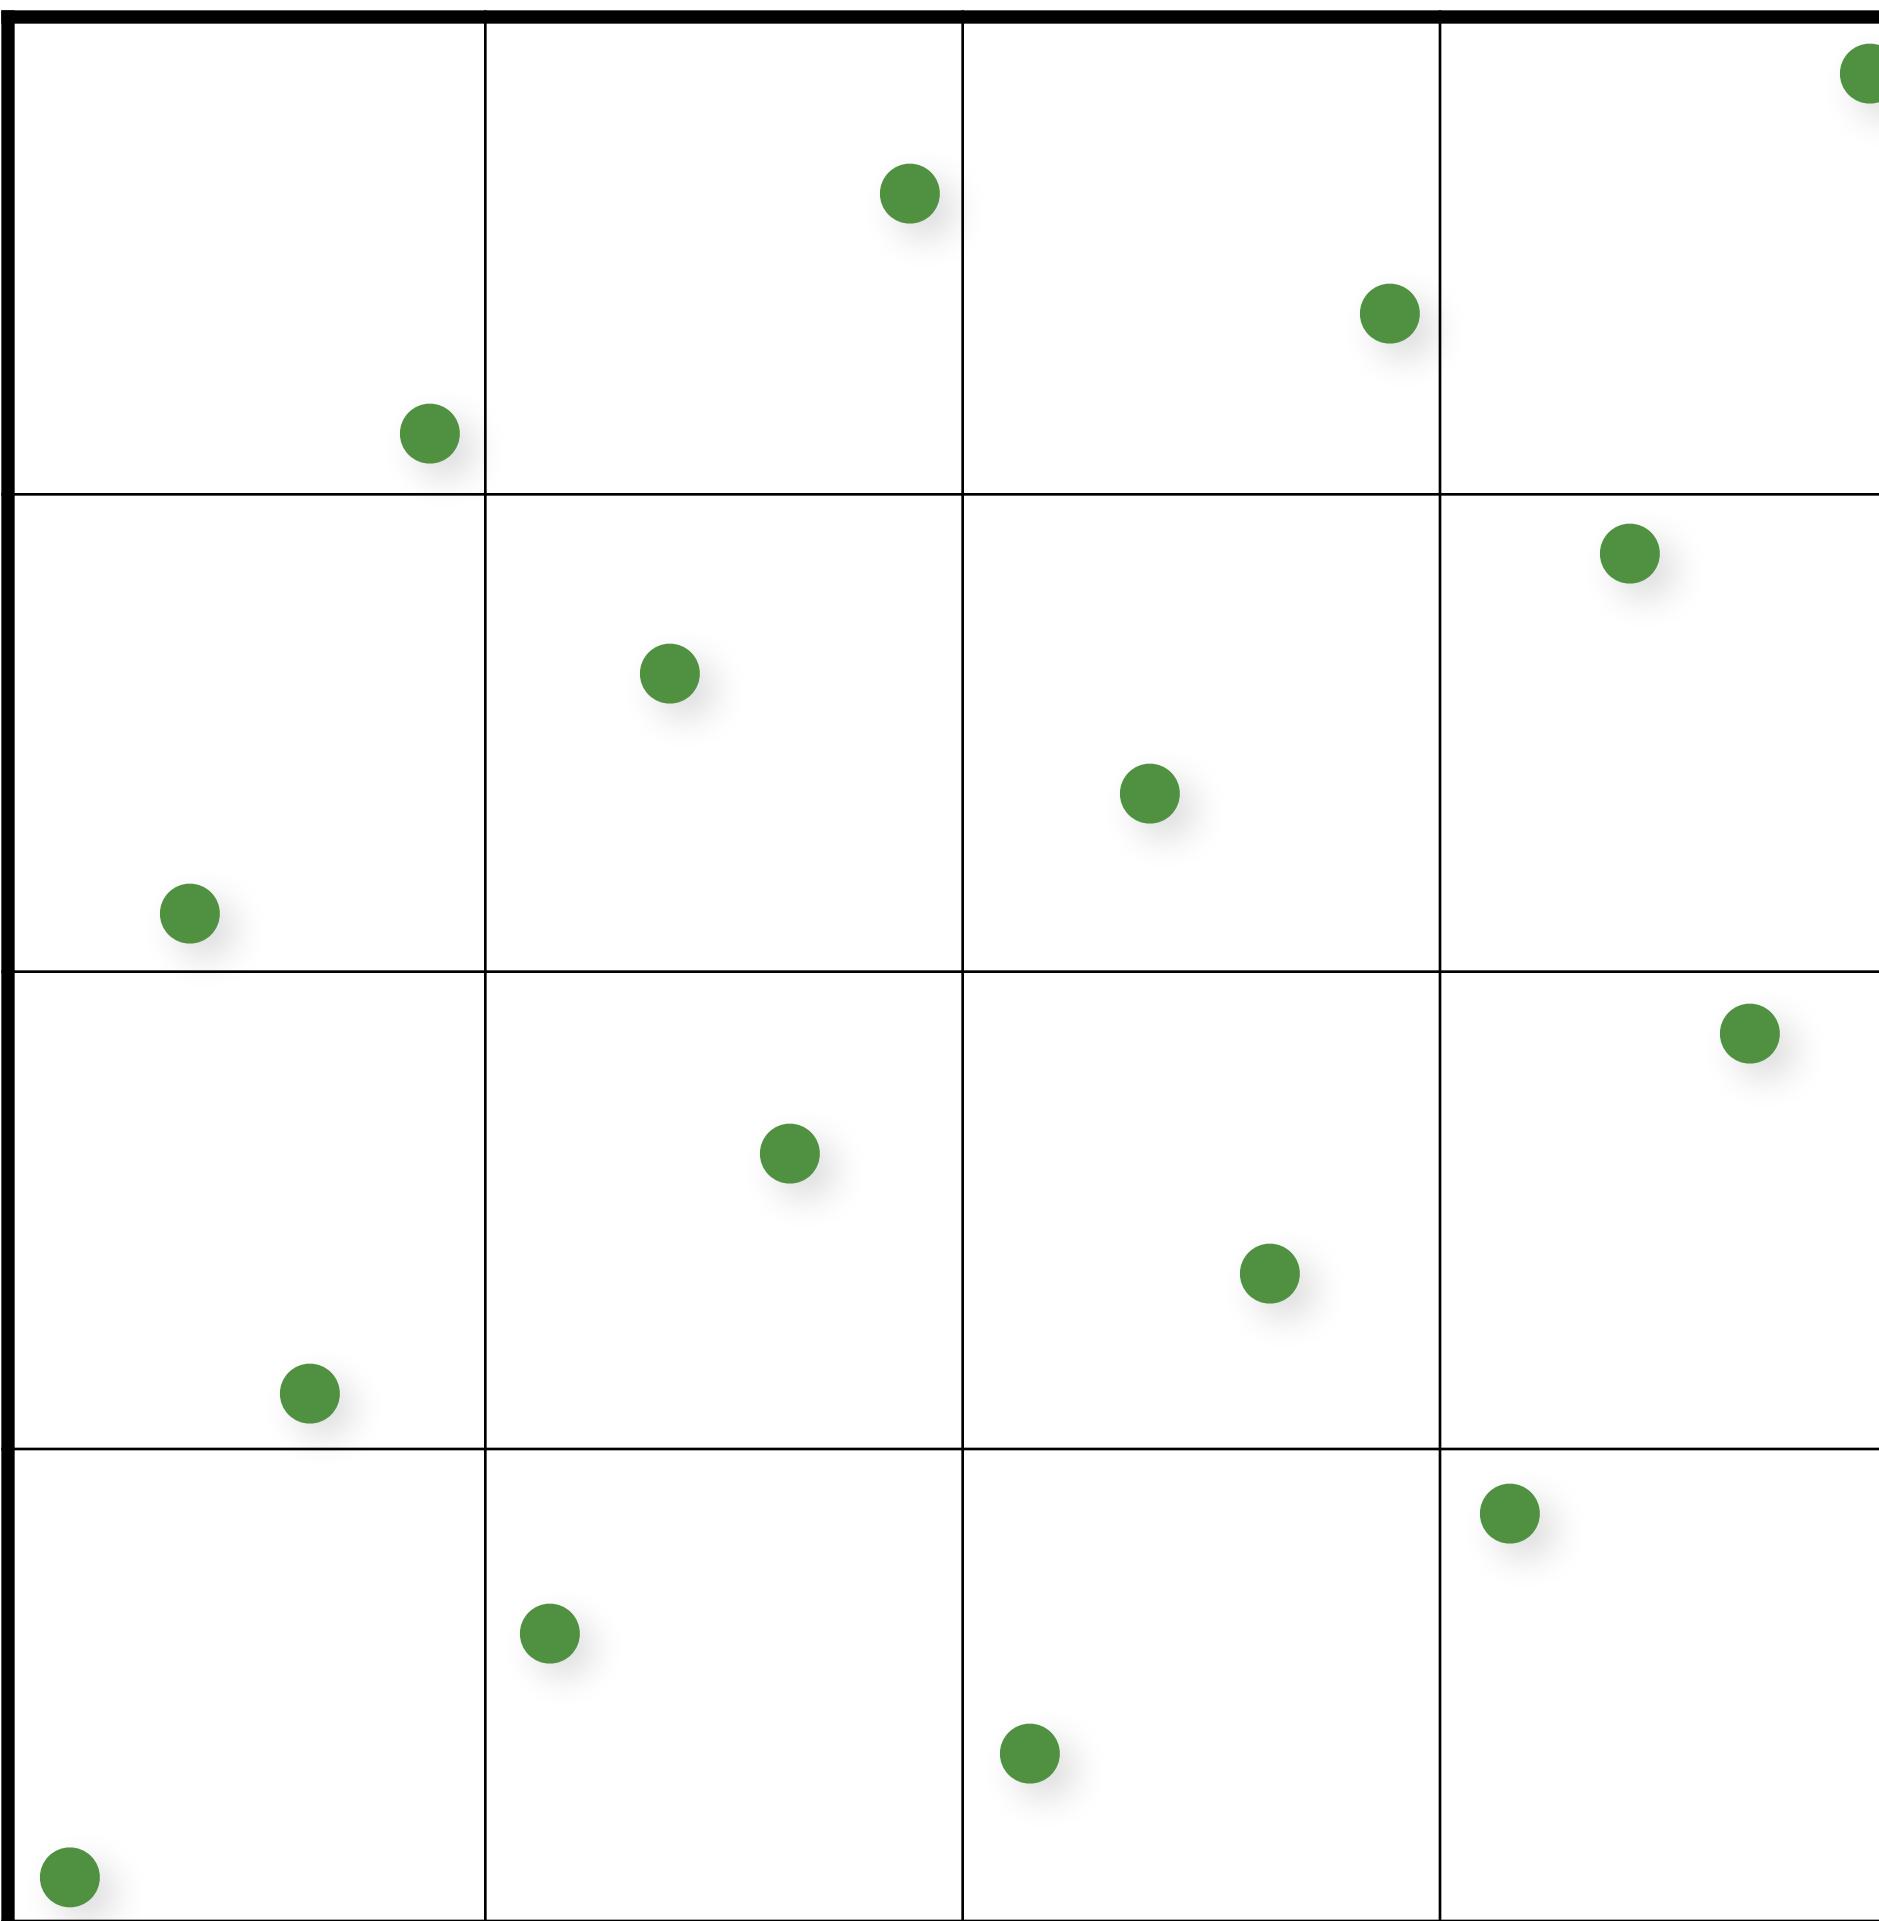
- Not incremental, need to know sample count, N , in advance

The Hammersley Sequence



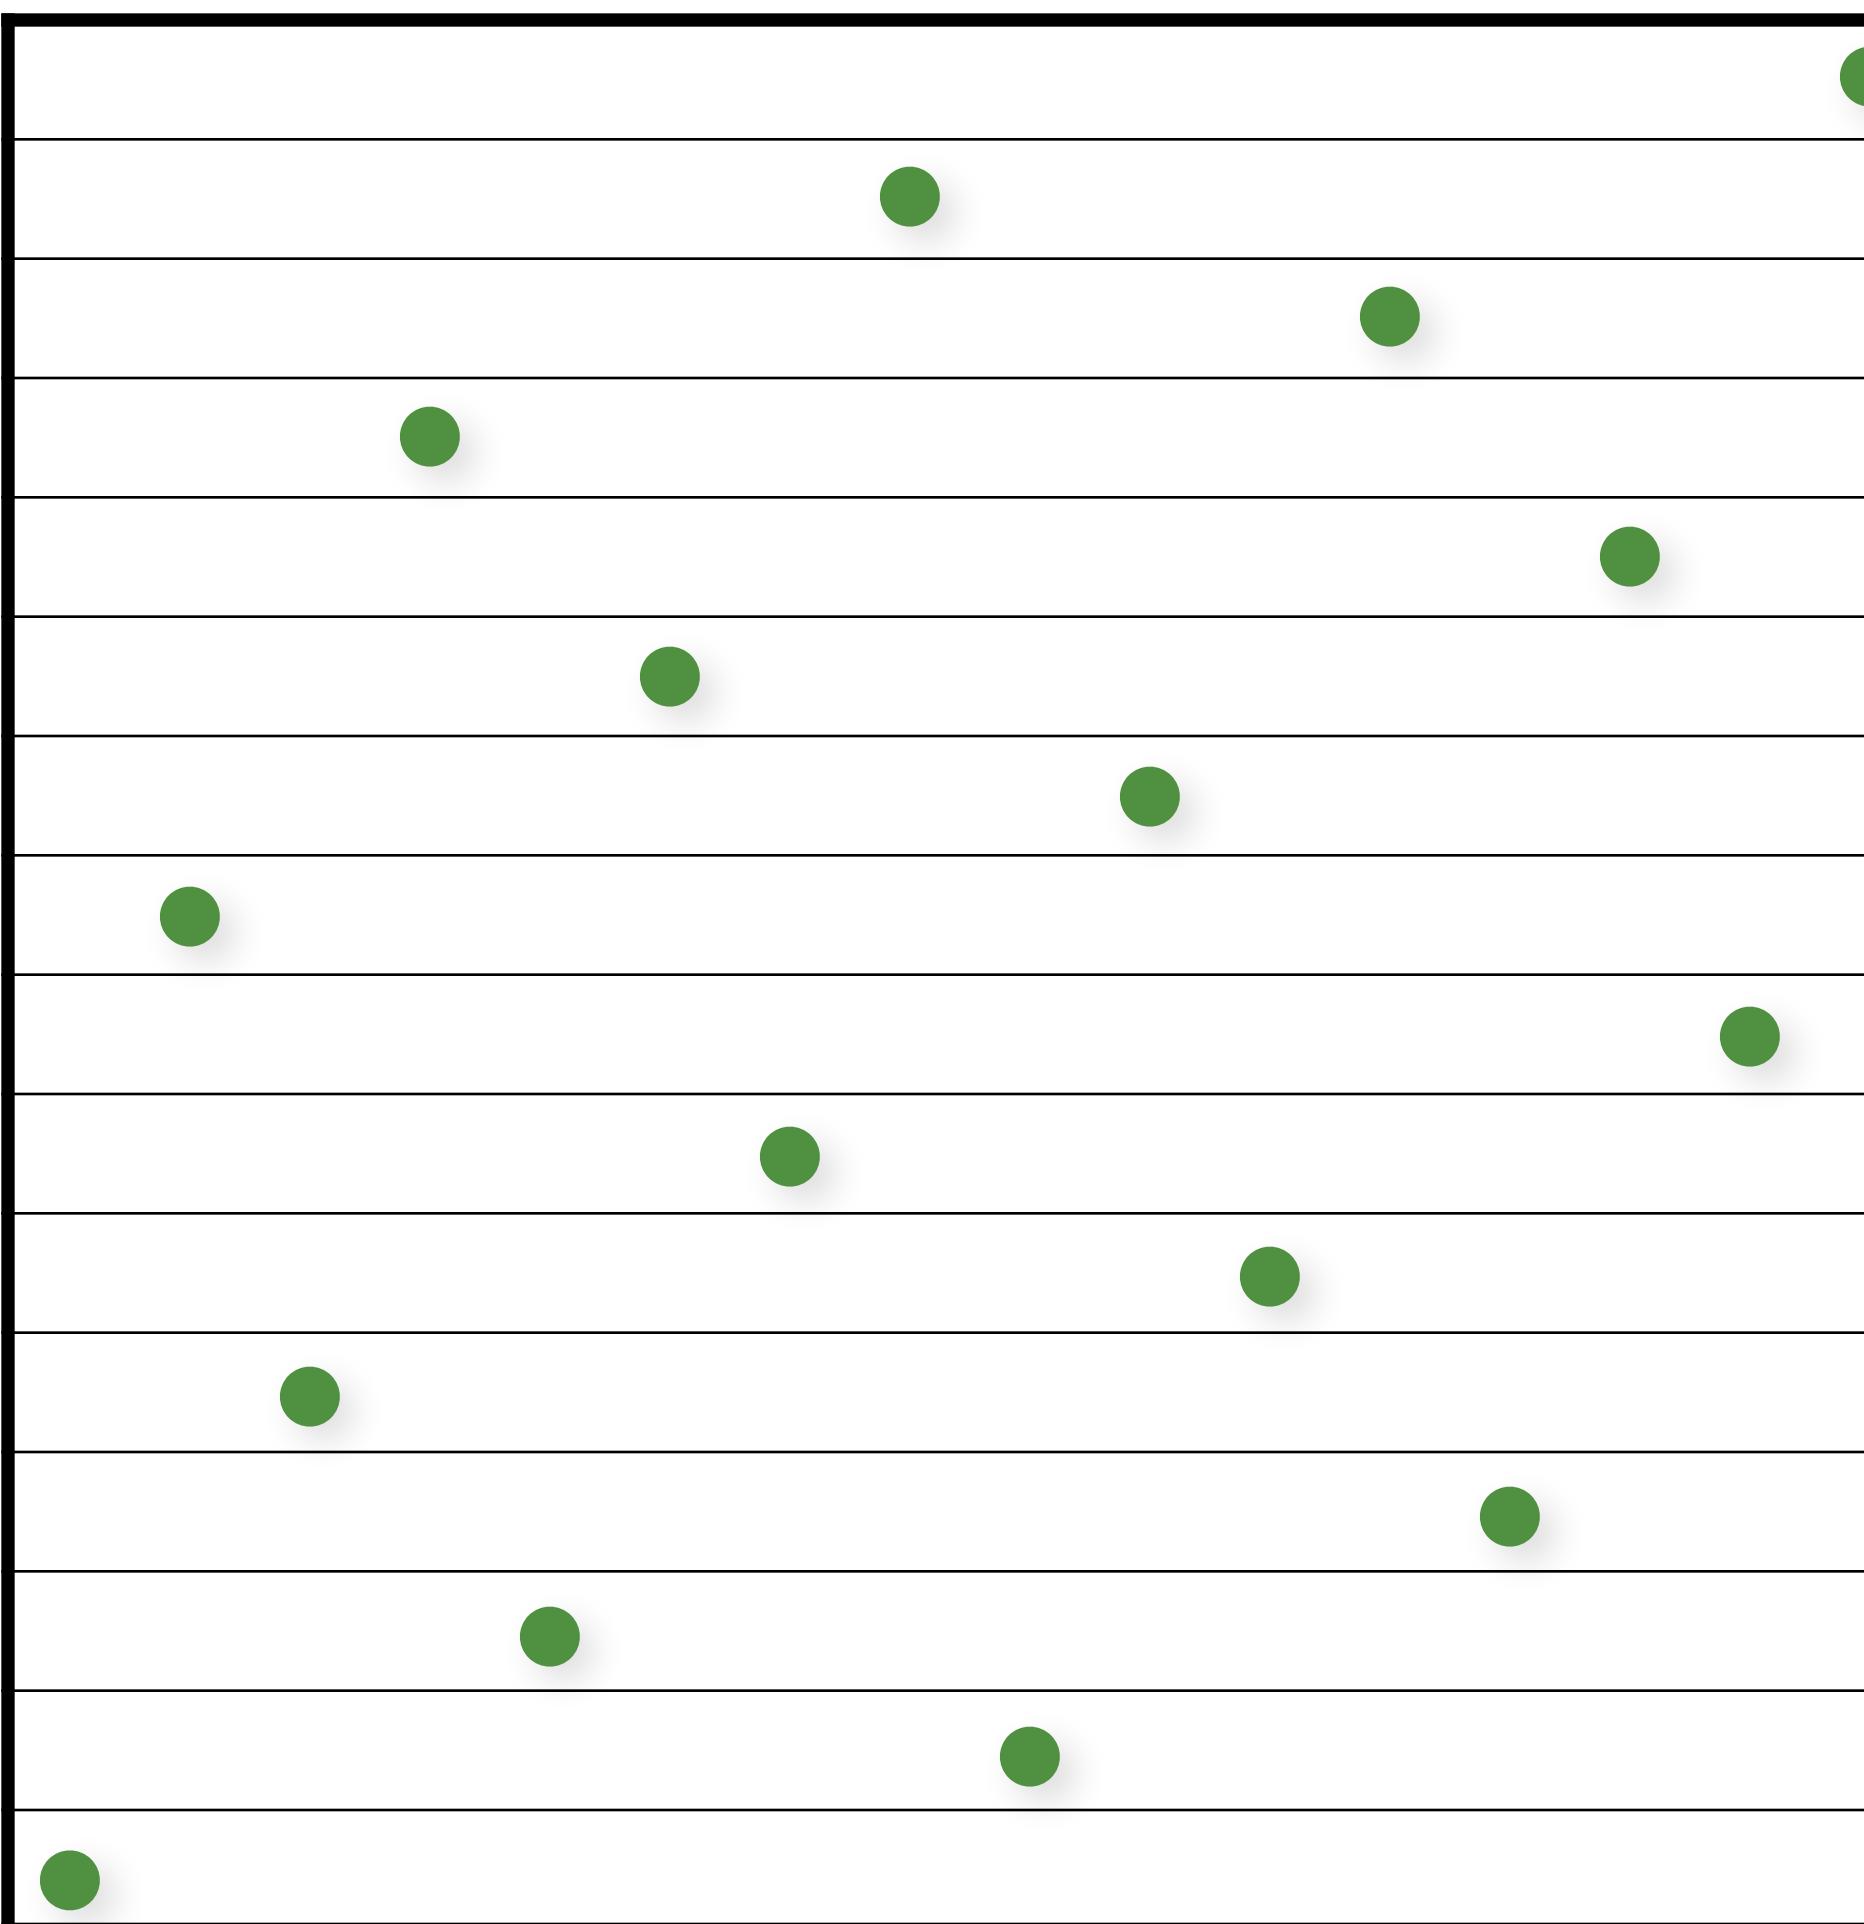
1 sample in each “elementary interval”

The Hammersley Sequence



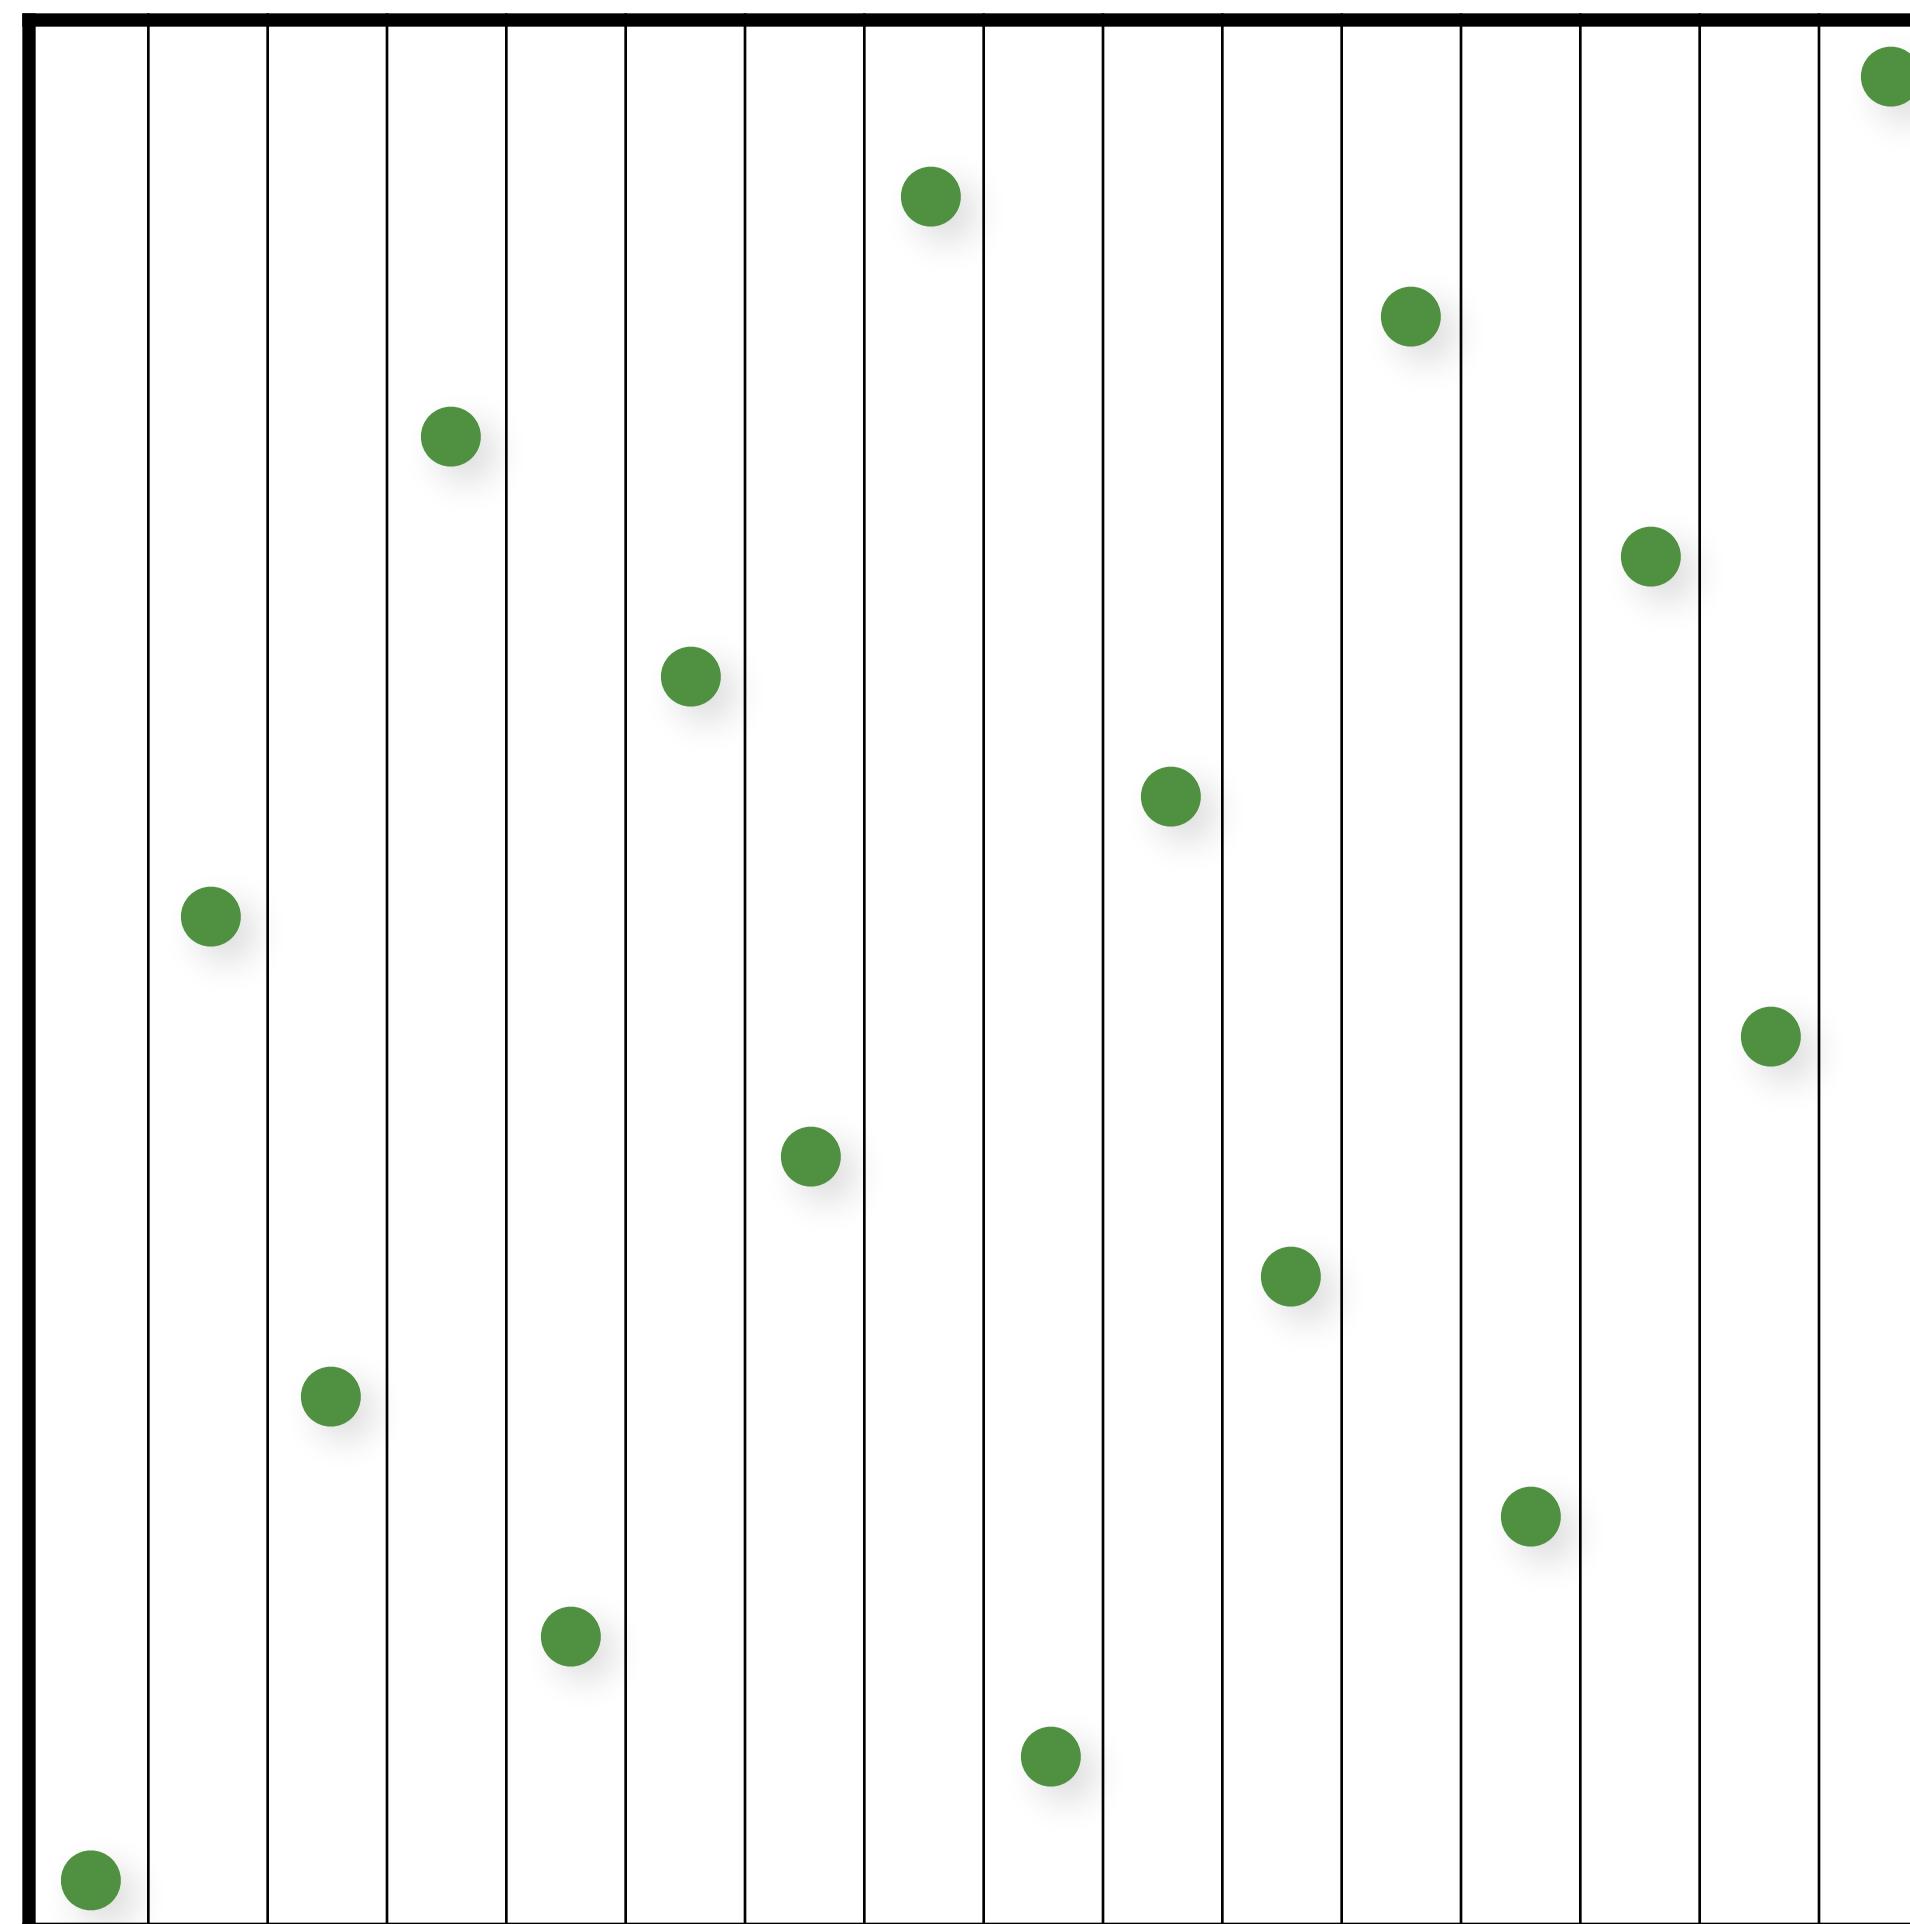
1 sample in each “elementary interval”

The Hammersley Sequence



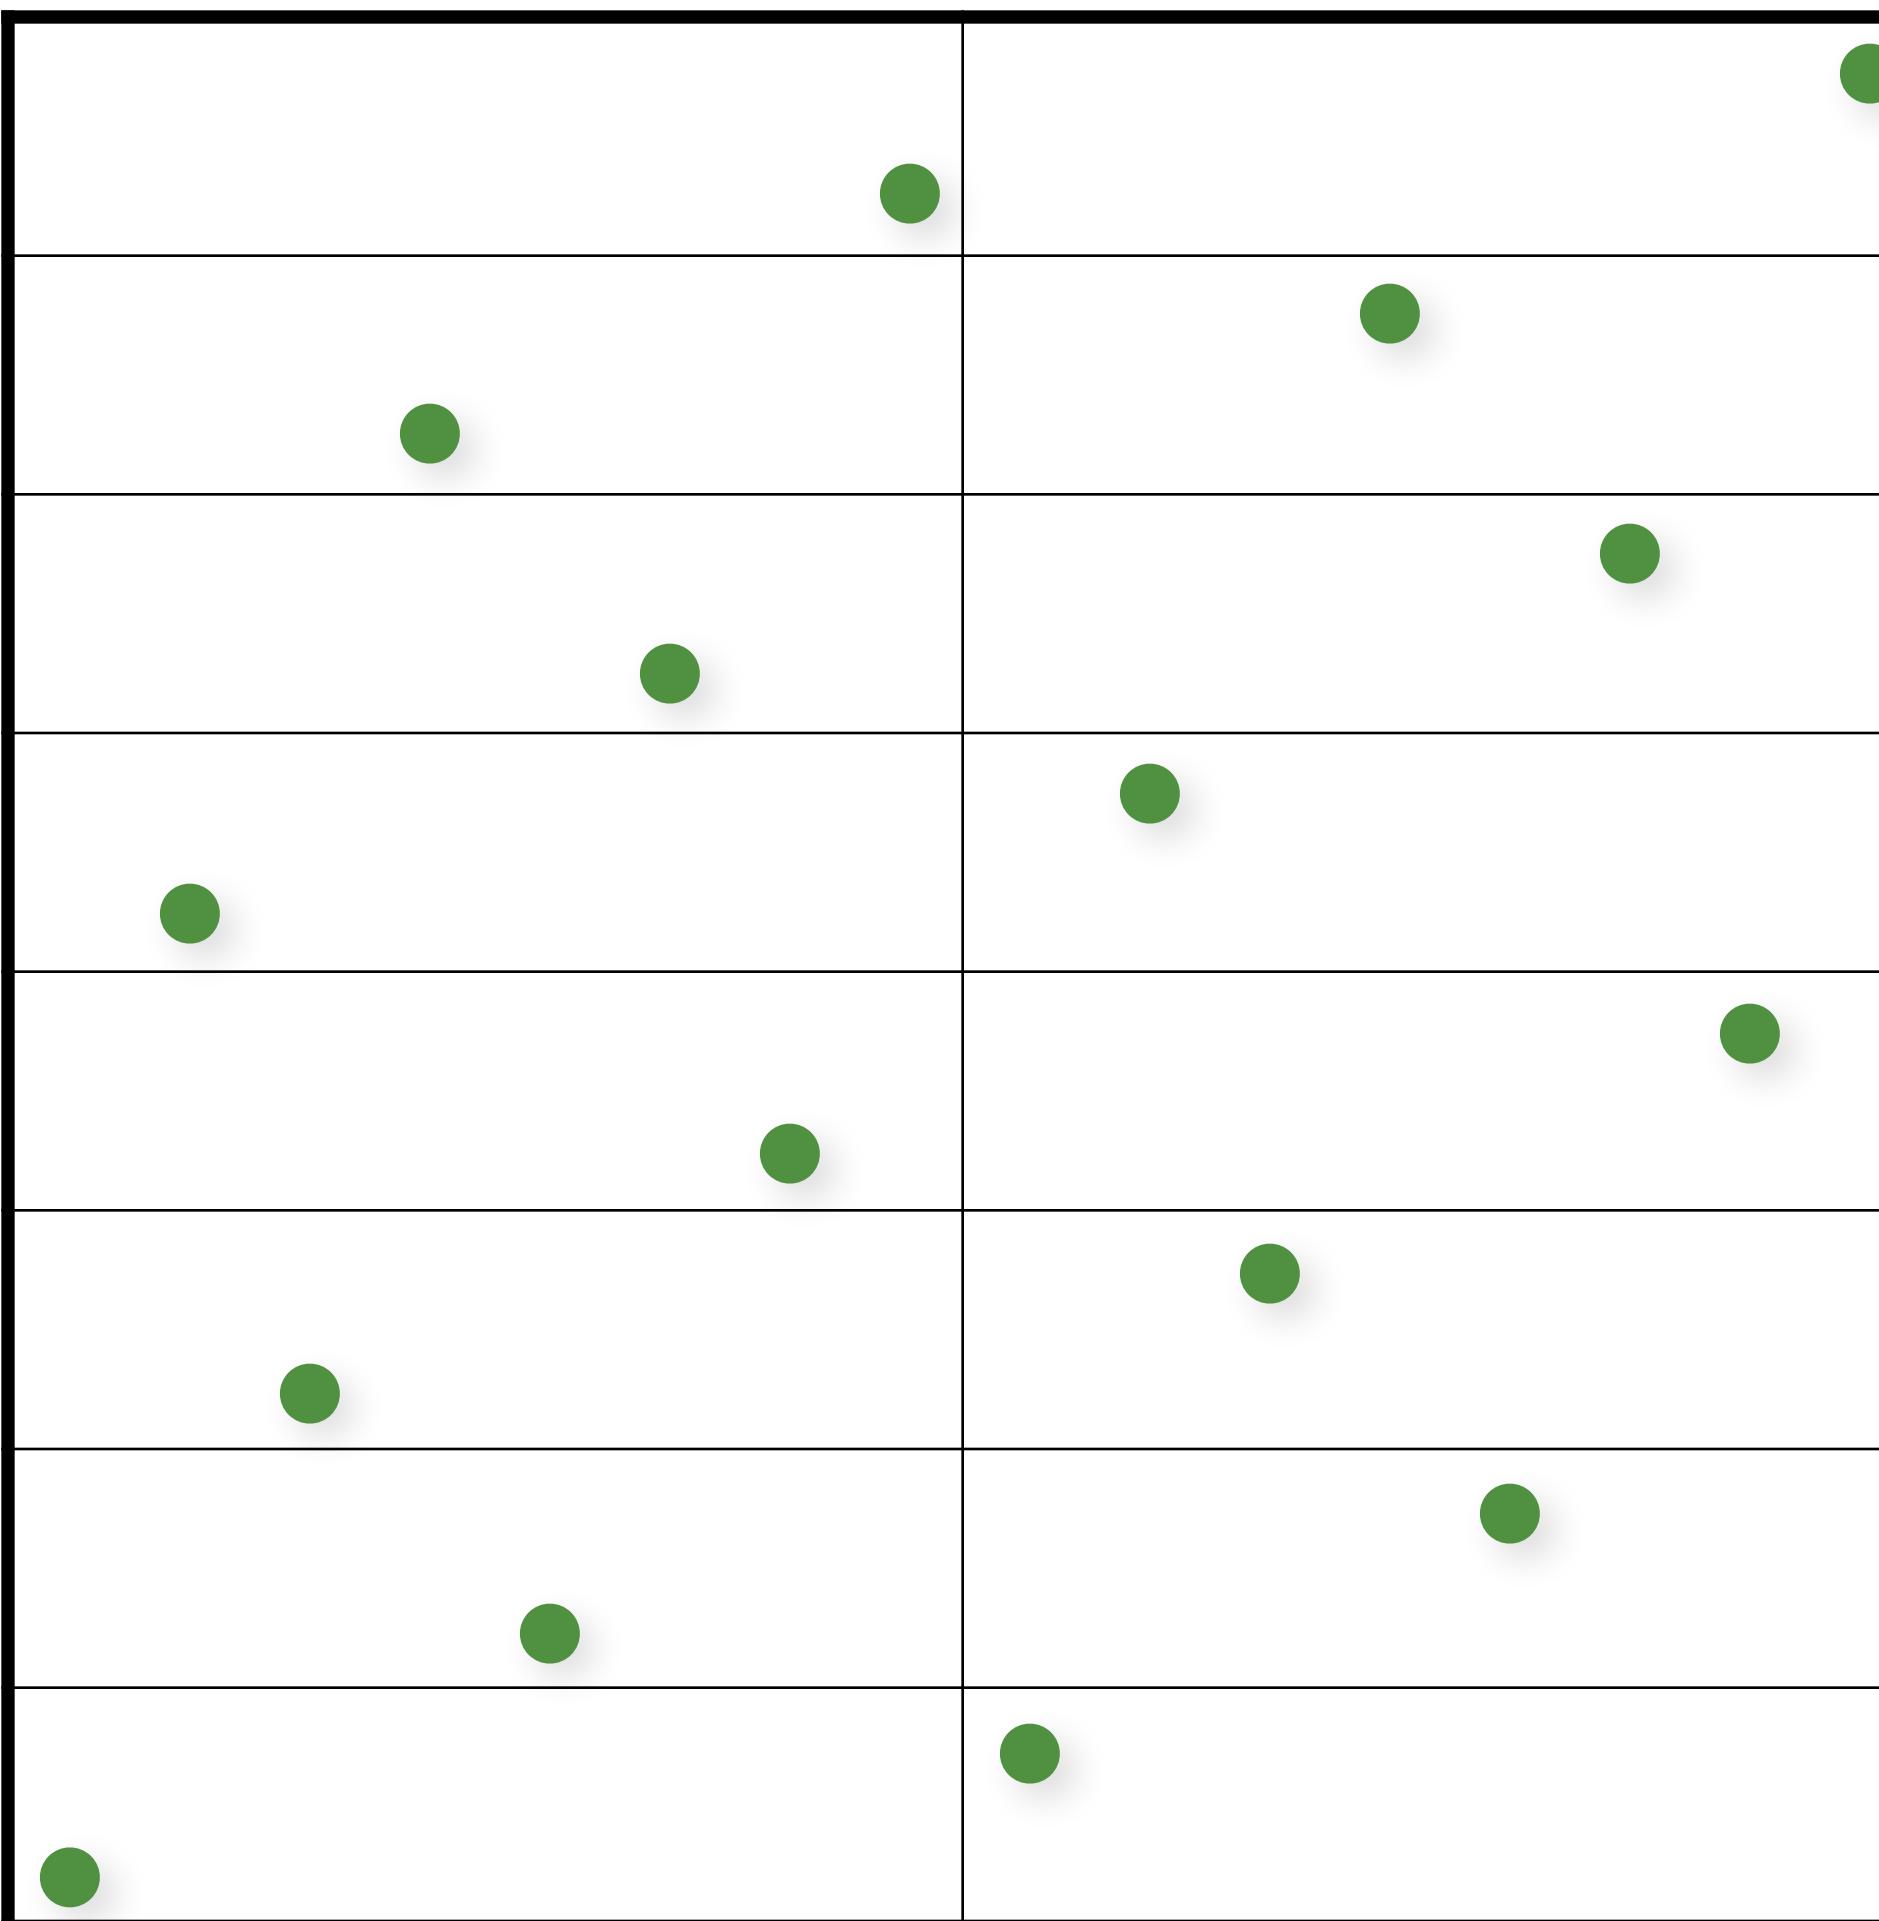
1 sample in each “elementary interval”

The Hammersley Sequence



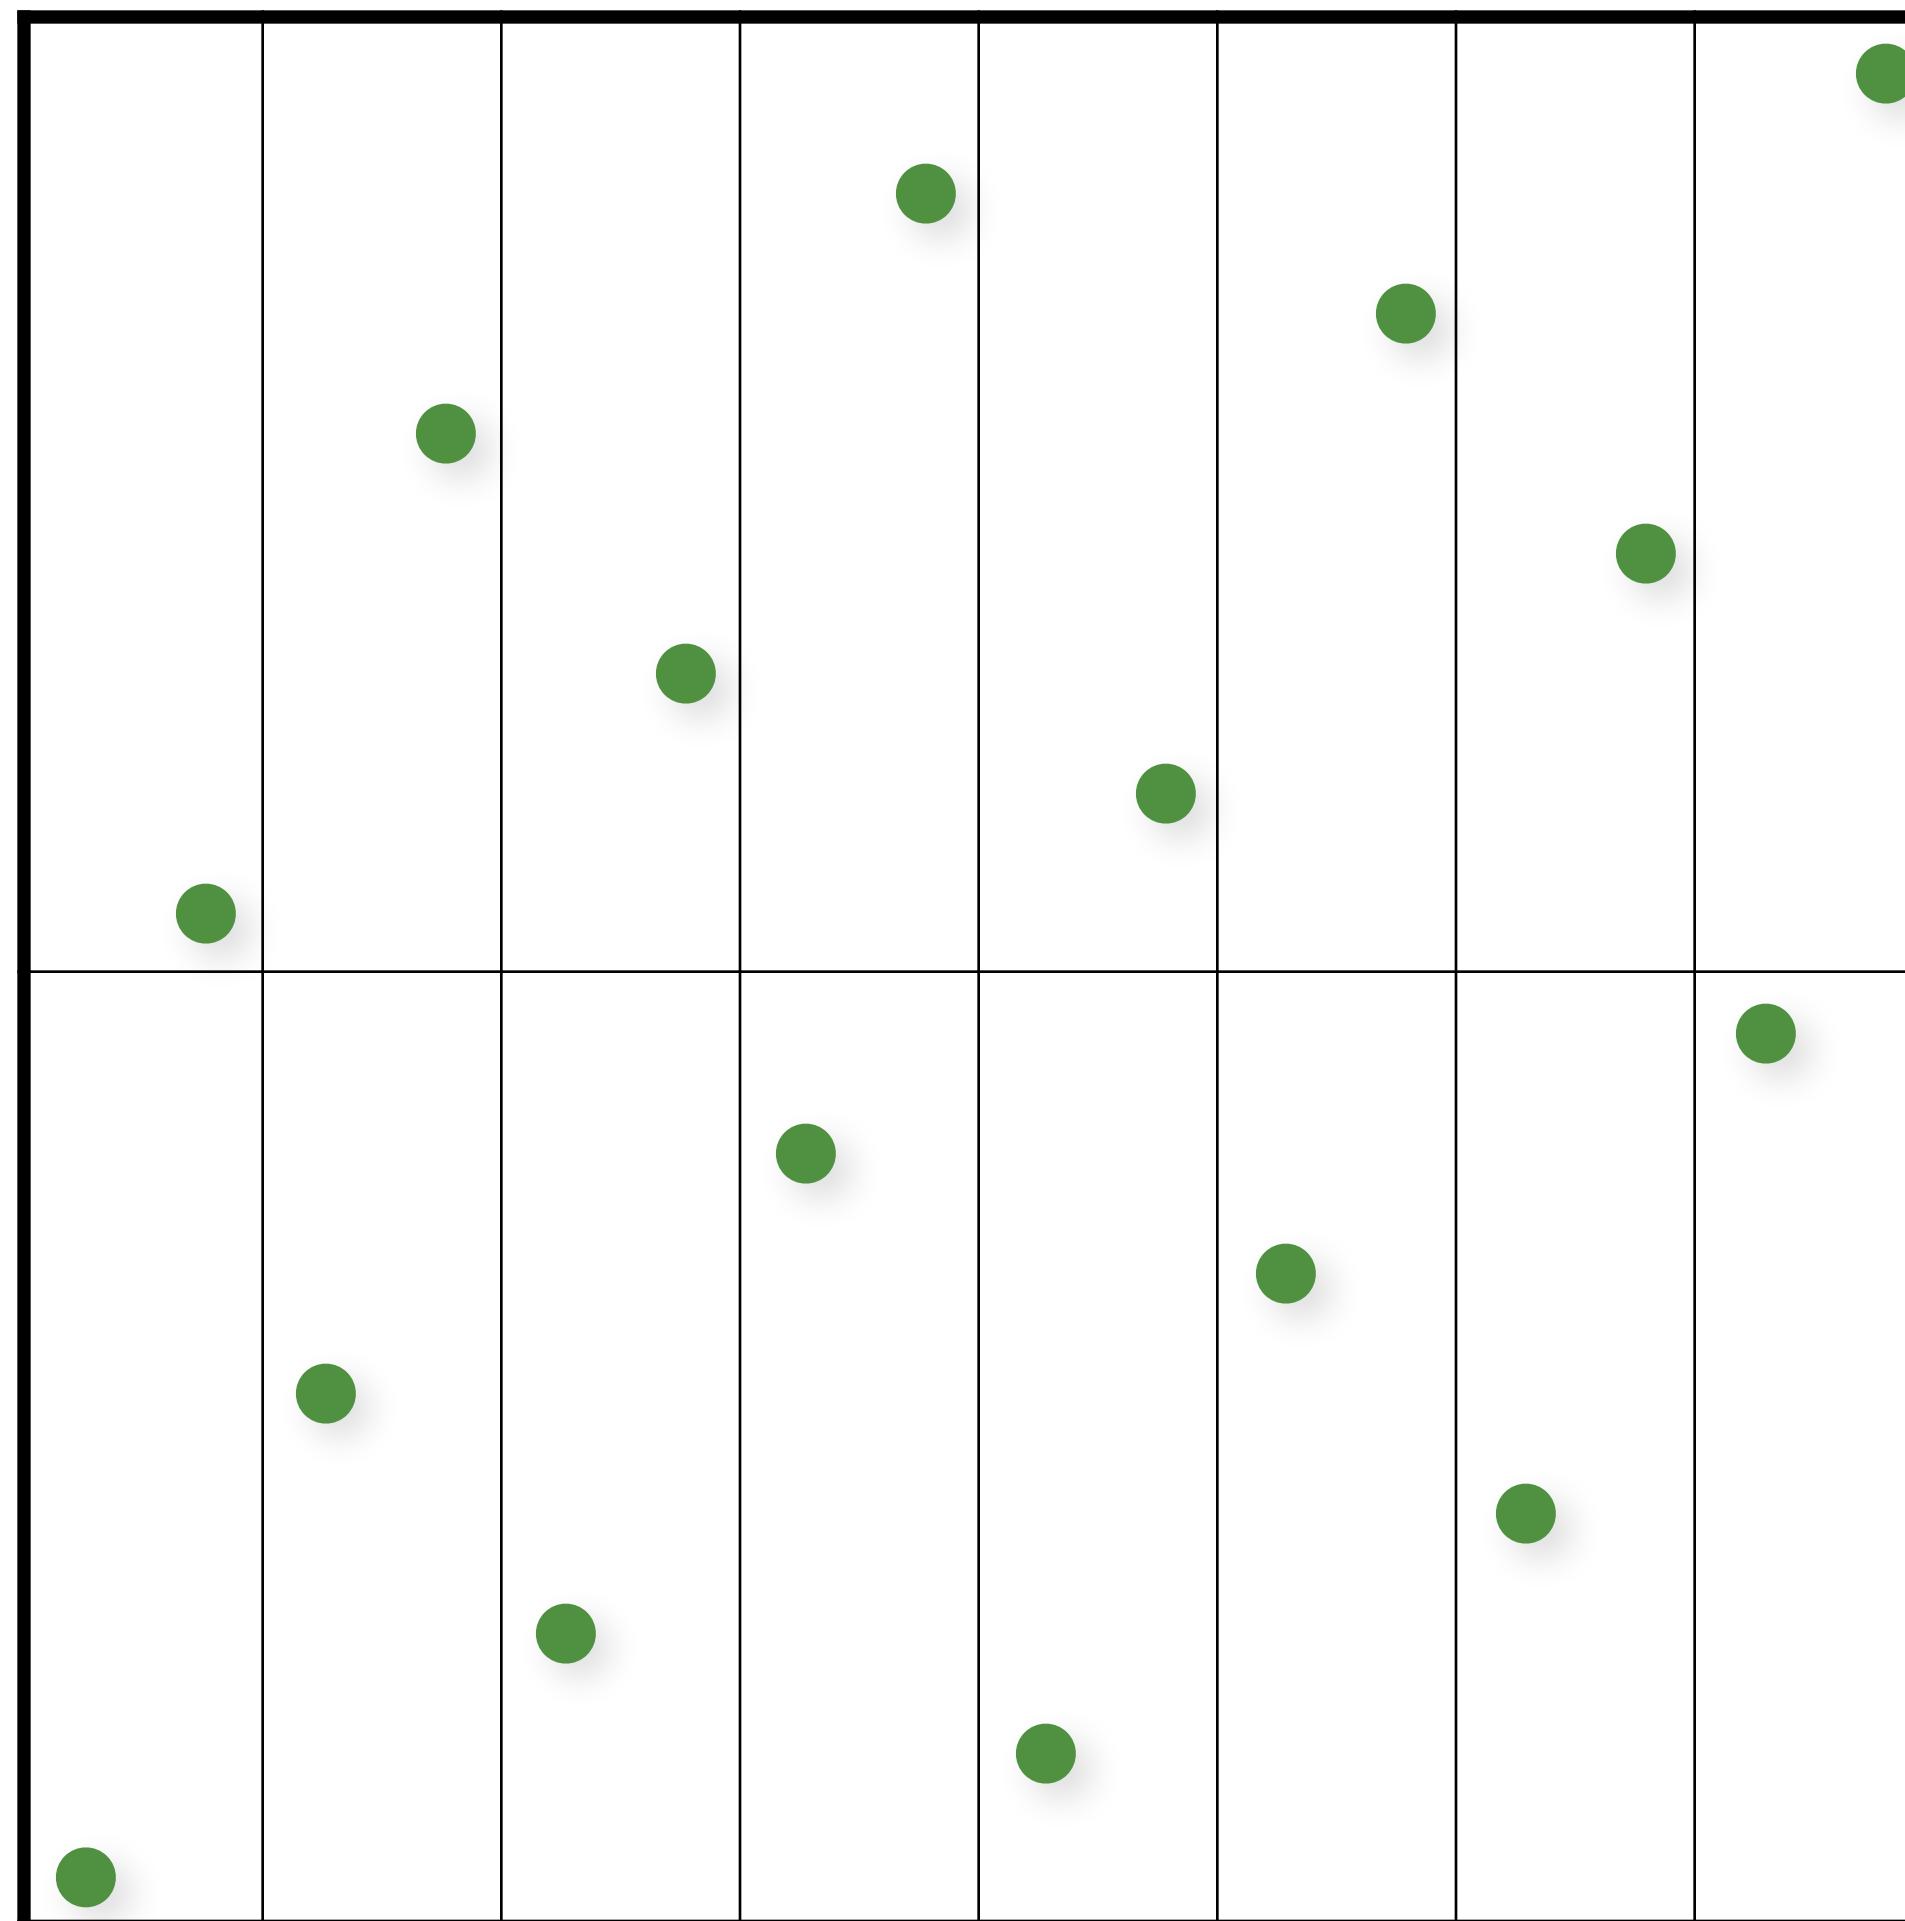
1 sample in each “elementary interval”

The Hammersley Sequence



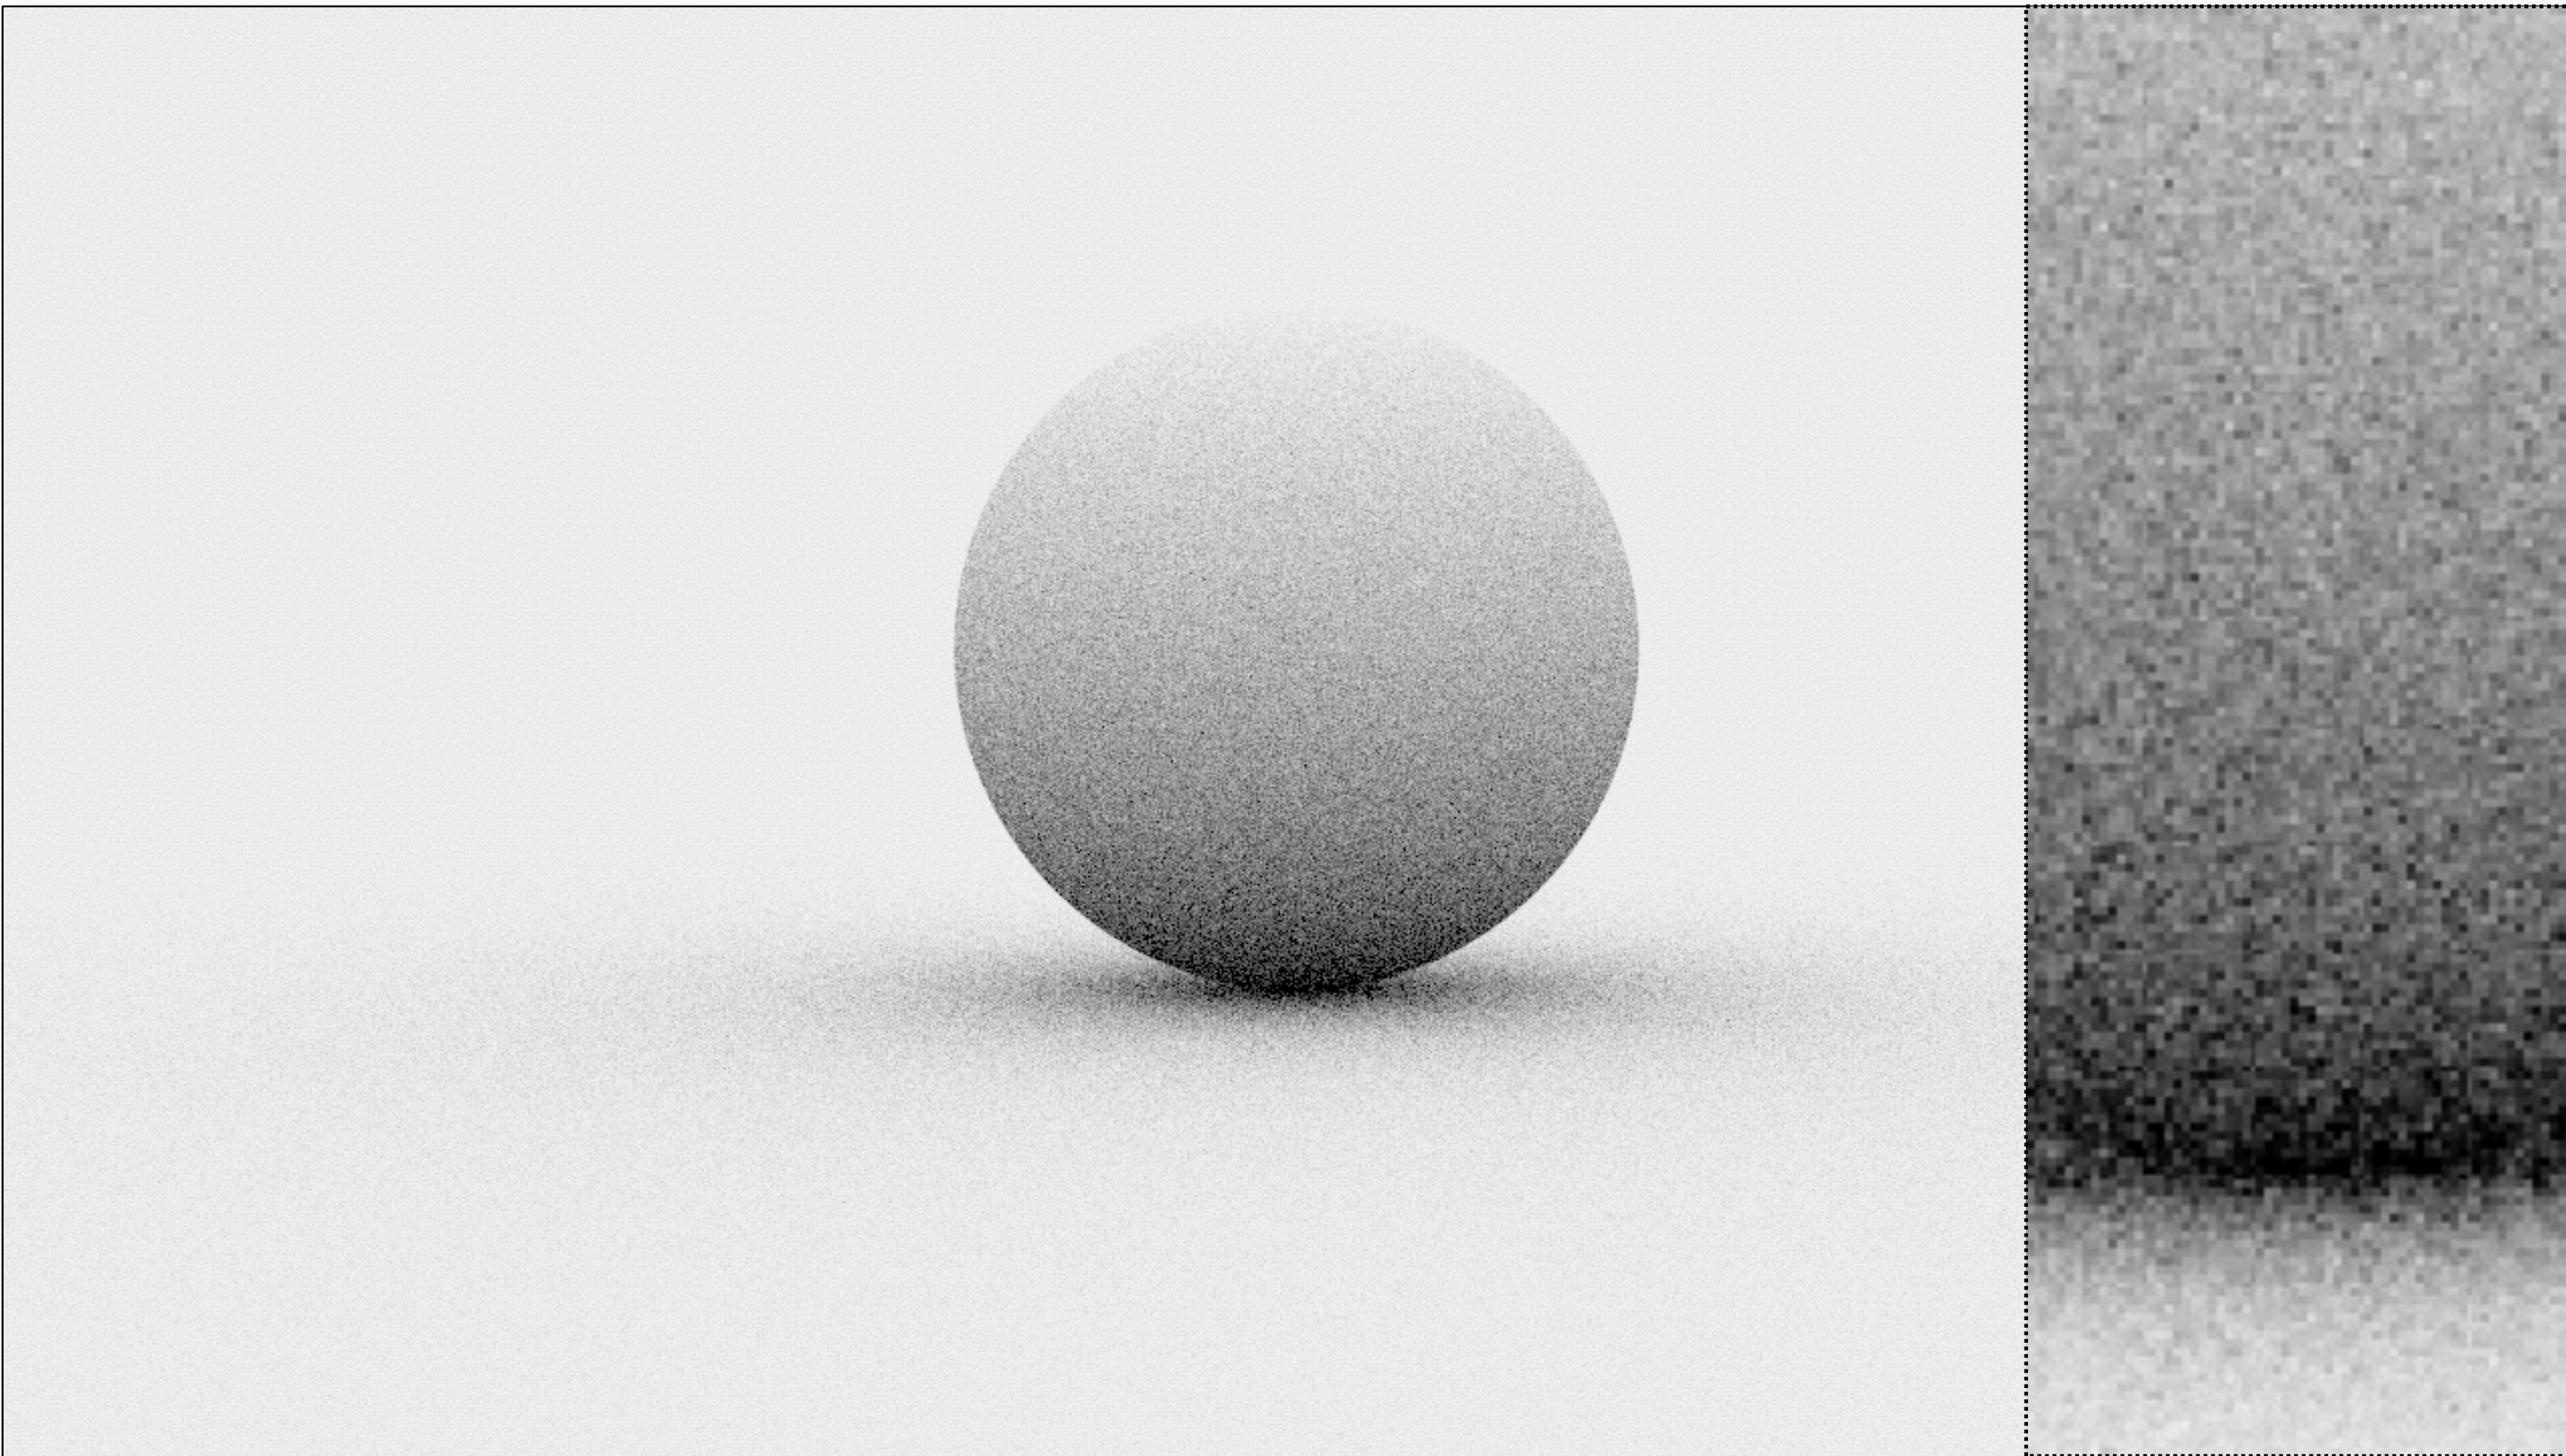
1 sample in each “elementary interval”

The Hammersley Sequence

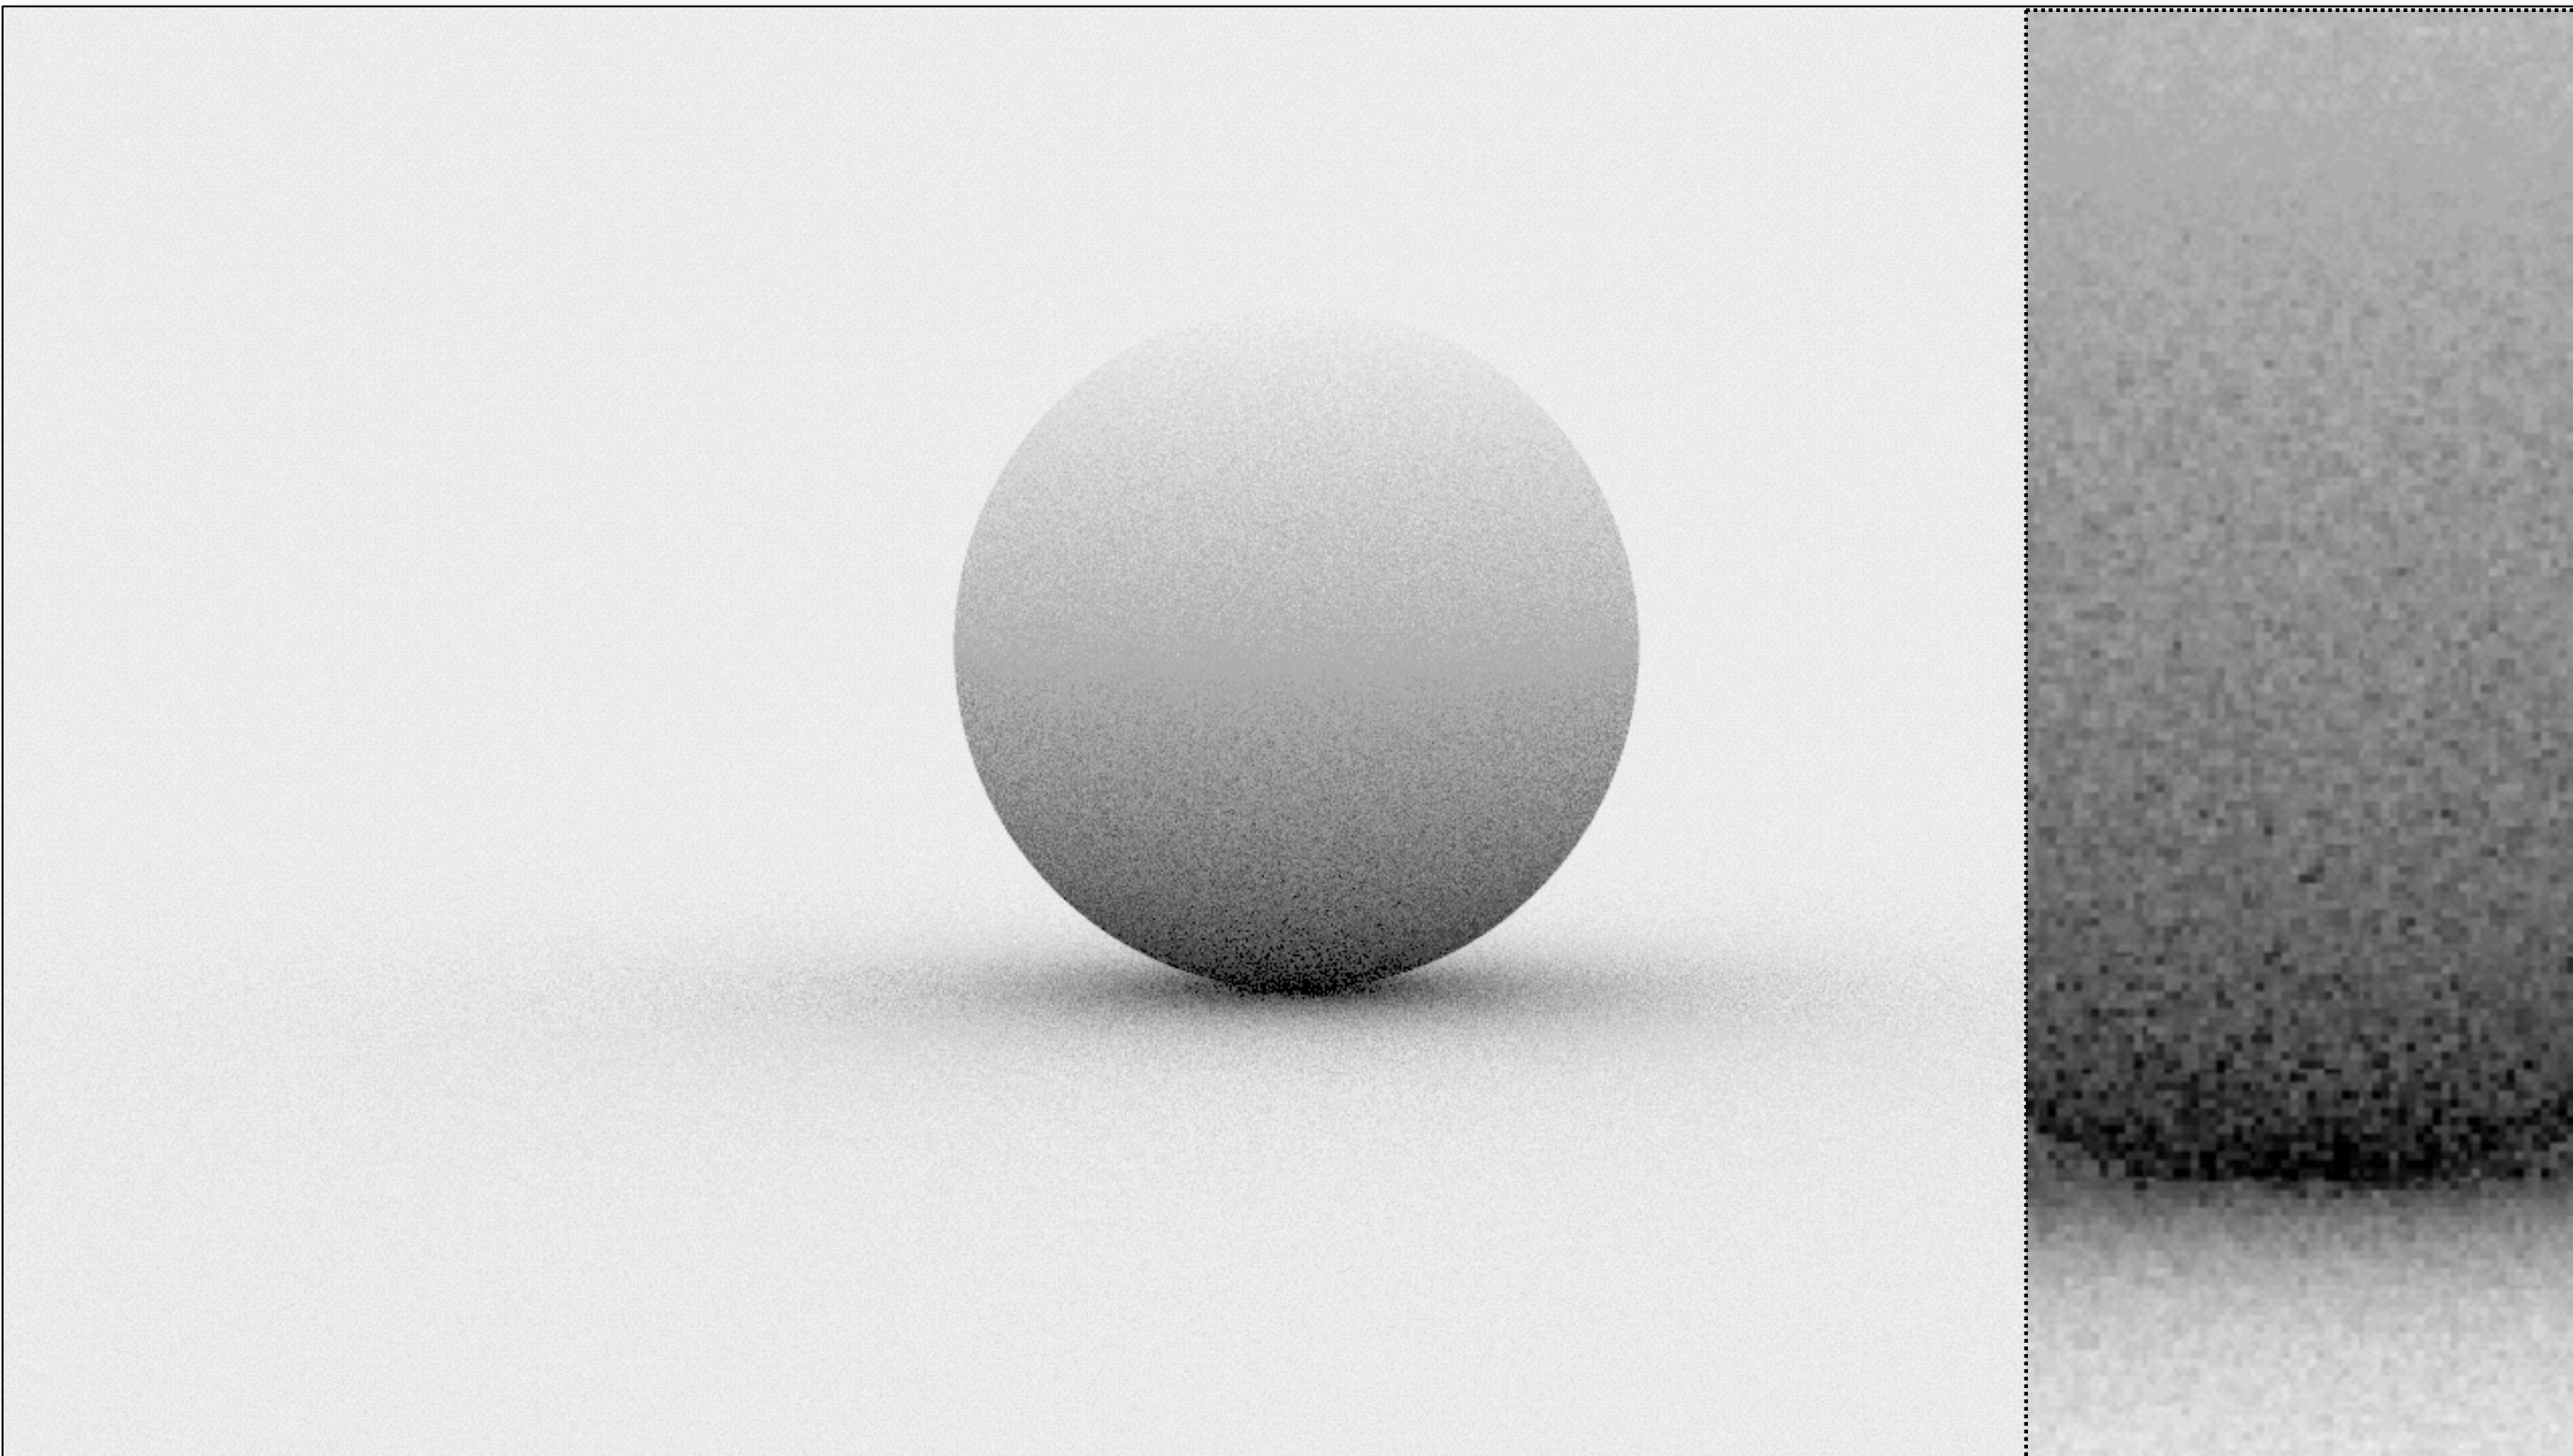


1 sample in each “elementary interval”

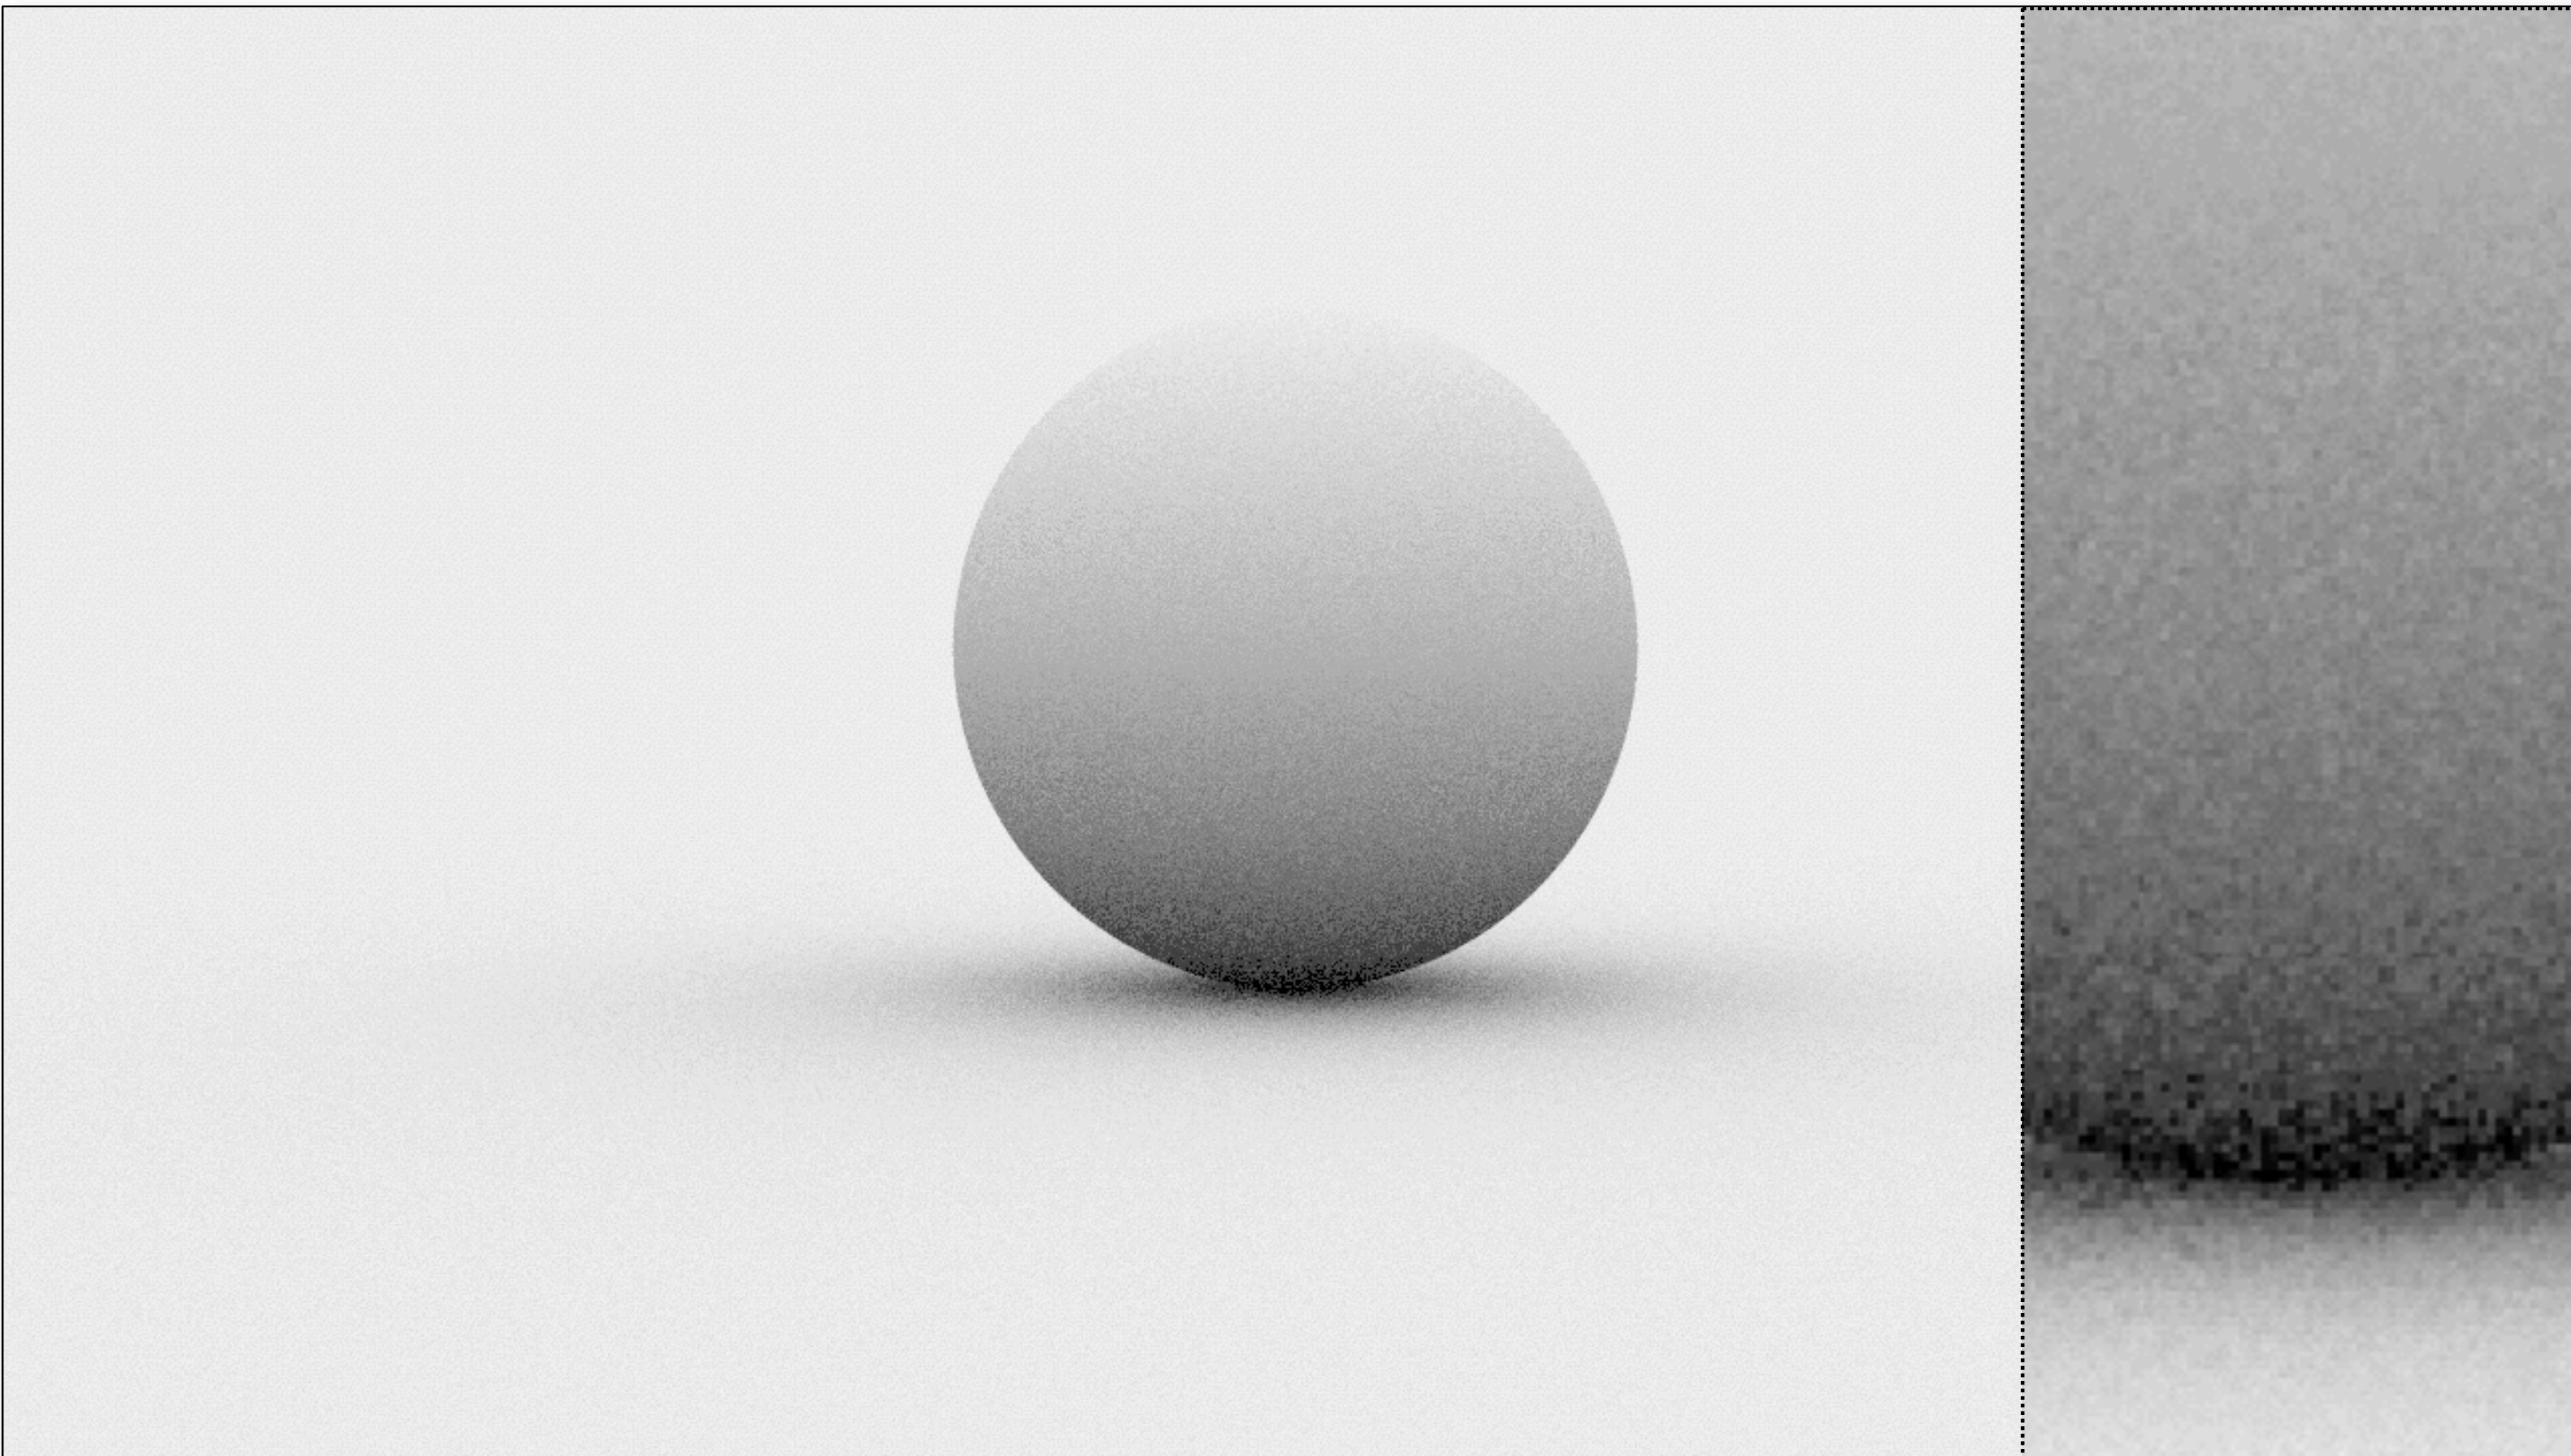
Monte Carlo (16 random samples)



Monte Carlo (16 stratified samples)



Scrambled Low-Discrepancy Sampling



More info on QMC in Rendering

S. Premoze, A. Keller, and M. Raab.

Advanced (Quasi-) Monte Carlo Methods for Image Synthesis. In SIGGRAPH 2012 courses.

How can we predict variance from these?

